

# System Overview

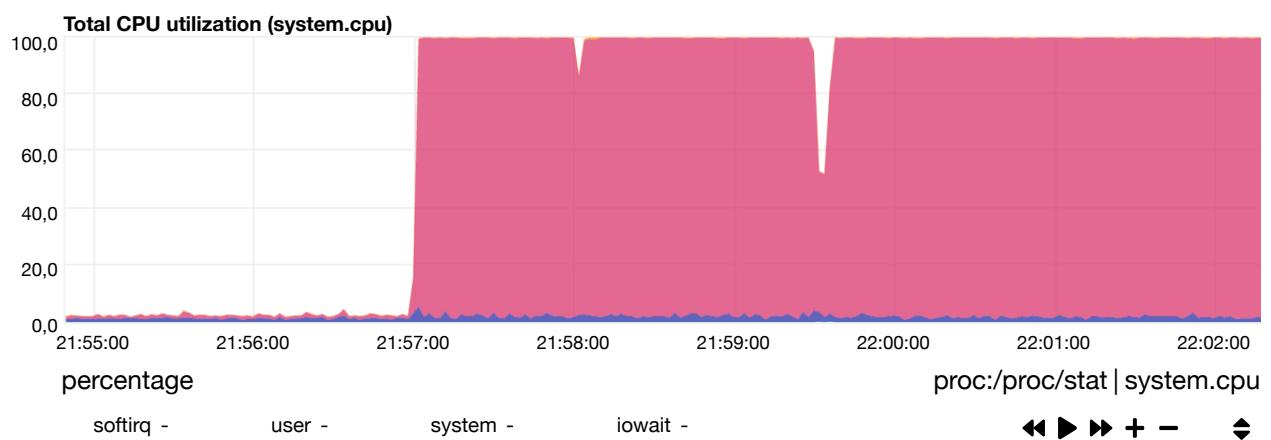
Overview of the key system metrics.

## cpu

Total CPU utilization (all cores). 100% here means there is no CPU idle time at all. You can get per core usage at the CPUs section and per application usage at the Applications Monitoring section.

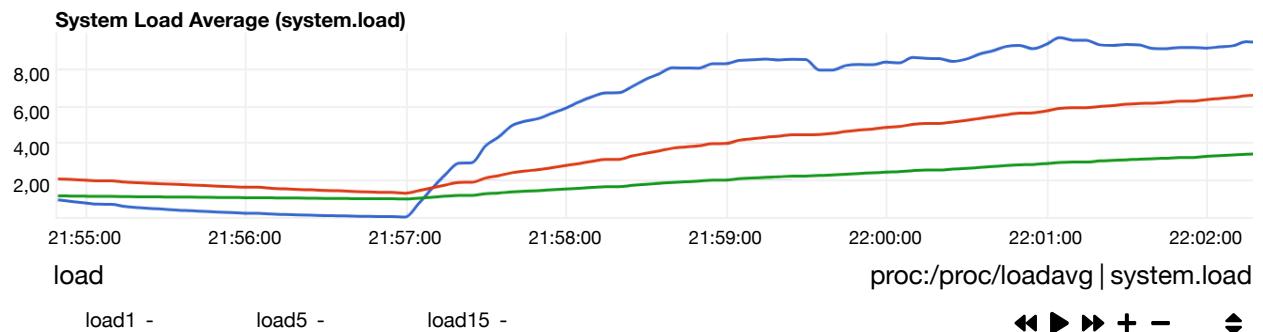
Keep an eye on **iowait**  (  %). If it is constantly high, your disks are a bottleneck and they slow your system down.

An important metric worth monitoring, is **softirq**  (  %). A constantly high percentage of softirq may indicate network driver issues.



## load

Current system load, i.e. the number of processes using CPU or waiting for system resources (usually CPU and disk). The 3 metrics refer to 1, 5 and 15 minute averages. The system calculates this once every 5 seconds. For more information check this wikipedia article ([https://en.wikipedia.org/wiki/Load\\_\(computing\)](https://en.wikipedia.org/wiki/Load_(computing))).

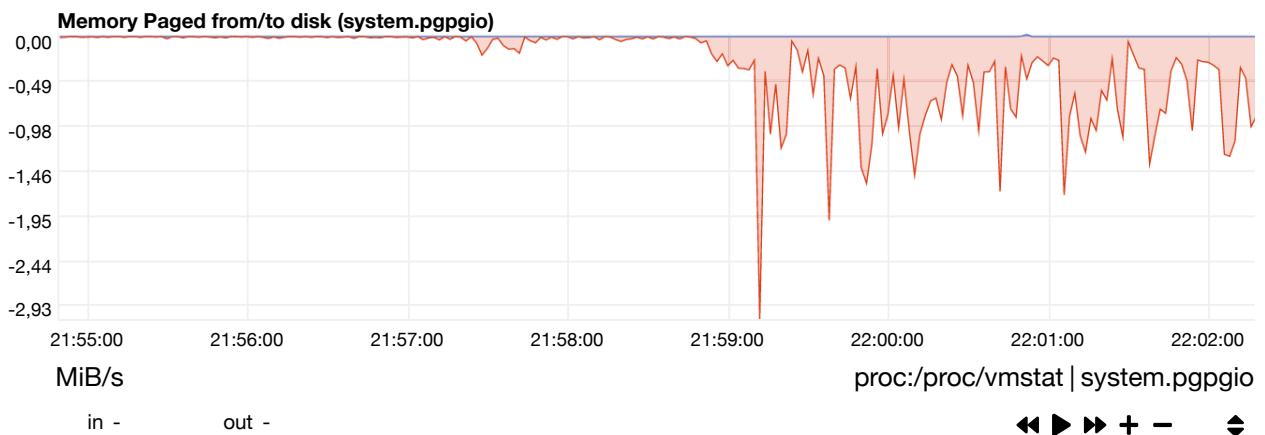


## disk

Total Disk I/O, for all physical disks. You can get detailed information about each disk at the Disks section and per application Disk usage at the Applications Monitoring section. Physical are all the disks that are listed in `/sys/block`, but do not exist in `/sys/devices/virtual/block`.

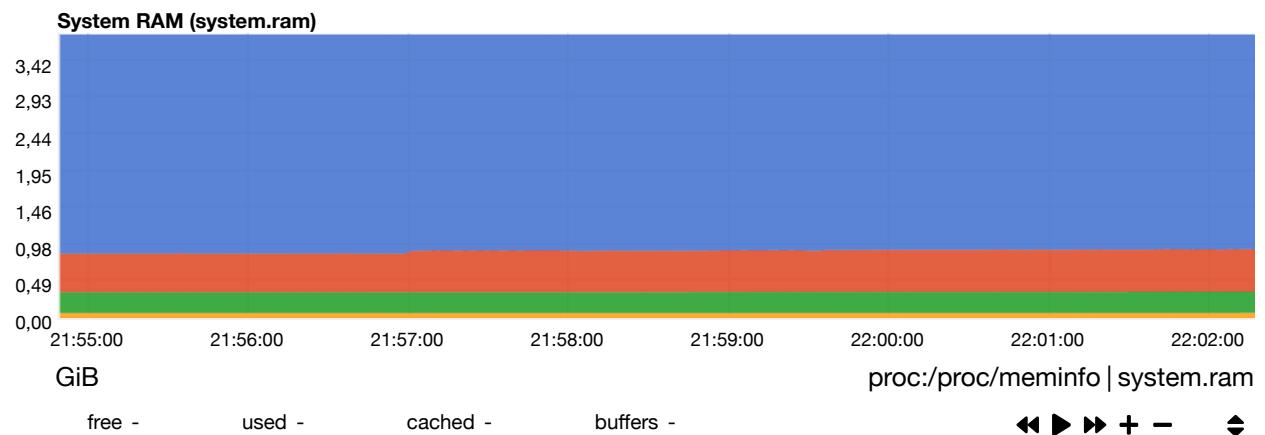


Memory paged from/to disk. This is usually the total disk I/O of the system.



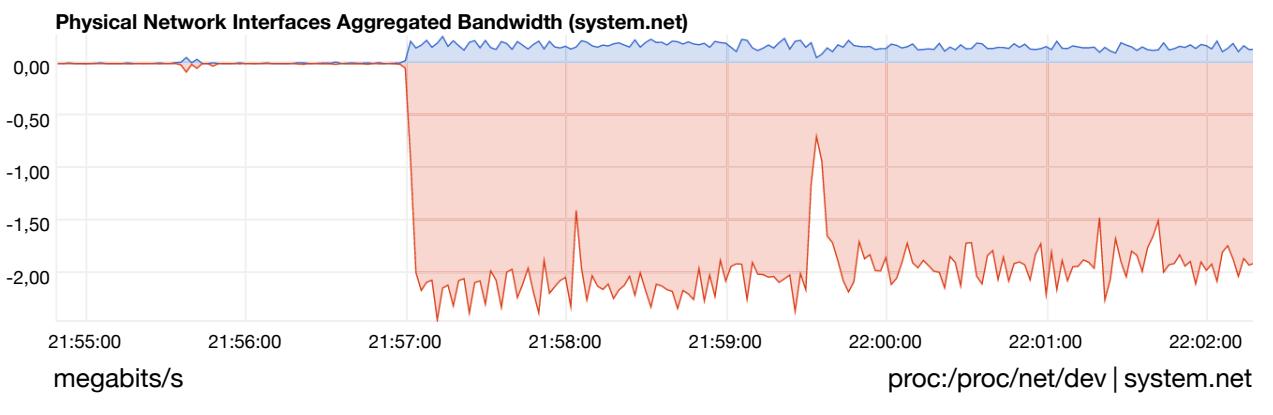
## ram

System Random Access Memory (i.e. physical memory) usage.

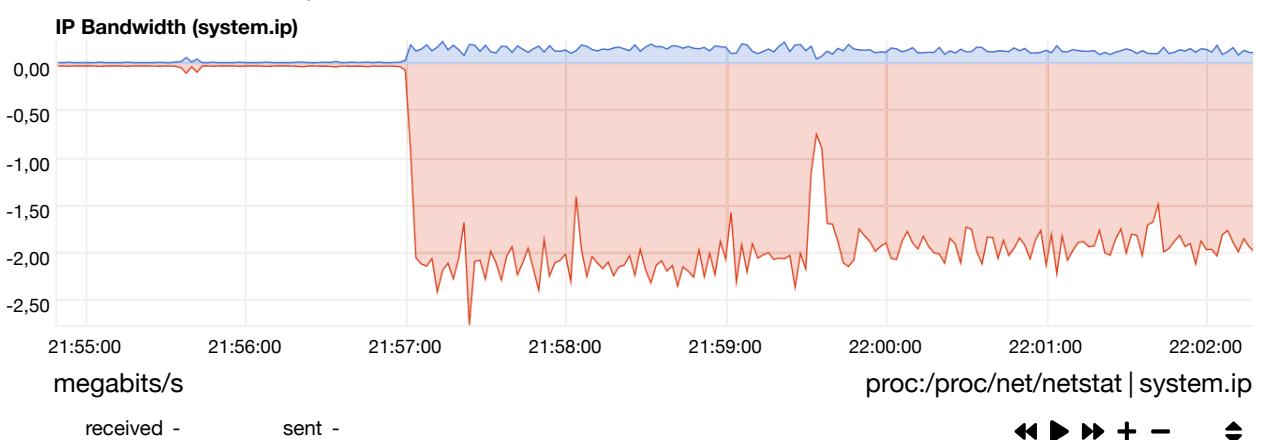


## network

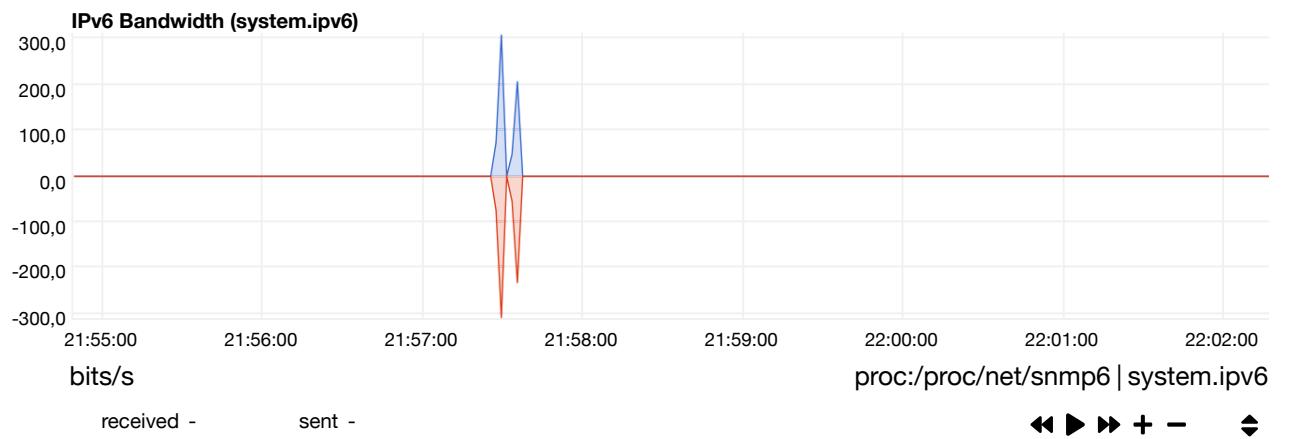
Total bandwidth of all physical network interfaces. This does not include `lo`, VPNs, network bridges, IFB devices, bond interfaces, etc. Only the bandwidth of physical network interfaces is aggregated. Physical are all the network interfaces that are listed in `/proc/net/dev`, but do not exist in `/sys/devices/virtual/net`.



Total IP traffic in the system.



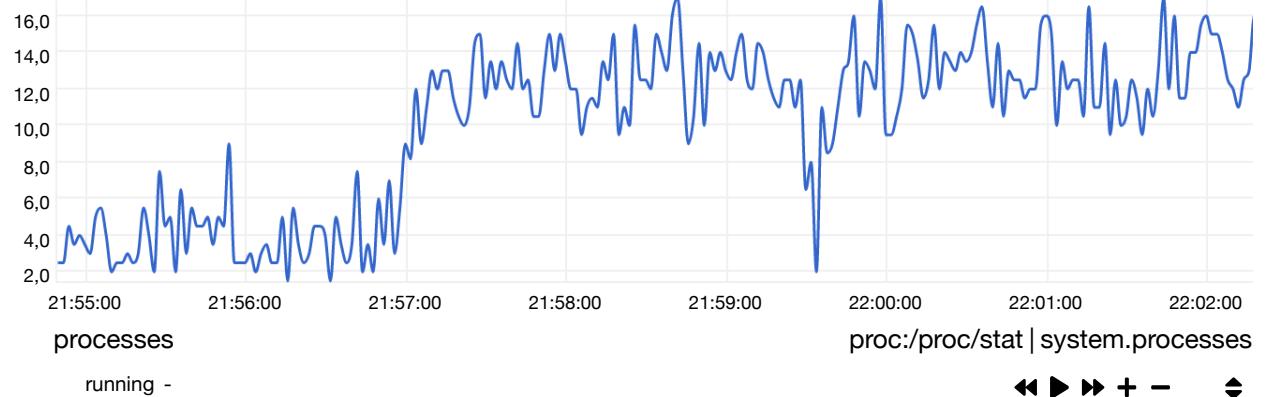
Total IPv6 Traffic.



## processes

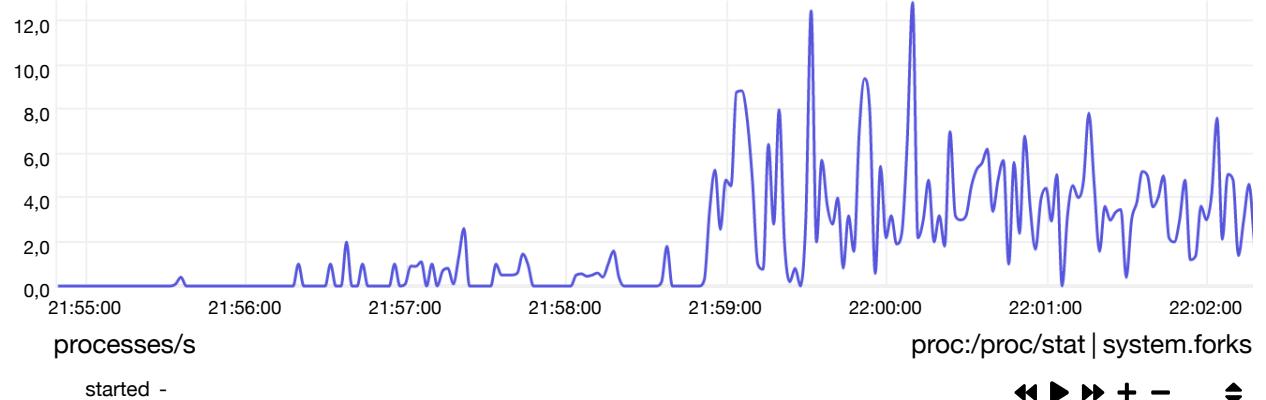
System processes. **Running** are the processes in the CPU. **Blocked** are processes that are willing to enter the CPU, but they cannot, e.g. because they wait for disk activity.

#### System Processes (system.processes)



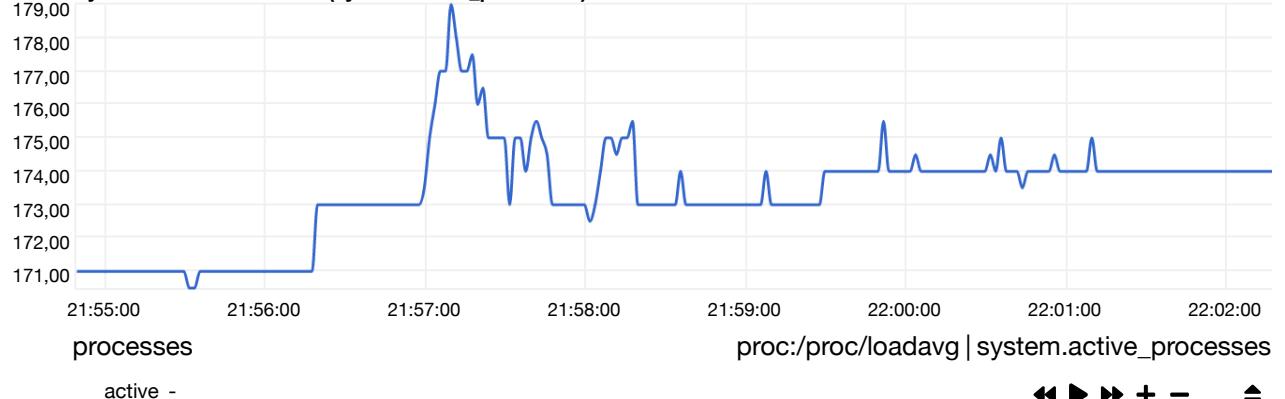
Number of new processes created.

#### Started Processes (system.forks)

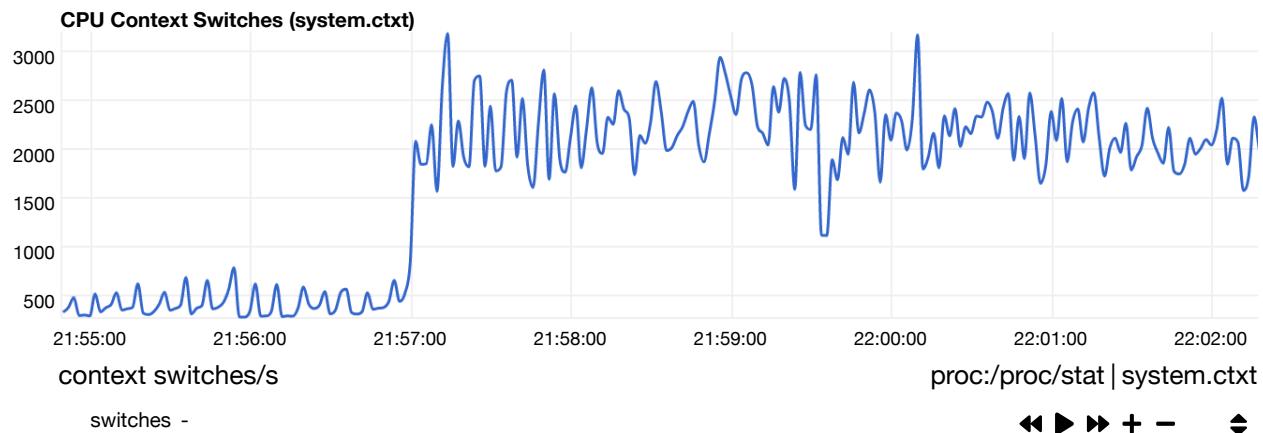


All system processes.

#### System Active Processes (system.active\_processes)

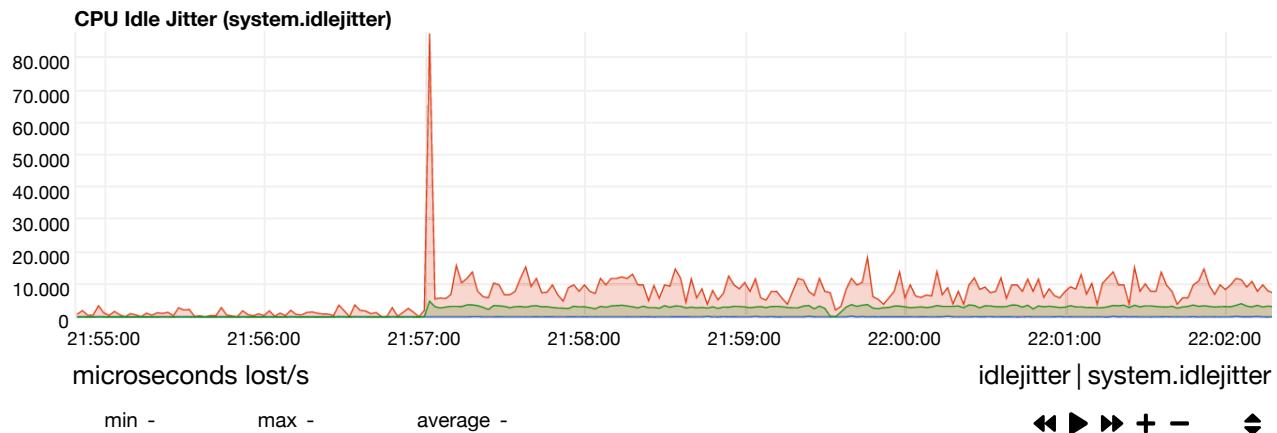


Context Switches ([https://en.wikipedia.org/wiki/Context\\_switch](https://en.wikipedia.org/wiki/Context_switch)), is the switching of the CPU from one process, task or thread to another. If there are many processes or threads willing to execute and very few CPU cores available to handle them, the system is making more context switching to balance the CPU resources among them. The whole process is computationally intensive. The more the context switches, the slower the system gets.



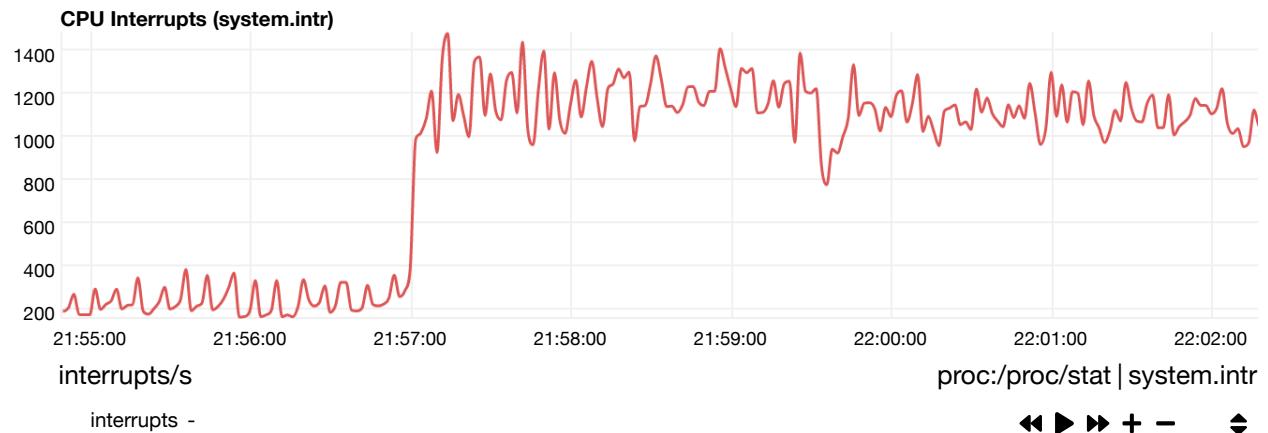
## idlejitter

Idle jitter is calculated by netdata. A thread is spawned that requests to sleep for a few microseconds. When the system wakes it up, it measures how many microseconds have passed. The difference between the requested and the actual duration of the sleep, is the **idle jitter**. This number is useful in real-time environments, where CPU jitter can affect the quality of the service (like VoIP media gateways).

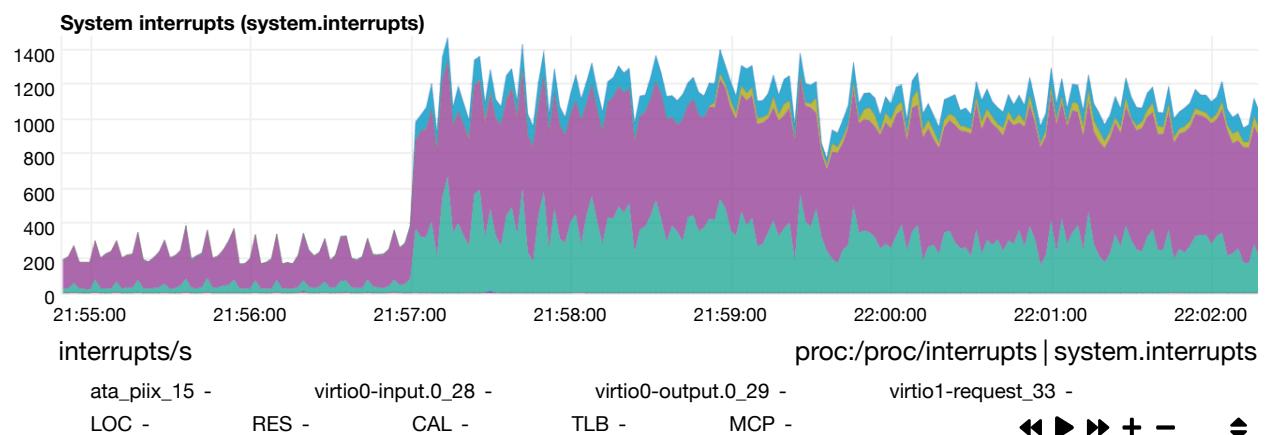


## interrupts

Total number of CPU interrupts. Check `system.interrupts` that gives more detail about each interrupt and also the CPUs section where interrupts are analyzed per CPU core.

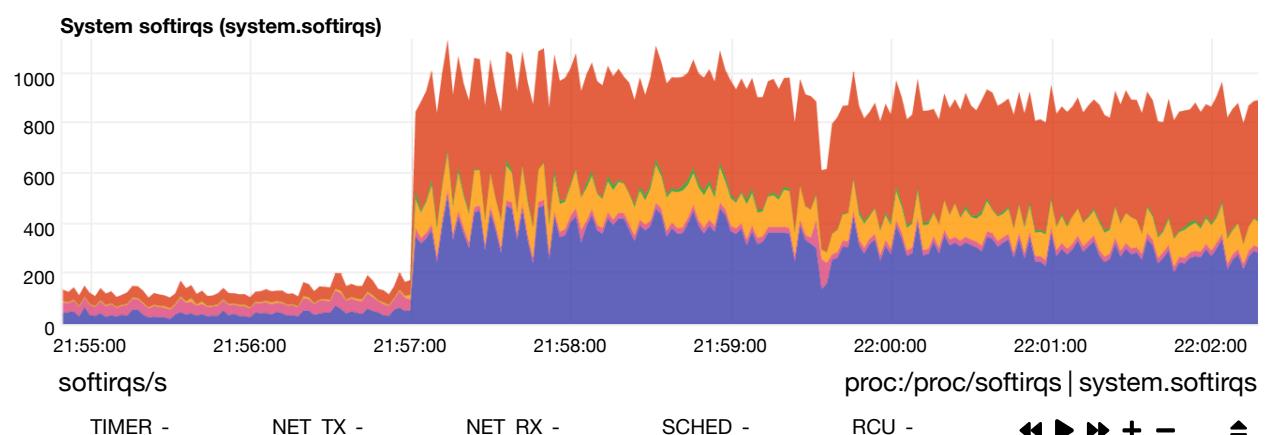


CPU interrupts in detail. At the CPUs section, interrupts are analyzed per CPU core.



## softirqs

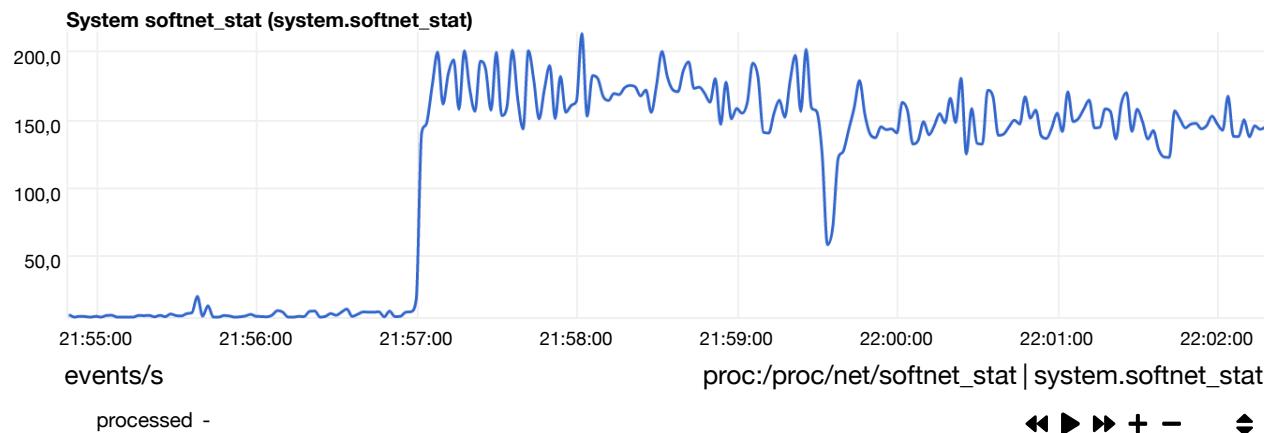
CPU softirqs in detail. At the CPUs section, softirqs are analyzed per CPU core.



## softnet

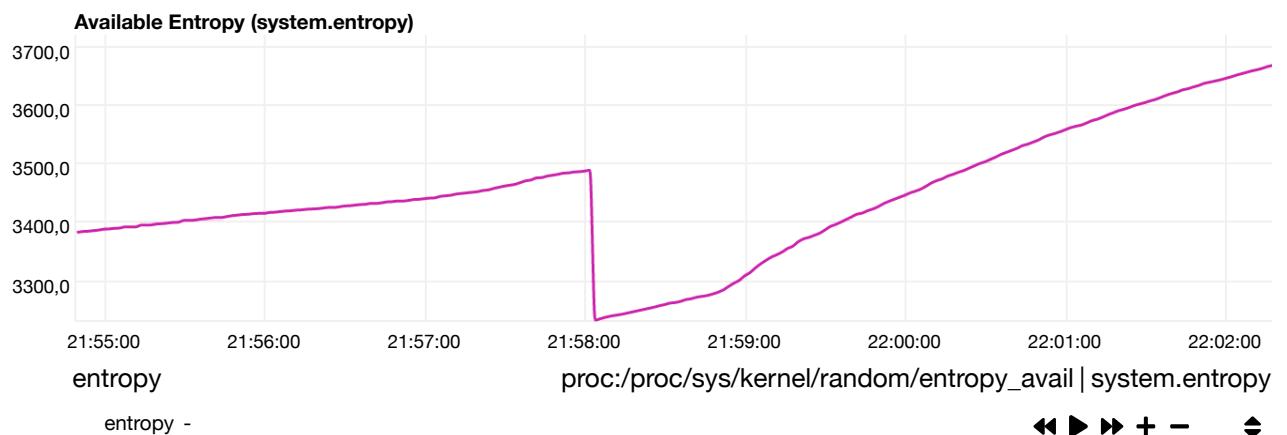
Statistics for CPUs SoftIRQs related to network receive work. Break down per CPU core can be found at CPU / softnet statistics. **processed** states the number of packets processed, **dropped** is the number packets dropped because the network device backlog was full (to fix them on Linux use `sysctl` to increase `net.core.netdev_max_backlog`), **squeezed** is the number of packets dropped

because the network device budget ran out (to fix them on Linux use `sysctl` to increase `net.core.netdev_budget` and/or `net.core.netdev_budget_usecs`). More information about identifying and troubleshooting network driver related issues can be found at Red Hat Enterprise Linux Network Performance Tuning Guide ([https://access.redhat.com/sites/default/files/attachments/20150325\\_network\\_performance\\_tuning.pdf](https://access.redhat.com/sites/default/files/attachments/20150325_network_performance_tuning.pdf)).

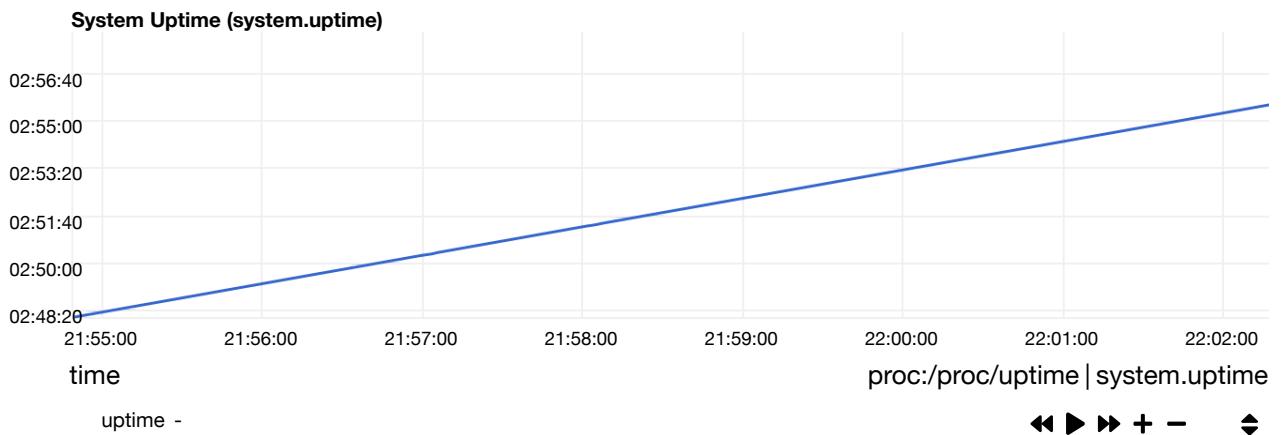


## entropy

Entropy ([https://en.wikipedia.org/wiki/Entropy\\_\(computing\)](https://en.wikipedia.org/wiki/Entropy_(computing))), is a pool of random numbers (`/dev/random` (<https://en.wikipedia.org/wiki//dev/random>)) that is mainly used in cryptography. If the pool of entropy gets empty, processes requiring random numbers may run a lot slower (it depends on the interface each program uses), waiting for the pool to be replenished. Ideally a system with high entropy demands should have a hardware device for that purpose (TPM is one such device). There are also several software-only options you may install, like `haveged` , although these are generally useful only in servers.



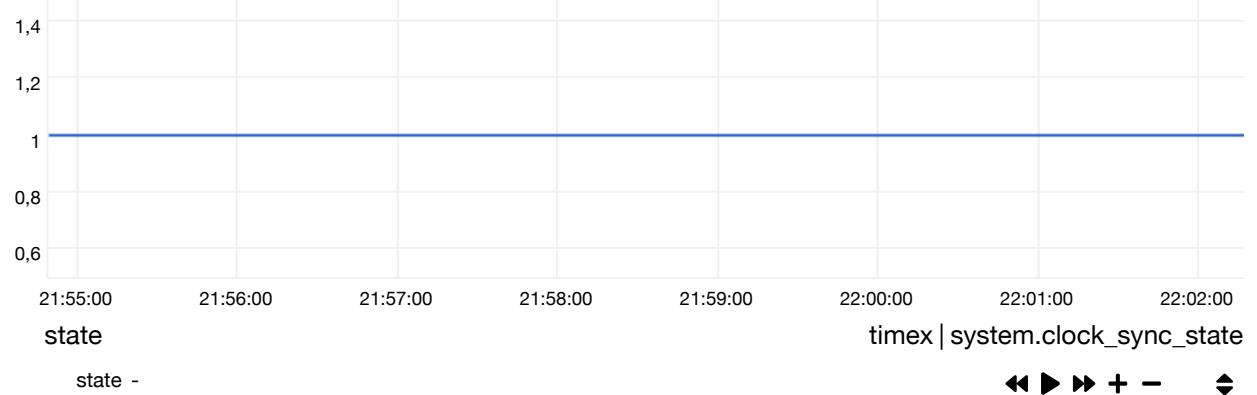
## uptime



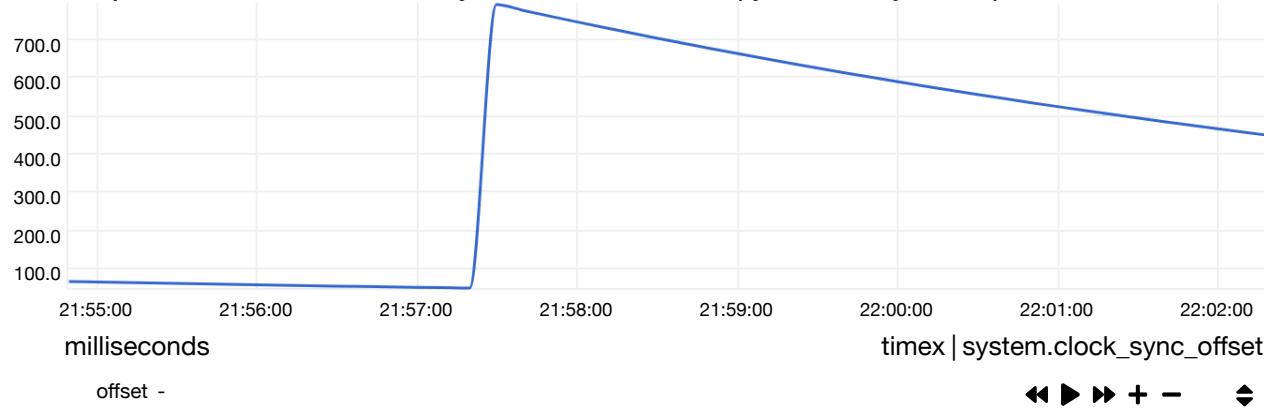
## clock synchronization

State map: 0 - not synchronized, 1 - synchronized

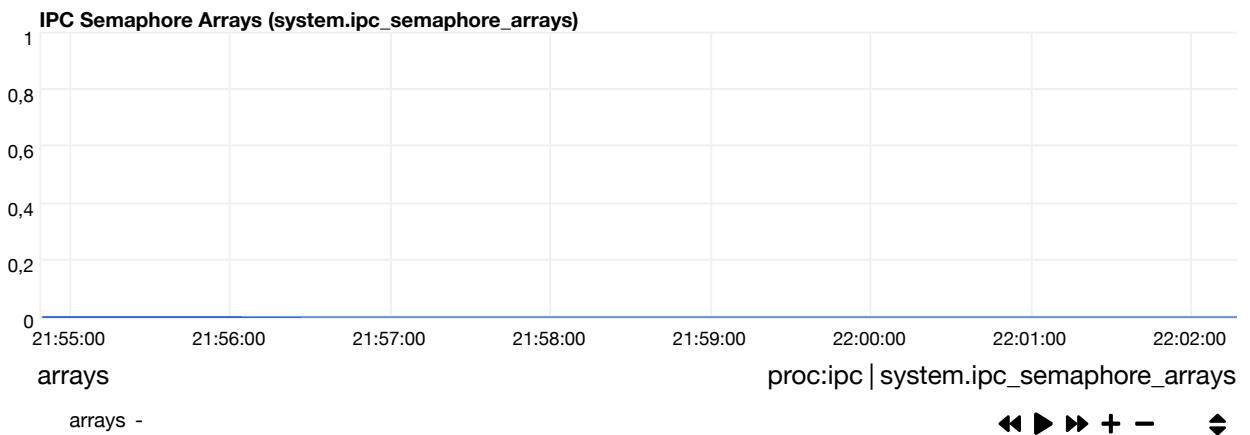
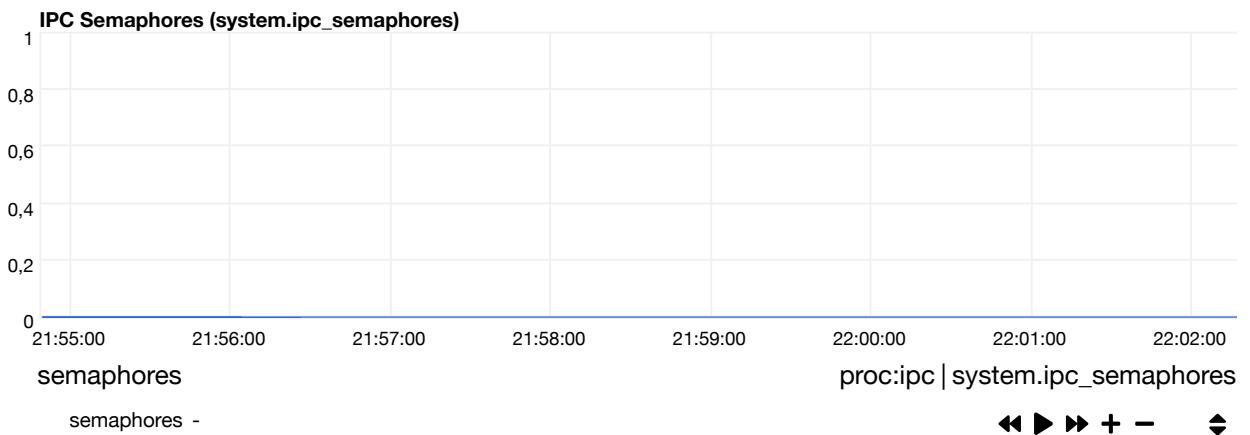
**System Clock Synchronization State (system.clock\_sync\_state)**



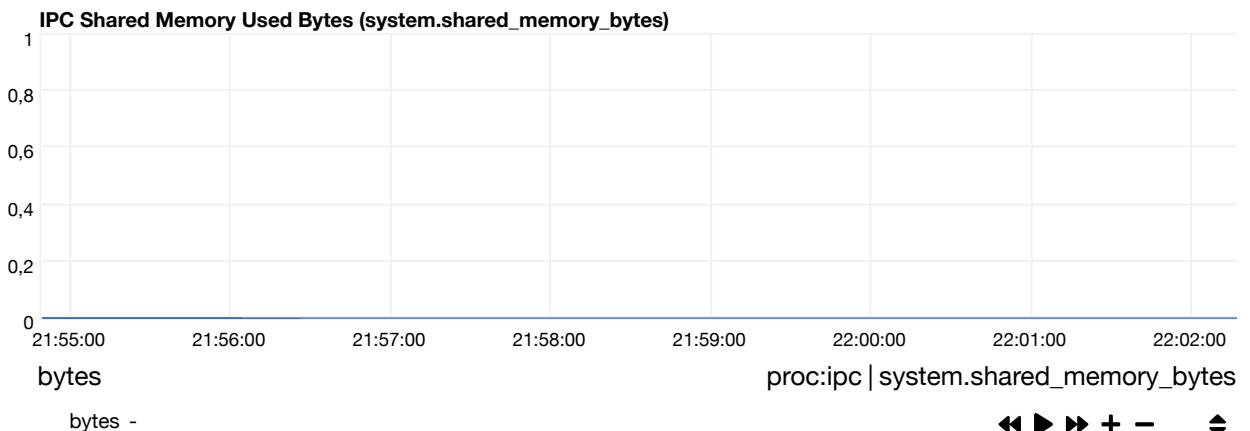
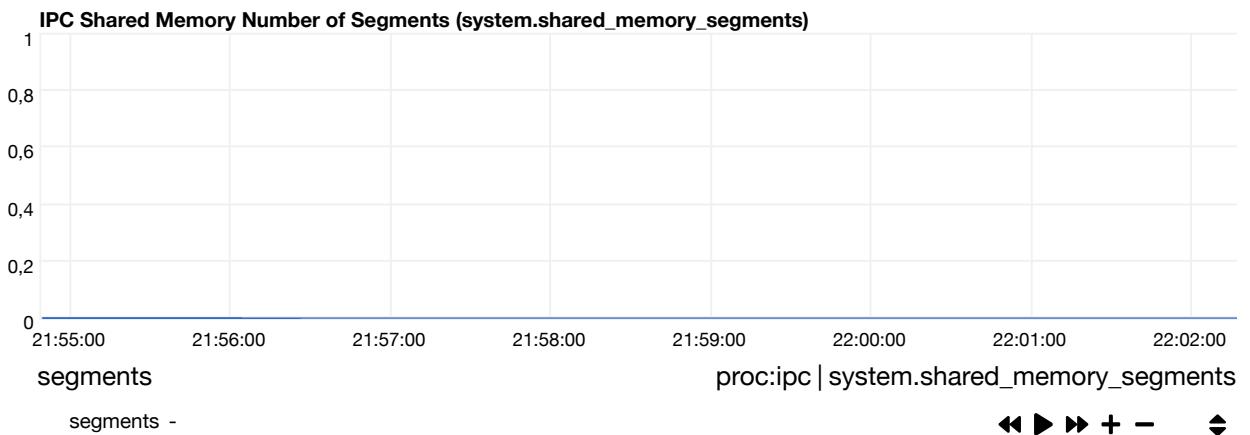
**Computed Time Offset Between Local System and Reference Clock (system.clock\_sync\_offset)**



## ipc semaphores



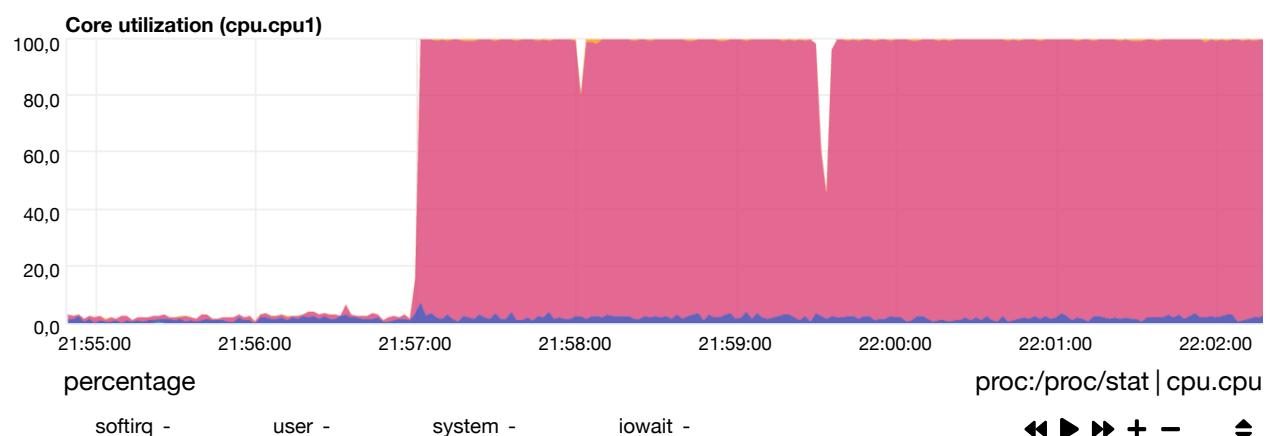
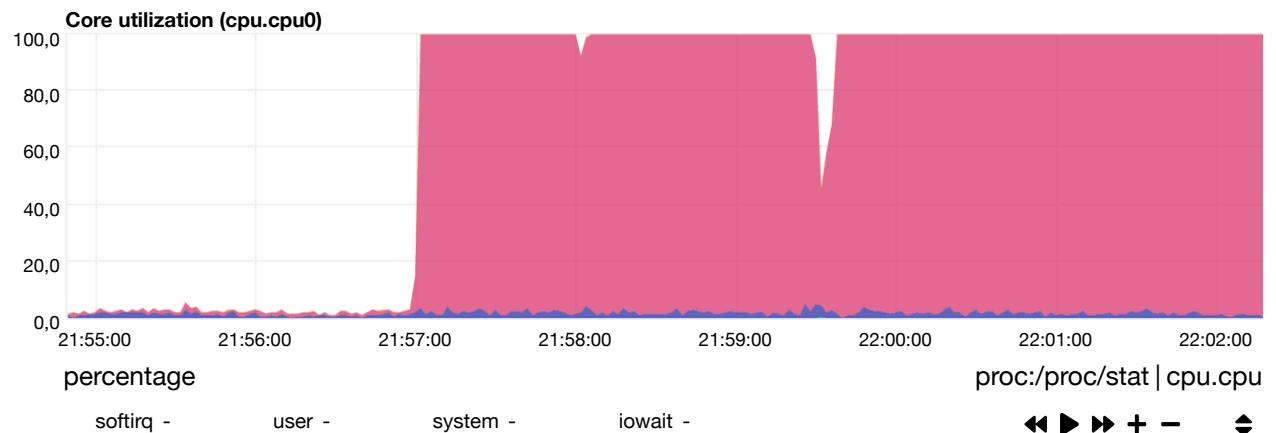
## ipc shared memory



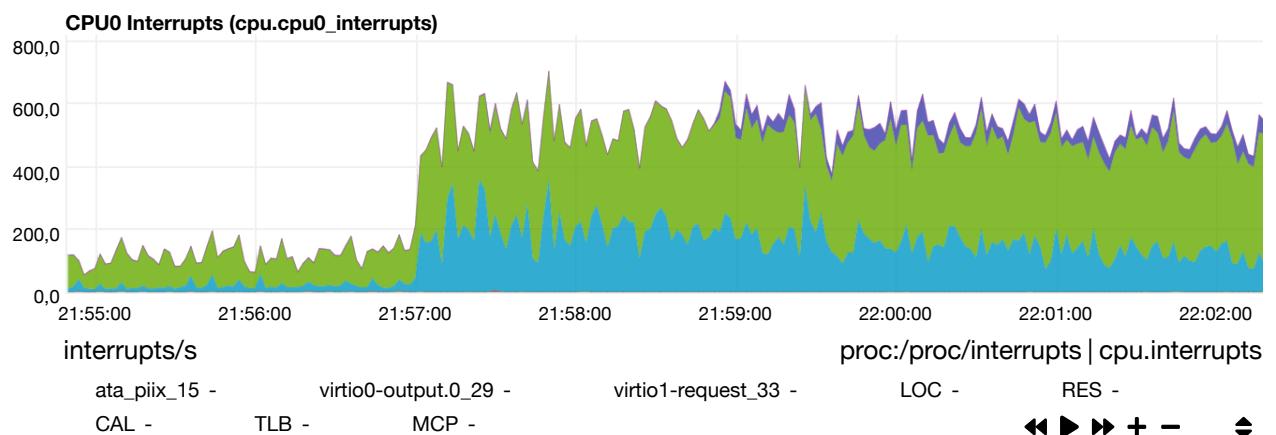
# ⚡ CPUs

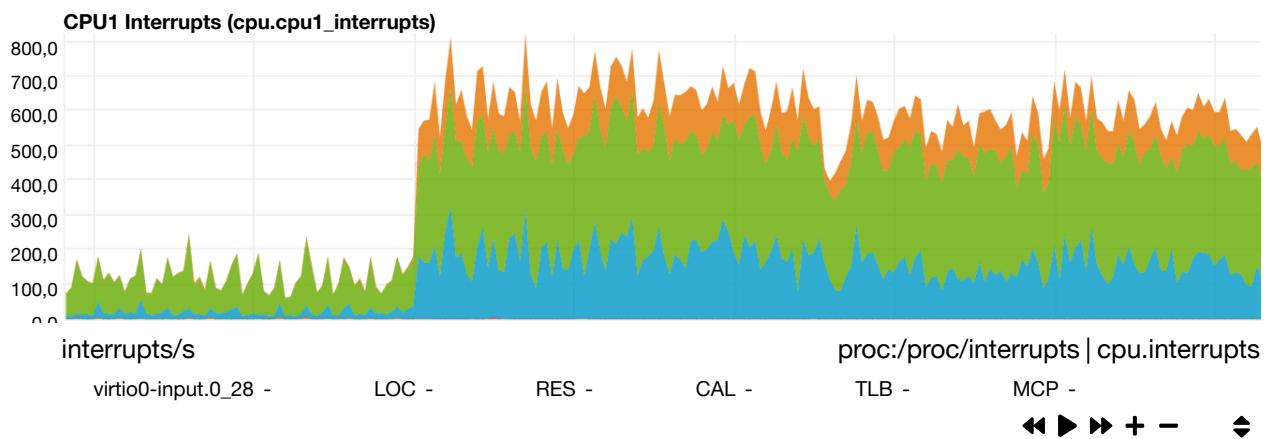
Detailed information for each CPU of the system. A summary of the system for all CPUs can be found at the System Overview section.

## utilization

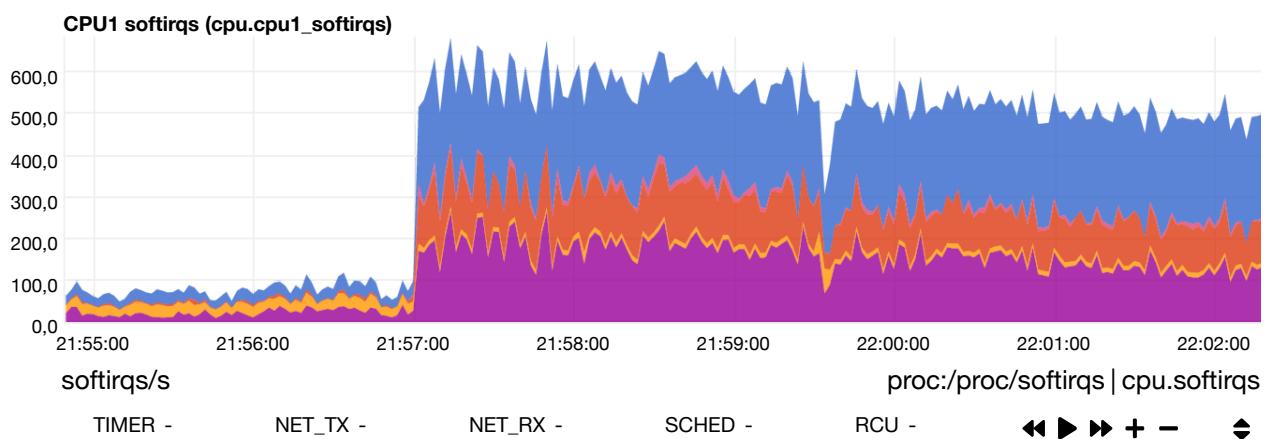
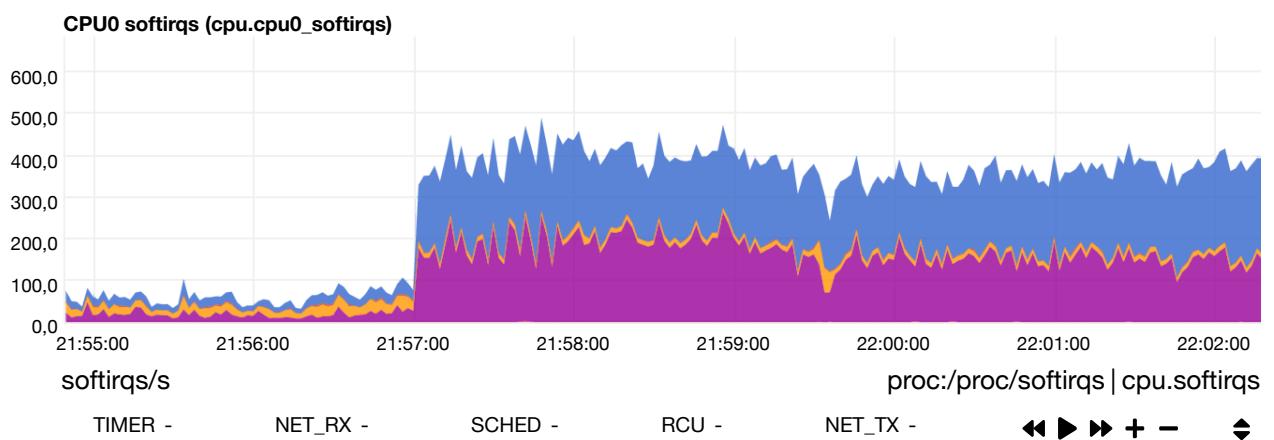


## interrupts



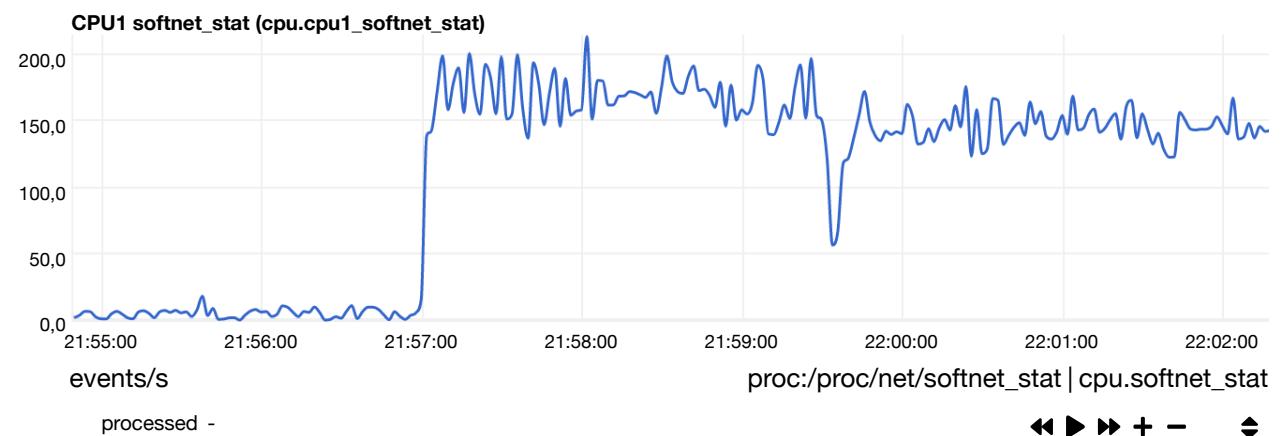
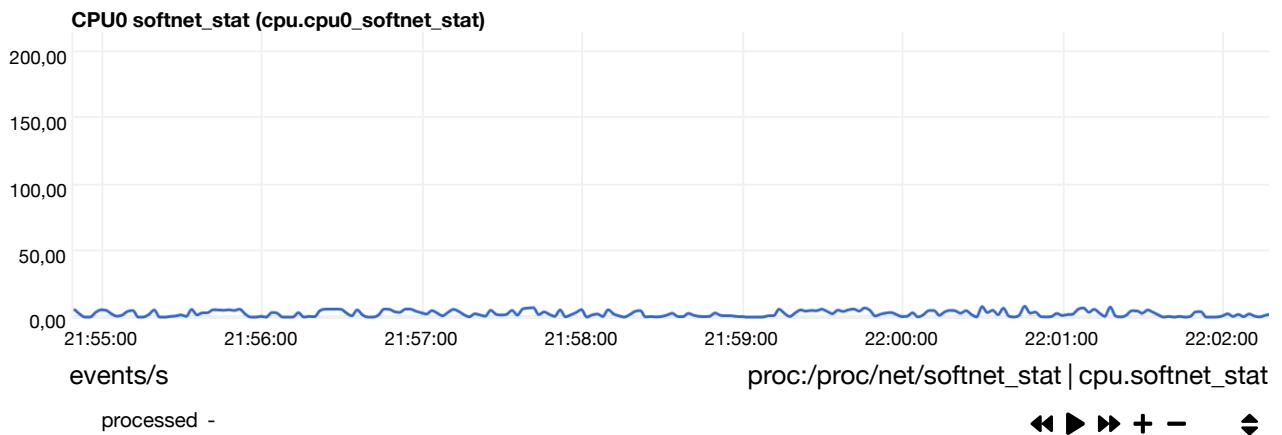


## softirqs

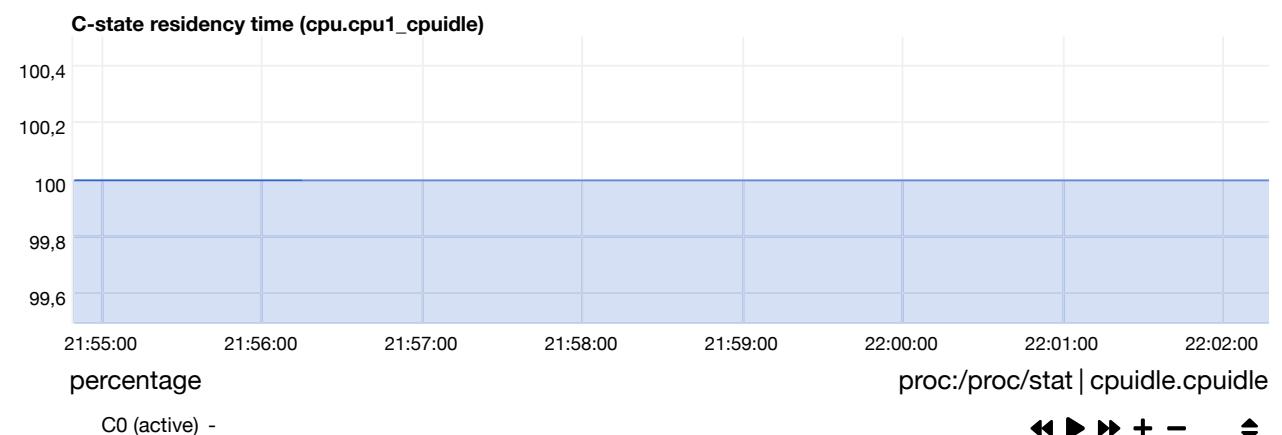
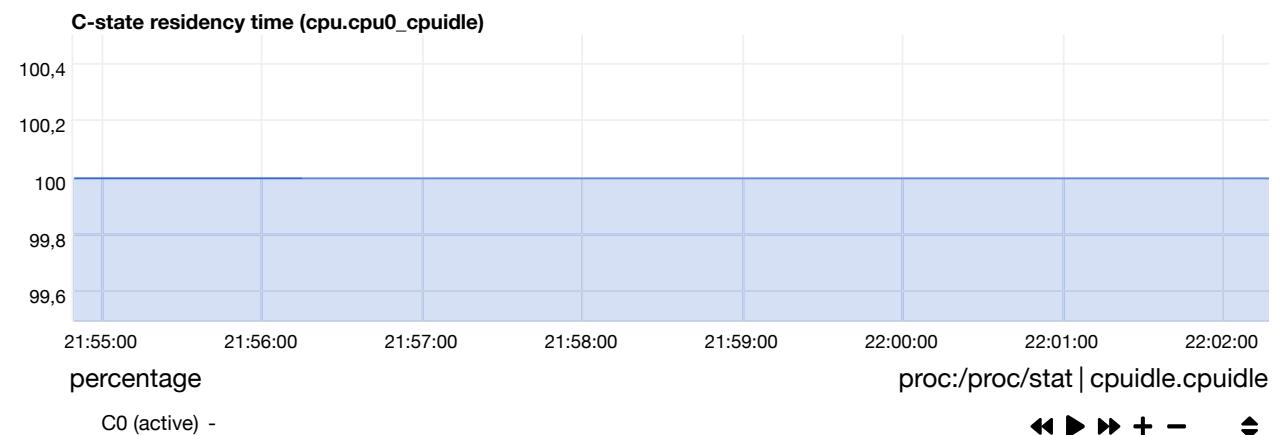


## softnet

Statistics for per CPUs core SoftIRQs related to network receive work. Total for all CPU cores can be found at System / softnet statistics. **processed** states the number of packets processed, **dropped** is the number packets dropped because the network device backlog was full (to fix them on Linux use `sysctl` to increase `net.core.netdev_max_backlog`), **squeezed** is the number of packets dropped because the network device budget ran out (to fix them on Linux use `sysctl` to increase `net.core.netdev_budget` and/or `net.core.netdev_budget_usecs`). More information about identifying and troubleshooting network driver related issues can be found at Red Hat Enterprise Linux Network Performance Tuning Guide ([https://access.redhat.com/sites/default/files/attachments/20150325\\_network\\_performance\\_tuning.pdf](https://access.redhat.com/sites/default/files/attachments/20150325_network_performance_tuning.pdf)).



## cpuidle

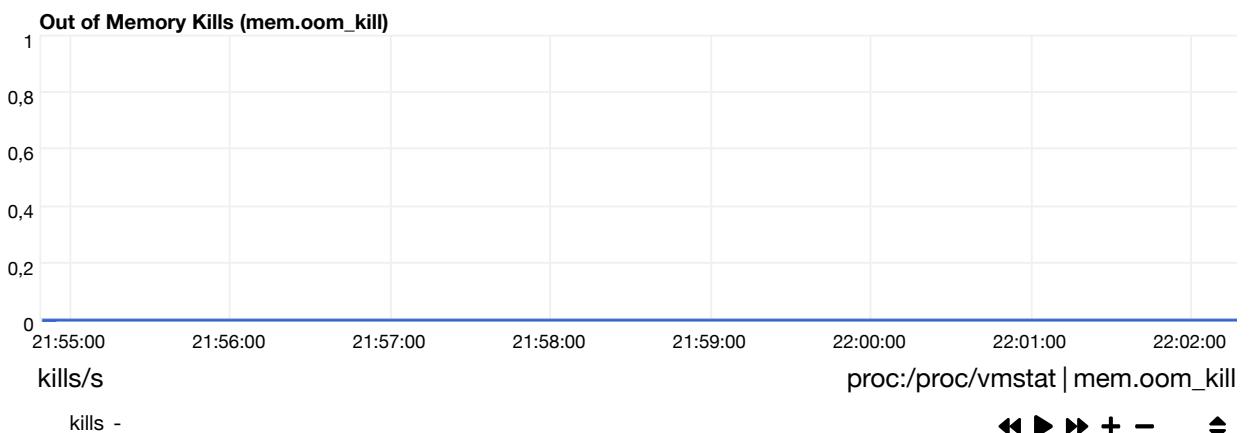
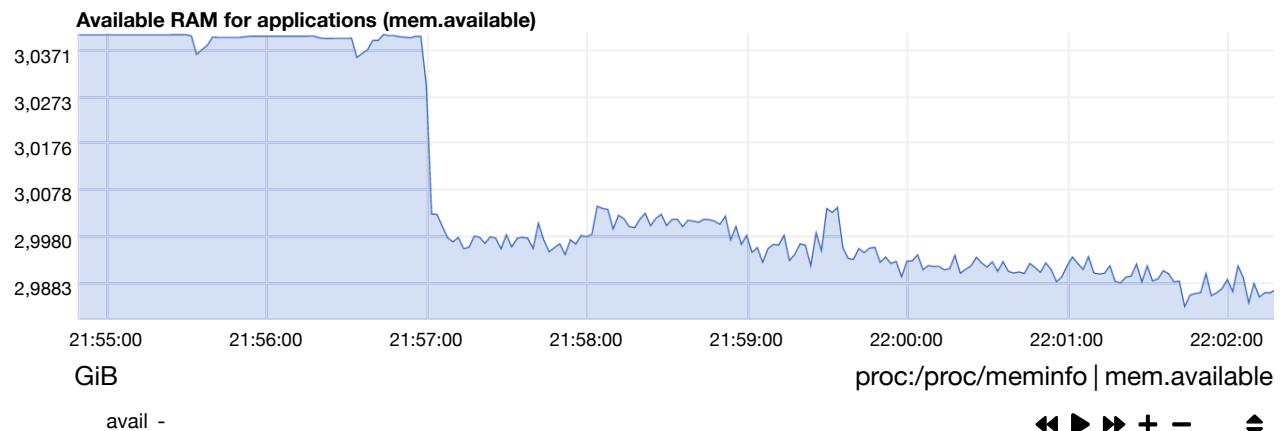


# Memory

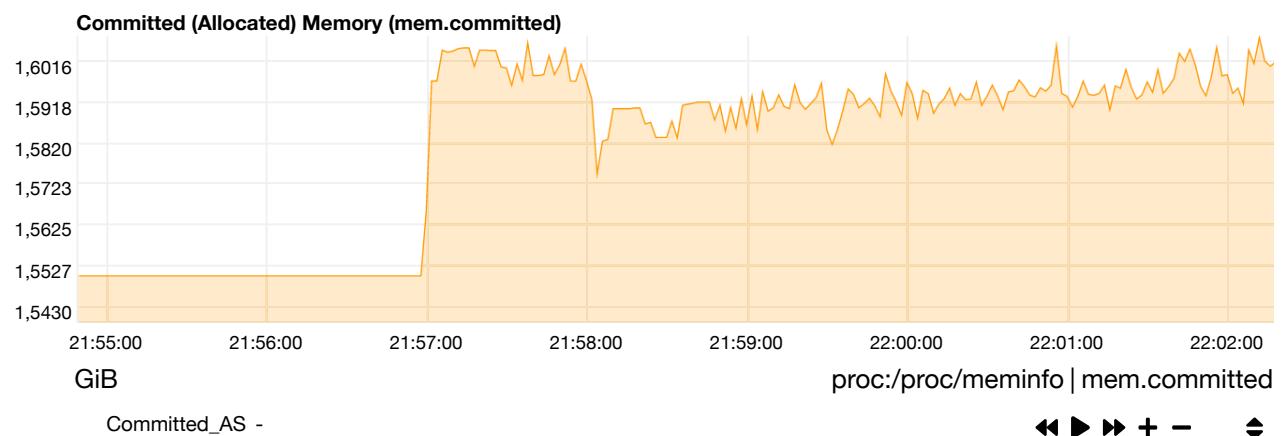
Detailed information about the memory management of the system.

## system

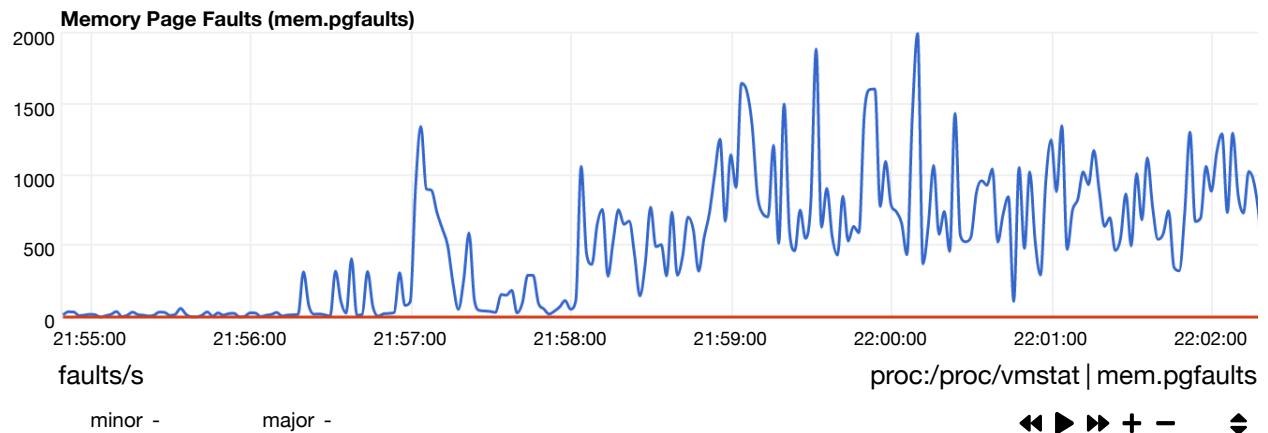
Available Memory is estimated by the kernel, as the amount of RAM that can be used by userspace processes, without causing swapping.



Committed Memory, is the sum of all memory which has been allocated by processes.

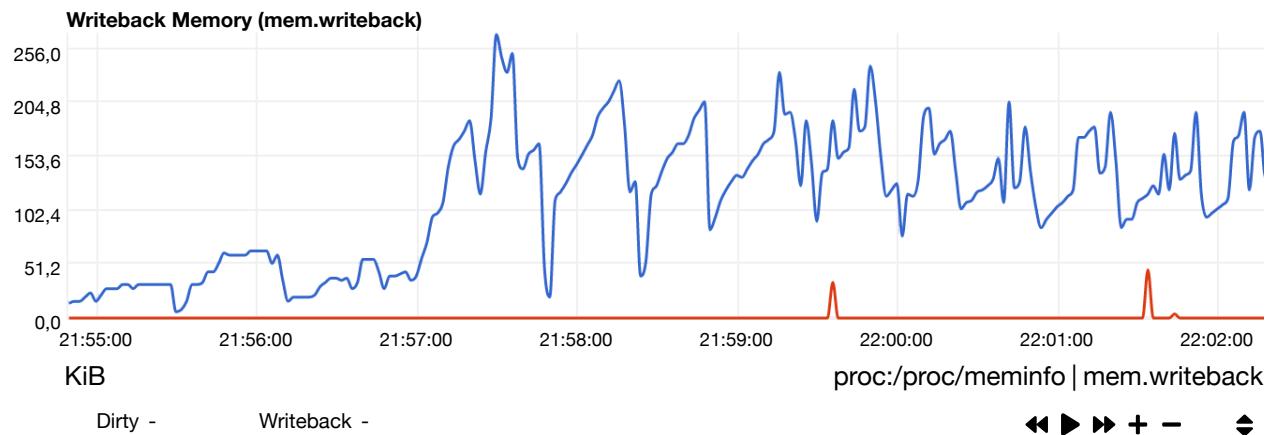


A page fault ([https://en.wikipedia.org/wiki/Page\\_fault](https://en.wikipedia.org/wiki/Page_fault)) is a type of interrupt, called trap, raised by computer hardware when a running program accesses a memory page that is mapped into the virtual address space, but not actually loaded into main memory. If the page is loaded in memory at the time the fault is generated, but is not marked in the memory management unit as being loaded in memory, then it is called a **minor** or soft page fault. A **major** page fault is generated when the system needs to load the memory page from disk or swap memory.

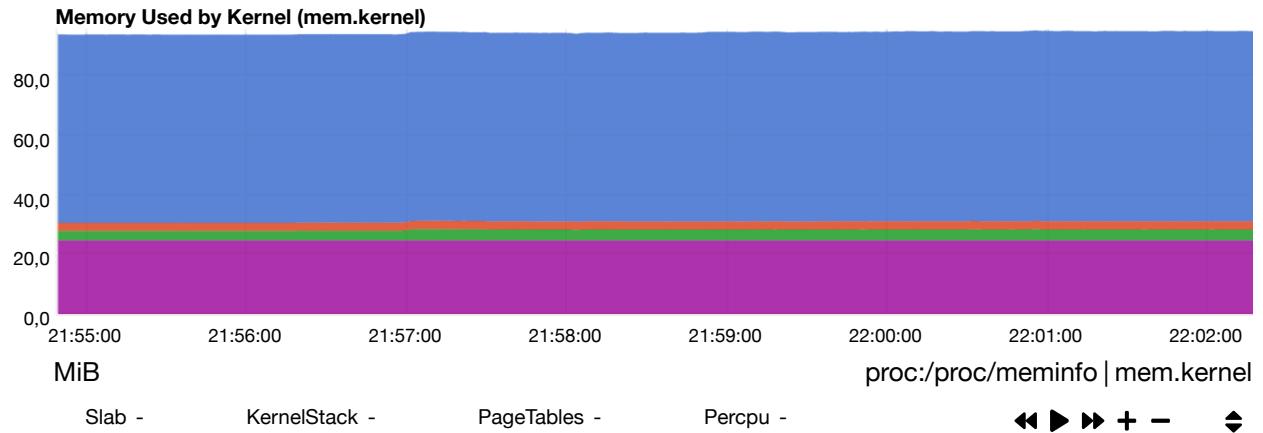


## kernel

**Dirty** is the amount of memory waiting to be written to disk. **Writeback** is how much memory is actively being written to disk.

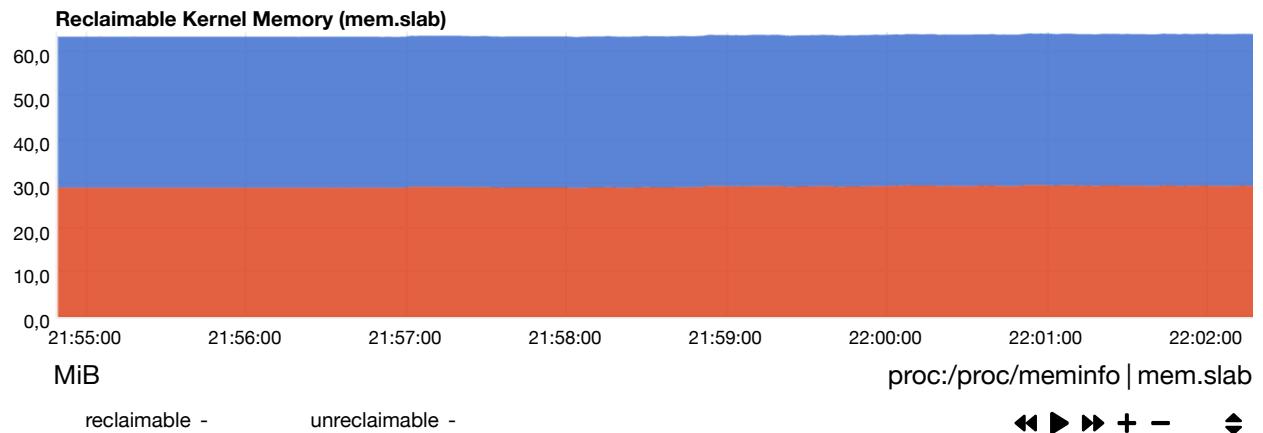


The total amount of memory being used by the kernel. **Slab** is the amount of memory used by the kernel to cache data structures for its own use. **KernelStack** is the amount of memory allocated for each task done by the kernel. **PageTables** is the amount of memory dedicated to the lowest level of page tables (A page table is used to turn a virtual address into a physical memory address). **VmallocUsed** is the amount of memory being used as virtual address space.



## slab

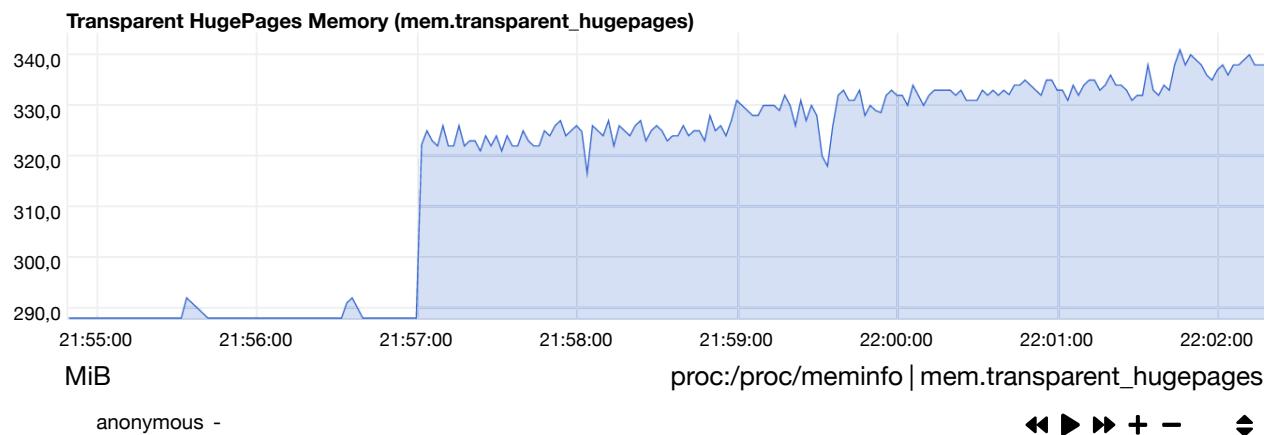
**Reclaimable** is the amount of memory which the kernel can reuse. **Unreclaimable** can not be reused even when the kernel is lacking memory.



## hugepages

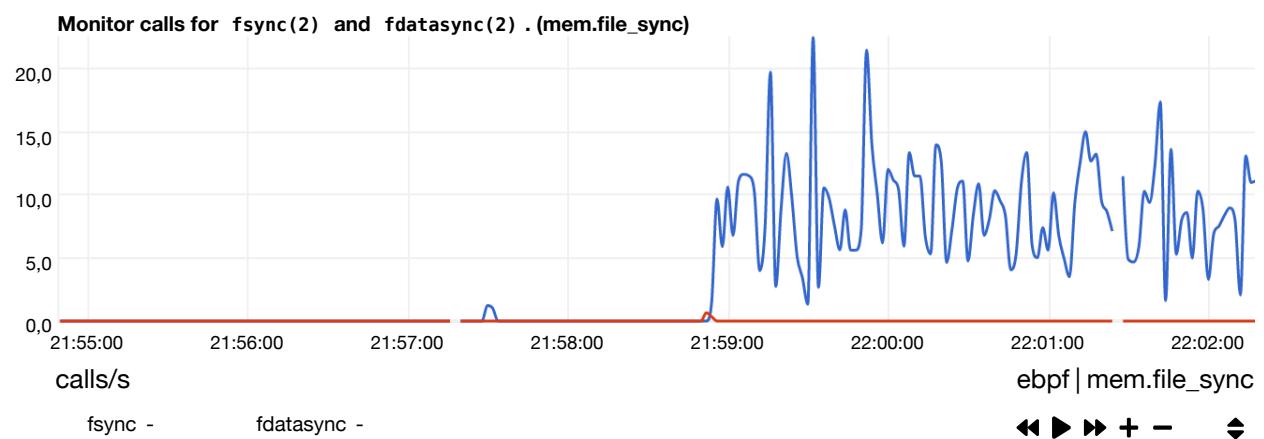
Hugepages is a feature that allows the kernel to utilize the multiple page size capabilities of modern hardware architectures. The kernel creates multiple pages of virtual memory, mapped from both physical RAM and swap. There is a mechanism in the CPU architecture called "Translation Lookaside Buffers" (TLB) to manage the mapping of virtual memory pages to actual physical memory addresses. The TLB is a limited hardware resource, so utilizing a large amount of physical memory with the default page size consumes the TLB and adds processing overhead. By utilizing Huge Pages, the kernel is able to create pages of much larger sizes, each page consuming a single resource in the TLB. Huge Pages are pinned to physical RAM and cannot be swapped/paged out.

Transparent HugePages (THP) is backing virtual memory with huge pages, supporting automatic promotion and demotion of page sizes. It works for all applications for anonymous memory mappings and tmpfs/shmem.

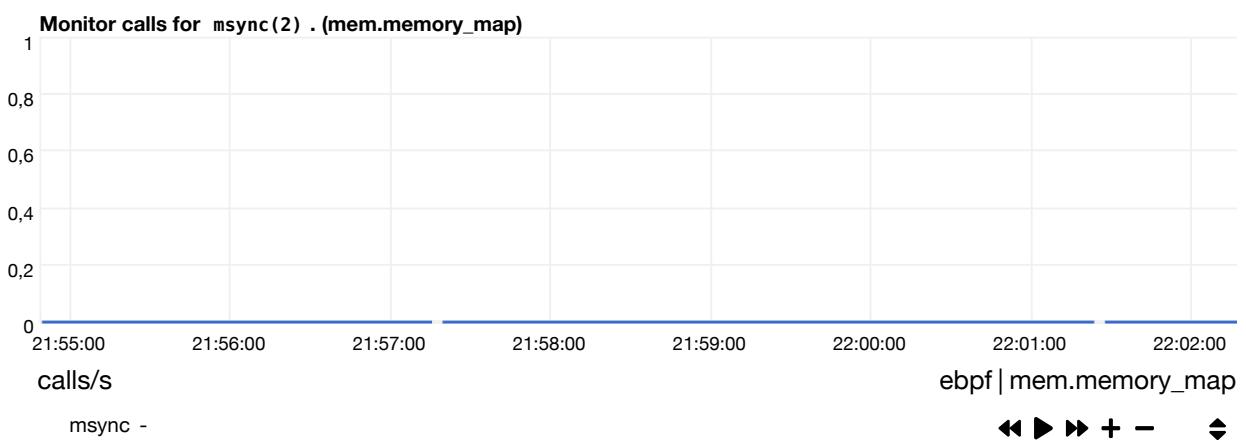


## synchronization (eBPF)

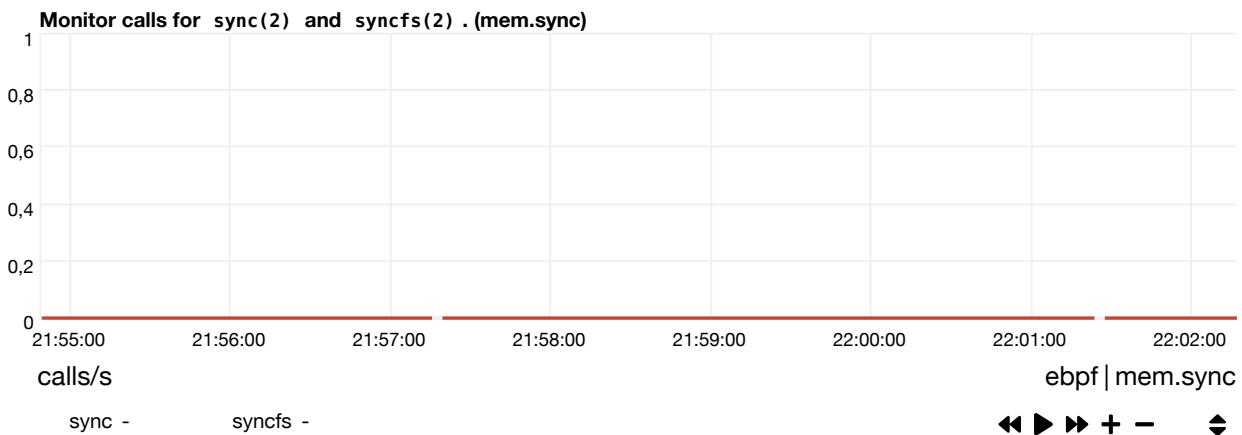
System calls for fsync() and fdatasync() (<https://man7.org/linux/man-pages/man2/fsync.2.html>) transfer all modified page caches for the files on disk devices. These calls block until the device reports that the transfer has been completed.



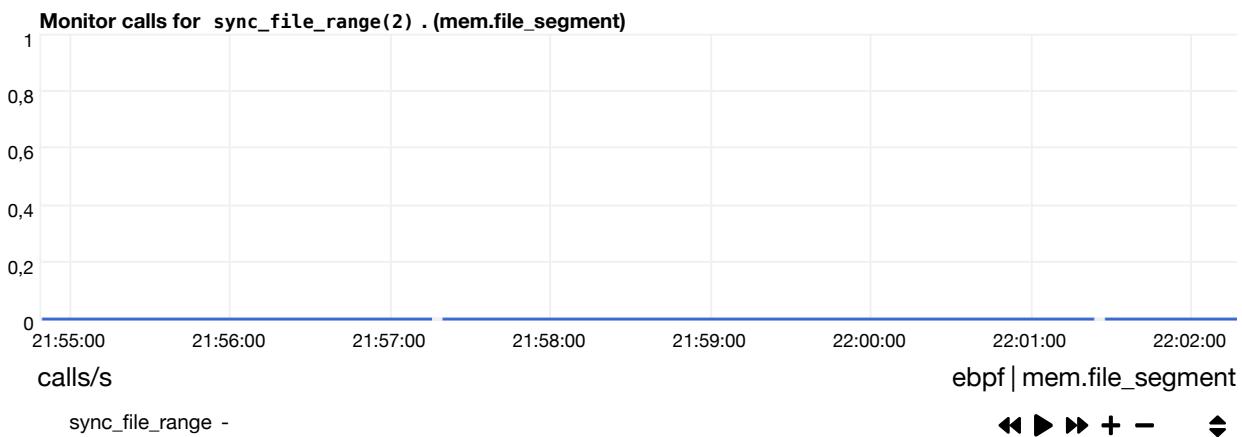
System calls for msync() (<https://man7.org/linux/man-pages/man2/msync.2.html>) which flushes changes made to the in-core copy of a file that was mapped.



System calls for sync() and syncfs() (<https://man7.org/linux/man-pages/man2/sync.2.html>) which flush the file system buffers to storage devices. Performance perturbations might be caused by these calls. The sync() calls are based on the eBPF syncsnoop (<https://github.com/iovisor/bcc/blob/master/tools/syncsnoop.py>) from BCC tools.



System calls for sync\_file\_range() ([https://man7.org/linux/man-pages/man2/sync\\_file\\_range.2.html](https://man7.org/linux/man-pages/man2/sync_file_range.2.html)) permits fine control when synchronizing the open file referred to by the file descriptor fd with disk. This system call is extremely dangerous and should not be used in portable programs.



## Disks

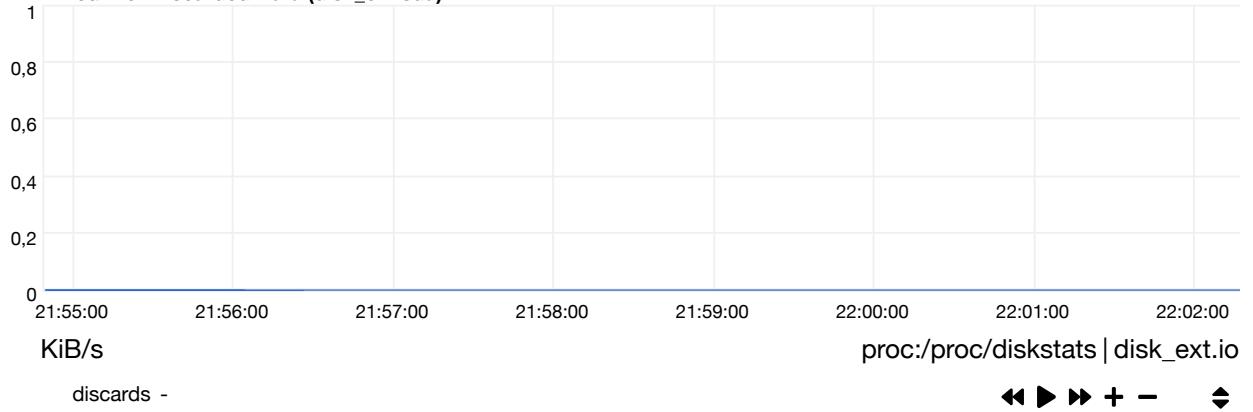
Charts with performance information for all the system disks. Special care has been given to present disk performance metrics in a way compatible with `iostat -x`. netdata by default prevents rendering performance charts for individual partitions and unmounted virtual disks. Disabled charts can still be enabled by configuring the relative settings in the netdata configuration file.

### sda

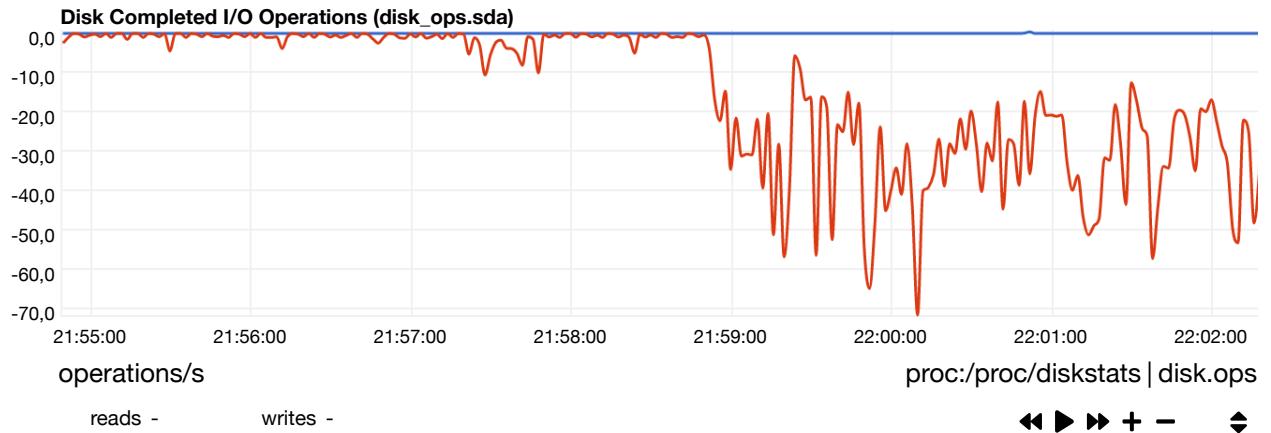
Amount of data transferred to and from disk.

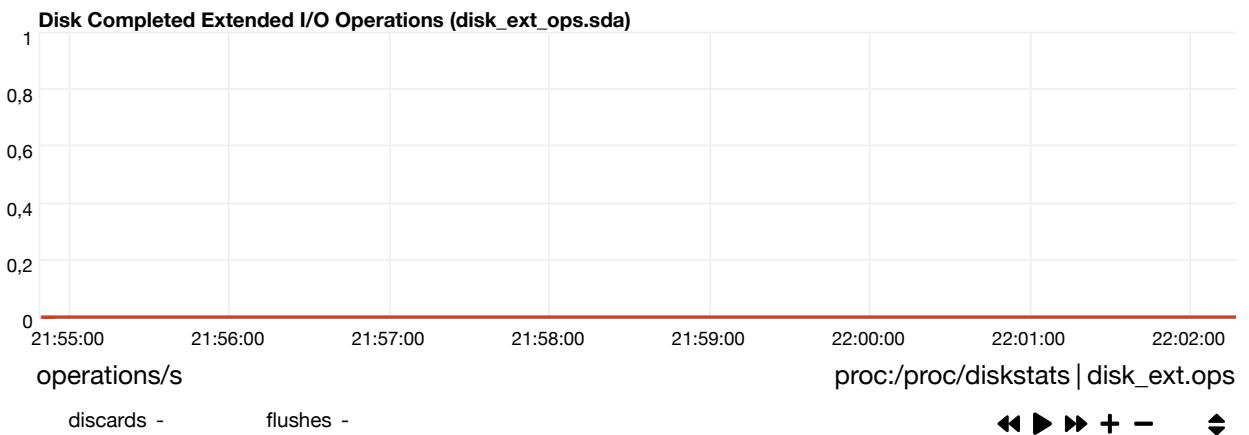


Amount of Discarded Data (disk\_ext.sda)

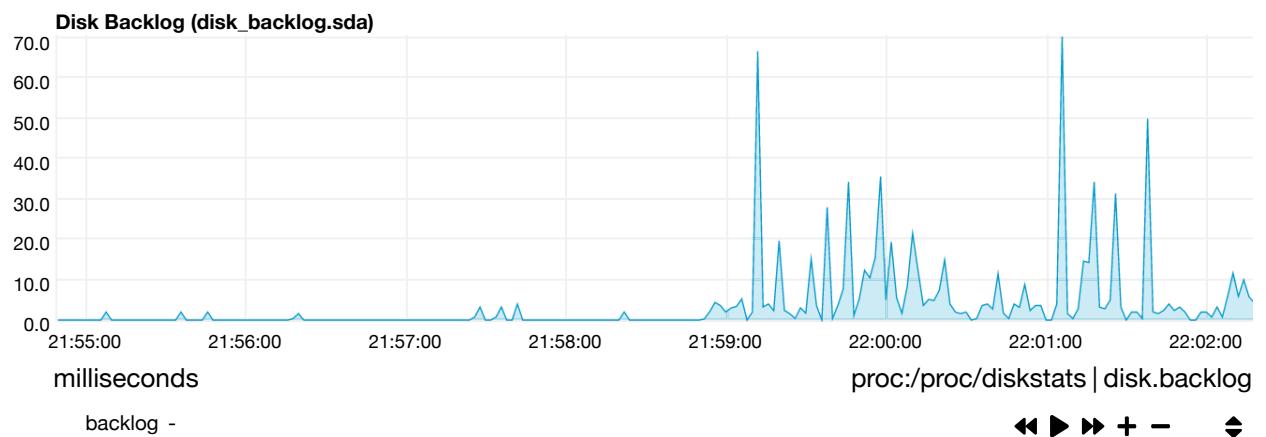


Completed disk I/O operations. Keep in mind the number of operations requested might be higher, since the system is able to merge adjacent to each other (see merged operations chart).

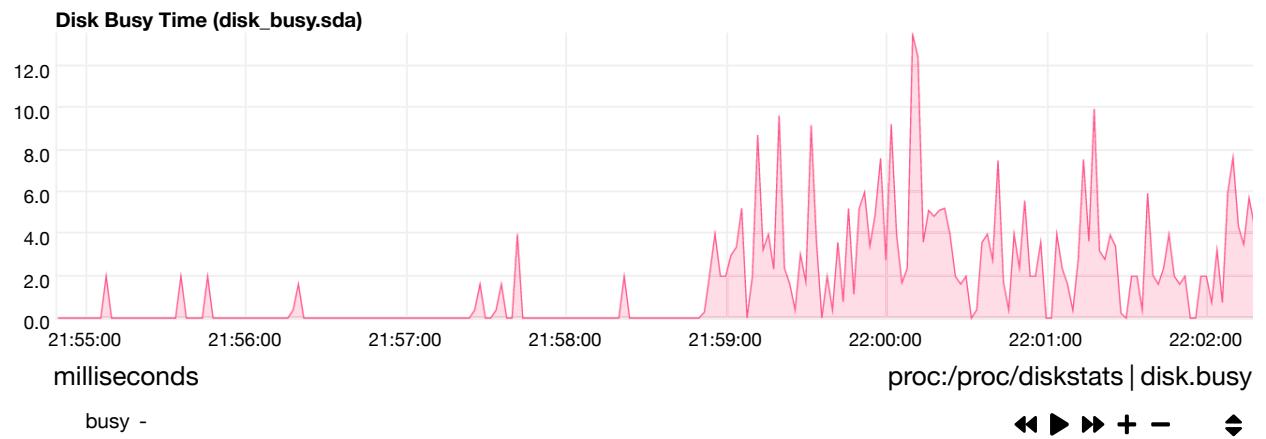




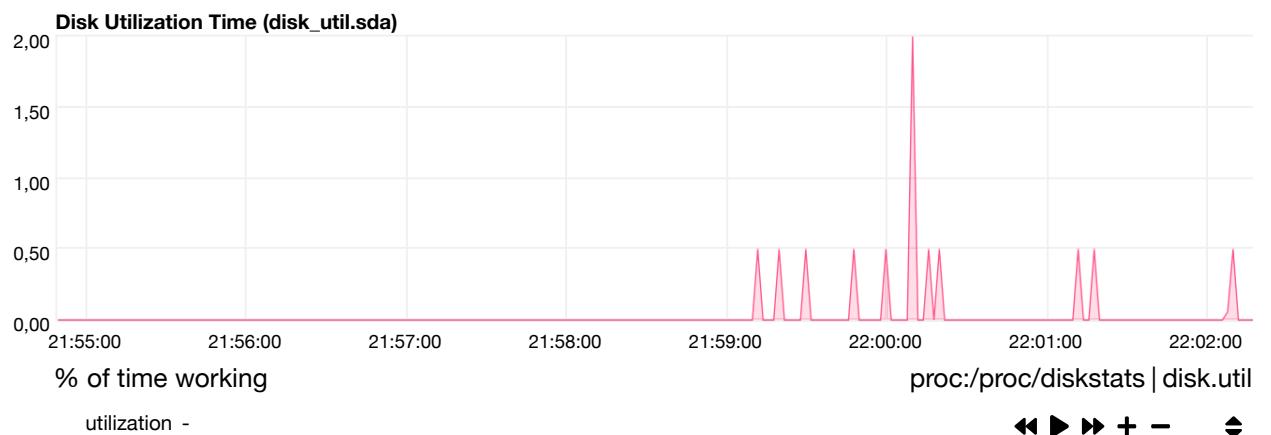
Backlog is an indication of the duration of pending disk operations. On every I/O event the system is multiplying the time spent doing I/O since the last update of this field with the number of pending operations. While not accurate, this metric can provide an indication of the expected completion time of the operations in progress.



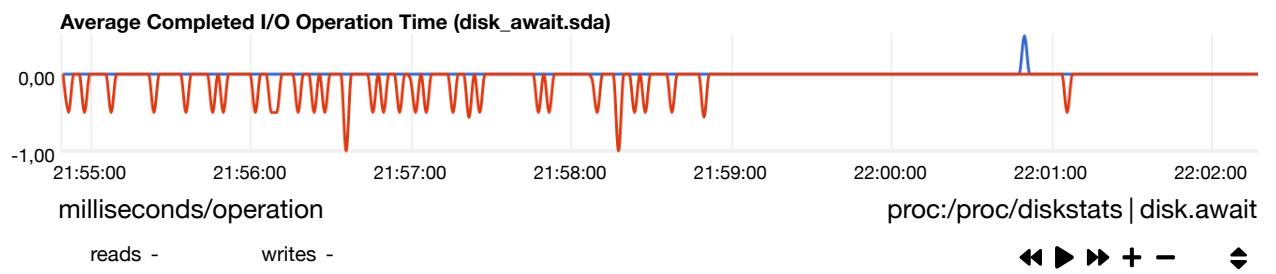
Disk Busy Time measures the amount of time the disk was busy with something.



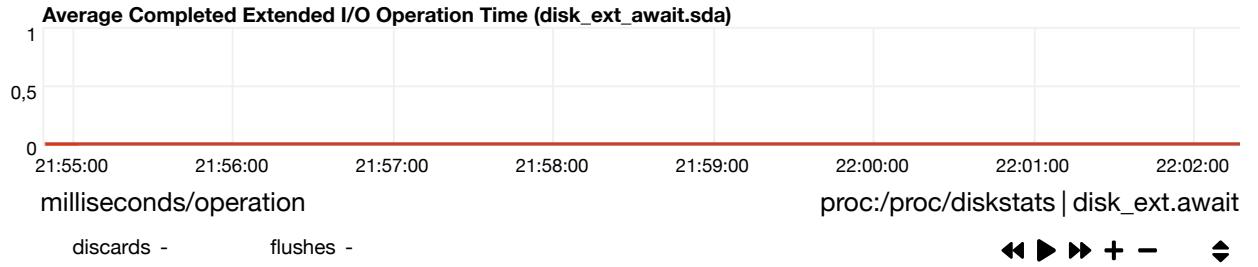
Disk Utilization measures the amount of time the disk was busy with something. This is not related to its performance. 100% means that the system always had an outstanding operation on the disk. Keep in mind that depending on the underlying technology of the disk, 100% here may or may not be an indication of congestion.



The average time for I/O requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.

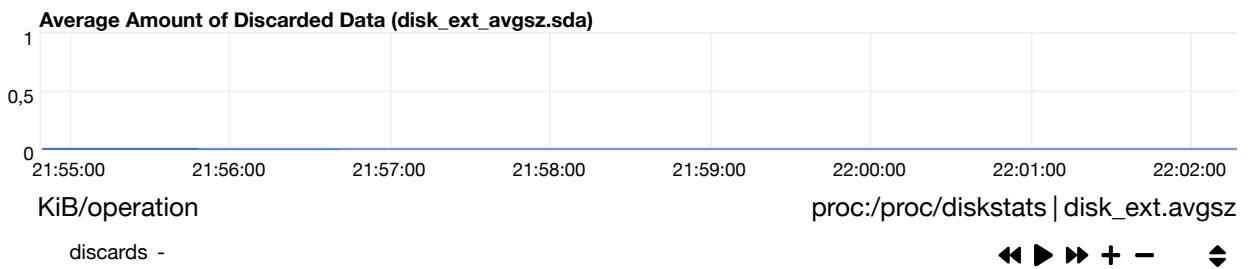


The average time for extended I/O requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.



The average I/O operation size.

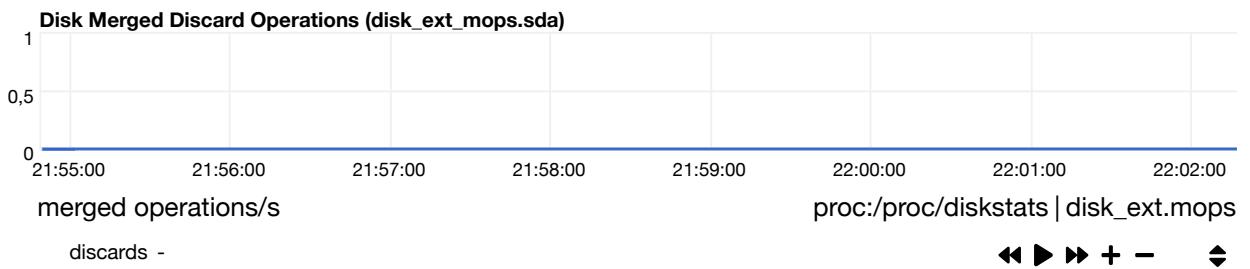




The average service time for completed I/O operations. This metric is calculated using the total busy time of the disk and the number of completed operations. If the disk is able to execute multiple parallel operations the reporting average service time will be misleading.

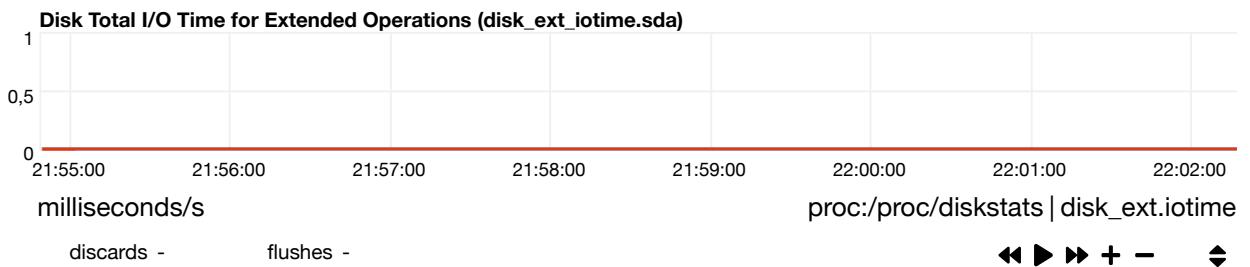


The number of merged disk operations. The system is able to merge adjacent I/O operations, for example two 4KB reads can become one 8KB read before given to disk.



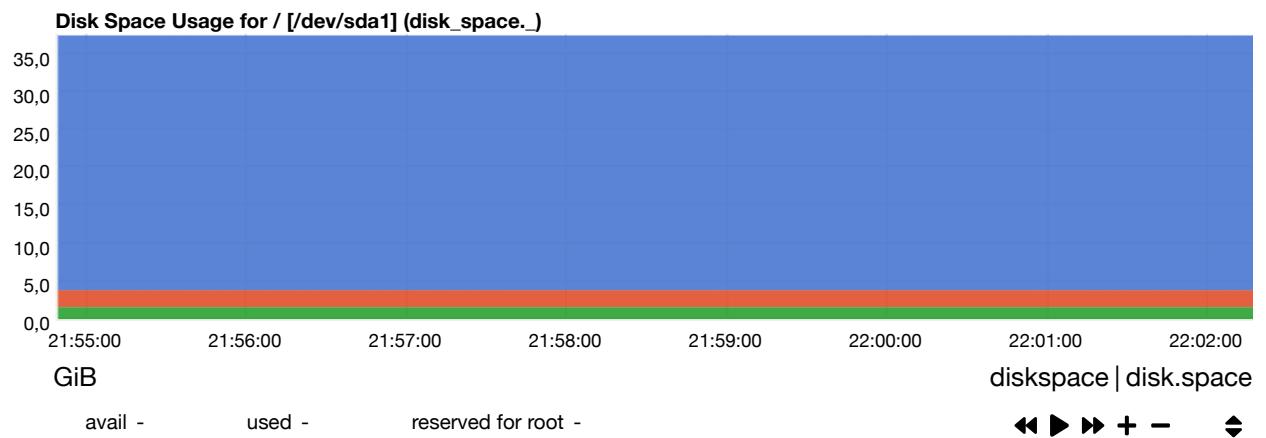
The sum of the duration of all completed I/O operations. This number can exceed the interval if the disk is able to execute I/O operations in parallel.



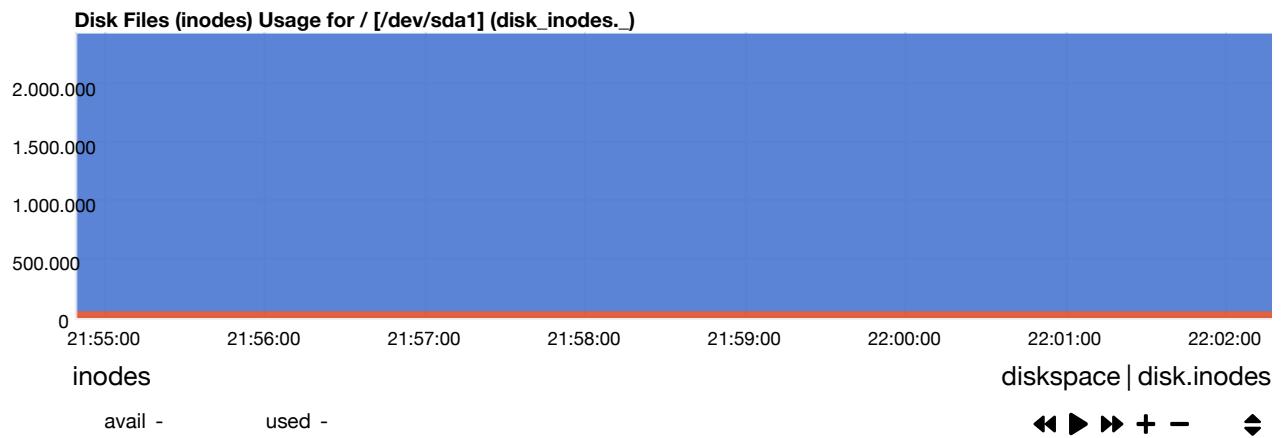


/

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

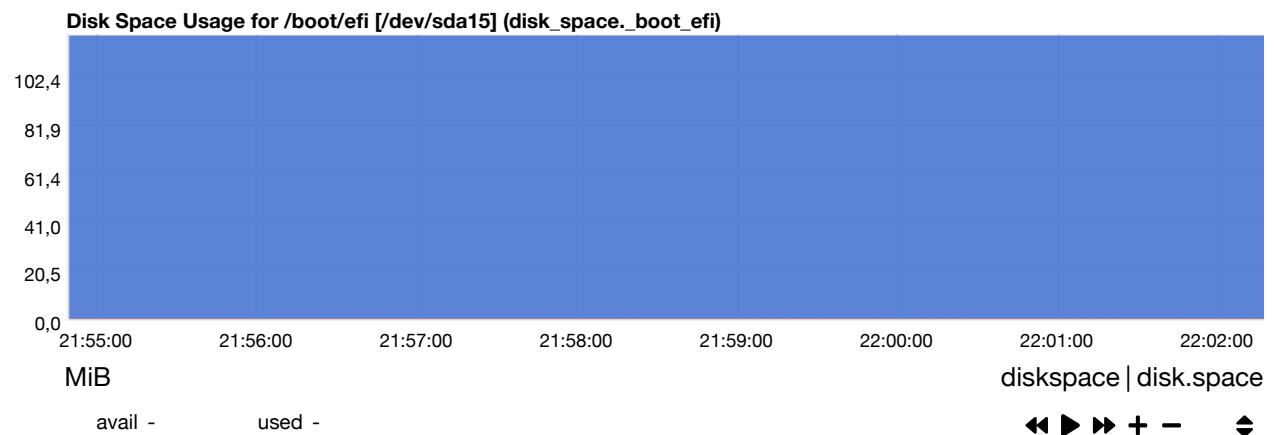


inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.



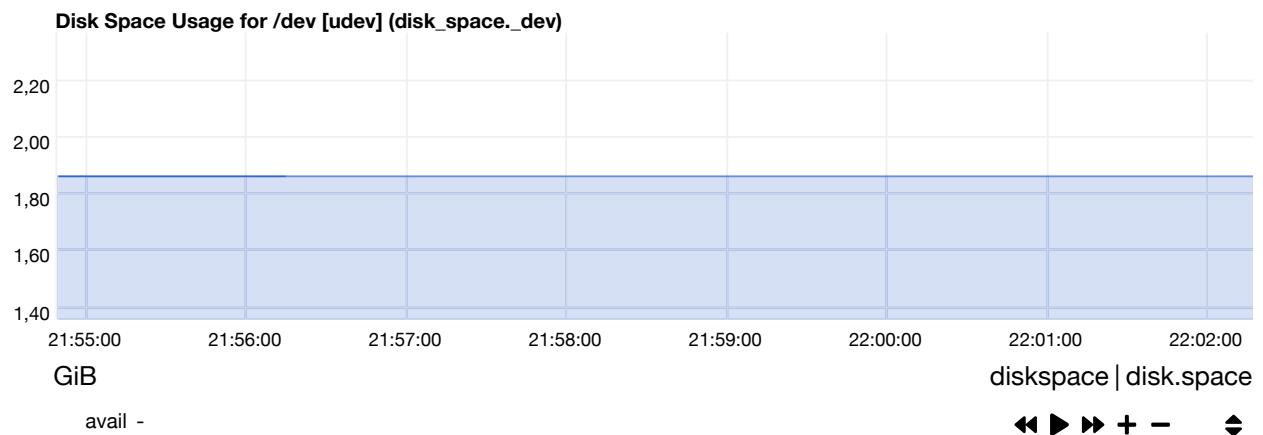
# /boot/efi

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

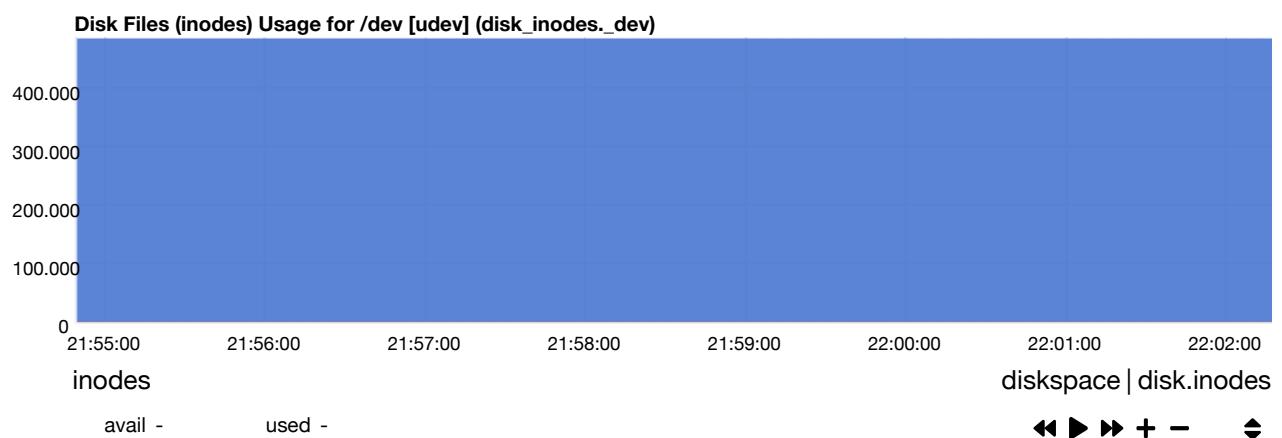


## /dev

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

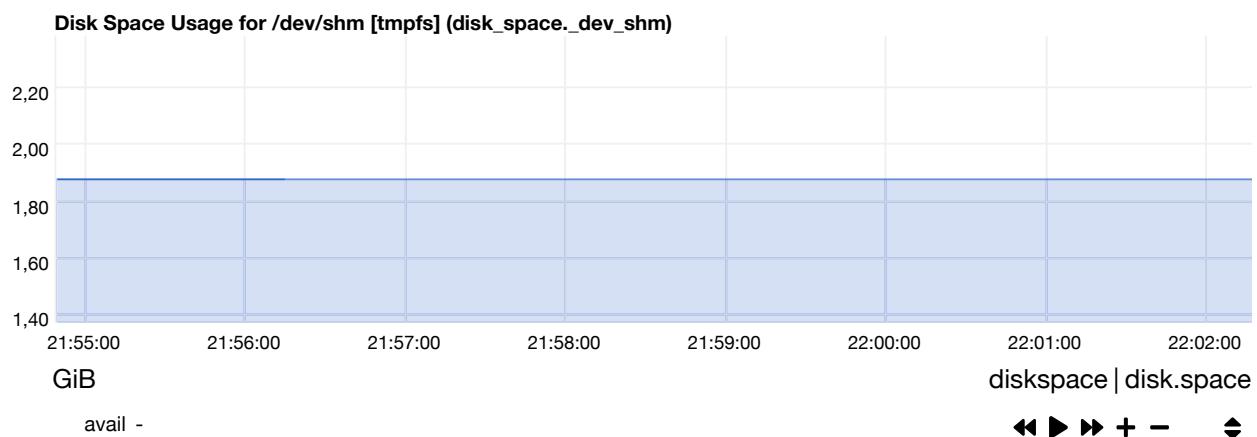


inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.

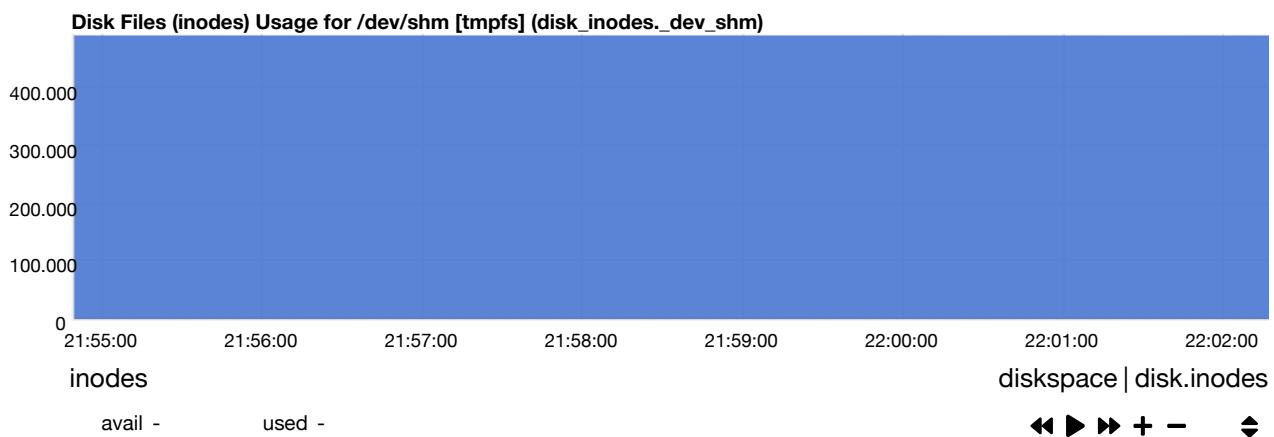


## /dev/shm

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

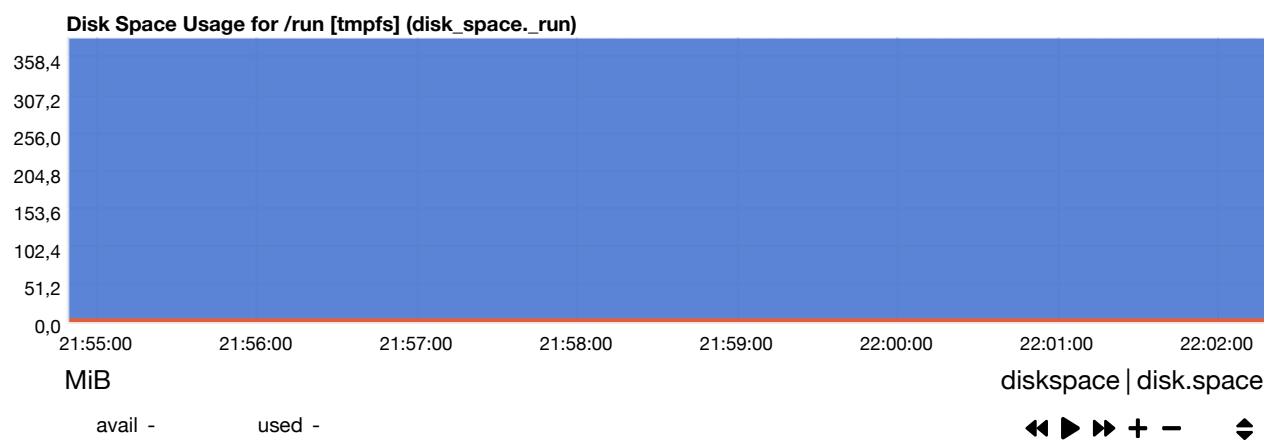


inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.

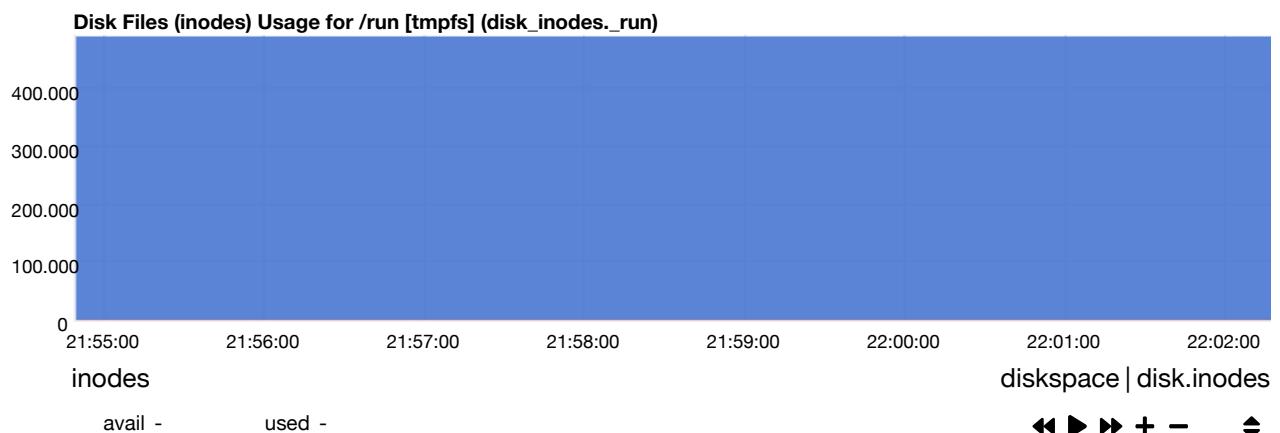


## /run

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

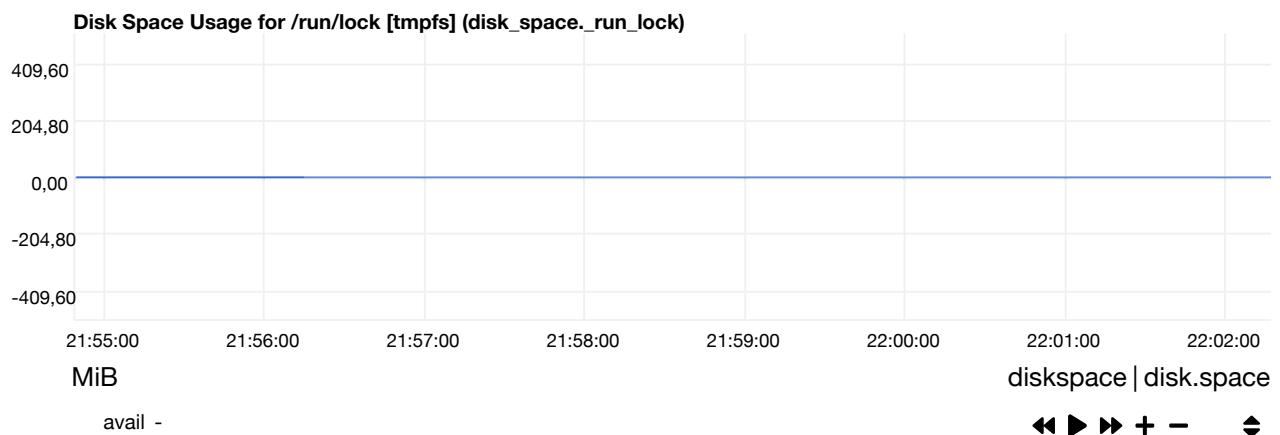


inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.

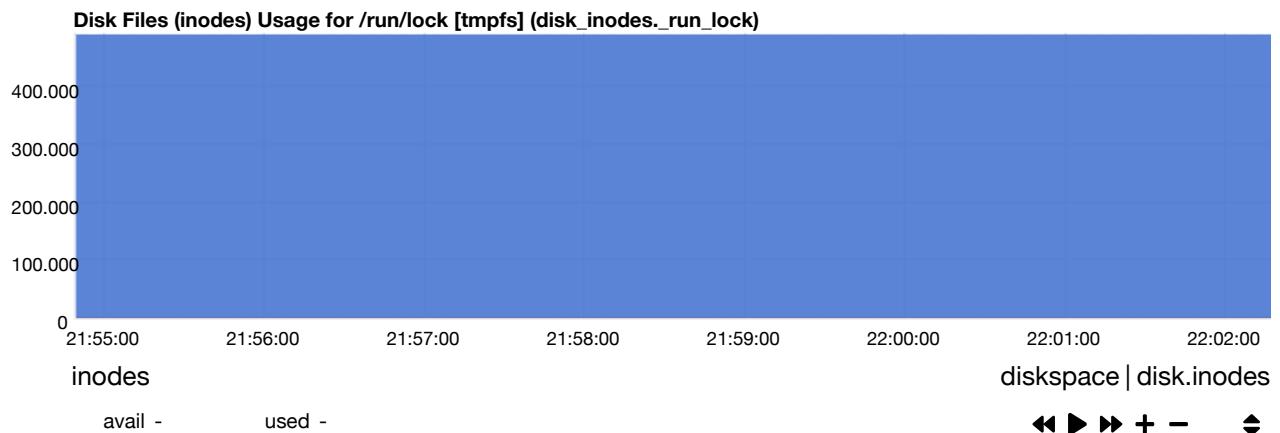


## /run/lock

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

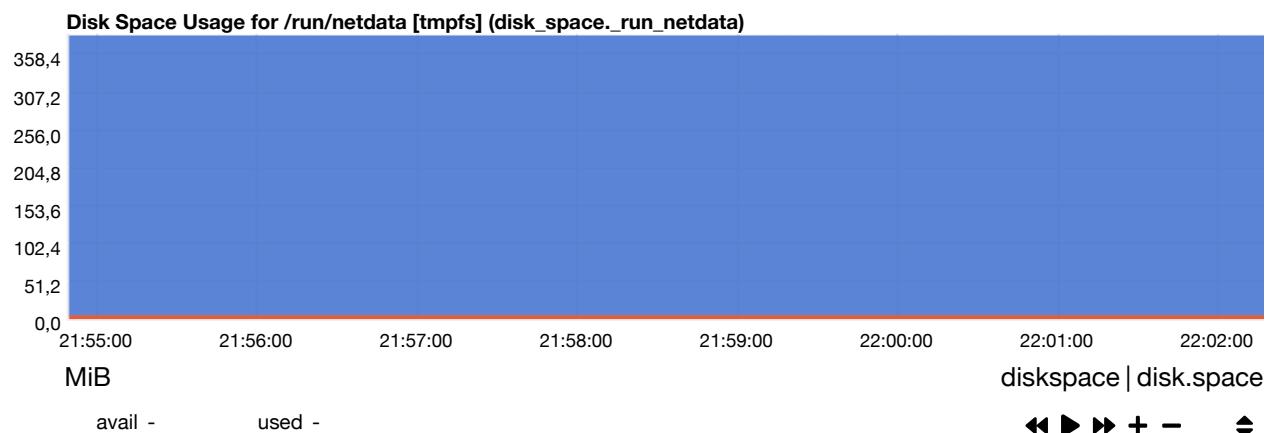


inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.

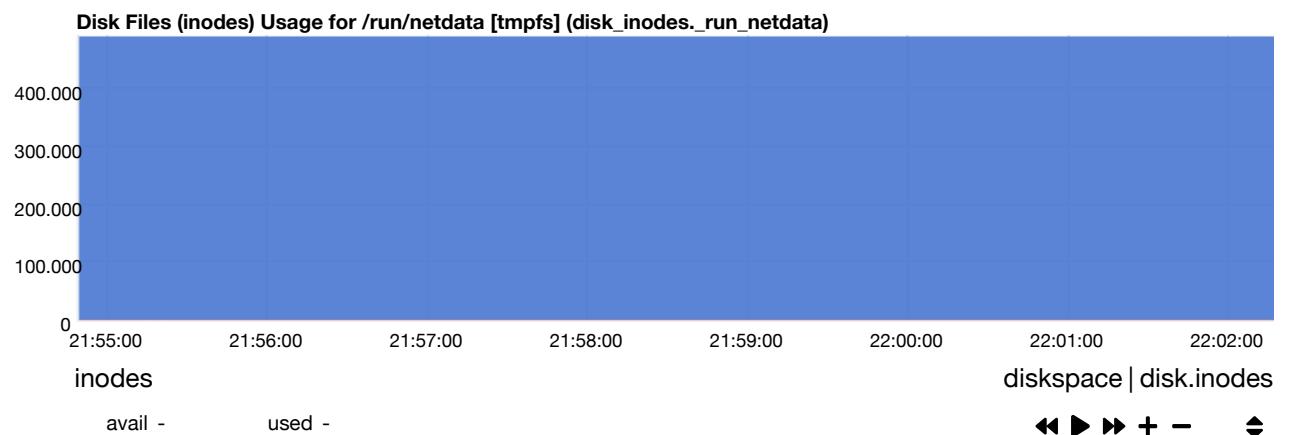


# /run/netdata

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.



inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.

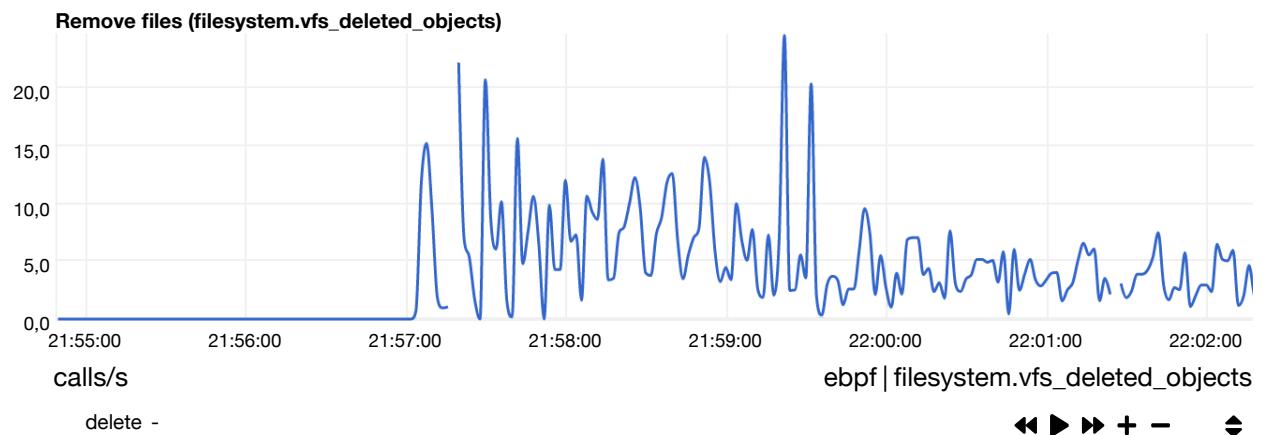



---

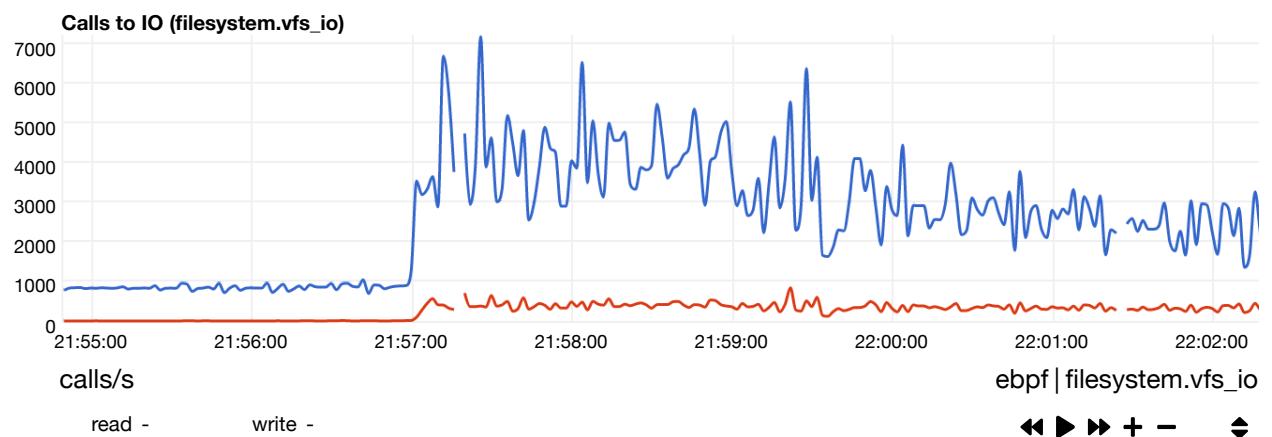
## Filesystem

### VFS (eBPF)

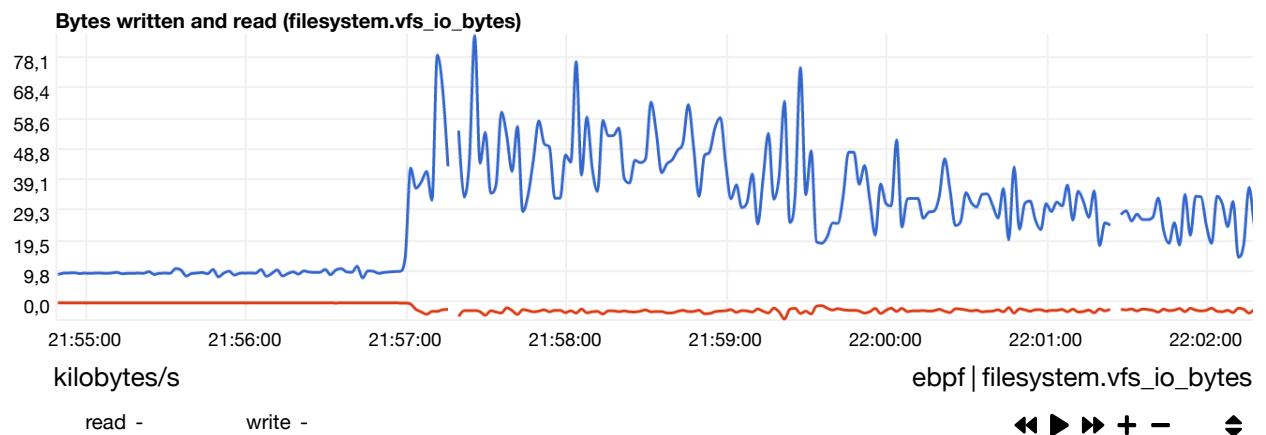
This chart does not show all events that remove files from the file system, because file systems can create their own functions to remove files, it shows calls for the function `vfs_unlink`.



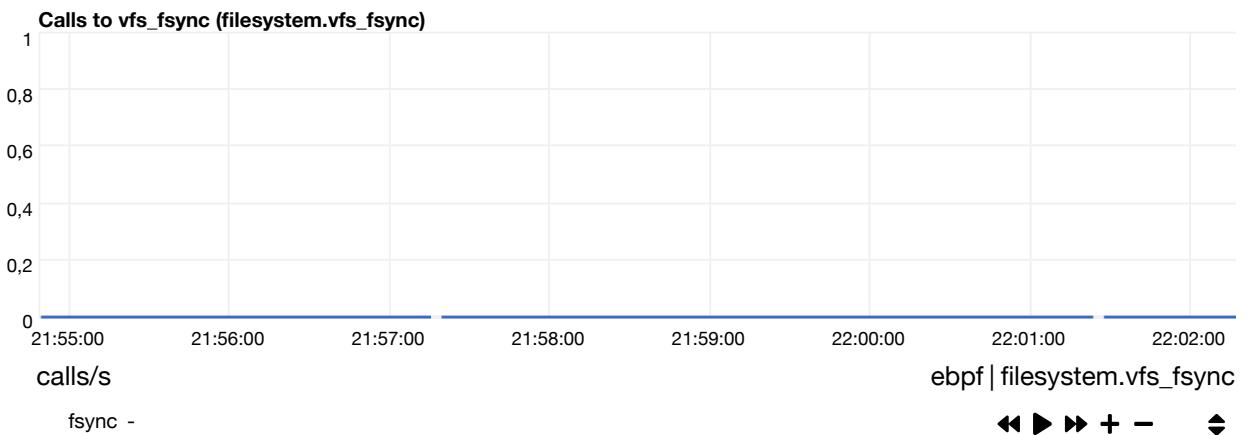
Successful or failed calls to functions `vfs_read` and `vfs_write`. This chart may not show all file system events if it uses other functions to store data on disk.



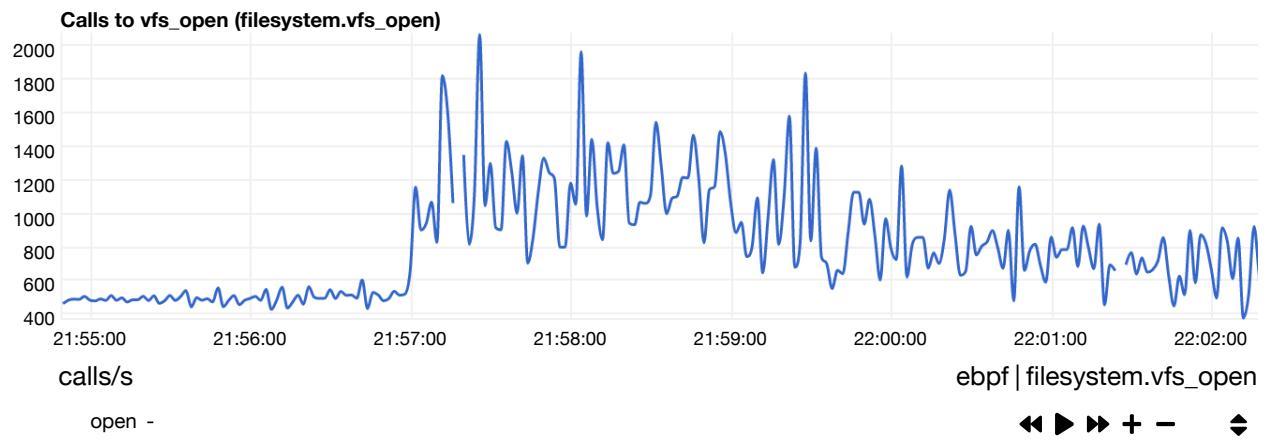
Total of bytes read or written with success using the functions `vfs_read` and `vfs_write`.



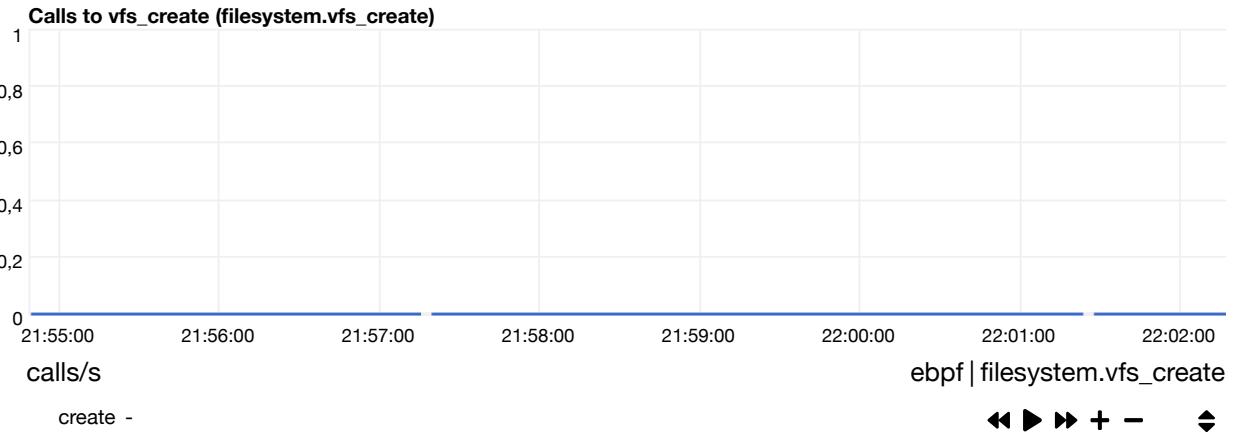
Successful or failed calls to functions `vfs_fsync` .



Successful or failed calls to functions `vfs_open` .



Successful or failed calls to functions `vfs_create` .

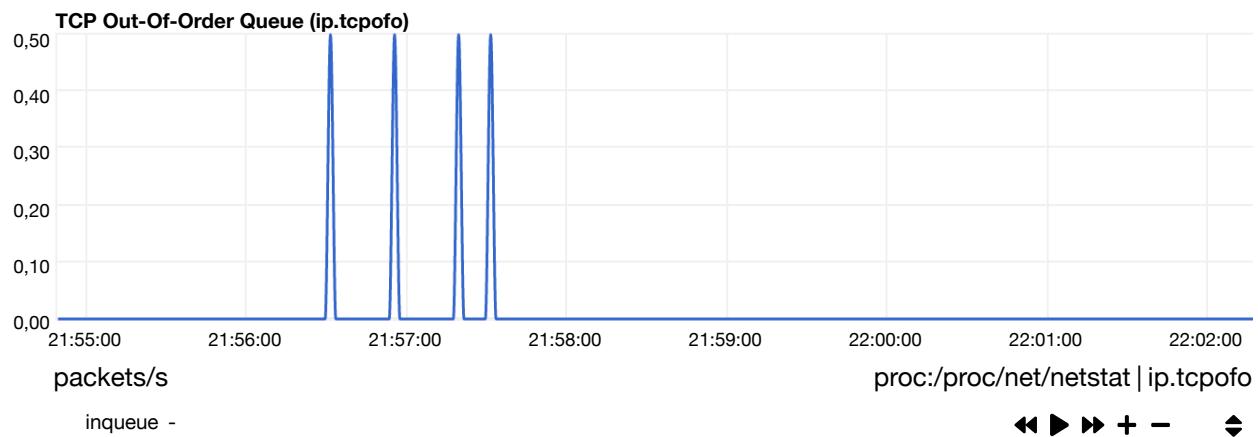
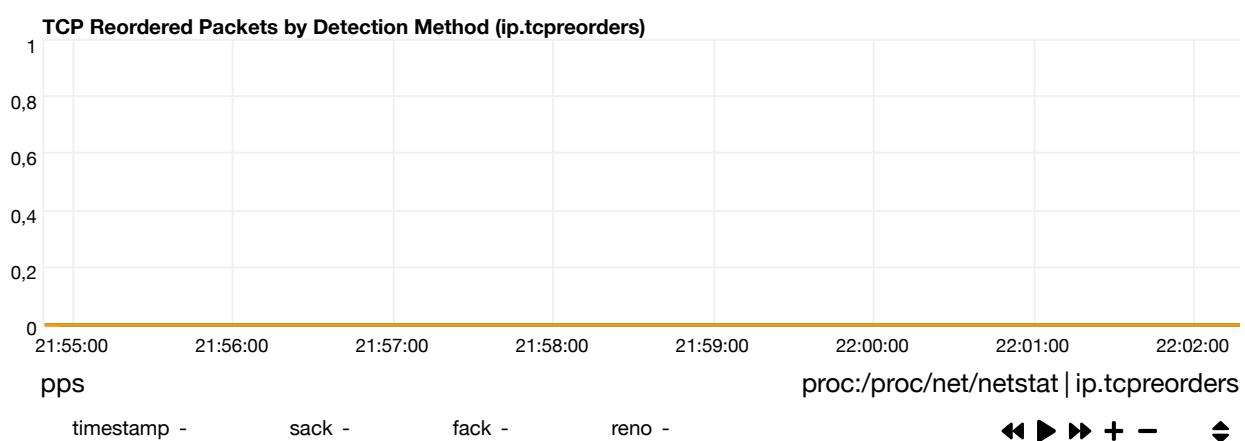
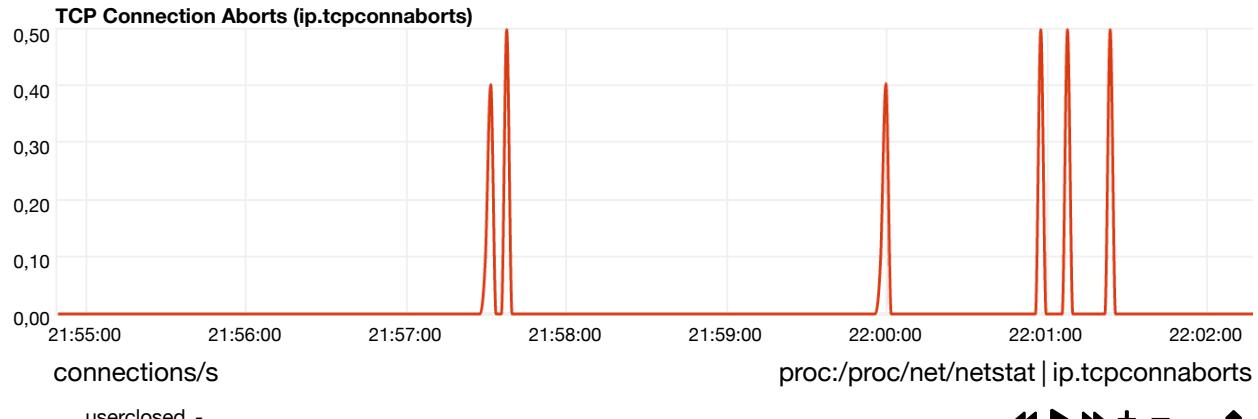


## Networking Stack

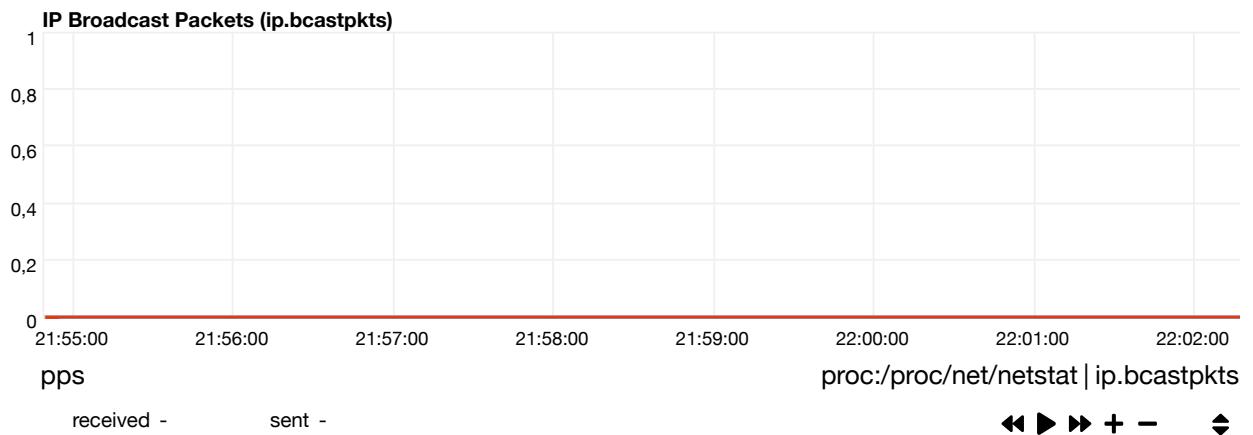
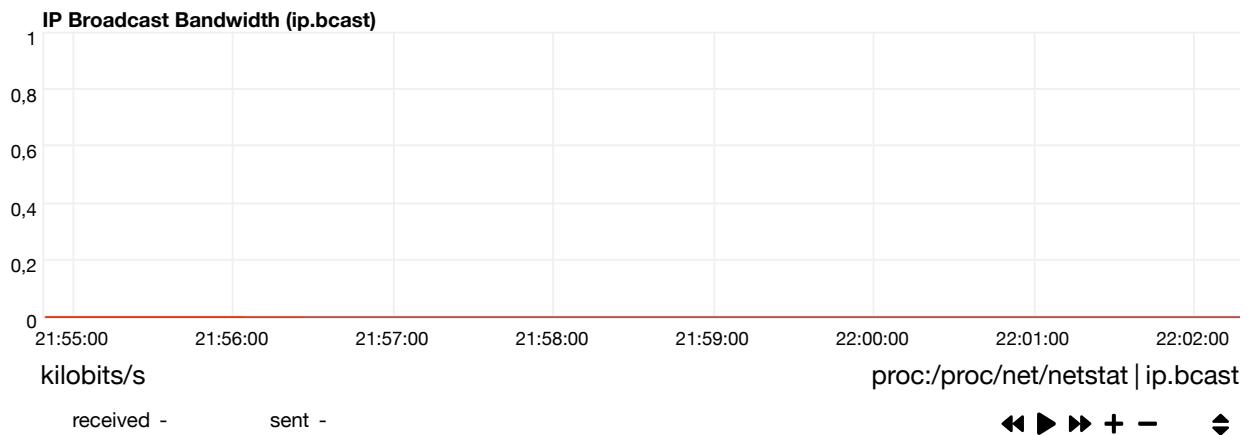
Metrics for the networking stack of the system. These metrics are collected from `/proc/net/netstat` , apply to both IPv4 and IPv6 traffic and are related to operation of the kernel networking stack.

### tcp

TCP connection aborts. **baddata** ( `TCPAbortOnData` ) happens while the connection is on `FIN_WAIT1` and the kernel receives a packet with a sequence number beyond the last one for this connection - the kernel responds with `RST` (closes the connection). **userclosed** ( `TCPAbortOnClose` ) happens when the kernel receives data on an already closed connection and responds with `RST` . **nomemory** ( `TCPAbortOnMemory` ) happens when there are too many orphaned sockets (not attached to an fd) and the kernel has to drop a connection - sometimes it will send an `RST` , sometimes it won't. **timeout** ( `TCPAbortOnTimeout` ) happens when a connection times out. **linger** ( `TCPAbortOnLinger` ) happens when the kernel killed a socket that was already closed by the application and lingered around for long enough. **failed** ( `TCPAbortFailed` ) happens when the kernel attempted to send an `RST` but failed because there was no memory available.



# broadcast



## ecn

Explicit Congestion Notification (ECN) ([https://en.wikipedia.org/wiki/Explicit\\_Congestion\\_Notation](https://en.wikipedia.org/wiki/Explicit_Congestion_Notation)) is a TCP extension that allows end-to-end notification of network congestion without dropping packets. ECN is an optional feature that may be used between two ECN-enabled endpoints when the underlying network infrastructure also supports it.

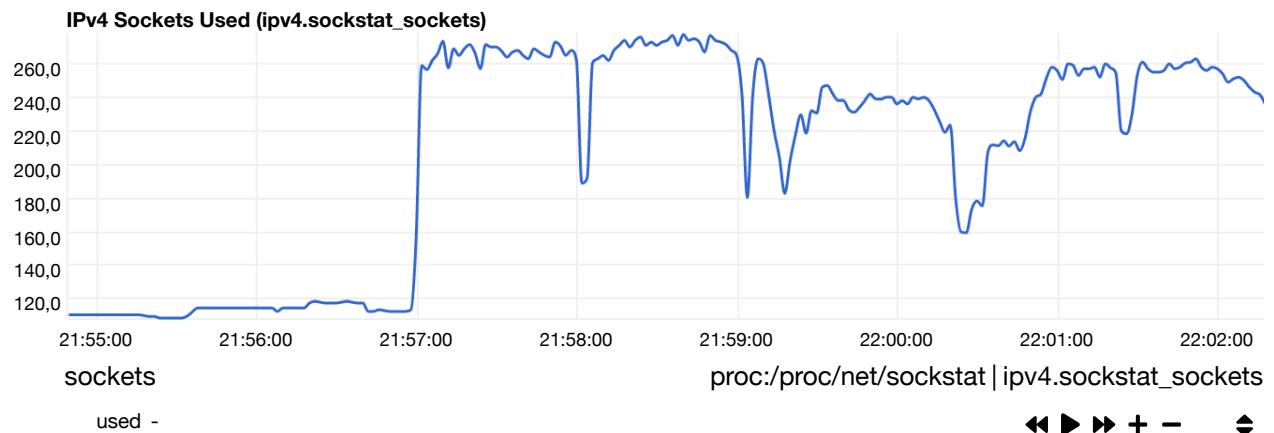


## IPv4 Networking

Metrics for the IPv4 stack of the system. Internet Protocol version 4 (IPv4) (<https://en.wikipedia.org/wiki/IPv4>) is the fourth version of the Internet Protocol (IP). It is one of the core protocols of standards-based internetworking methods in the Internet. IPv4 is a connectionless protocol for use on packet-switched networks. It operates on a best effort delivery model, in that it does not guarantee

delivery, nor does it assure proper sequencing or avoidance of duplicate delivery. These aspects, including data integrity, are addressed by an upper layer transport protocol, such as the Transmission Control Protocol (TCP).

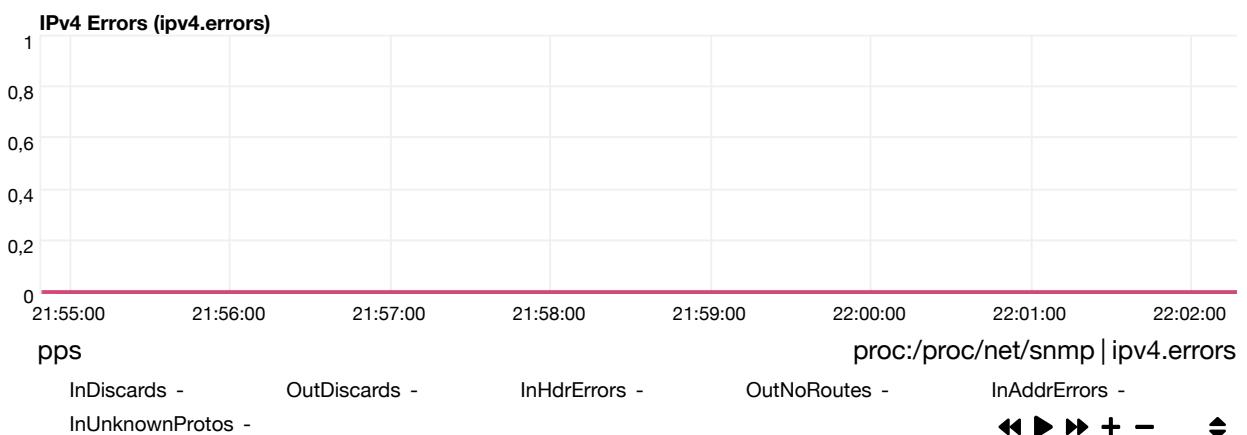
## sockets



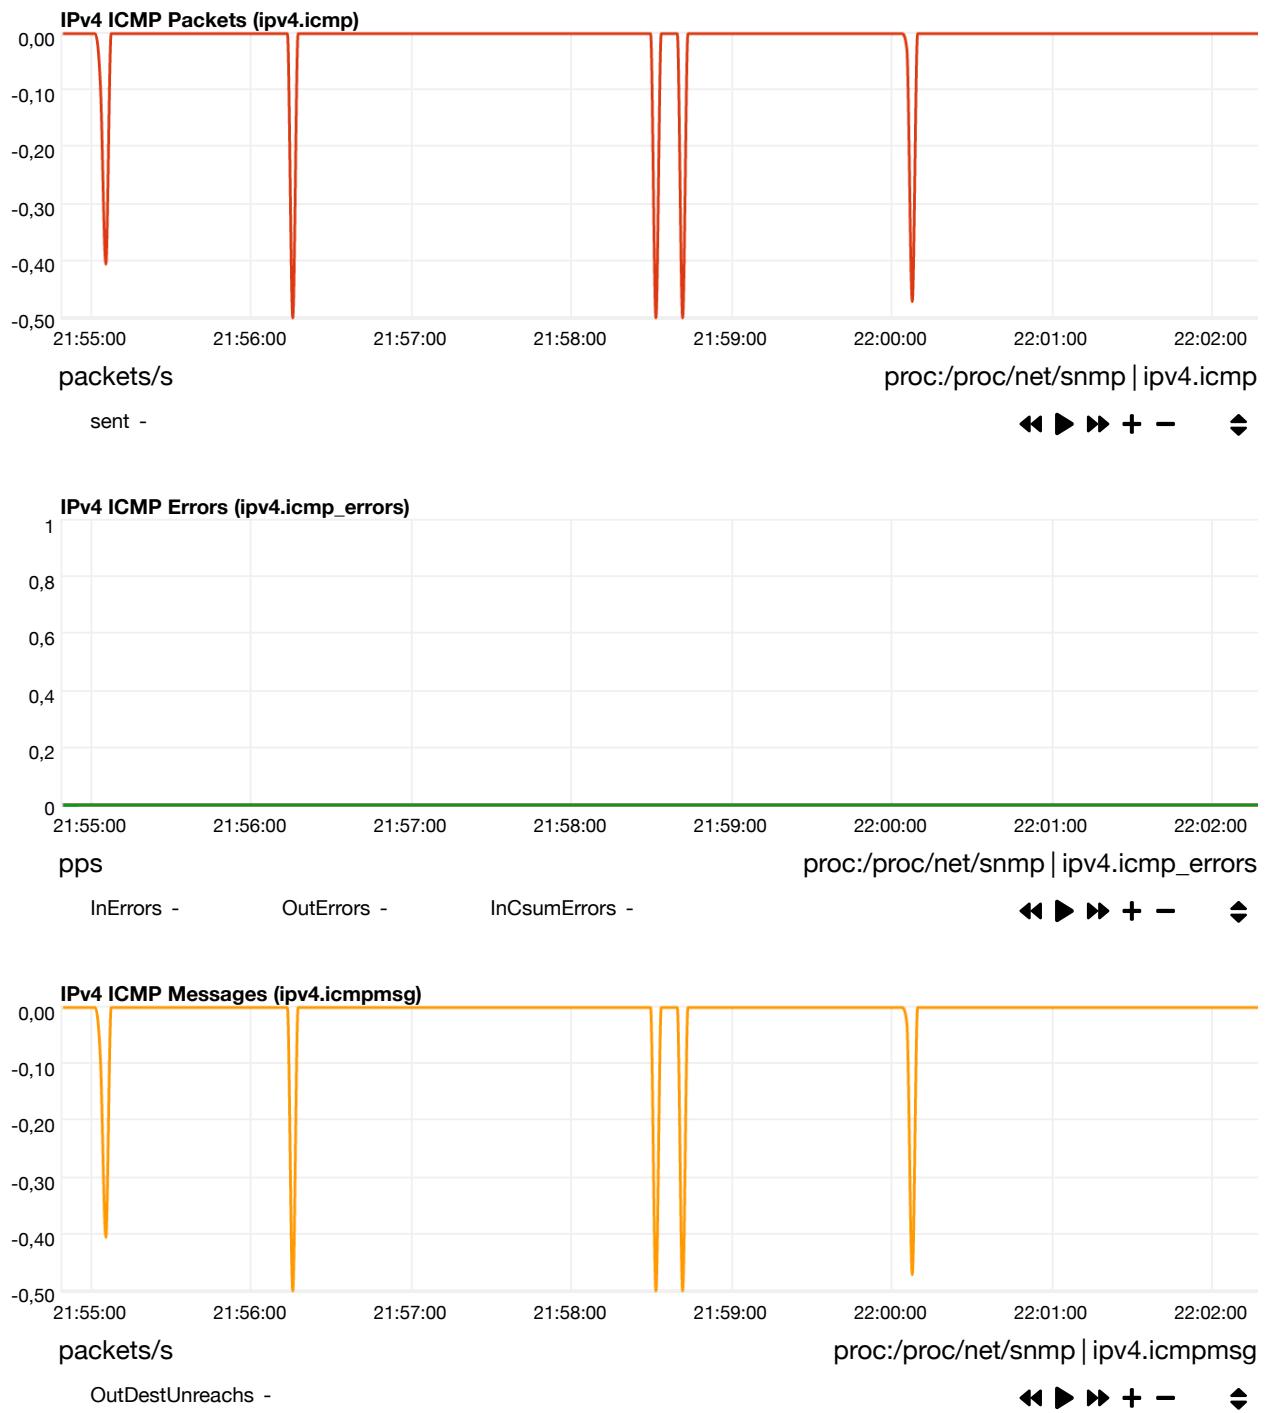
## packets



## errors



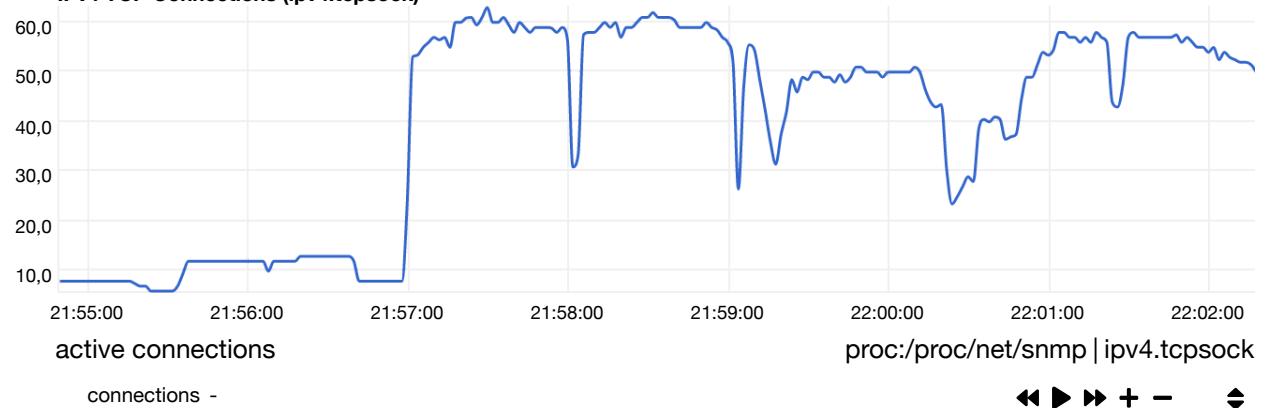
# icmp



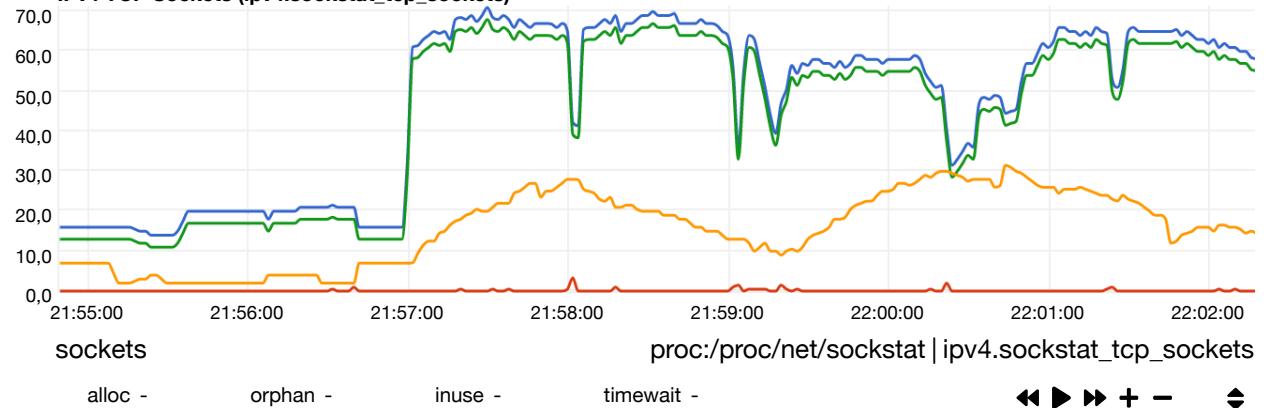
# tcp

The number of established TCP connections (known as `CurrEstab`). This is a snapshot of the established connections at the time of measurement (i.e. a connection established and a connection disconnected within the same iteration will not affect this metric).

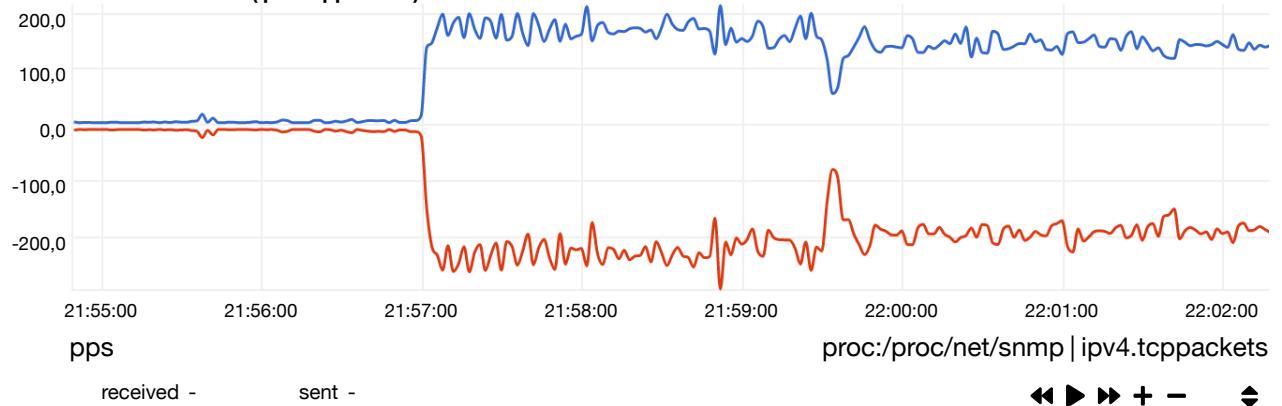
IPv4 TCP Connections (ipv4.tcpsock)



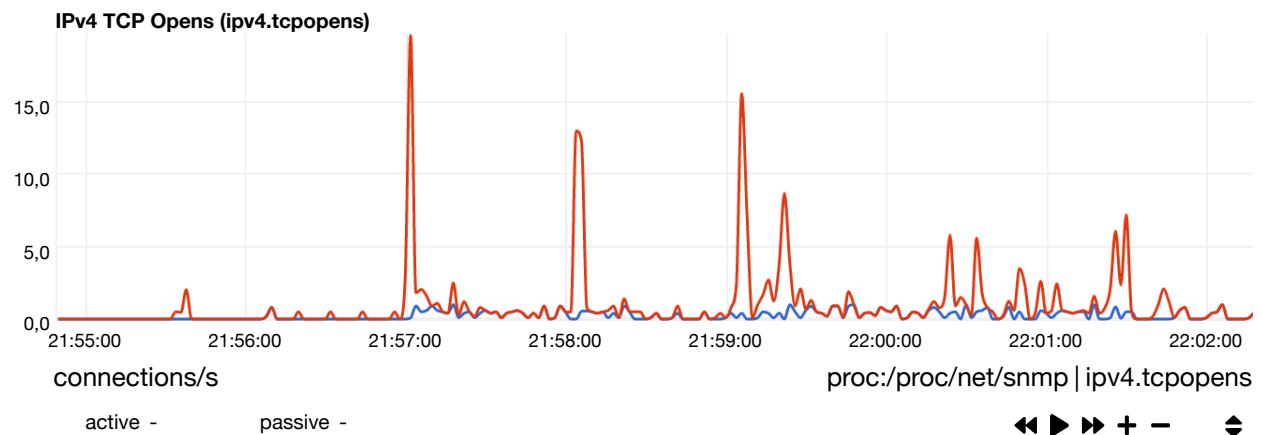
IPv4 TCP Sockets (ipv4.sockstat\_tcp\_sockets)



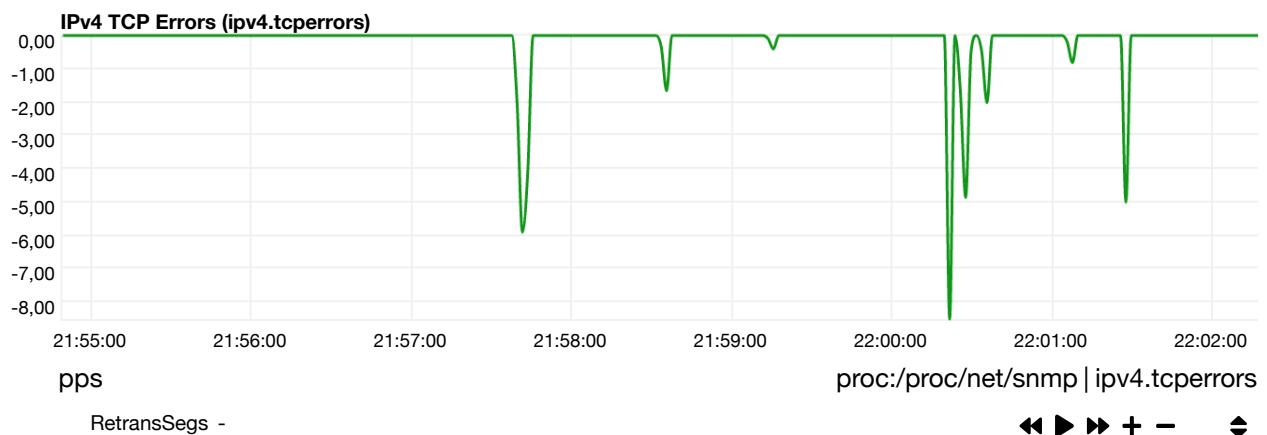
IPv4 TCP Packets (ipv4.tcppingackets)



**active** or **ActiveOpens** is the number of outgoing TCP **connections attempted** by this host.  
**passive** or **PassiveOpens** is the number of incoming TCP **connections accepted** by this host.

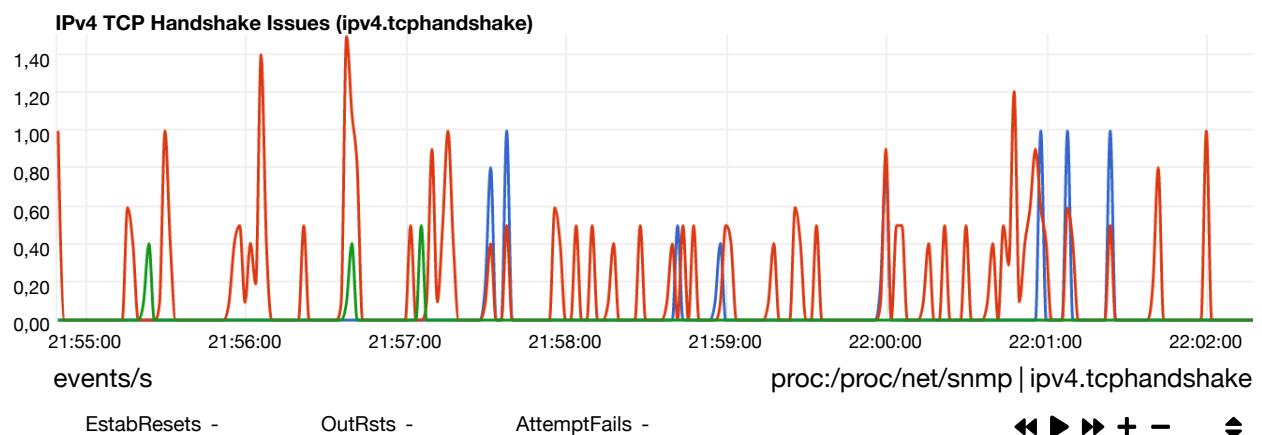


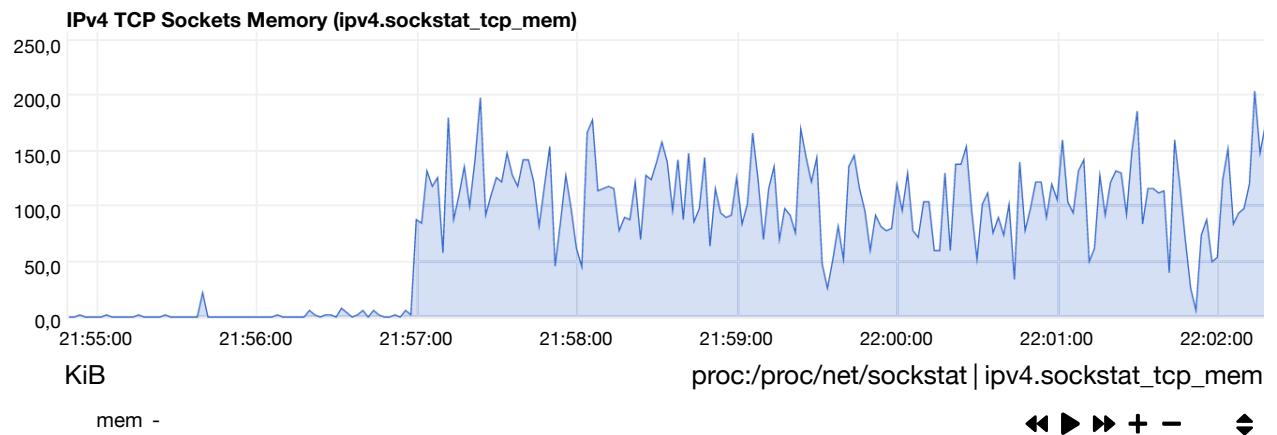
**InErrs** is the number of TCP segments received in error (including header too small, checksum errors, sequence errors, bad packets - for both IPv4 and IPv6). **InCsumErrors** is the number of TCP segments received with checksum errors (for both IPv4 and IPv6). **RetransSegs** is the number of TCP segments retransmitted.



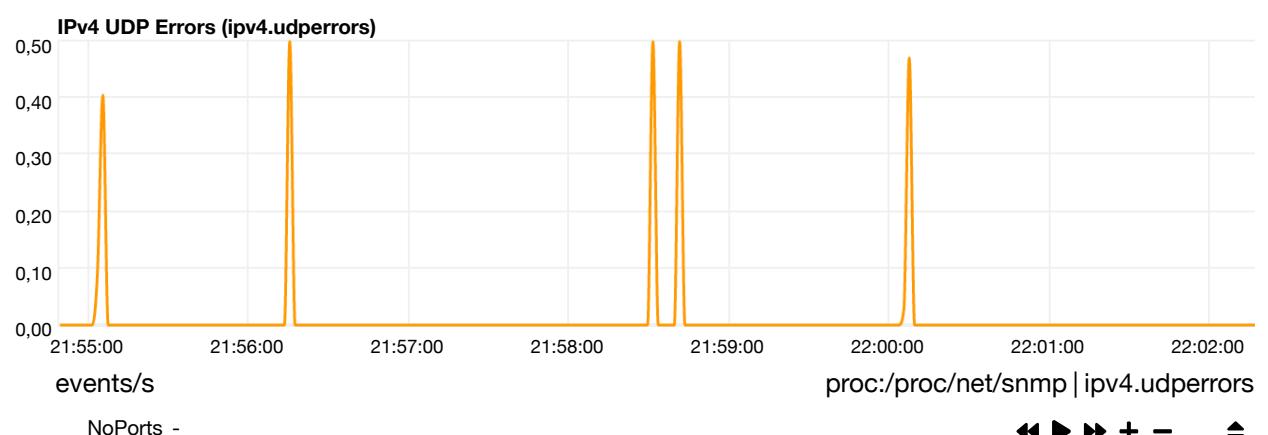
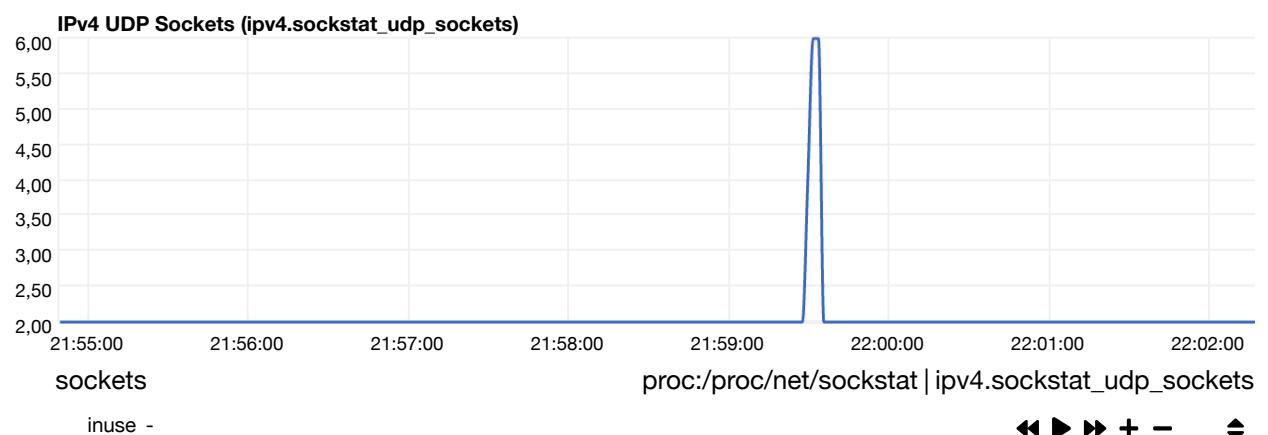
**EstabResets** is the number of established connections resets (i.e. connections that made a direct transition from `ESTABLISHED` or `CLOSE_WAIT` to `CLOSED` ). **OutRsts** is the number of TCP segments sent, with the `RST` flag set (for both IPv4 and IPv6). **AttemptFails** is the number of times TCP connections made a direct transition from either `SYN_SENT` or `SYN_RECV` to `CLOSED` , plus the number of times TCP connections made a direct transition from the `SYN_RECV` to `LISTEN` .

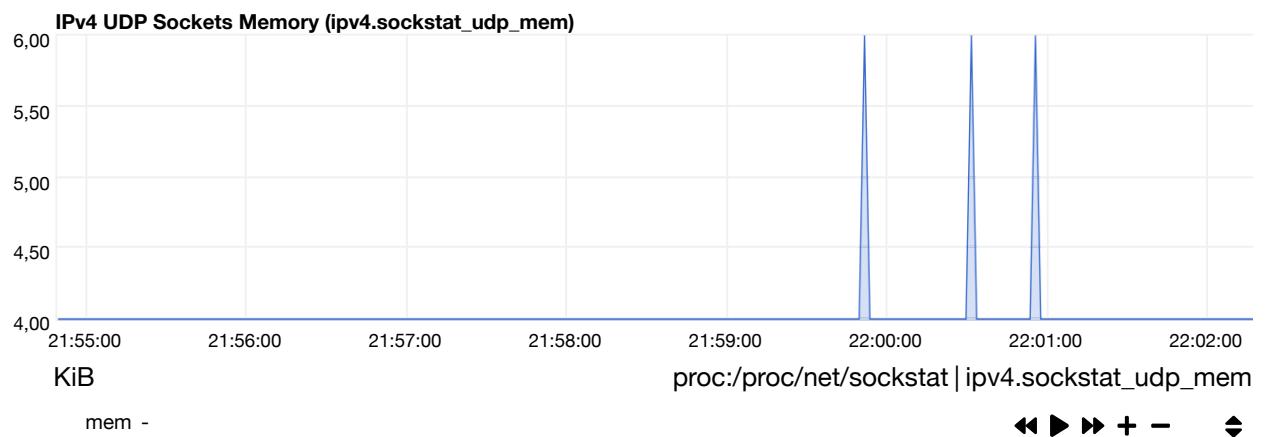
**TCPSynRetrans** shows retries for new outbound TCP connections, which can indicate general connectivity issues or backlog on the remote host.





## udp





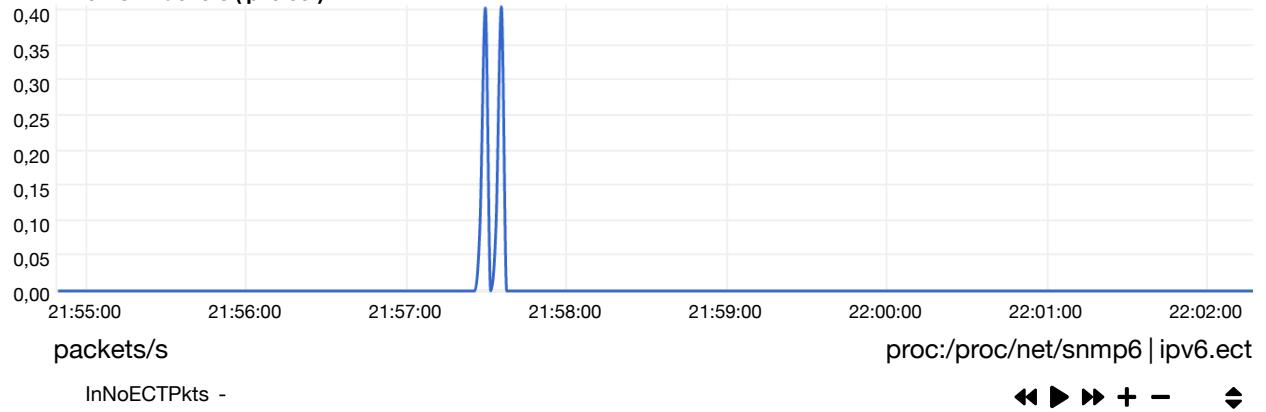
## IPv6 Networking

Metrics for the IPv6 stack of the system. Internet Protocol version 6 (IPv6) (<https://en.wikipedia.org/wiki/IPv6>) is the most recent version of the Internet Protocol (IP), the communications protocol that provides an identification and location system for computers on networks and routes traffic across the Internet. IPv6 was developed by the Internet Engineering Task Force (IETF) to deal with the long-anticipated problem of IPv4 address exhaustion. IPv6 is intended to replace IPv4.

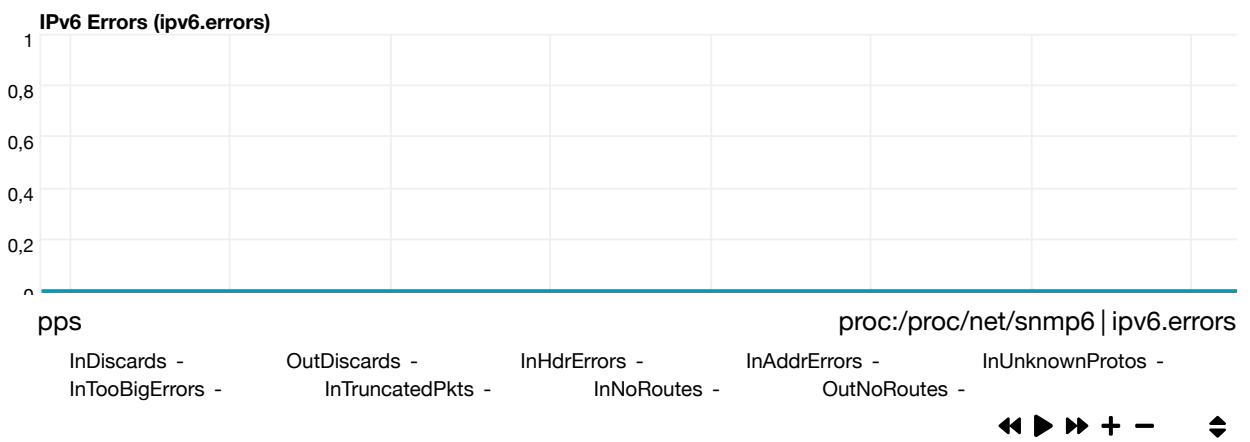
### packets



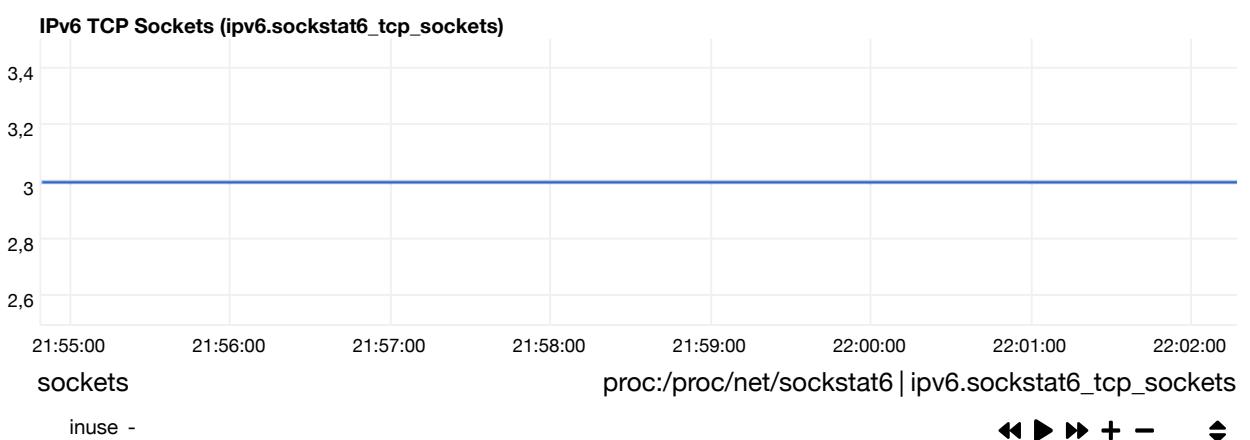
### IPv6 ECT Packets (ipv6.ect)



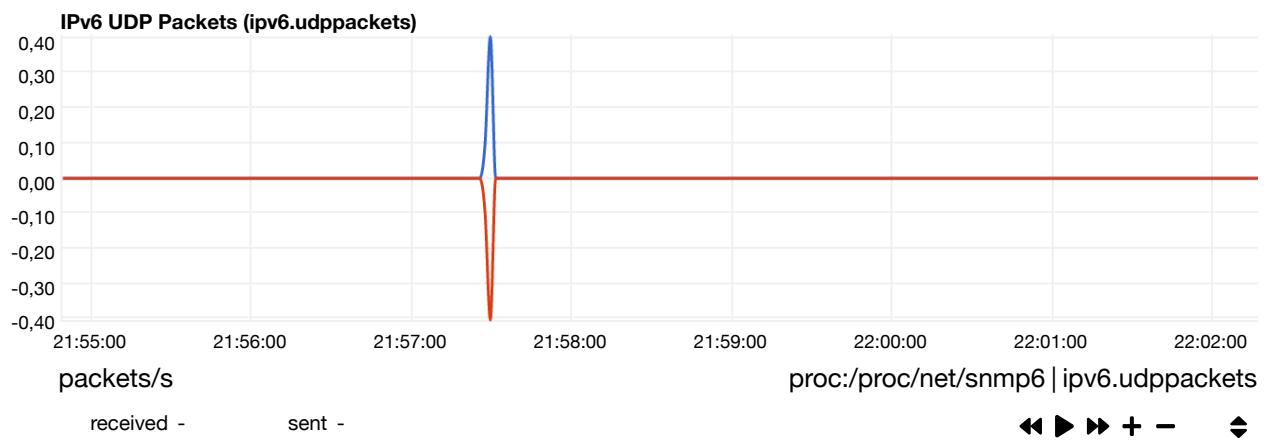
### errors

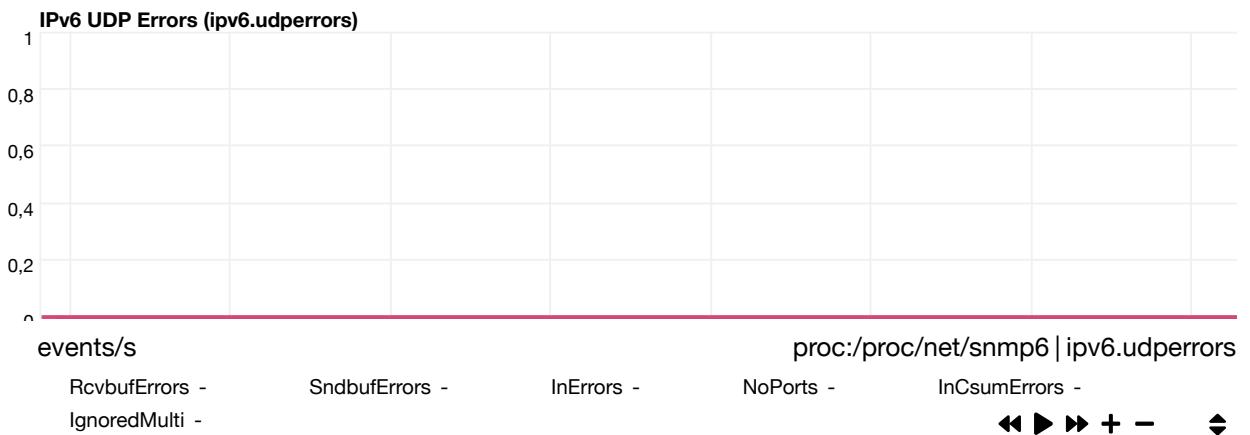


## tcp6

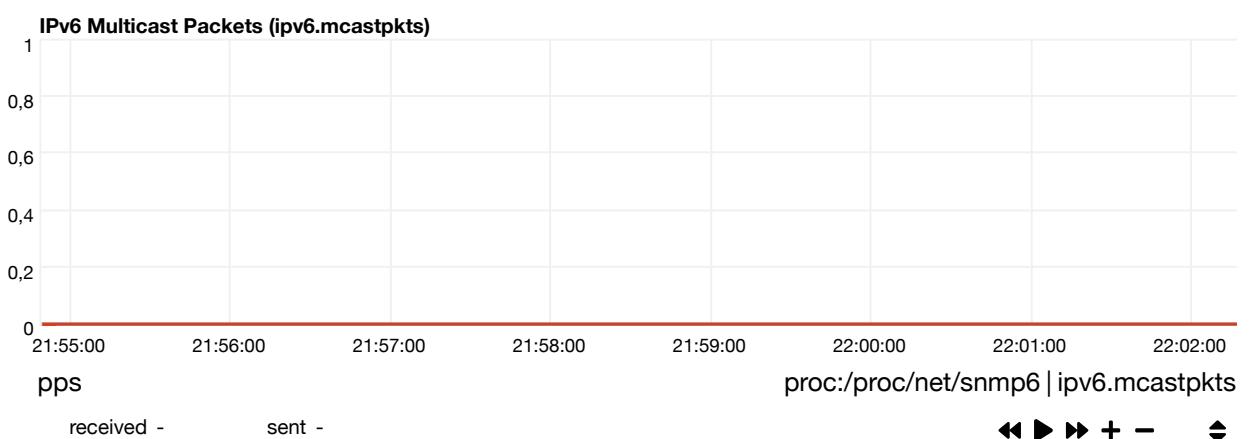
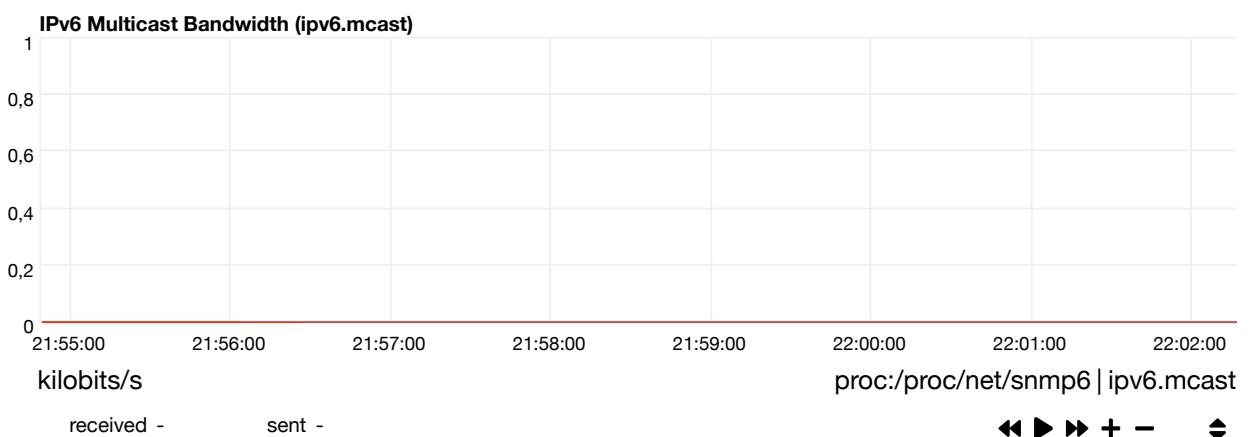


## udp6

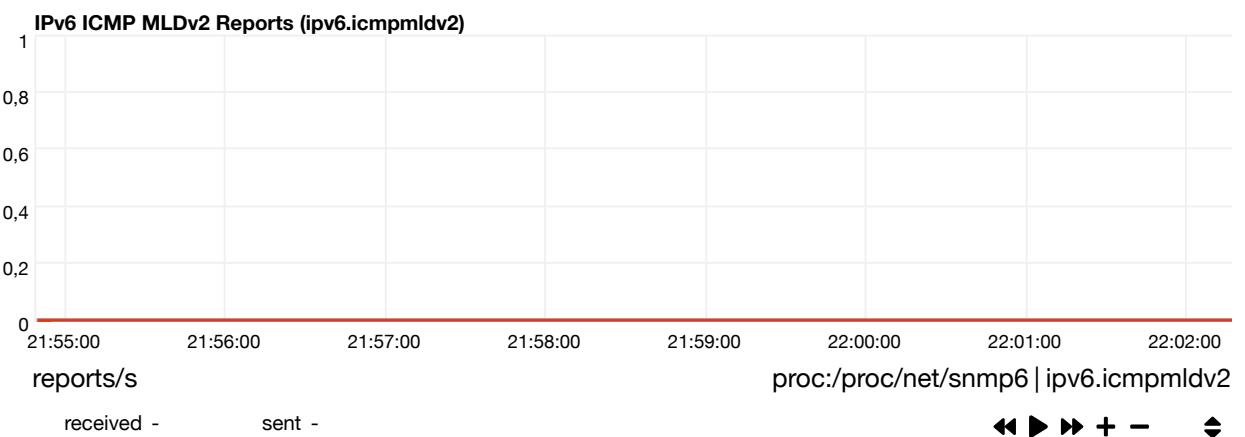
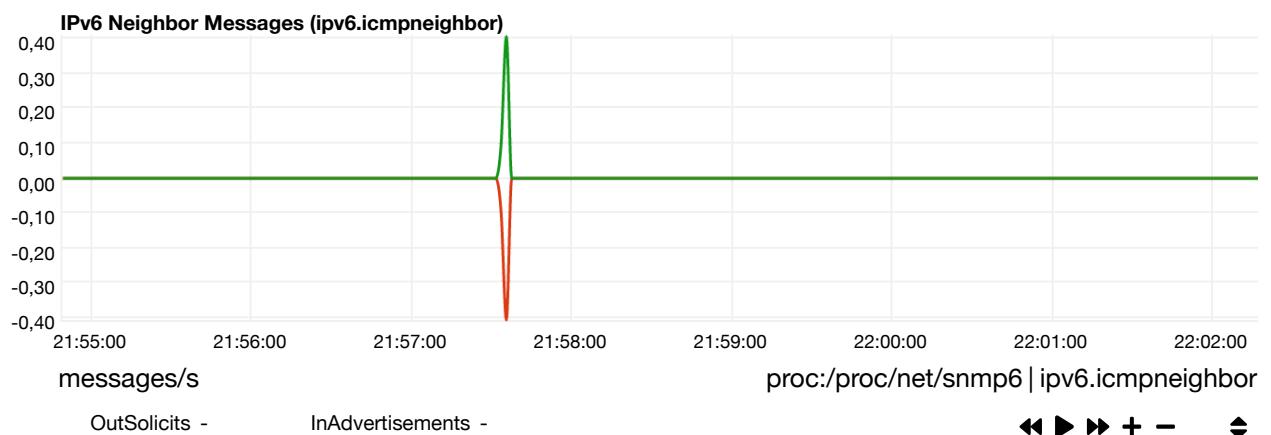
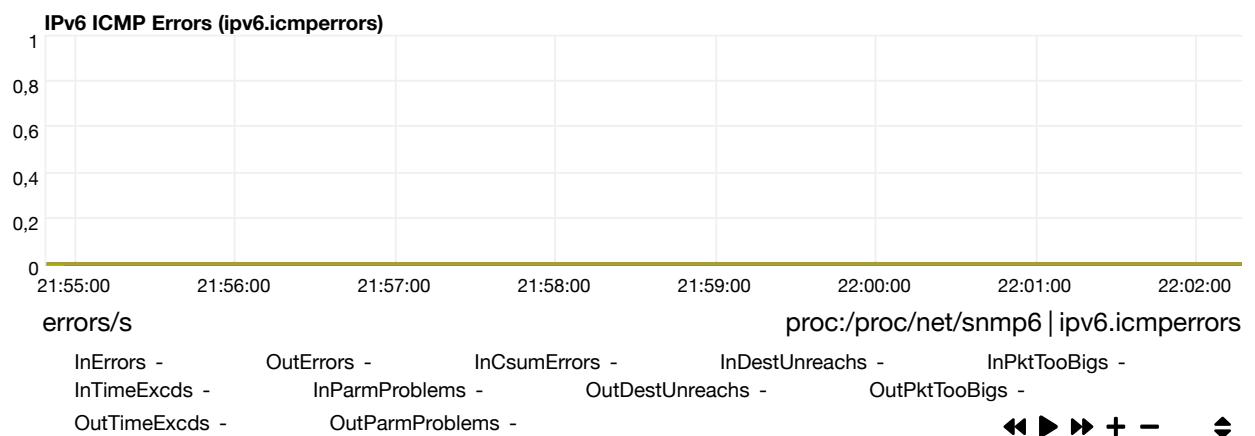


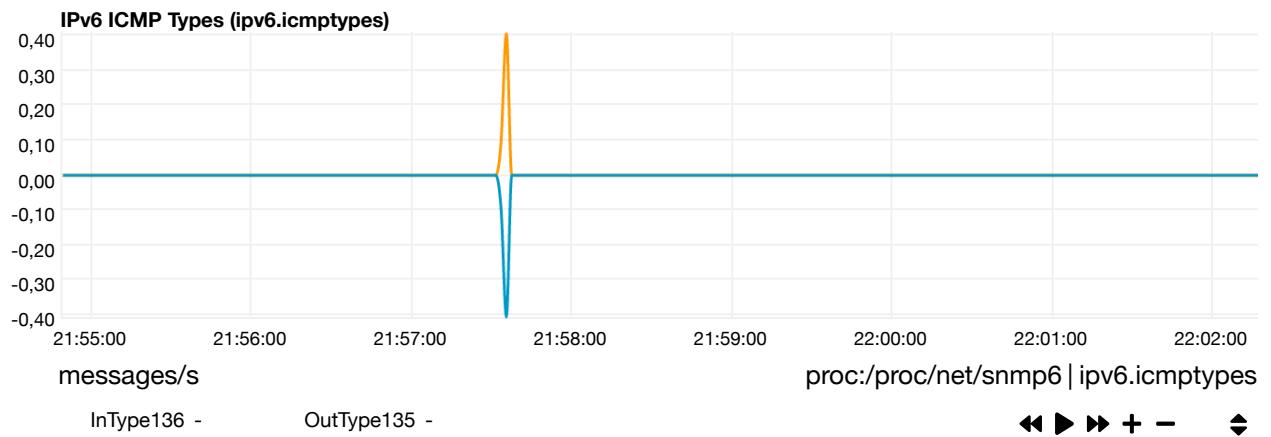


## multicast6



## icmp6

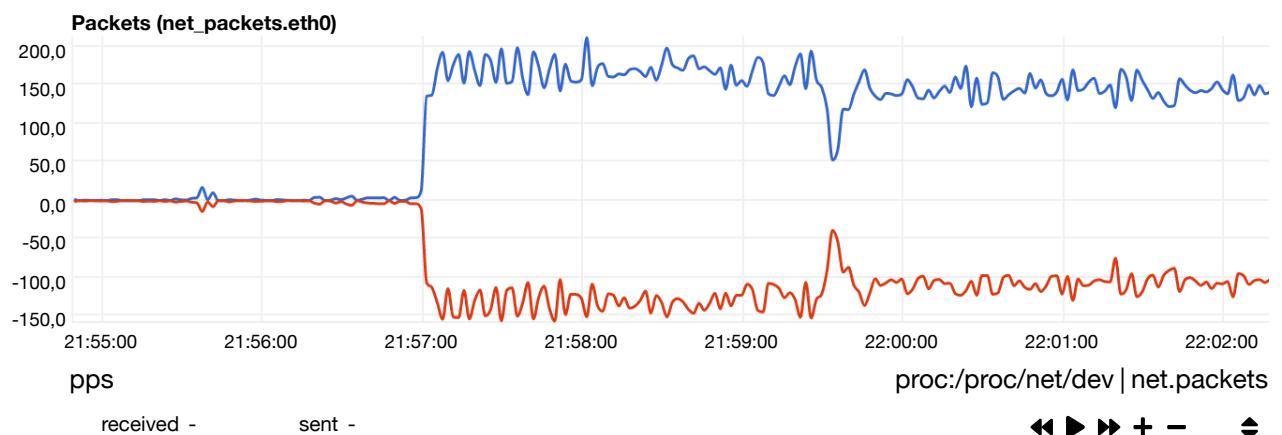
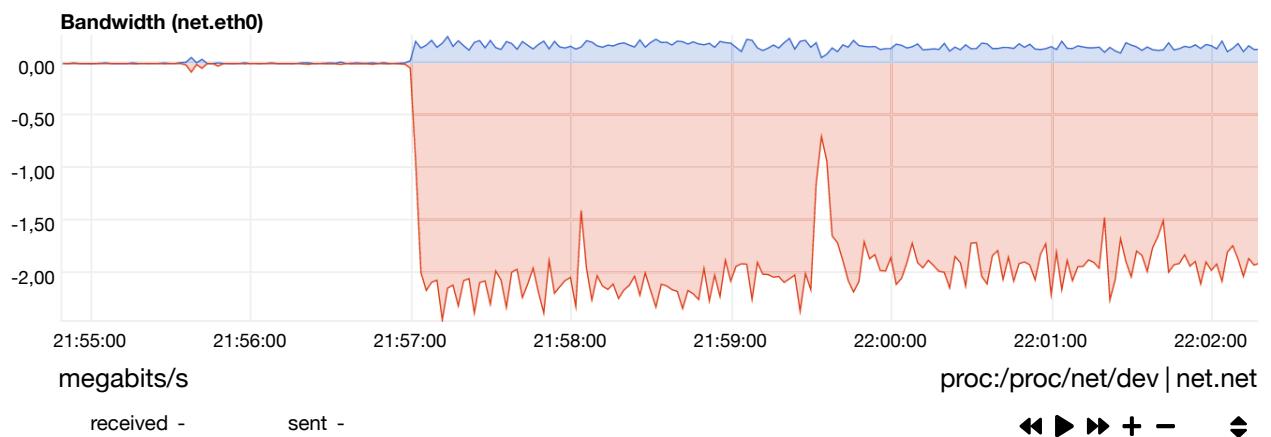


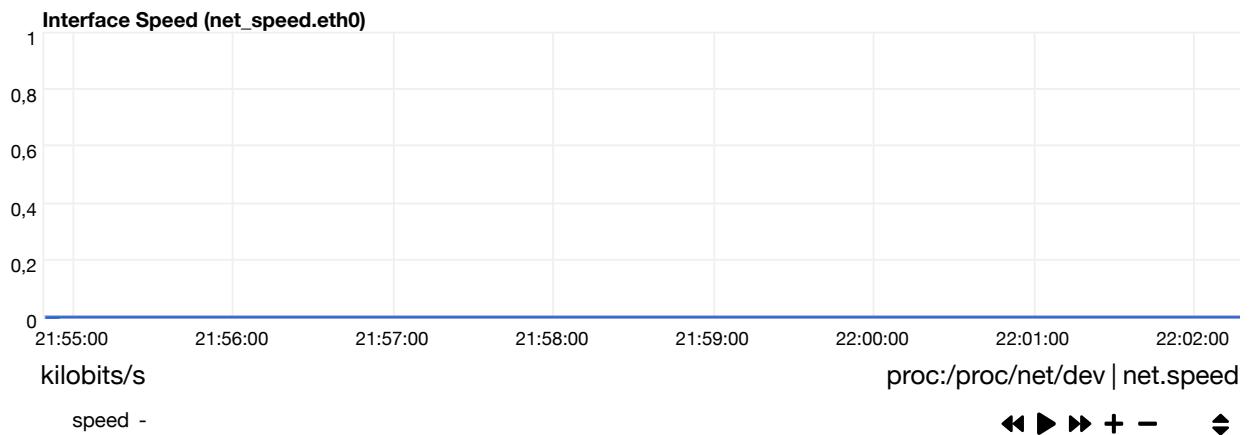


## Network Interfaces

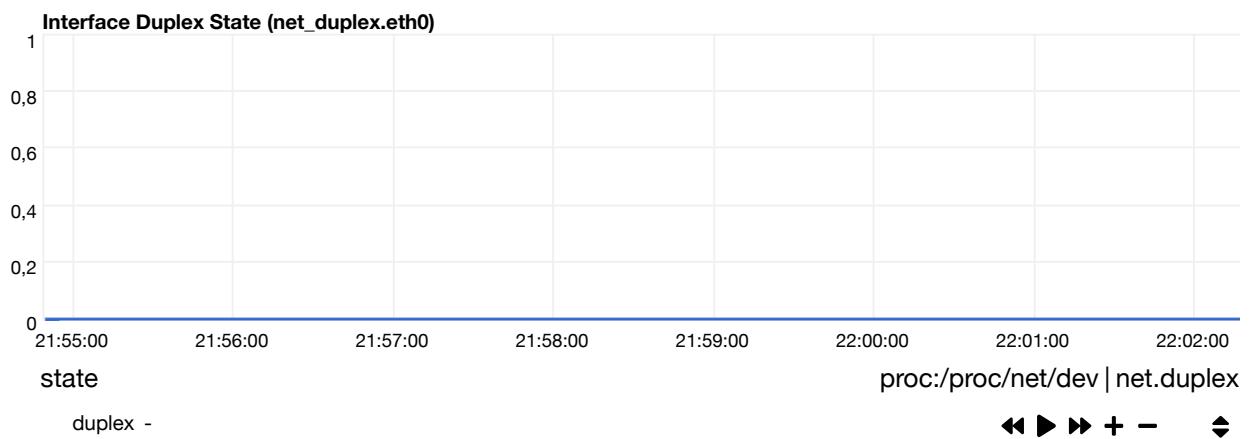
Performance metrics for network interfaces.

### eth0

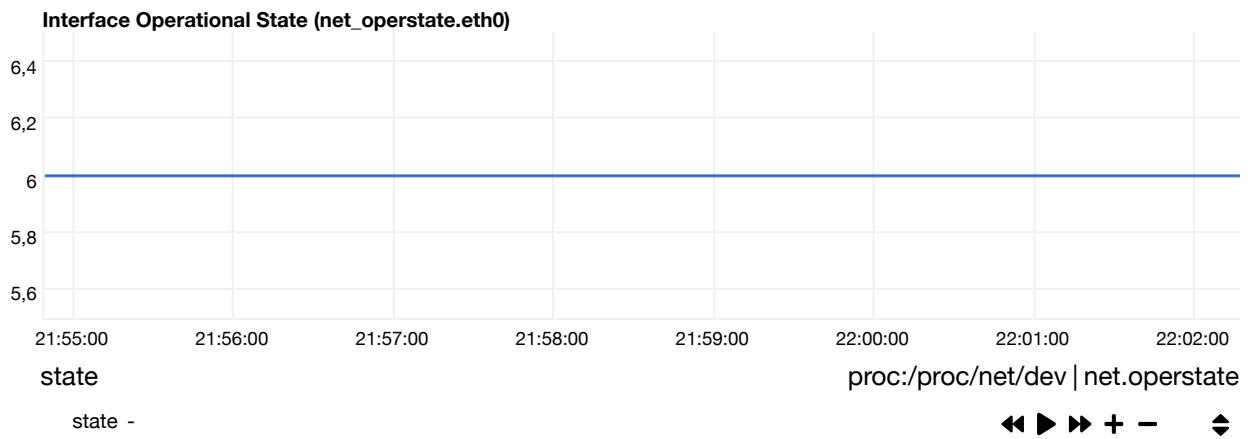




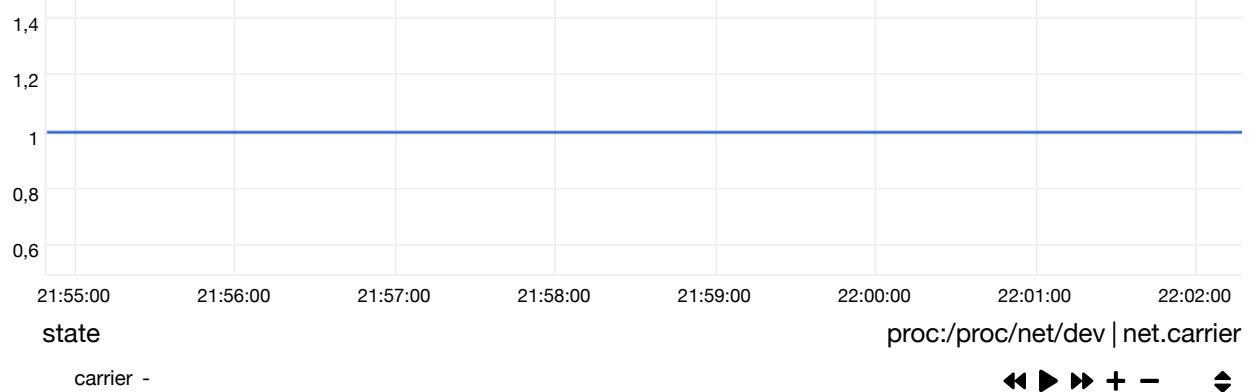
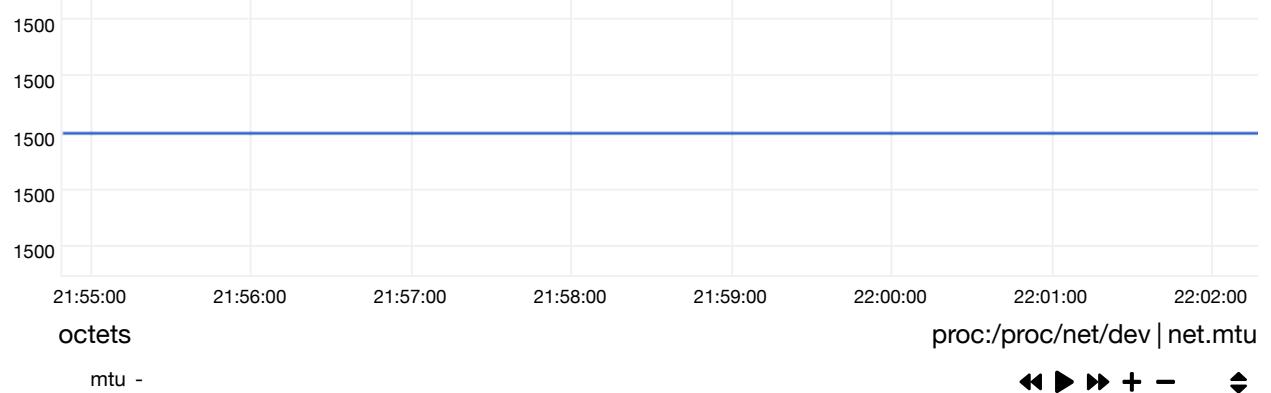
State map: 0 - unknown, 1 - half duplex, 2 - full duplex



State map: 0 - unknown, 1 - notpresent, 2 - down, 3 - lowerlayerdown, 4 - testing, 5 - dormant, 6 - up



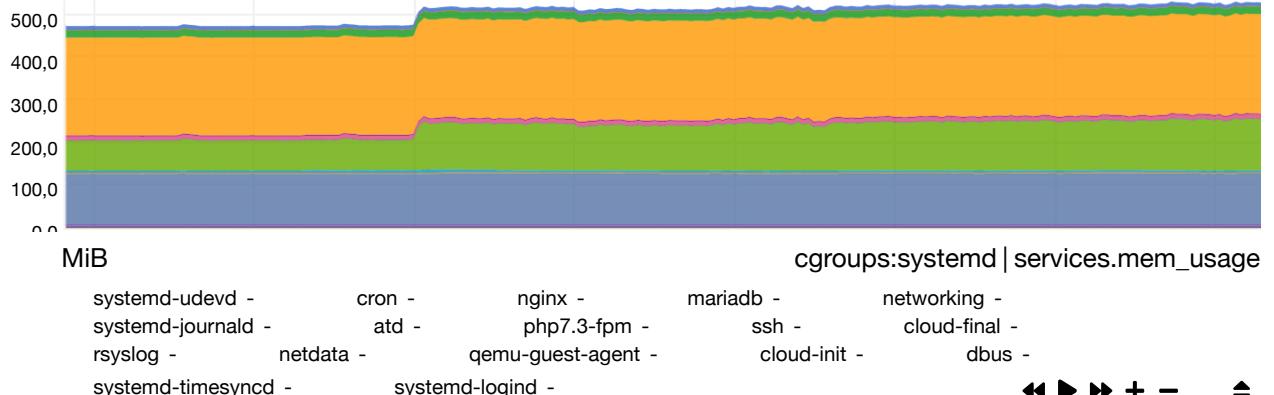
State map: 0 - down, 1 - up

**Interface Physical Link State (net\_carrier.eth0)****Interface MTU (net\_mtu.eth0)**

## ⚙️ systemd Services

Resources utilization of systemd services. netdata monitors all systemd services via CGROUPS (the resources accounting used by containers).

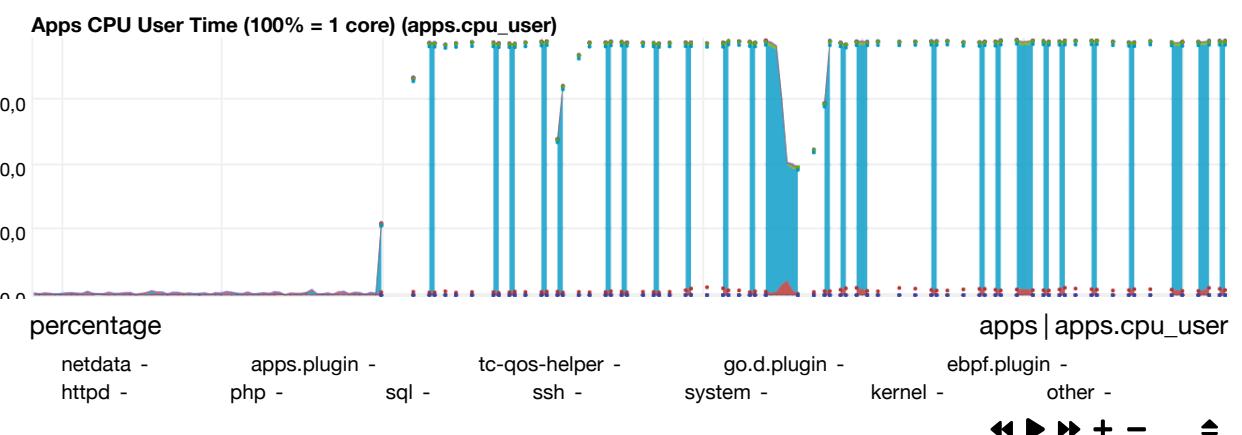
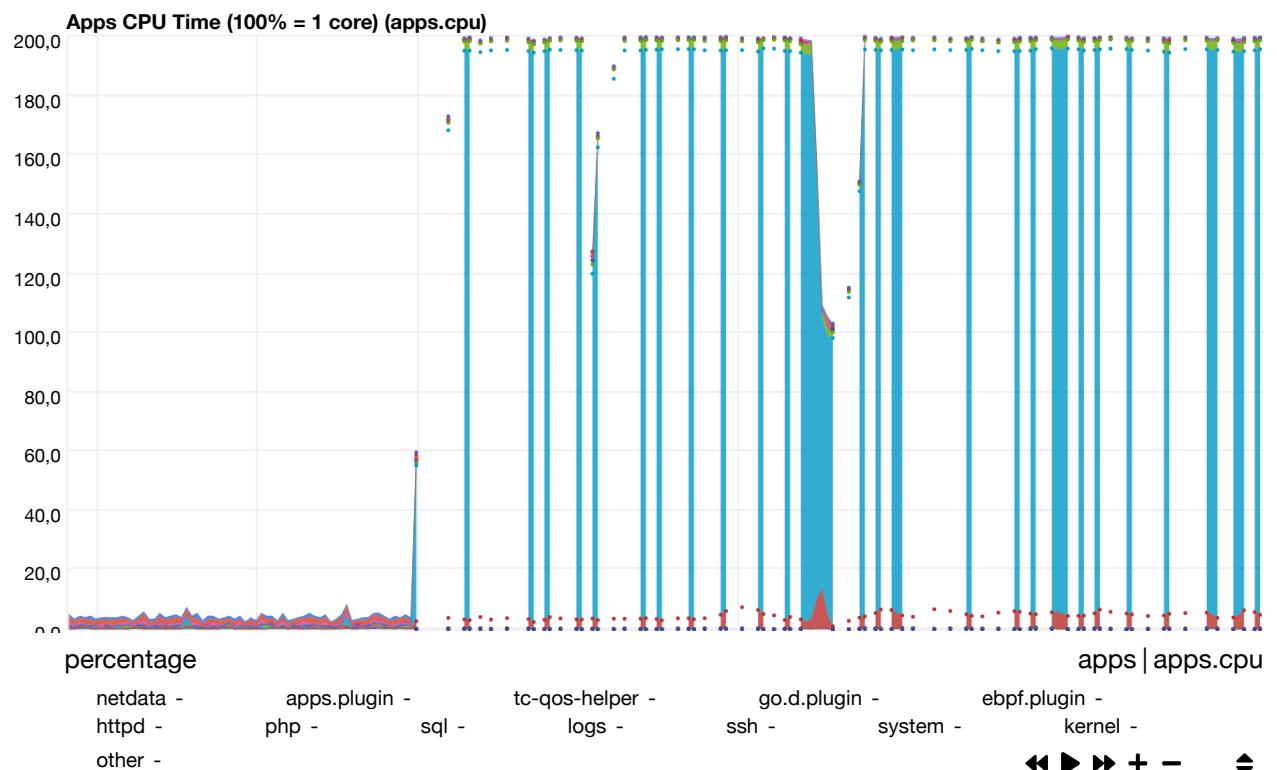
## mem

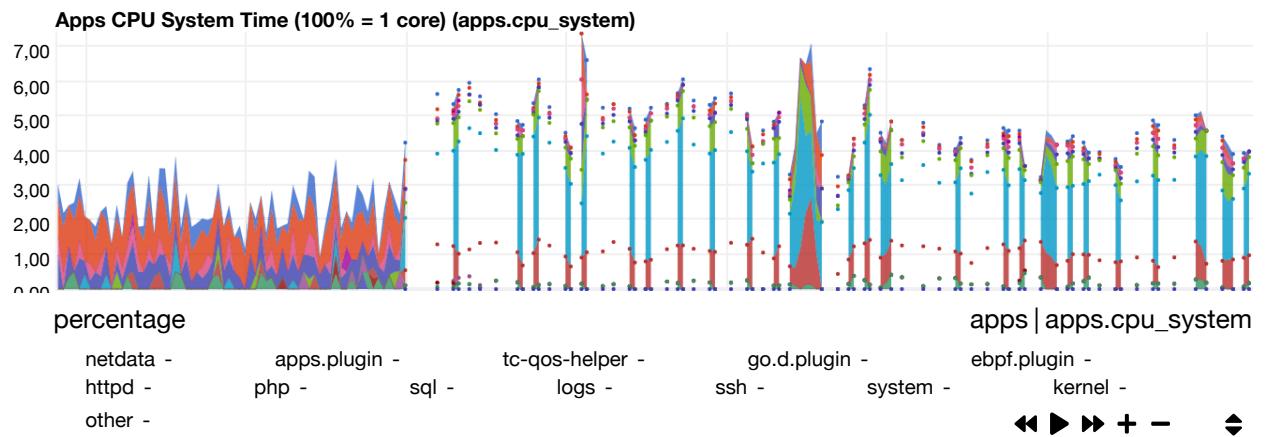
**Systemd Services Used Memory (services.mem\_usage)**

## ❤️ Applications

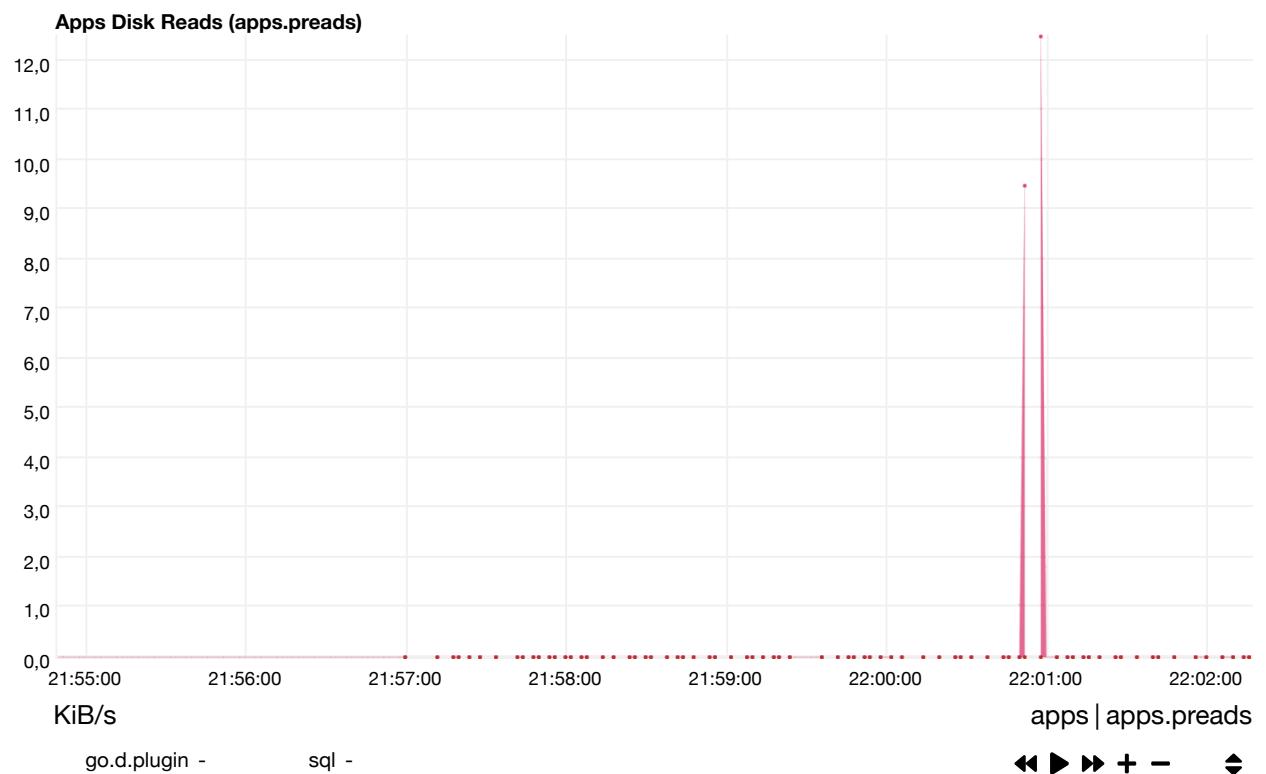
Per application statistics are collected using netdata's `apps.plugin`. This plugin walks through all processes and aggregates statistics for applications of interest, defined in `/etc/netdata/apps_groups.conf`, which can be edited by running `$ /etc/netdata/edit-config apps_groups.conf` (the default is here ([https://github.com/netdata/netdata/blob/master/collectors/apps.plugin/apps\\_groups.conf](https://github.com/netdata/netdata/blob/master/collectors/apps.plugin/apps_groups.conf))). The plugin internally builds a process tree (much like `ps fax` does), and groups processes together (evaluating both child and parent processes) so that the result is always a chart with a predefined set of dimensions (of course, only application groups found running are reported). The reported values are compatible with `top`, although the netdata plugin counts also the resources of exited children (unlike `top` which shows only the resources of the currently running processes). So for processes like shell scripts, the reported values include the resources used by the commands these scripts run within each timeframe.

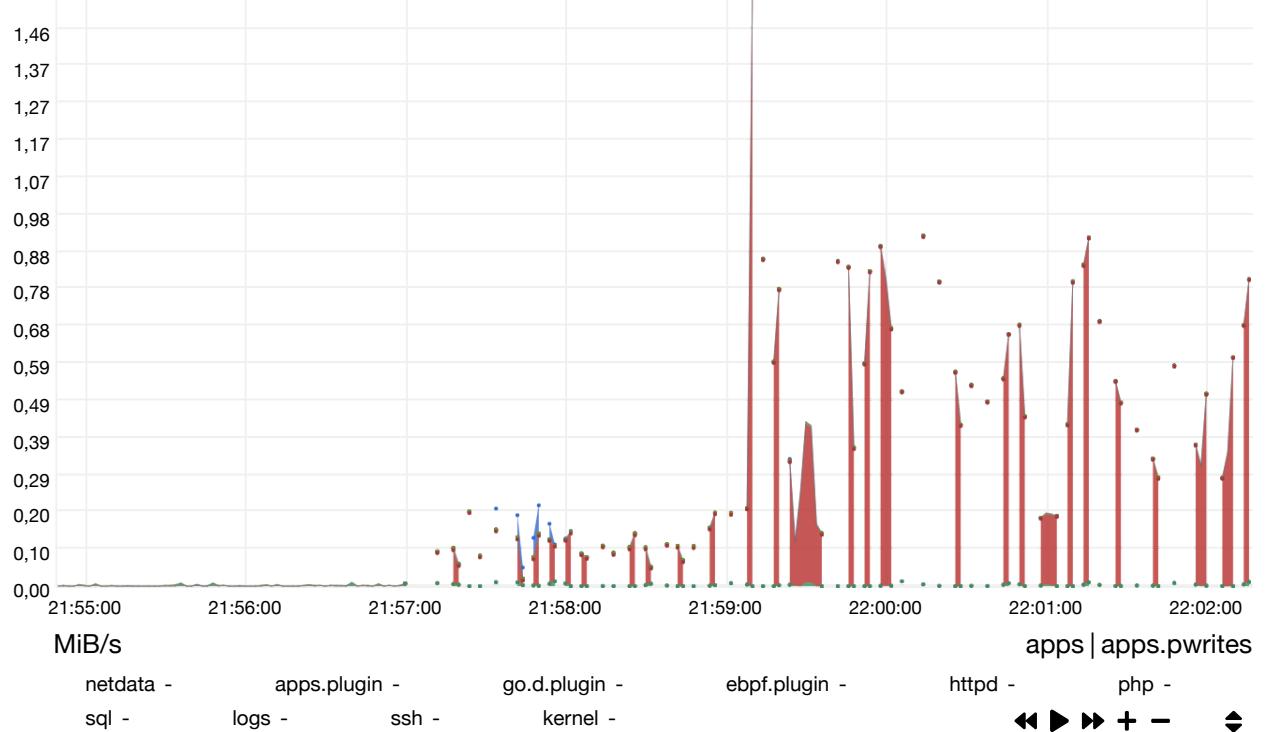
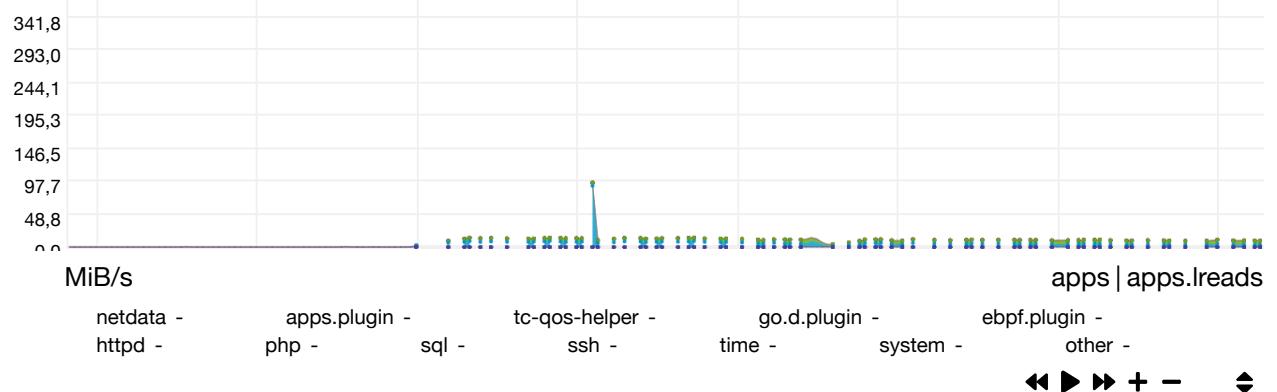
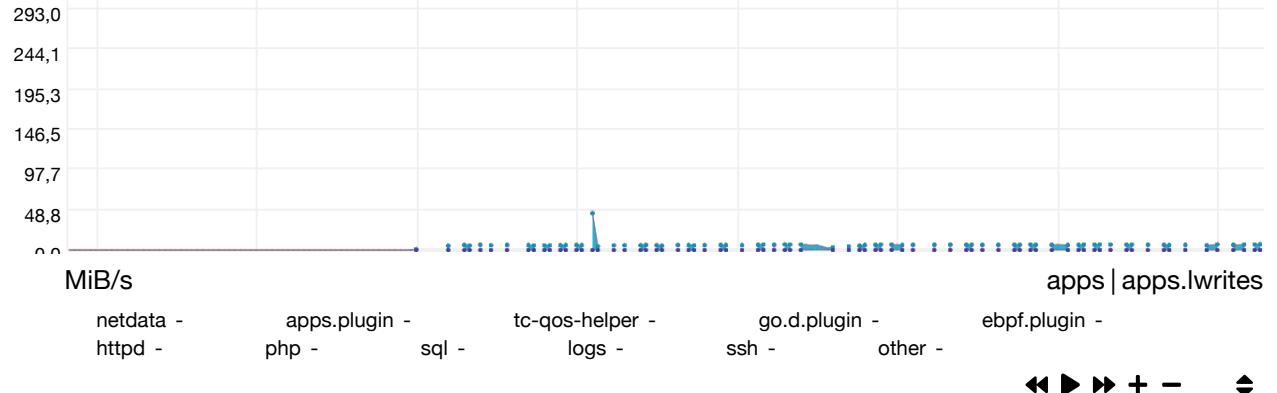
## Cpu

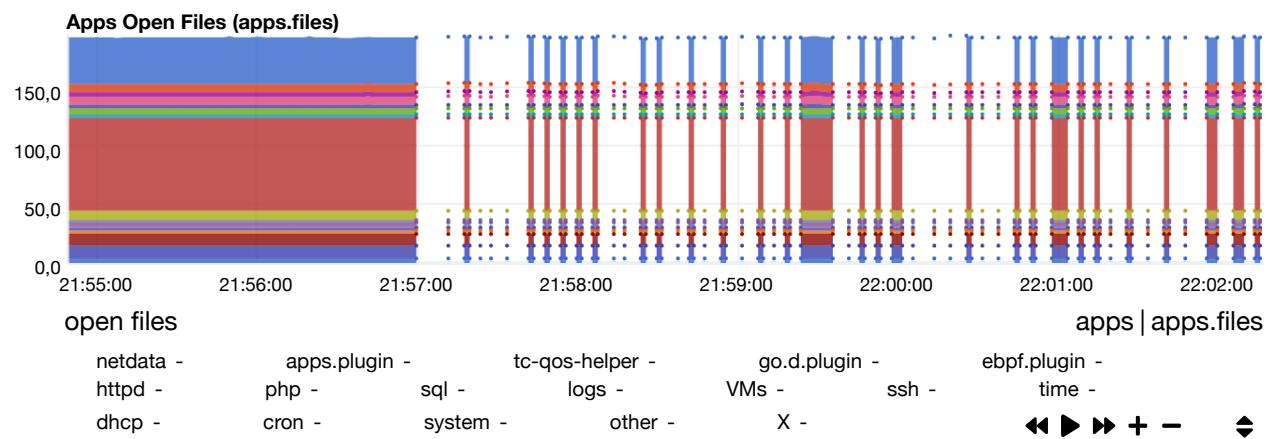




## disk

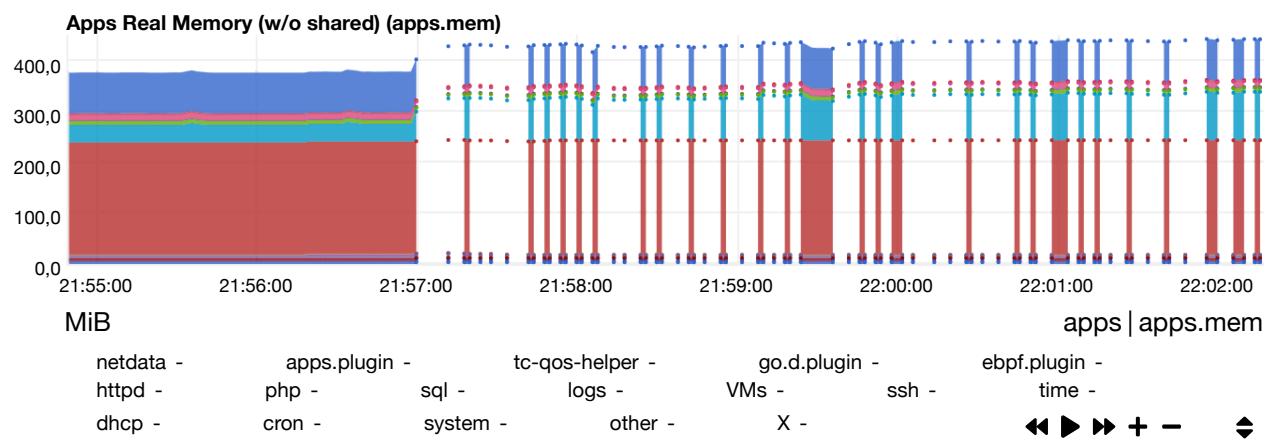


**Apps Disk Writes (apps.pwrites)****Apps Disk Logical Reads (apps.lreads)****Apps I/O Logical Writes (apps.lwrites)**



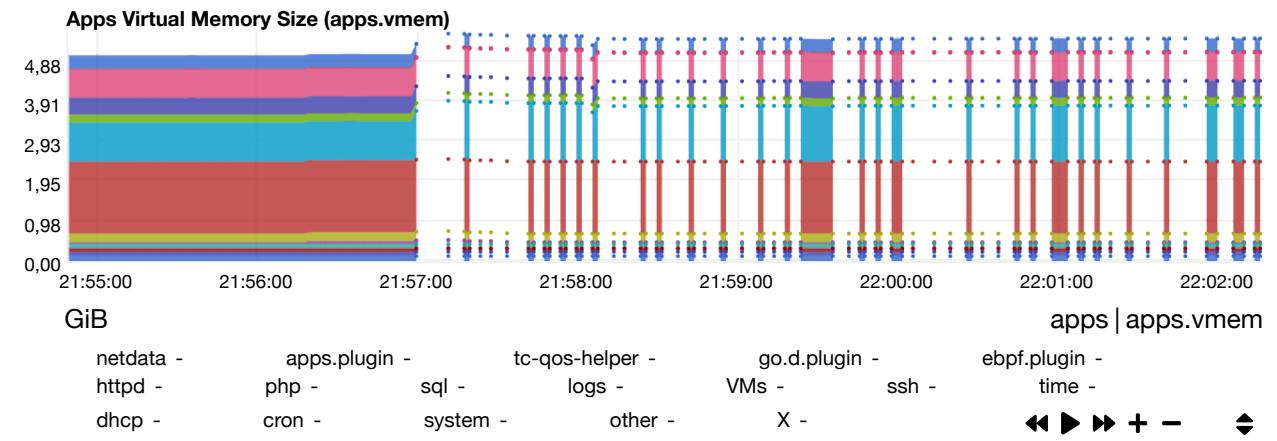
## mem

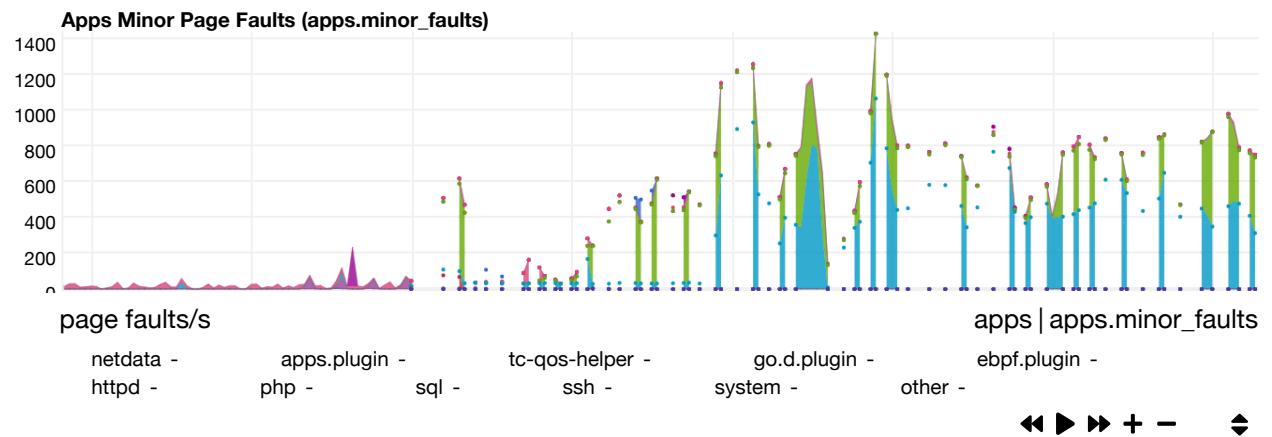
Real memory (RAM) used by applications. This does not include shared memory.



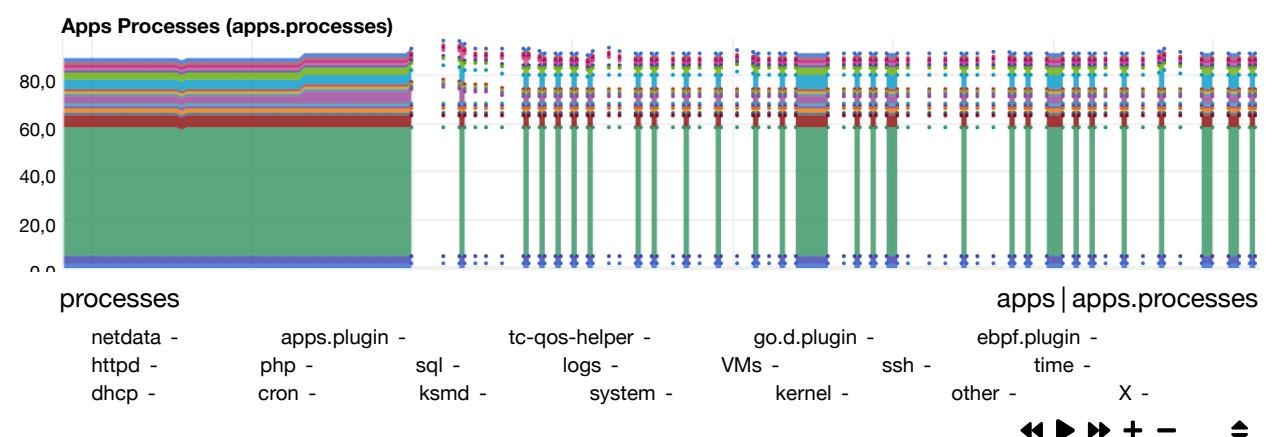
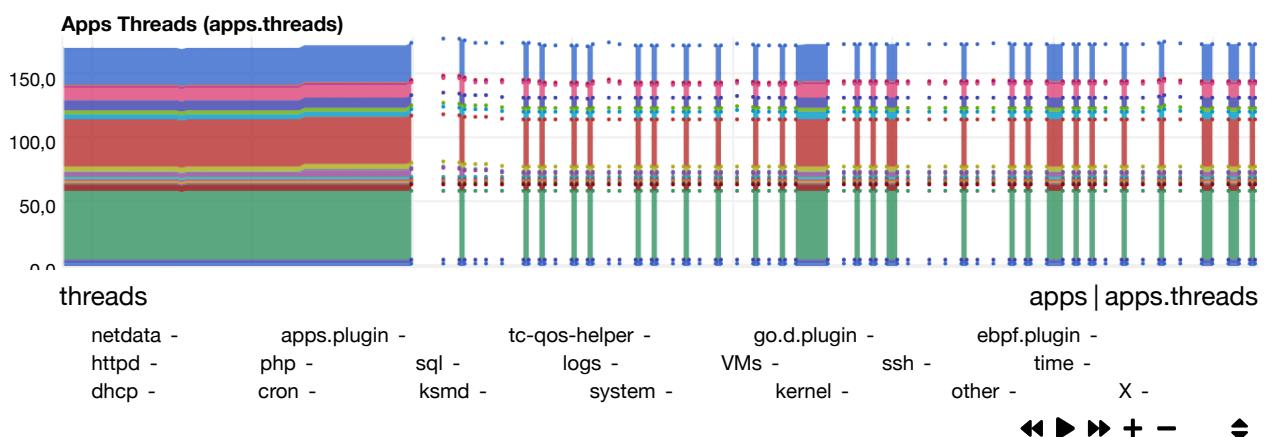
Virtual memory allocated by applications. Please check this article

(<https://github.com/netdata/netdata/tree/master/daemon#virtual-memory>) for more information.

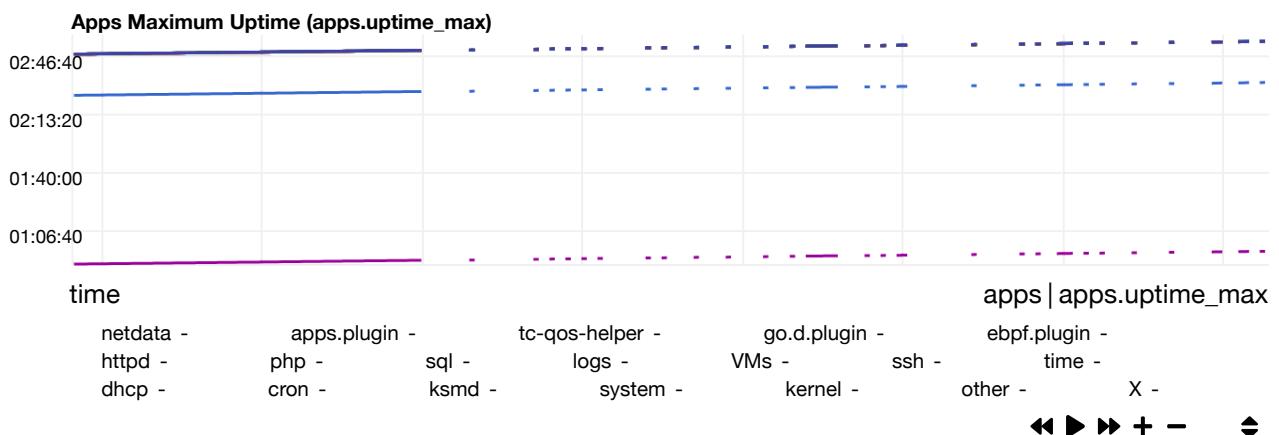
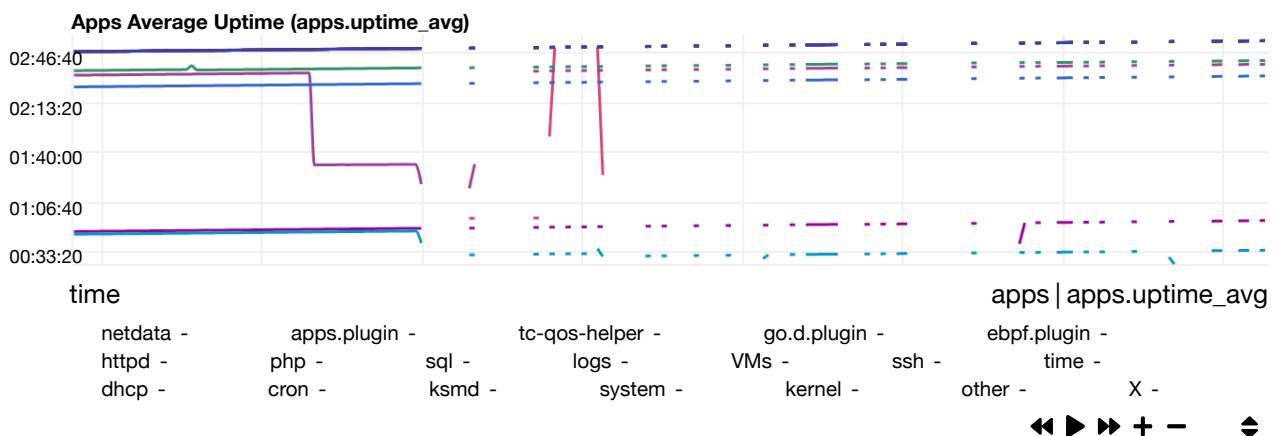
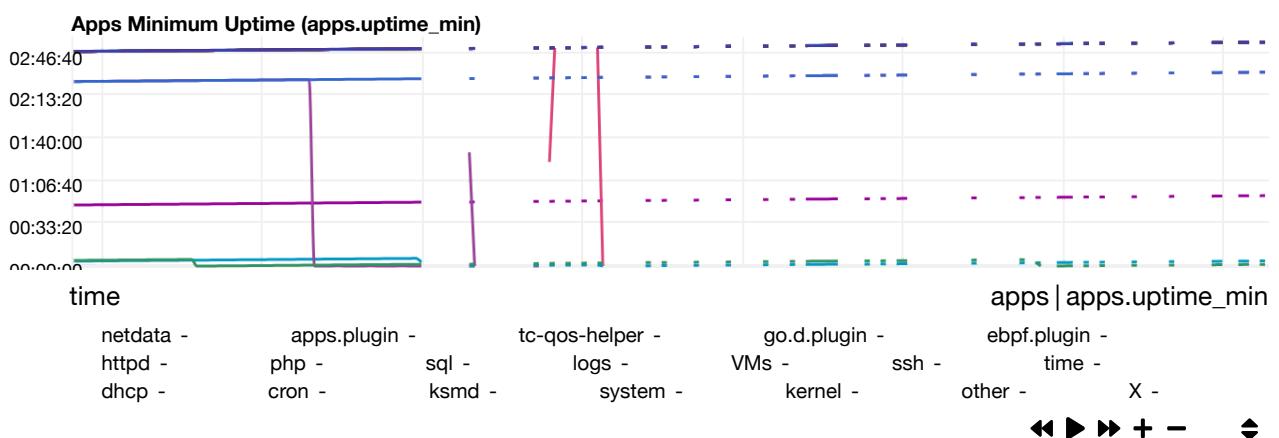
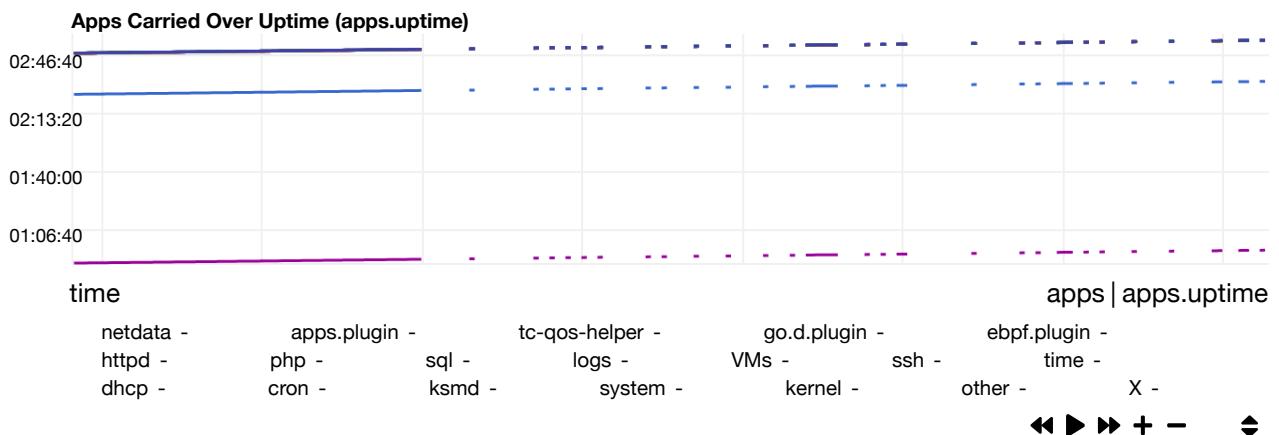


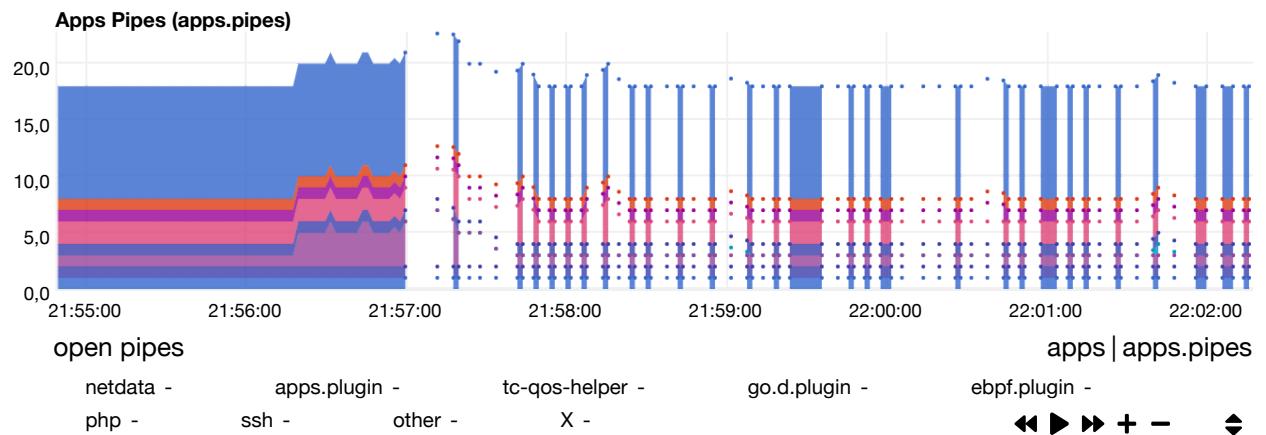


## processes

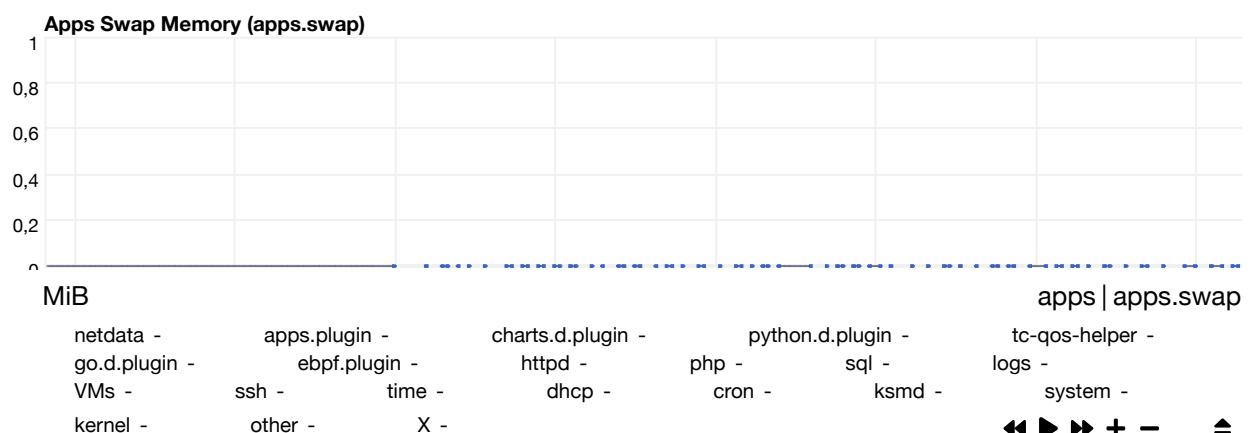


Carried over process group uptime since the Netdata restart. The period of time within which at least one process in the group was running.





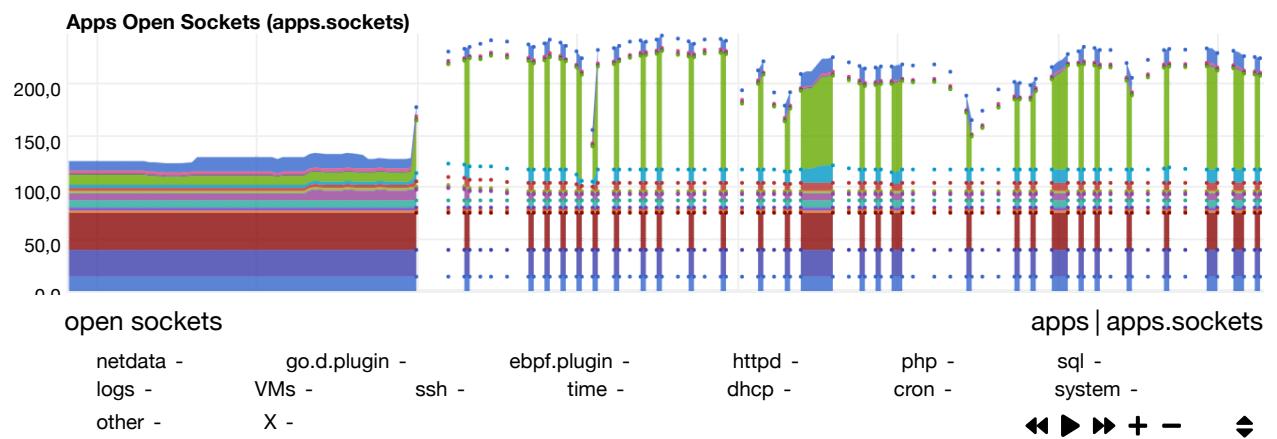
## swap



## Apps Major Page Faults (swap read) (apps.major\_faults)

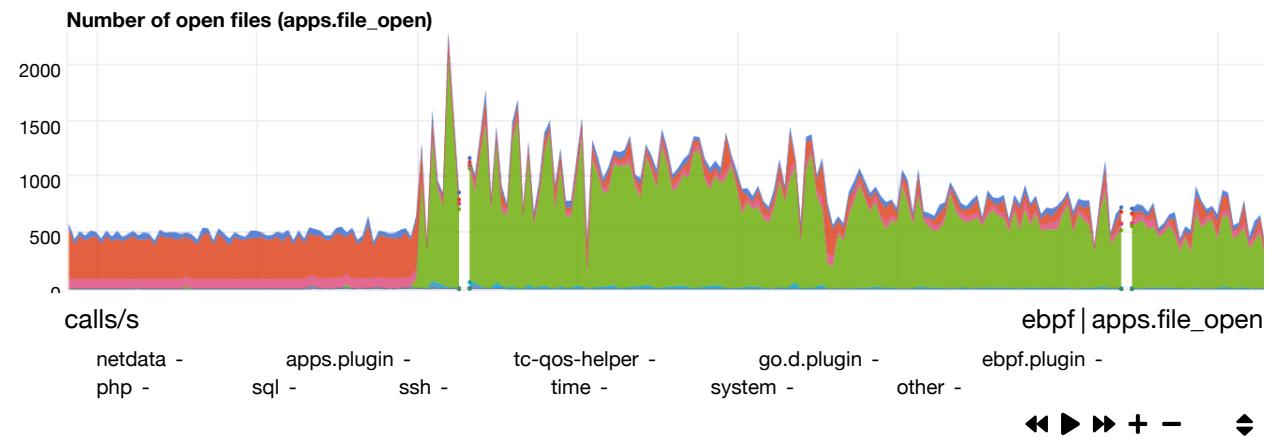


## net

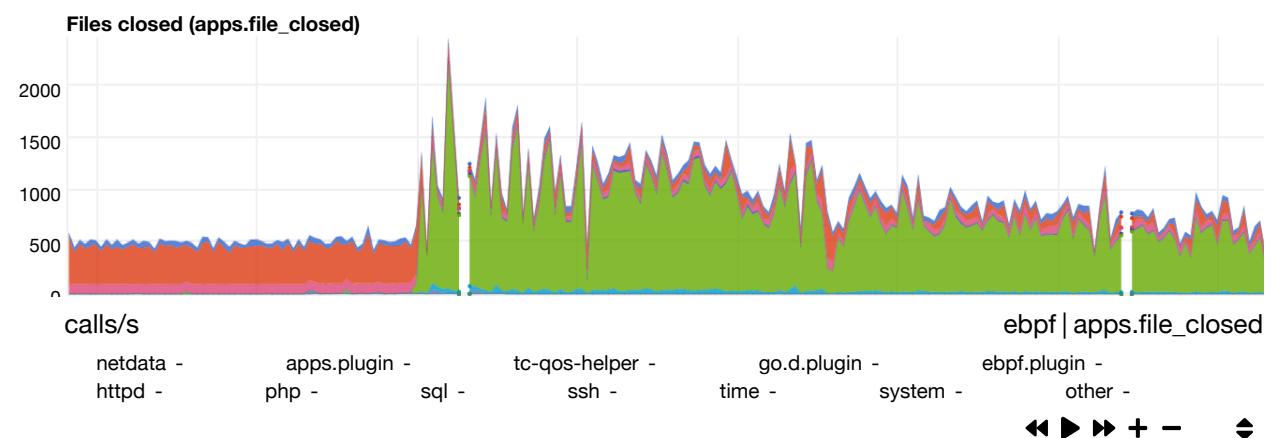


## file (eBPF)

Calls to the internal function `do_sys_open` ( For kernels newer than 5.5.19 we add a kprobe to `do_sys_openat2` ), which is the common function called from `open(2)` (<https://www.man7.org/linux/man-pages/man2/open.2.html>) and `openat(2)` (<https://www.man7.org/linux/man-pages/man2/openat.2.html>).



Calls to the internal function `__close_fd` (<https://elixir.bootlin.com/linux/v5.10/source/fs/file.c#L665>) or `close_fd` (<https://elixir.bootlin.com/linux/v5.11/source/fs/file.c#L617>) according to your kernel version, which is called from `close(2)` (<https://www.man7.org/linux/man-pages/man2/close.2.html>).

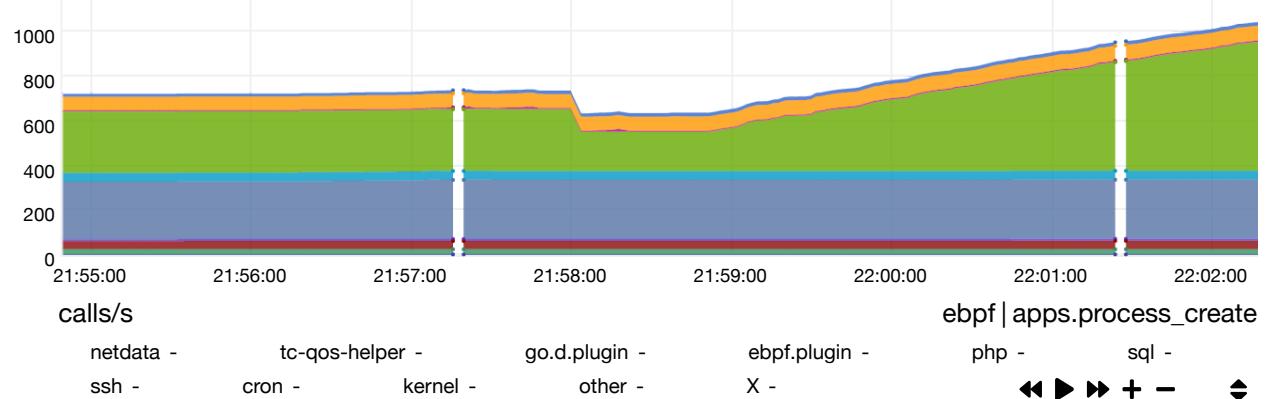


## process (eBPF)

## Calls to either do\_fork

(<https://www.ece.uic.edu/~yshi1/linux/lkse/node4.html#SECTION004210000000000000000000>), or kernel\_clone if you are running kernel newer than 5.9.16, to create a new task, which is the common name used to define process and tasks inside the kernel. Netdata identifies the process by counting the number of calls to sys\_clone (<https://linux.die.net/man/2/clone>) that do not have the flag CLONE\_THREAD set.

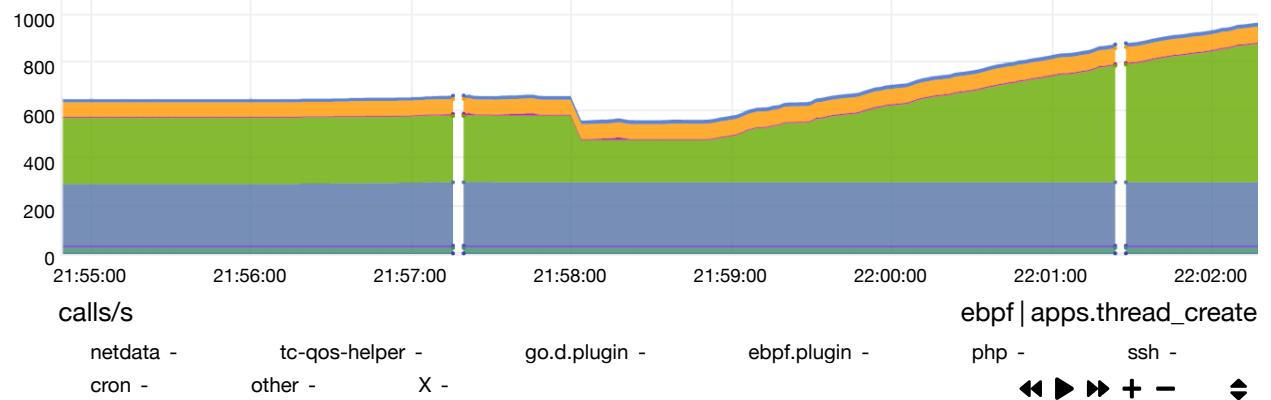
## Process started (apps.process\_create)



## Calls to either do\_fork

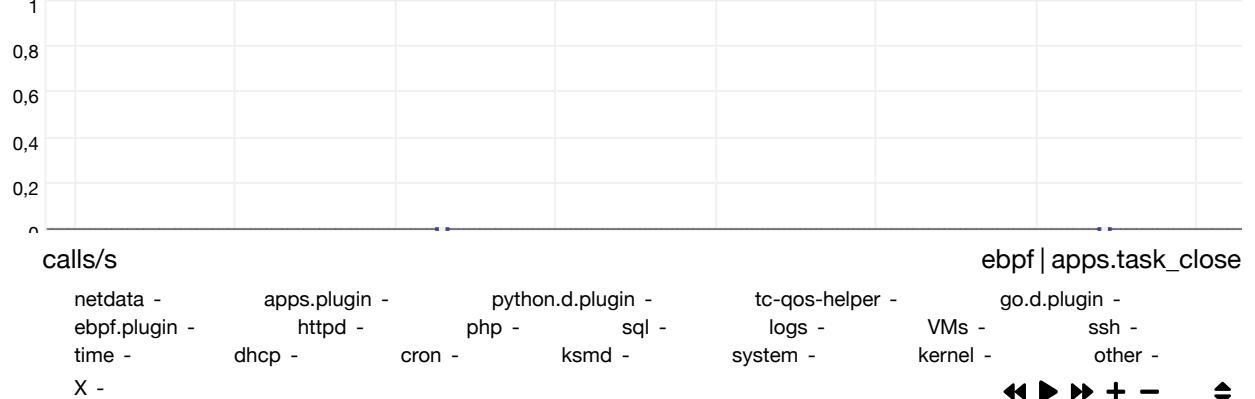
(<https://www.ece.uic.edu/~yshi1/linux/lkse/node4.html#SECTION004210000000000000000000>), or kernel\_clone if you are running kernel newer than 5.9.16, to create a new task, which is the common name used to define process and tasks inside the kernel. Netdata identifies the threads by counting the number of calls to sys\_clone (<https://linux.die.net/man/2/clone>) that have the flag CLONE\_THREAD set.

## Threads started (apps.thread\_create)



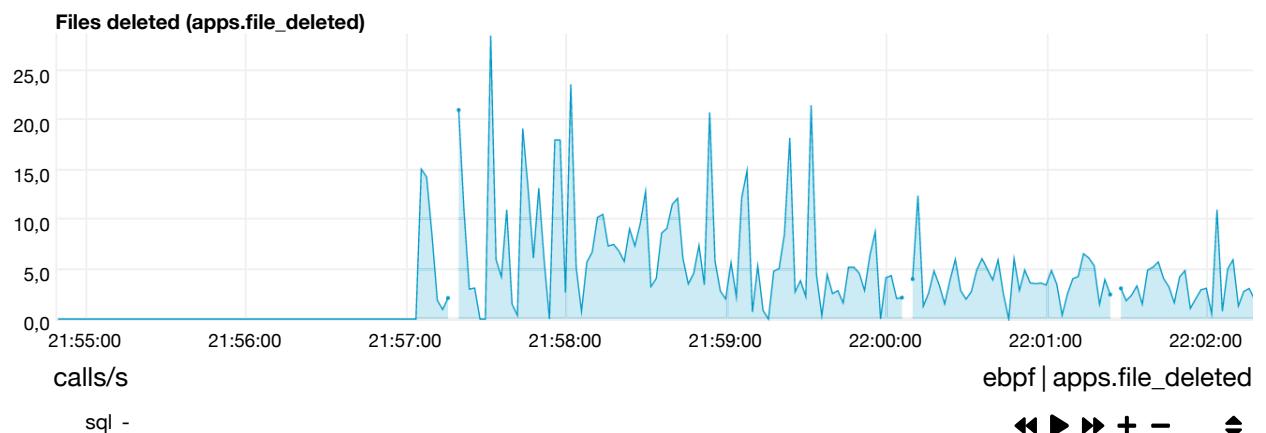
Calls to the functions responsible for closing (do\_exit (<https://www.informit.com/articles/article.aspx?p=370047&seqNum=4>)) and releasing (release\_task (<https://www.informit.com/articles/article.aspx?p=370047&seqNum=4>)) tasks.

## Tasks closed (apps.task\_close)

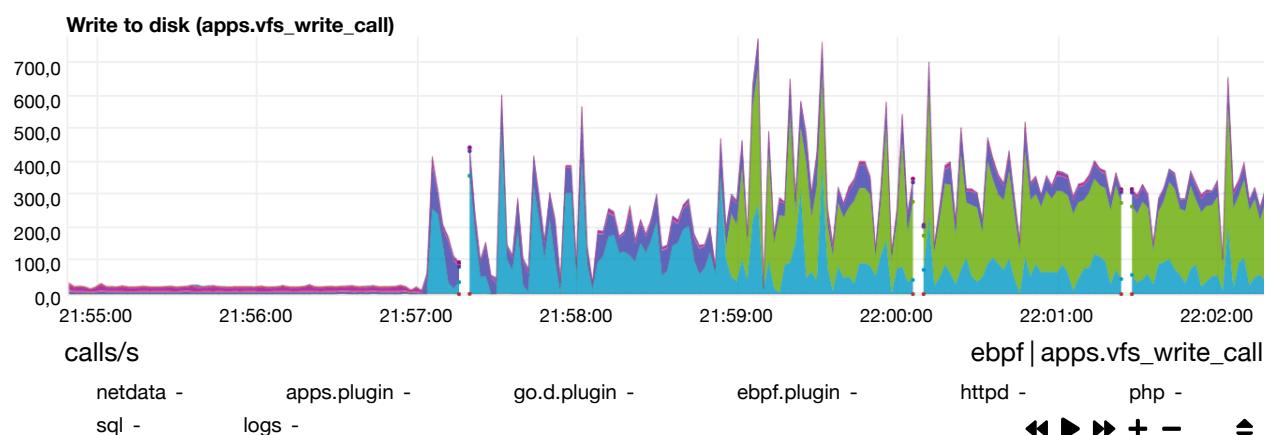


## vfs (eBPF)

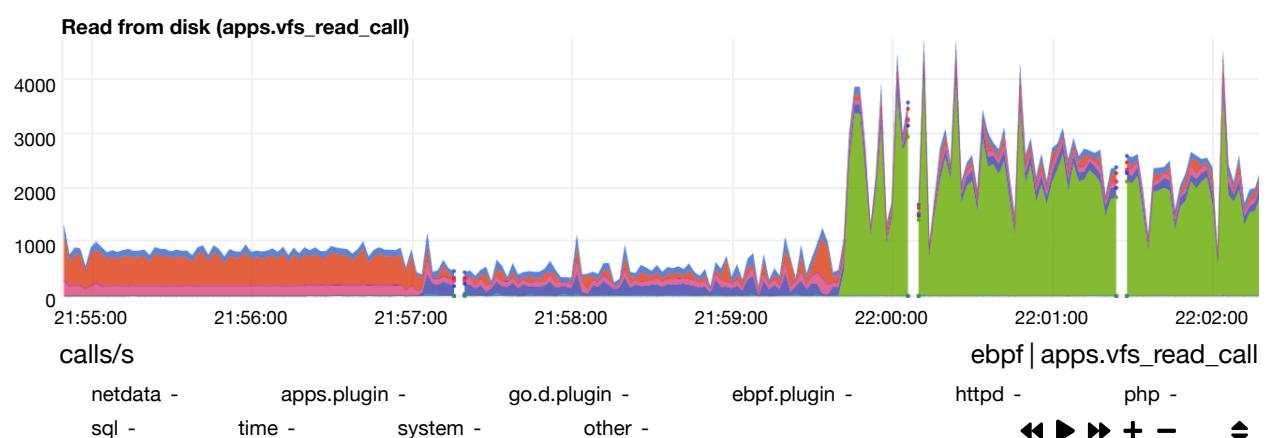
Calls to the function `vfs_unlink` (<https://www.kernel.org/doc/htmldocs/filesystems/API-vfs-unlink.html>). This chart does not show all events that remove files from the filesystem, because filesystems can create their own functions to remove files.



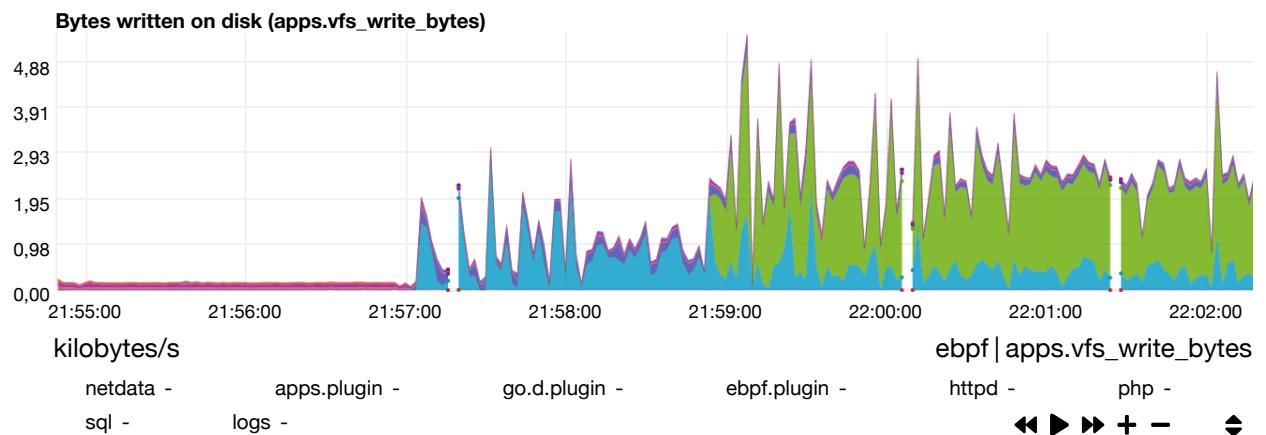
Successful calls to the function `vfs_write` ([https://topic.alibabacloud.com/a/kernel-state-file-operation-work-information-kernel\\_8\\_8\\_20287135.html](https://topic.alibabacloud.com/a/kernel-state-file-operation-work-information-kernel_8_8_20287135.html)). This chart may not show all filesystem events if it uses other functions to store data on disk.



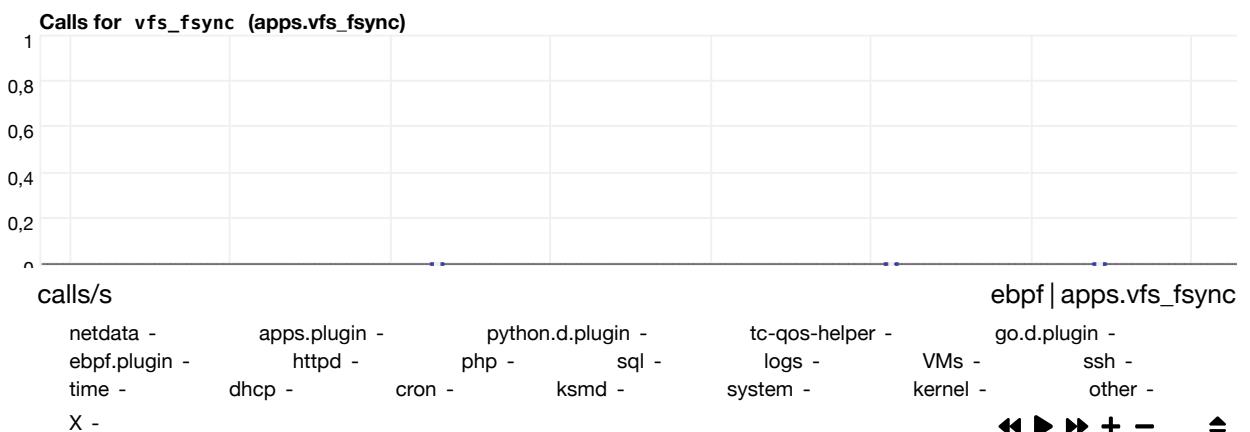
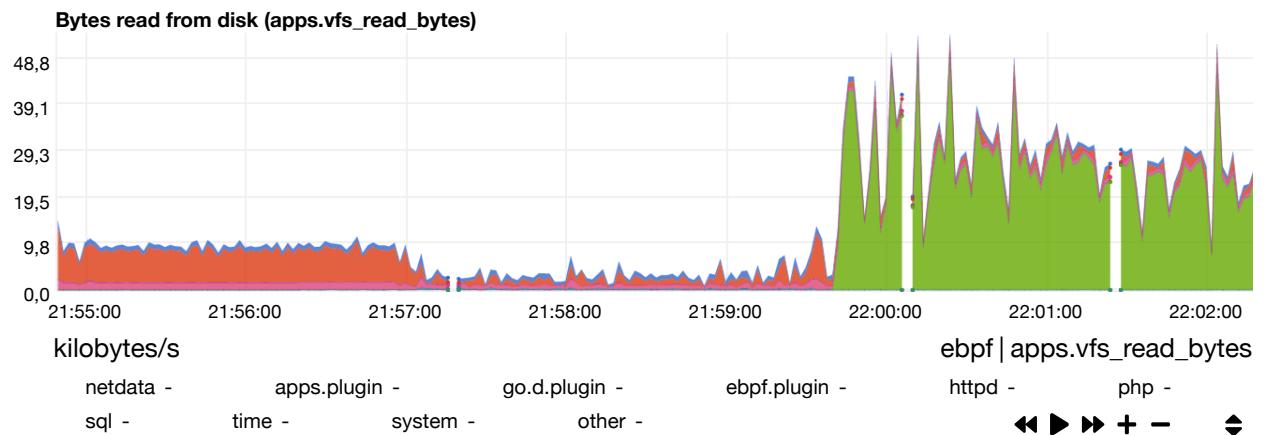
Successful calls to the function `vfs_read` ([https://topic.alibabacloud.com/a/kernel-state-file-operation-work-information-kernel\\_8\\_8\\_20287135.html](https://topic.alibabacloud.com/a/kernel-state-file-operation-work-information-kernel_8_8_20287135.html)). This chart may not show all filesystem events if it uses other functions to store data on disk.

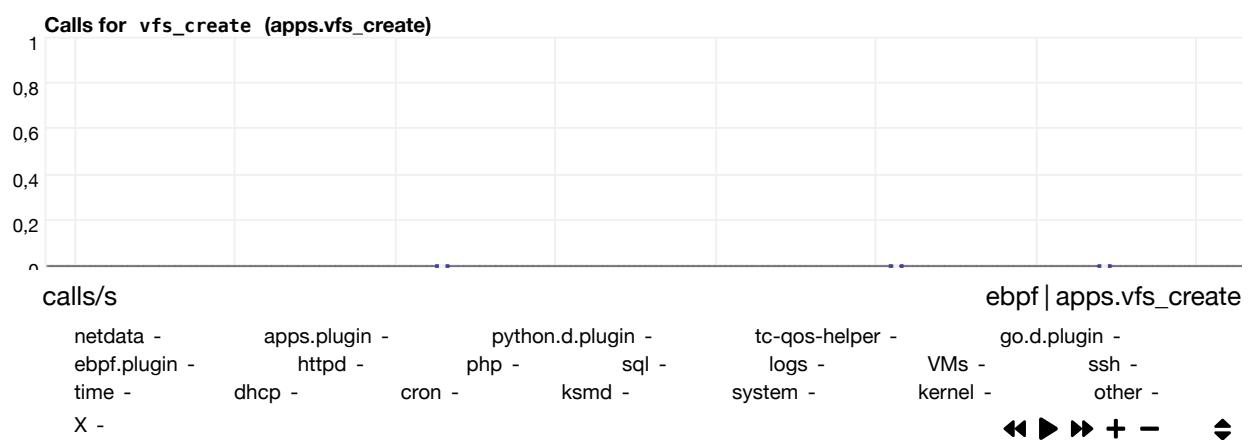
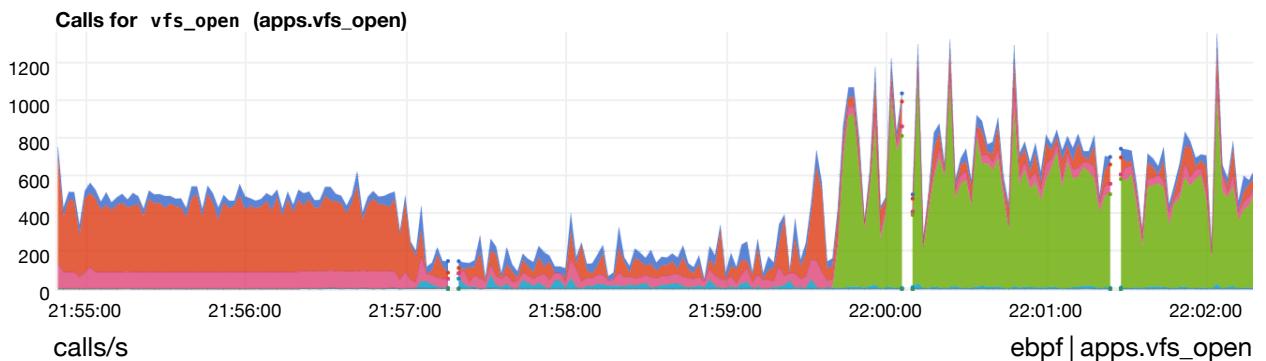


Total of bytes successfully written using the function vfs\_write  
[https://topic.alibabacloud.com/a/kernel-state-file-operation-\\_\\_-work-information-kernel\\_8\\_8\\_20287135.html](https://topic.alibabacloud.com/a/kernel-state-file-operation-__-work-information-kernel_8_8_20287135.html).

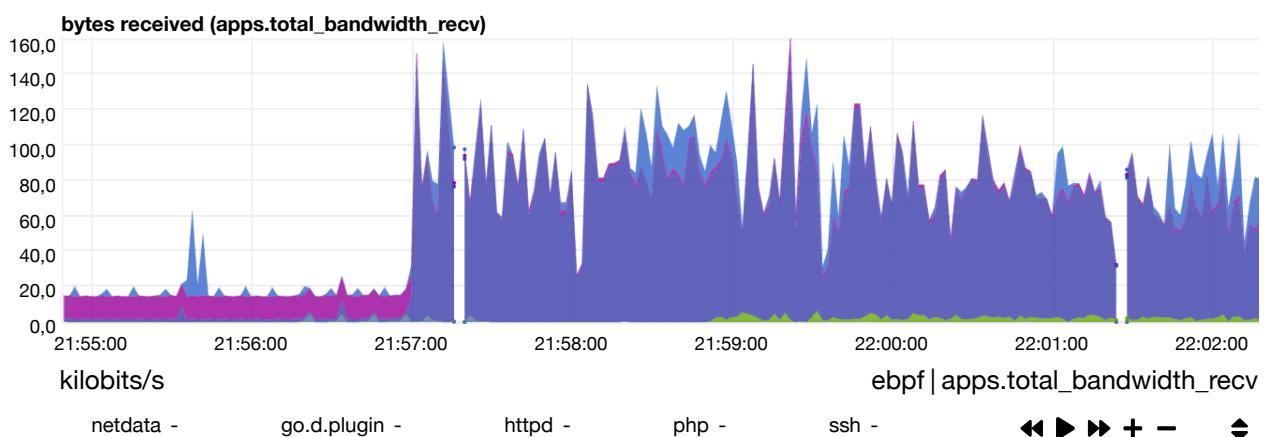
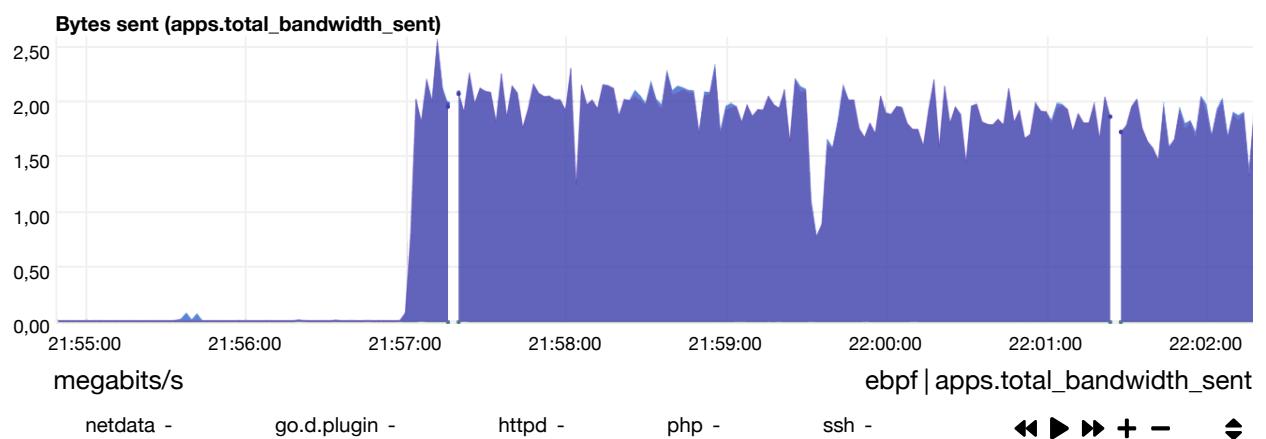


Total of bytes successfully read using the function vfs\_read ([https://topic.alibabacloud.com/a/kernel-state-file-operation-\\_\\_-work-information-kernel\\_8\\_8\\_20287135.html](https://topic.alibabacloud.com/a/kernel-state-file-operation-__-work-information-kernel_8_8_20287135.html)).





## net (eBPF)



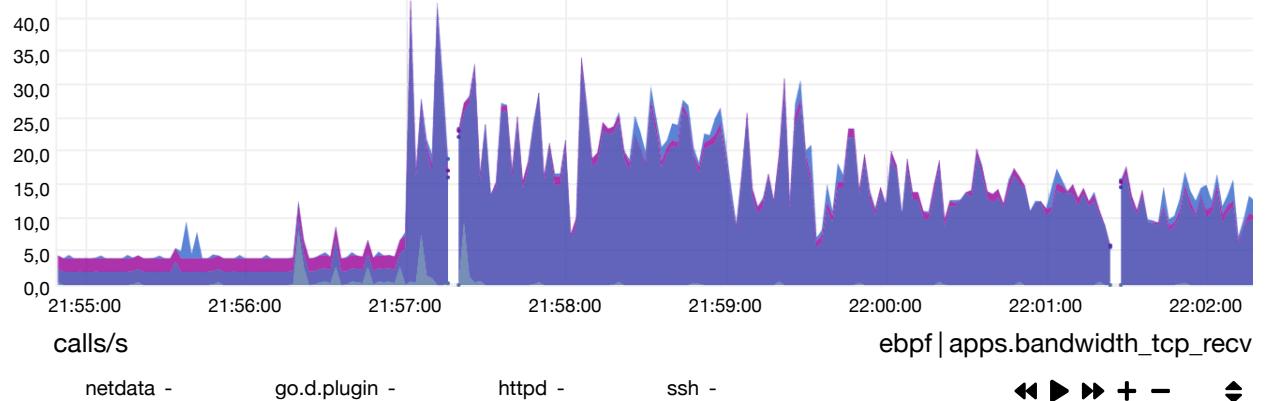
Calls for function `tcp_sendmsg` .

Calls for `tcp_sendmsg (apps.bandwidth_tcp_send)`



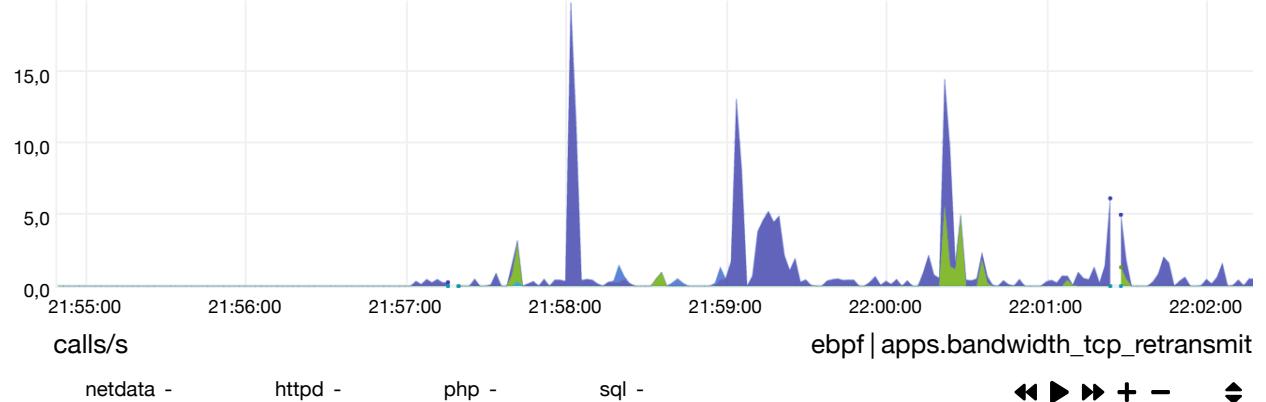
Calls for functions `tcp_cleanup_rbuf` . We use `tcp_cleanup_rbuf` instead `tcp_recvmsg` , because this last misses `tcp_read_sock()` traffic and we would also need to have more probes to get the socket and package size.

Calls for `tcp_cleanup_rbuf (apps.bandwidth_tcp_recv)`

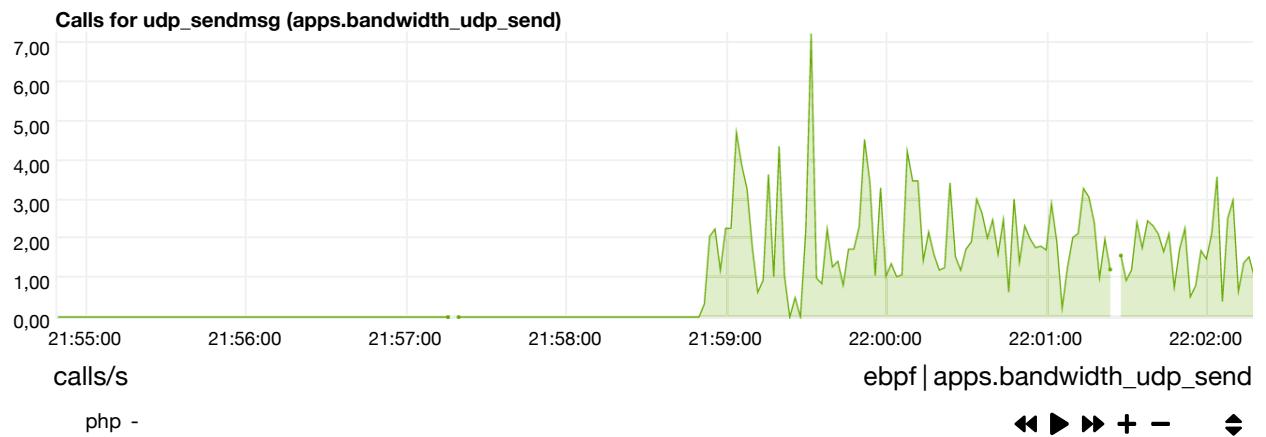


Calls for functions `tcp_retransmit_skb` .

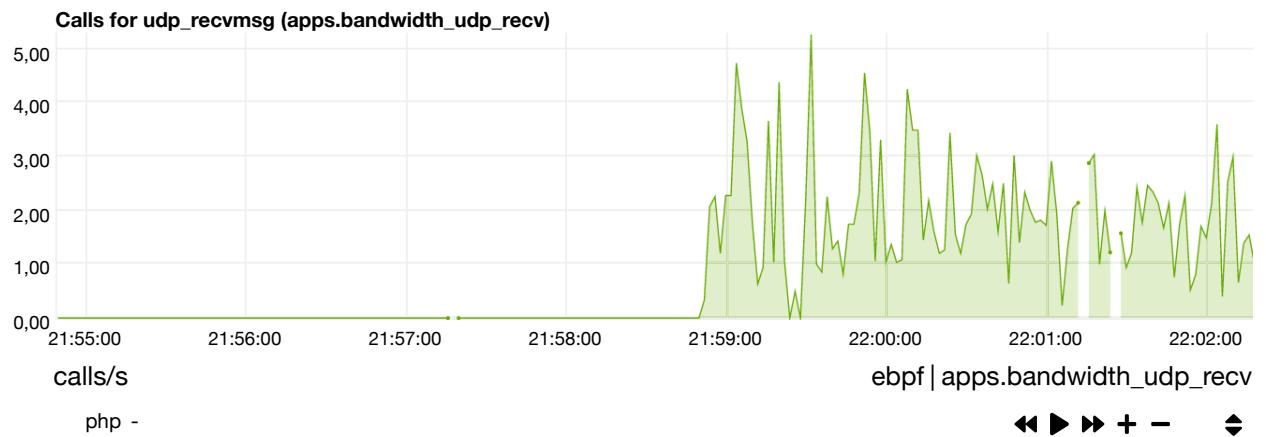
Calls for `tcp_retransmit (apps.bandwidth_tcp_retransmit)`



Calls for function `udp_sendmsg` .



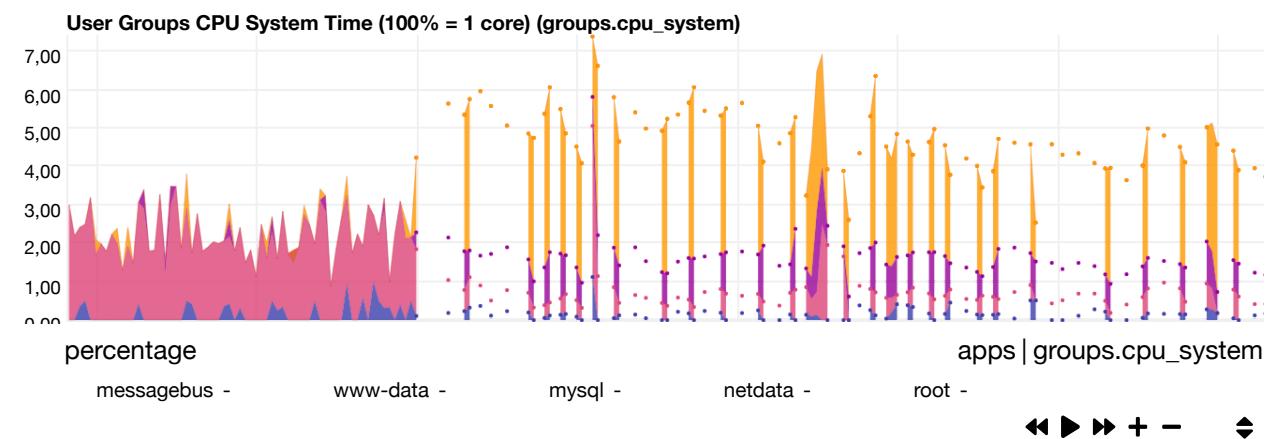
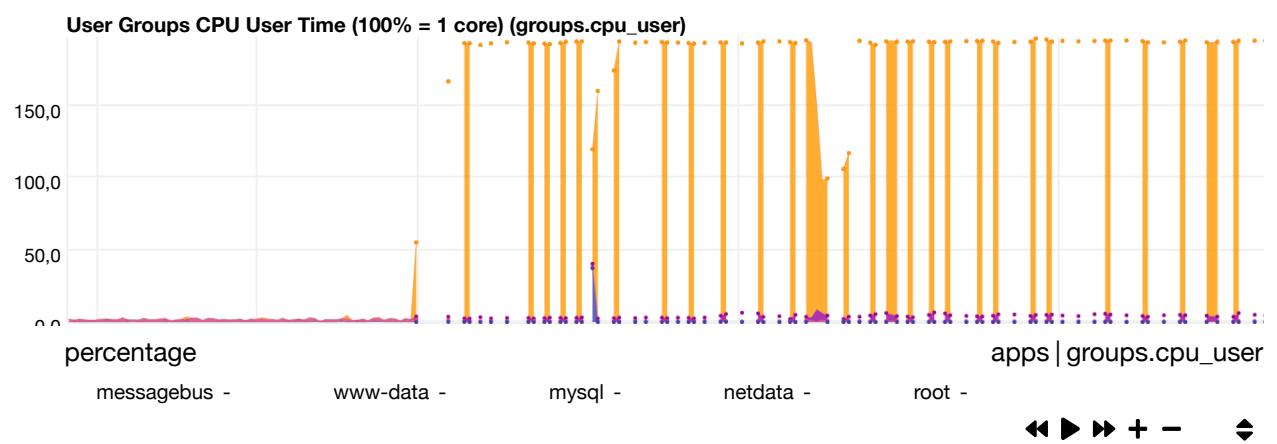
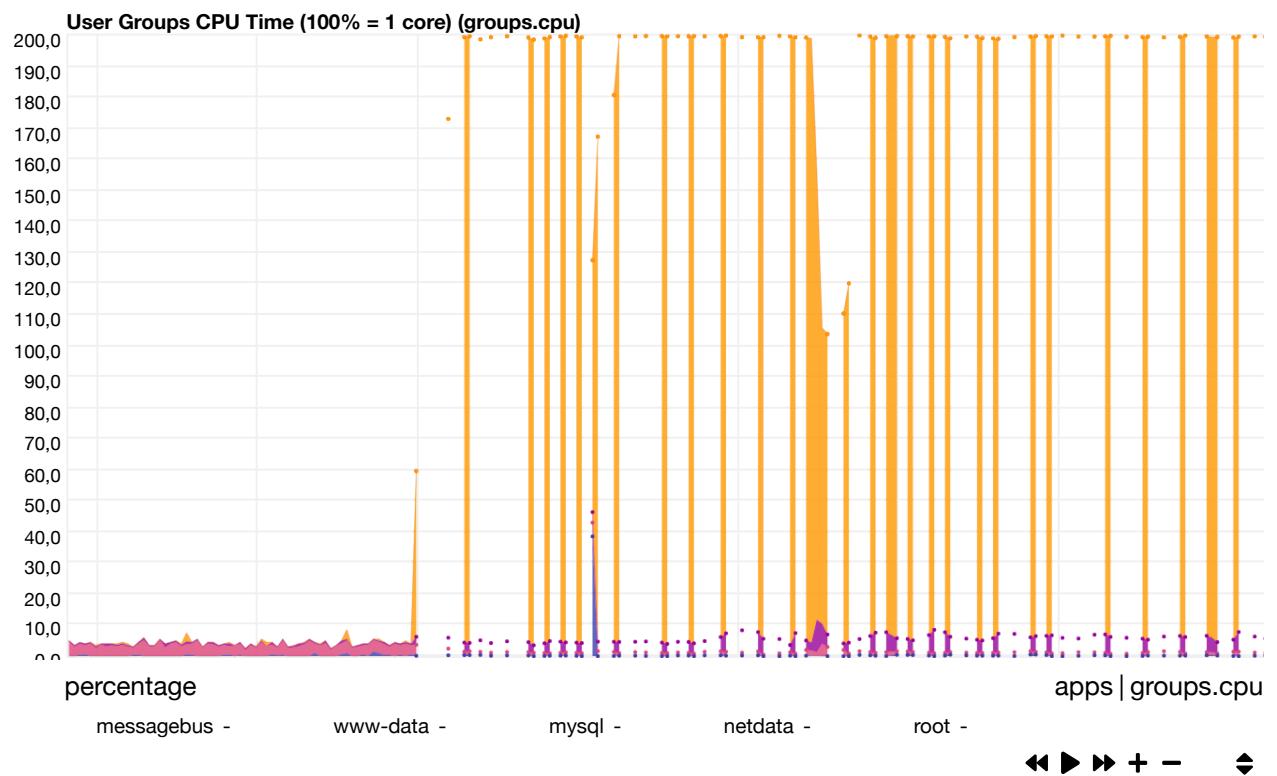
Calls for function `udp_recvmsg` .



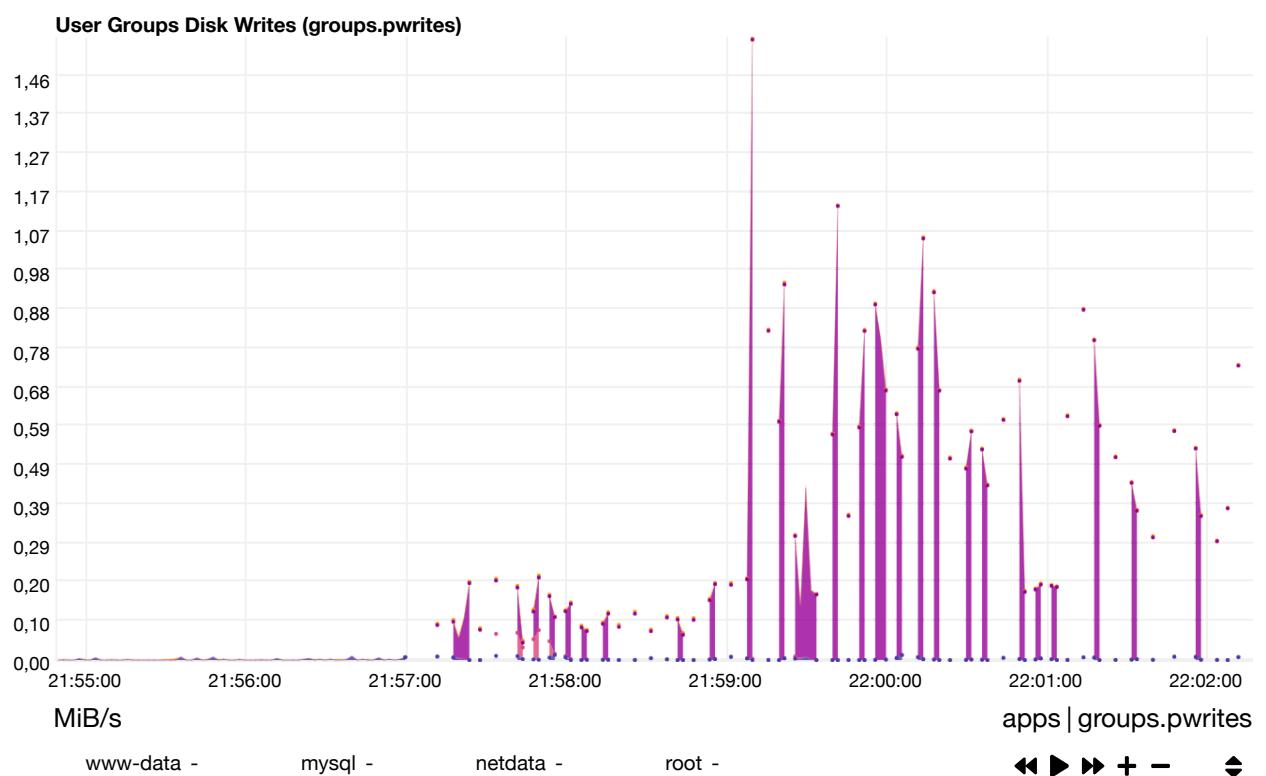
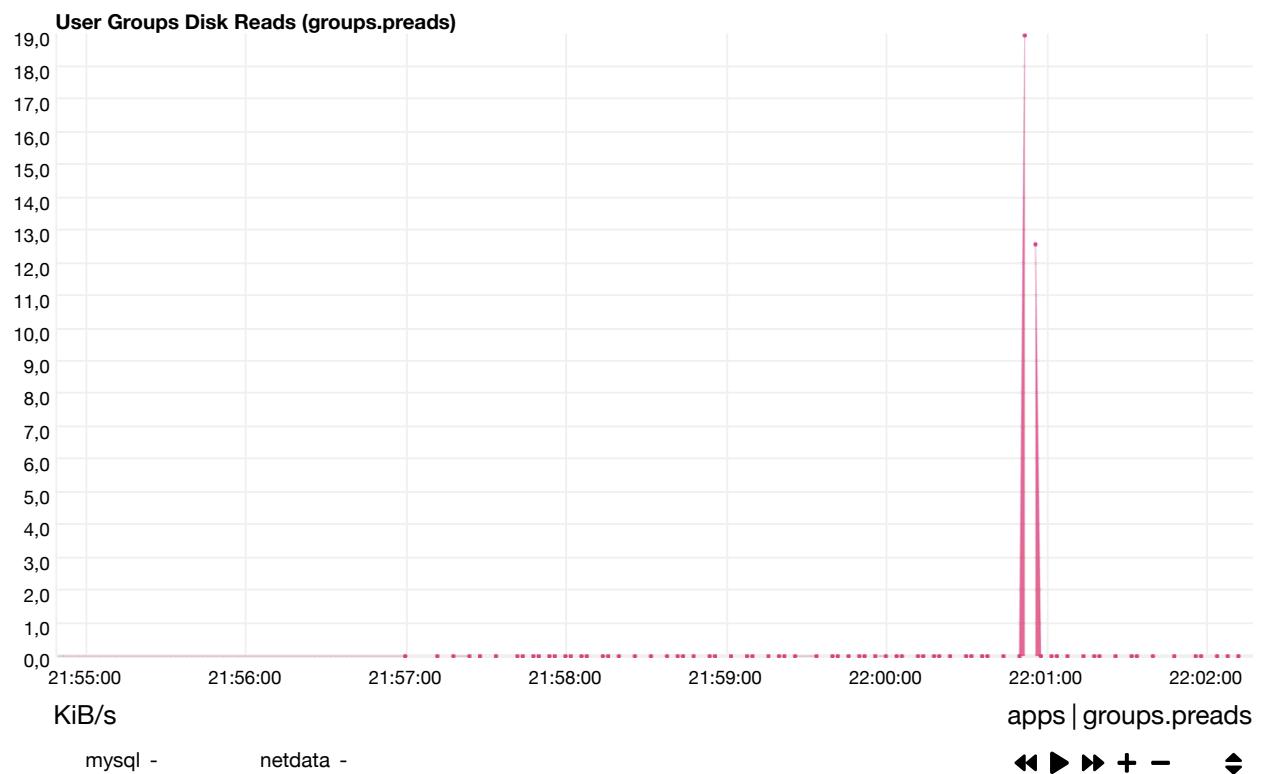
## >User Groups

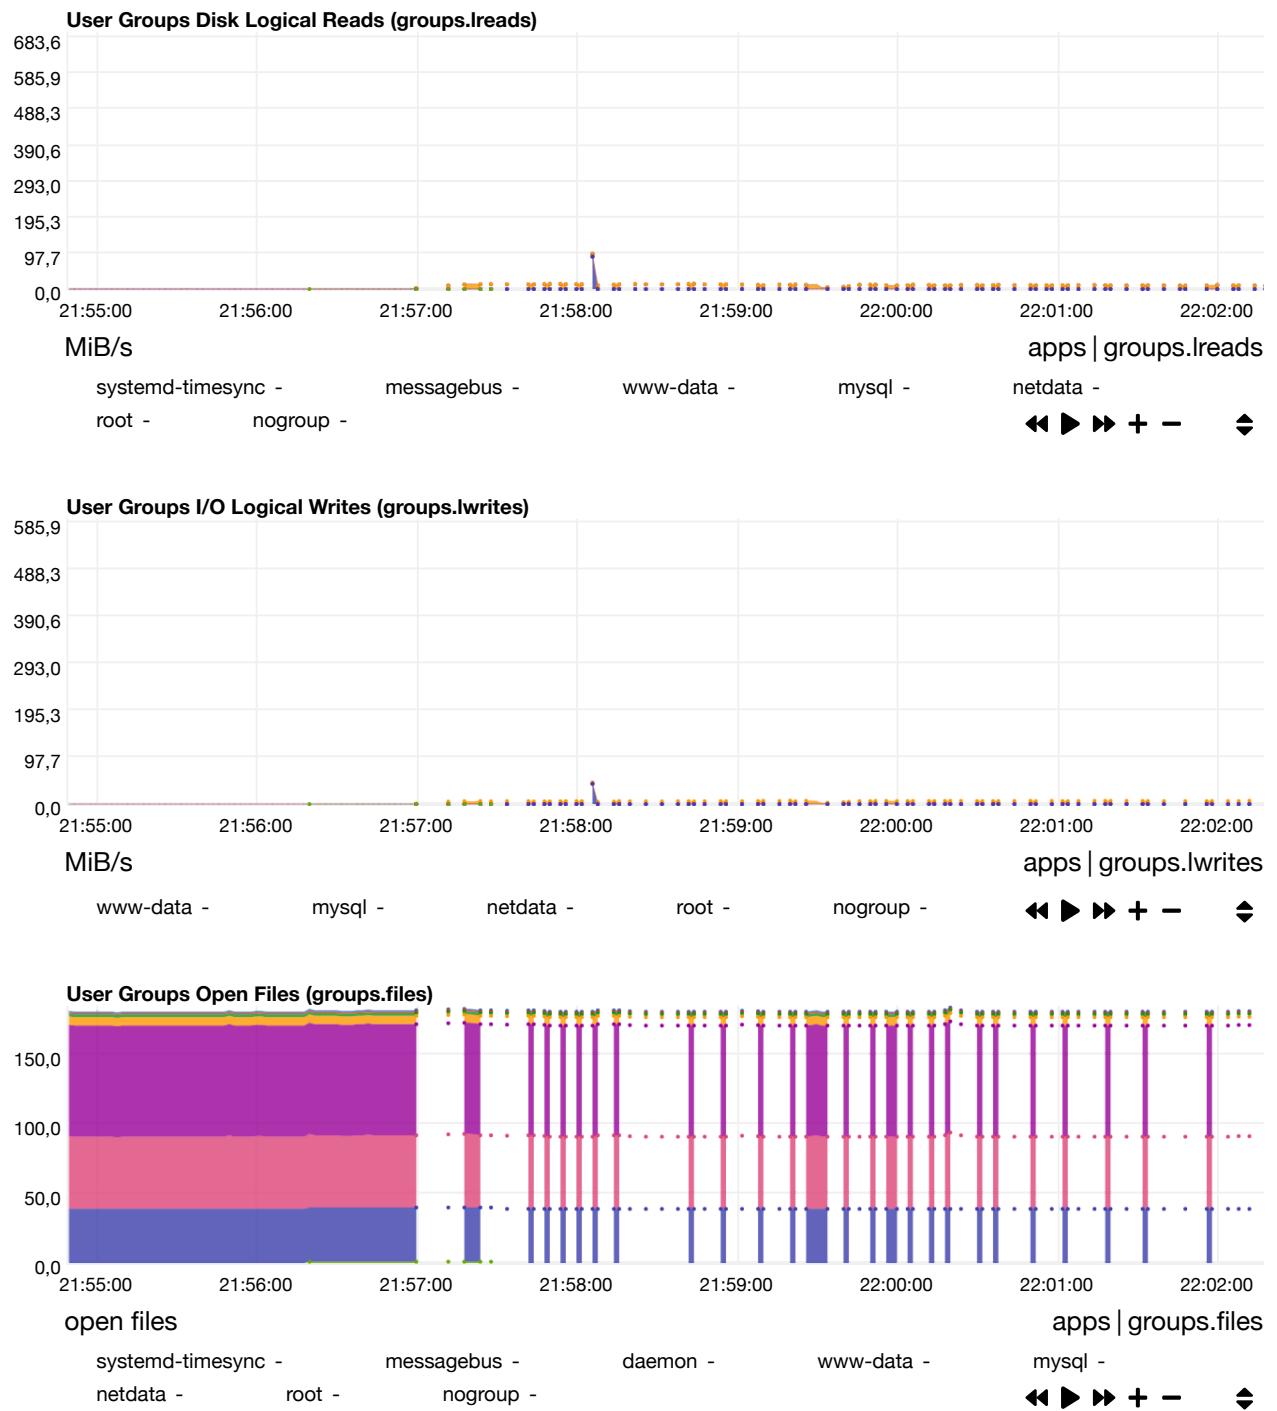
Per user group statistics are collected using netdata's `apps.plugin`. This plugin walks through all processes and aggregates statistics per user group. The reported values are compatible with `top`, although the netdata plugin counts also the resources of exited children (unlike `top` which shows only the resources of the currently running processes). So for processes like shell scripts, the reported values include the resources used by the commands these scripts run within each timeframe.

**cpu**



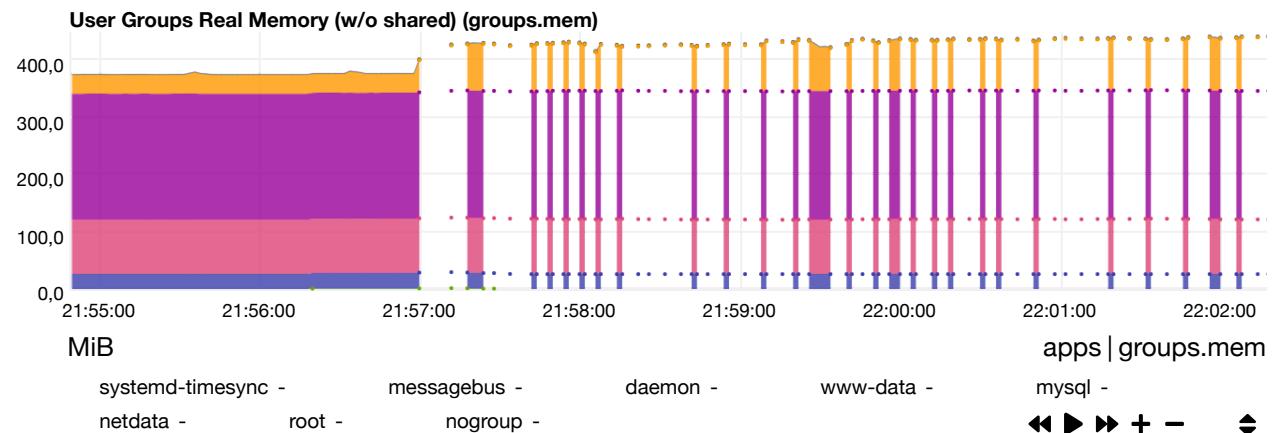
# disk



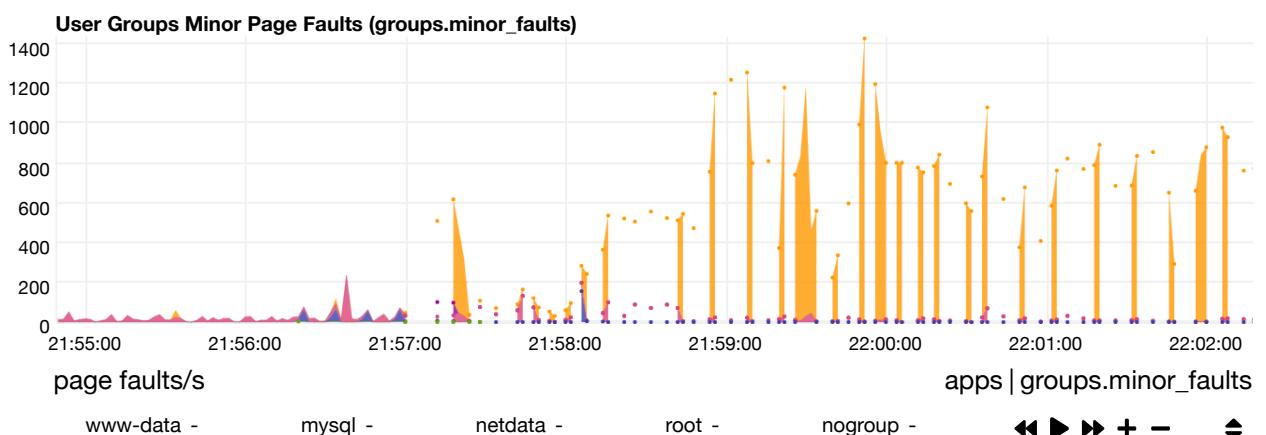
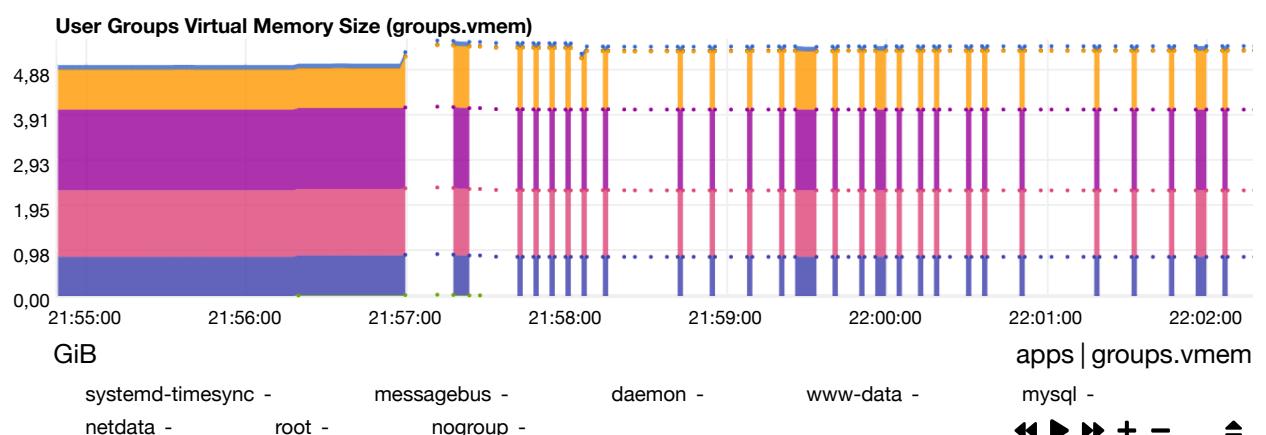


# mem

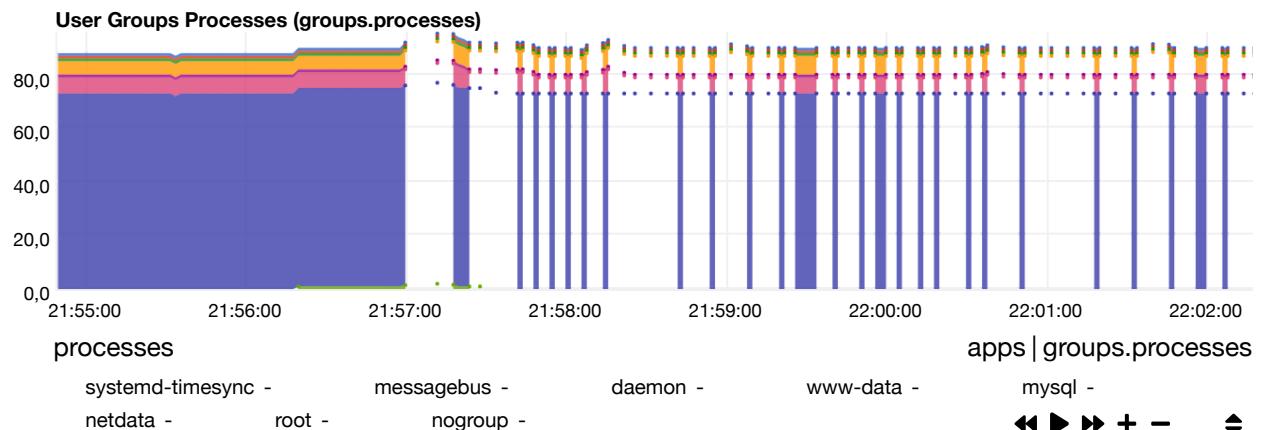
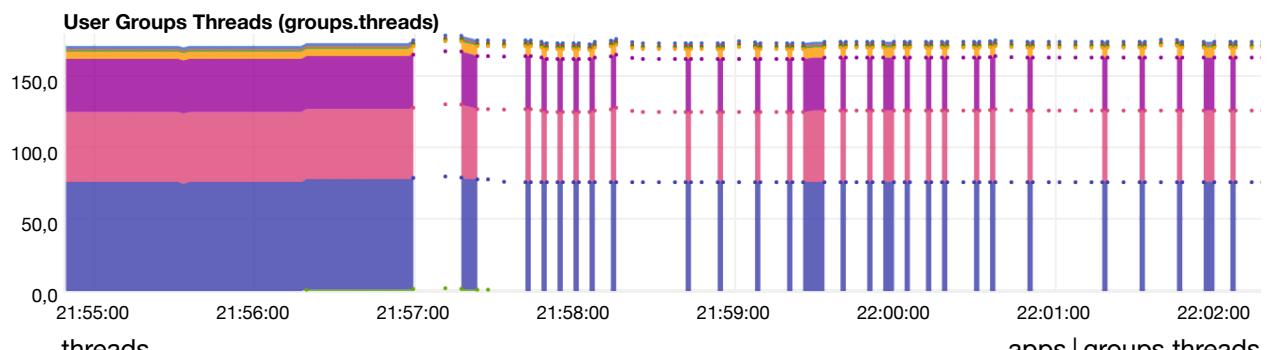
Real memory (RAM) used per user group. This does not include shared memory.



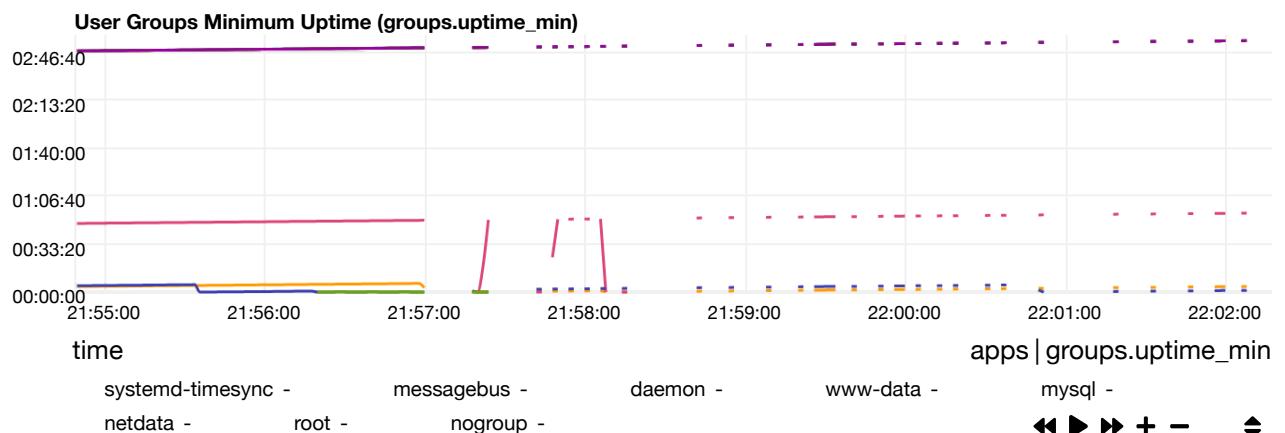
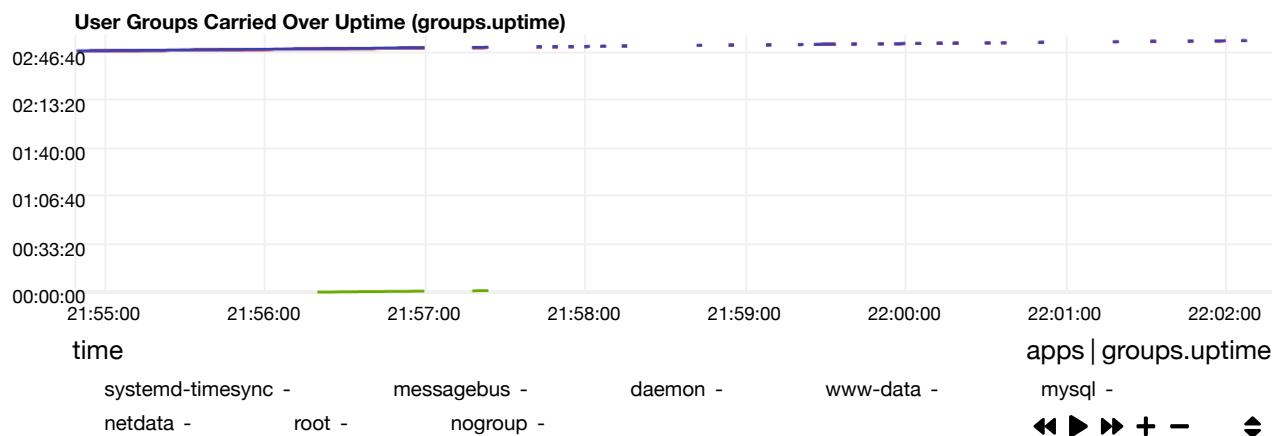
Virtual memory allocated per user group since the Netdata restart. Please check this article (<https://github.com/netdata/netdata/tree/master/daemon#virtual-memory>) for more information.

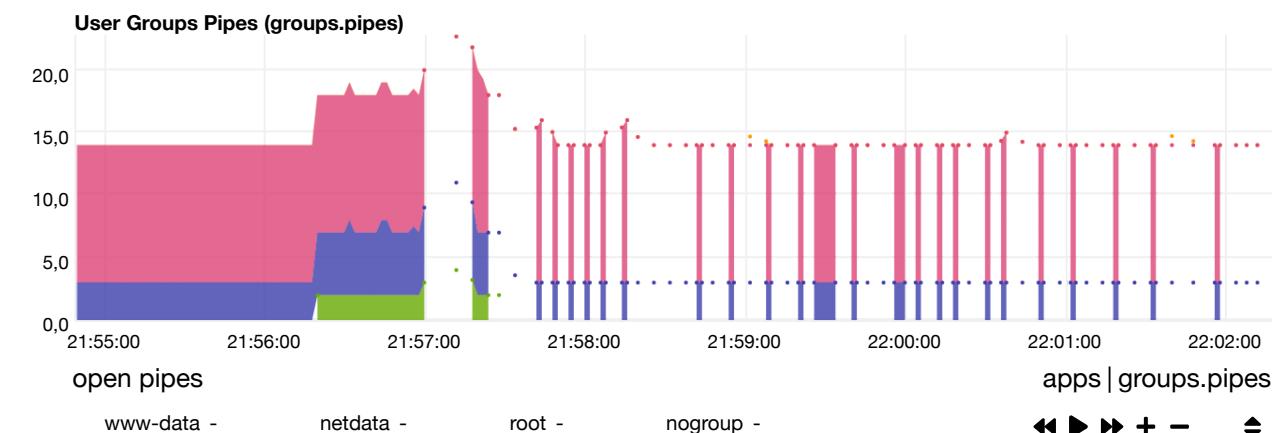
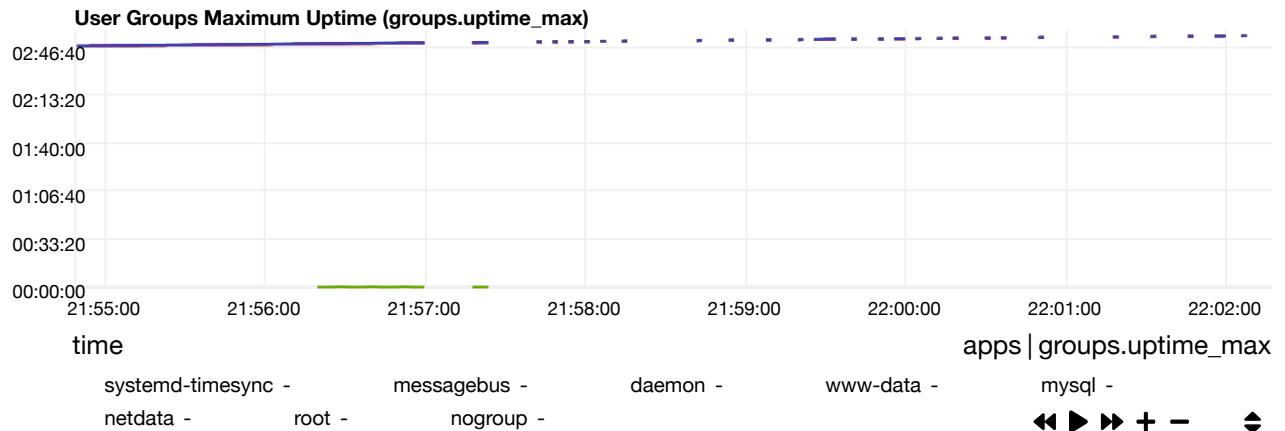
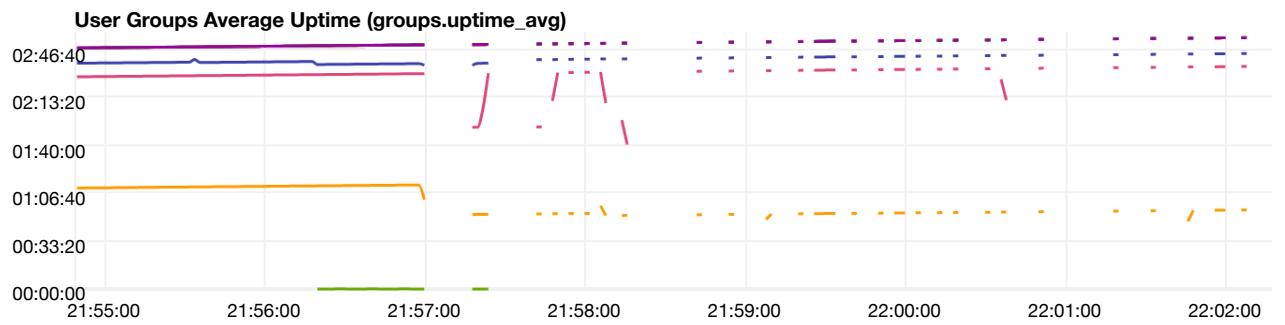


# processes

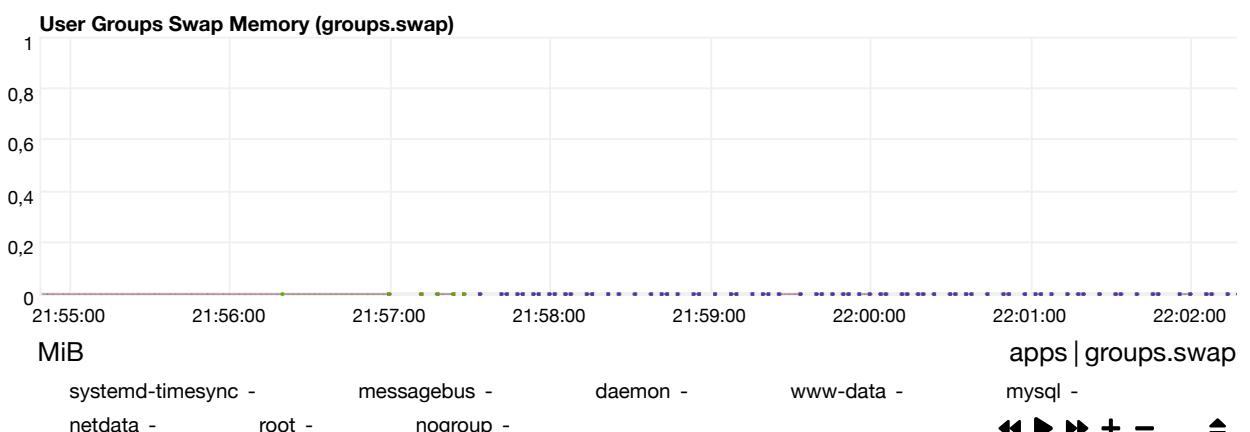


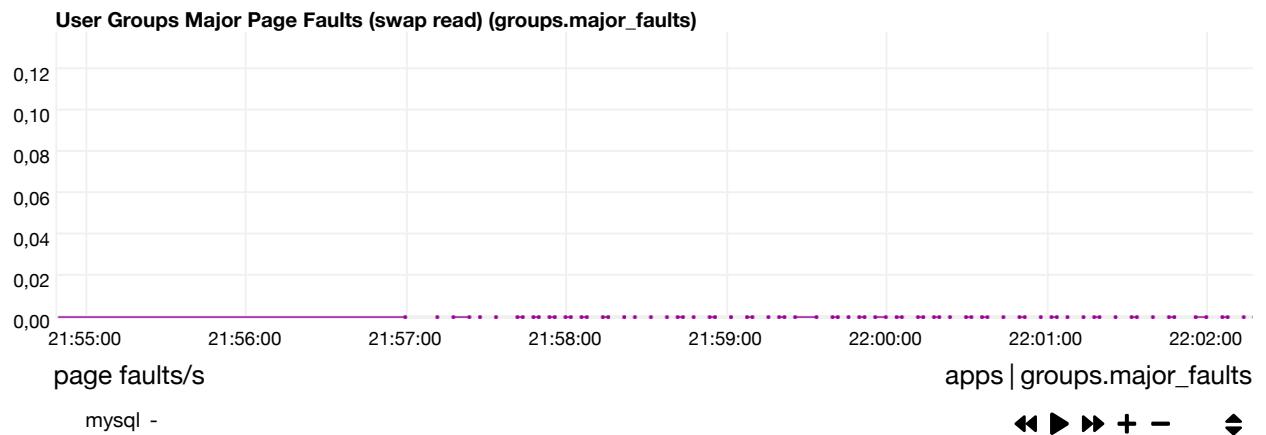
Carried over process group uptime. The period of time within which at least one process in the group was running.



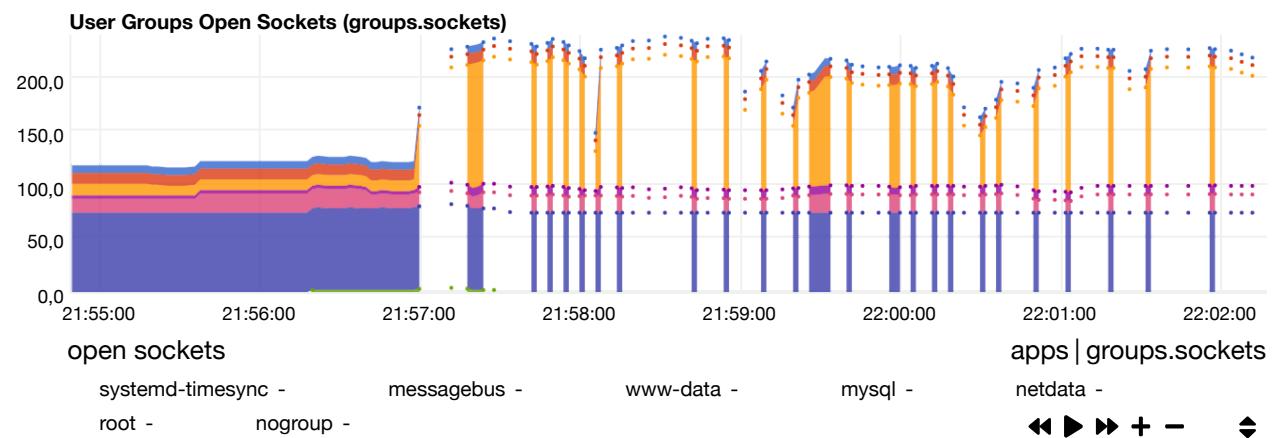


## swap





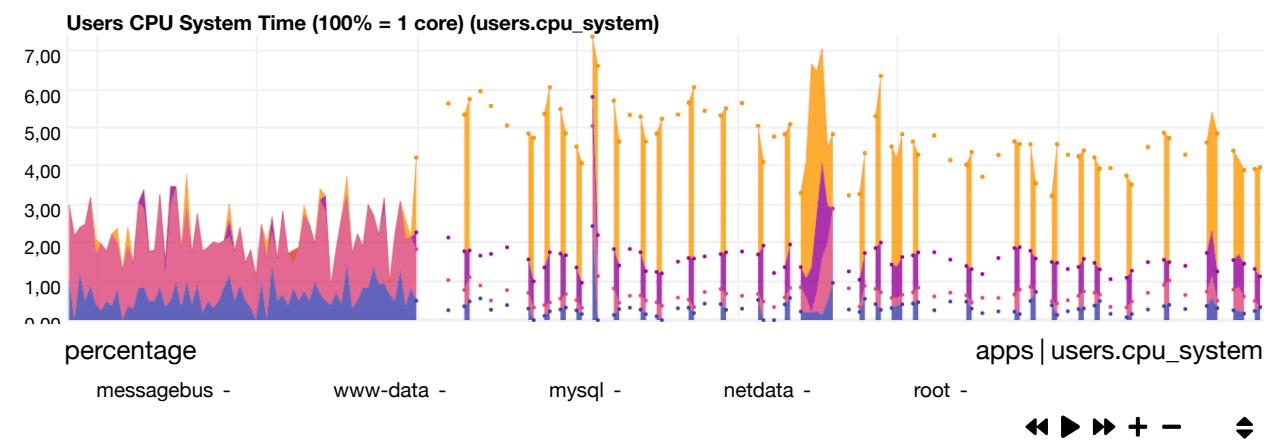
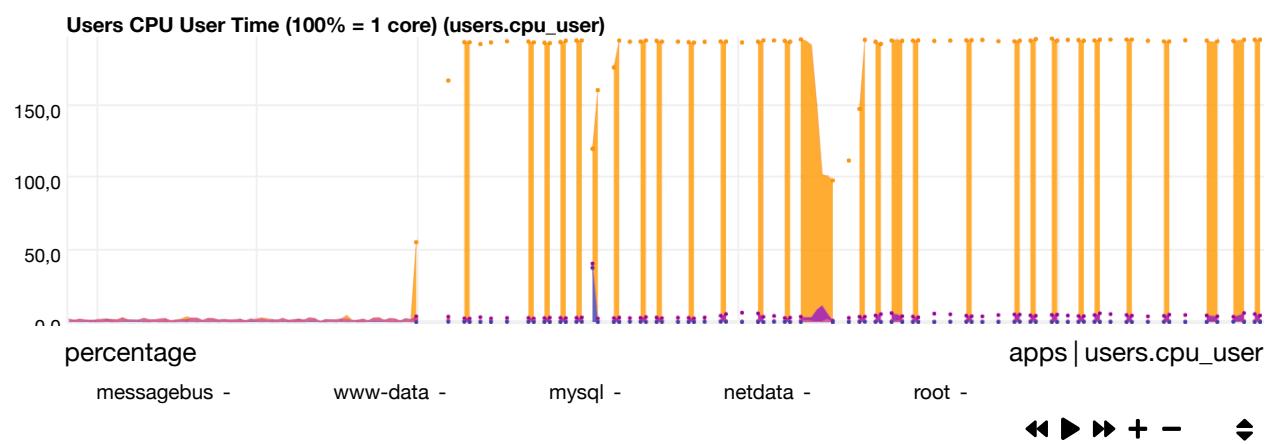
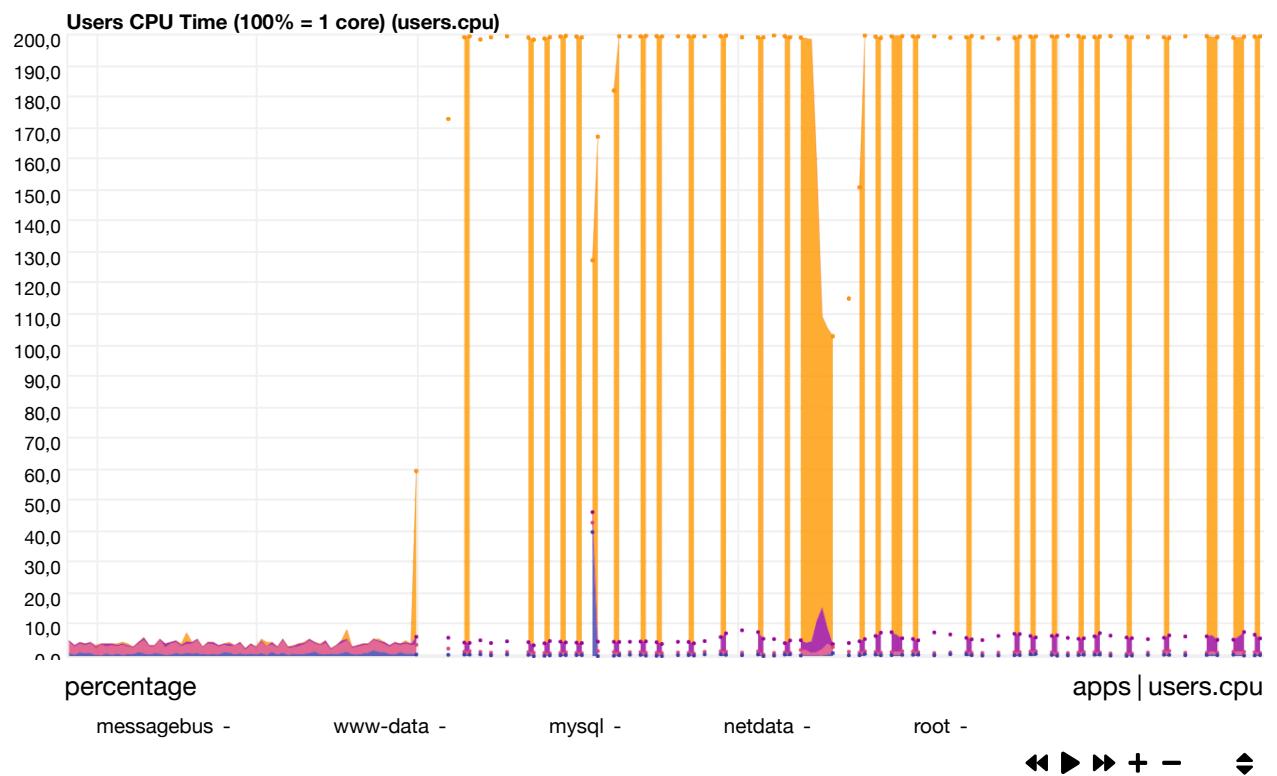
## net



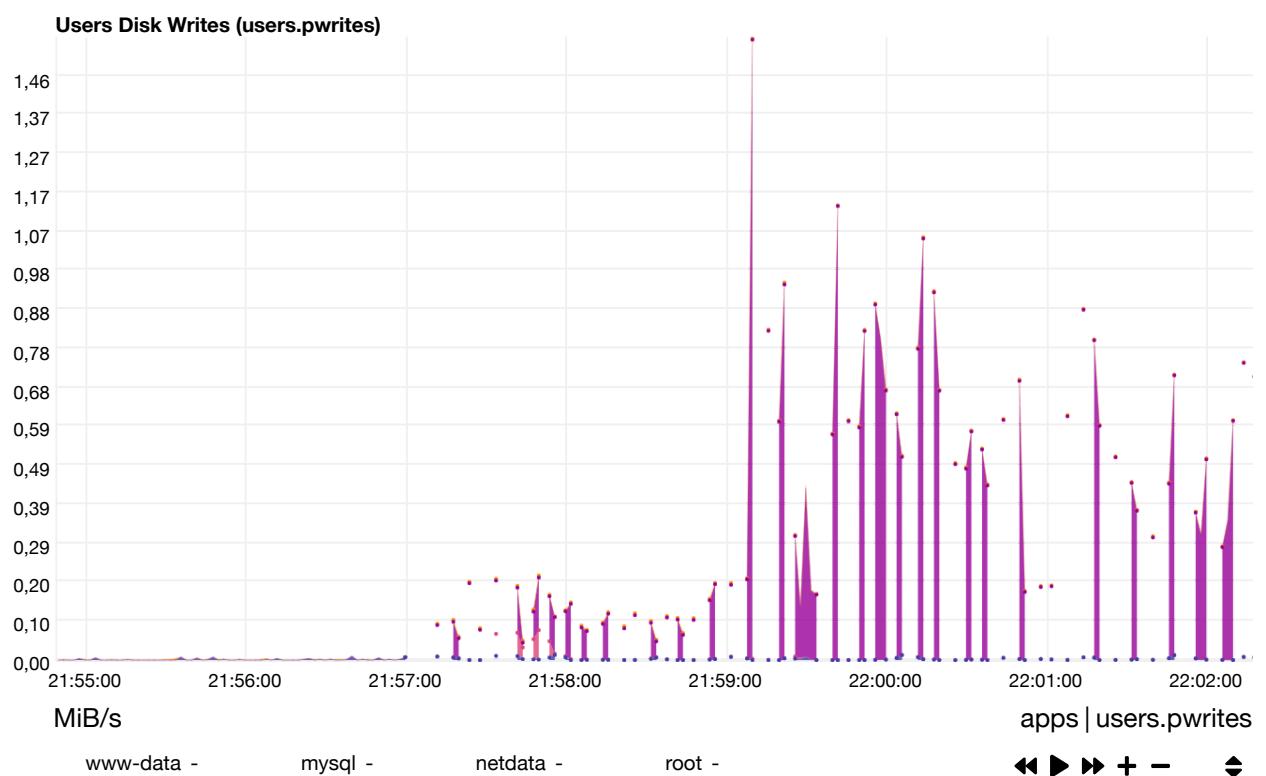
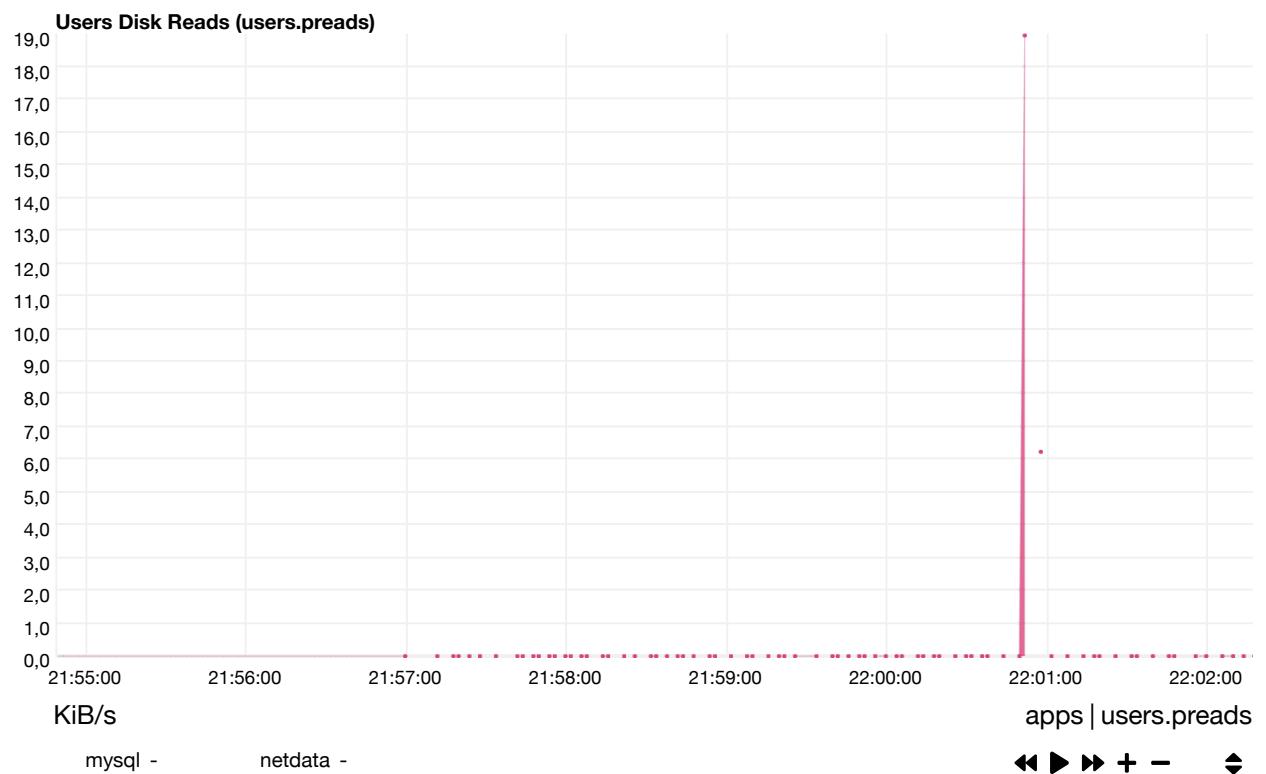
## Users

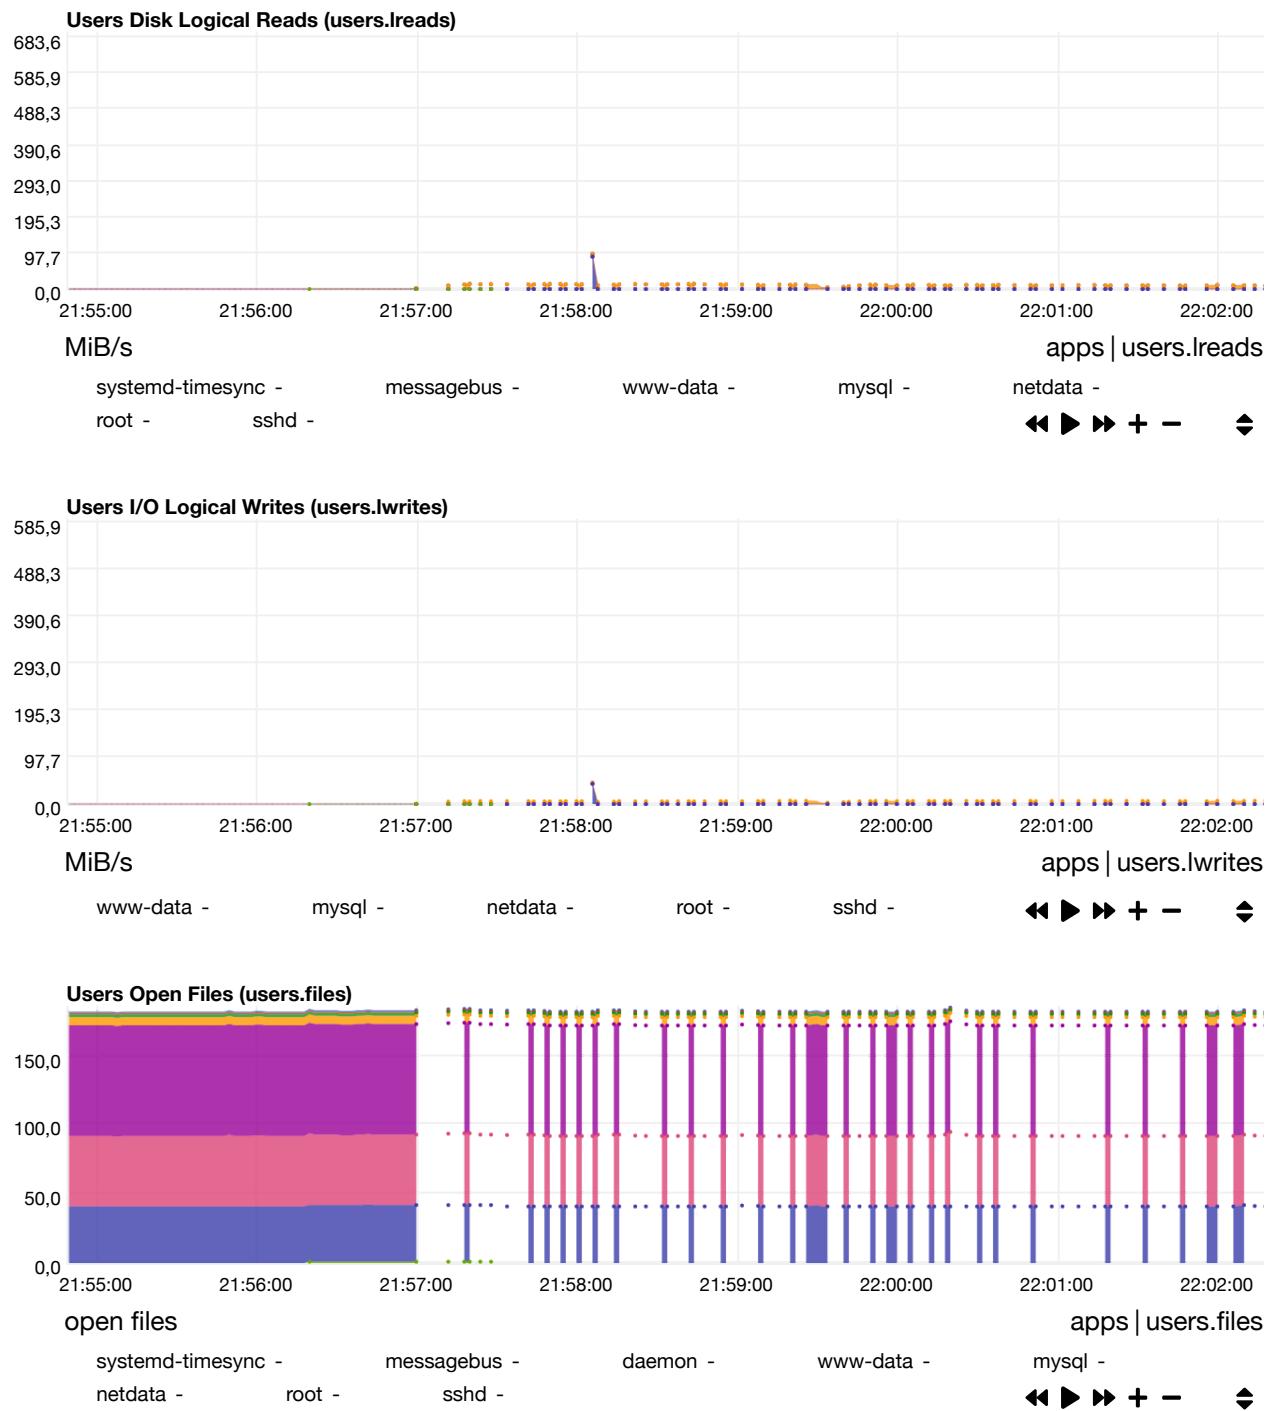
Per user statistics are collected using netdata's `apps.plugin`. This plugin walks through all processes and aggregates statistics per user. The reported values are compatible with `top`, although the netdata plugin counts also the resources of exited children (unlike `top` which shows only the resources of the currently running processes). So for processes like shell scripts, the reported values include the resources used by the commands these scripts run within each timeframe.

## cpu



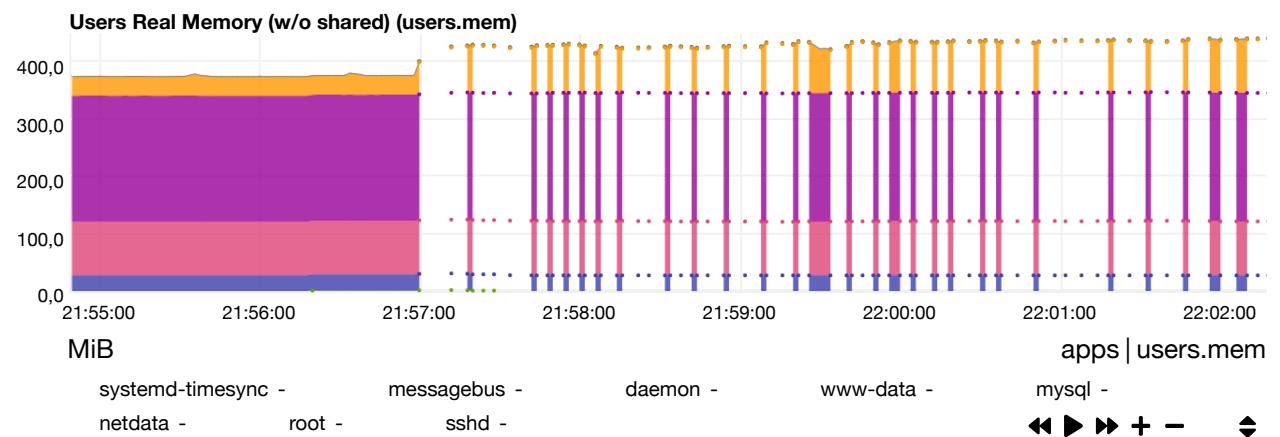
# disk



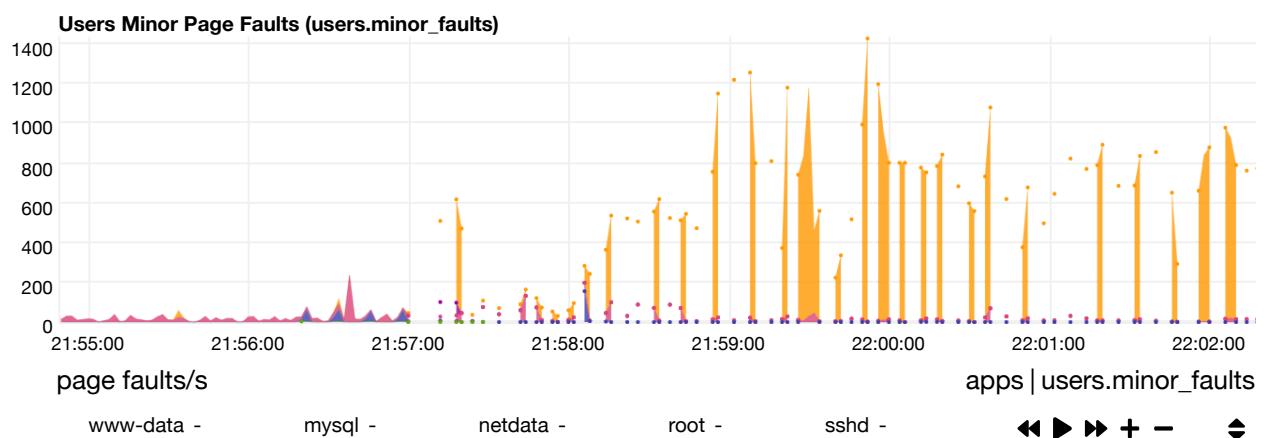
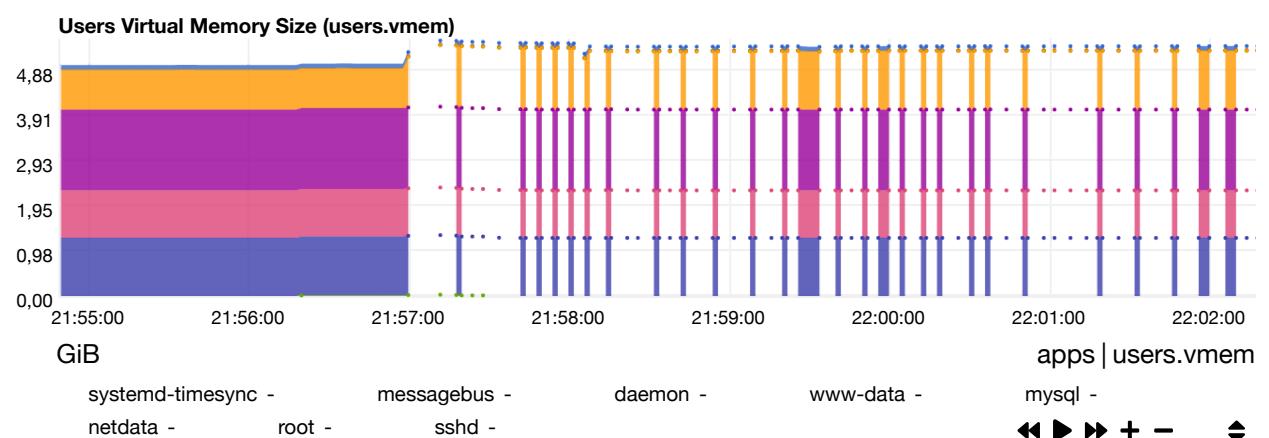


# mem

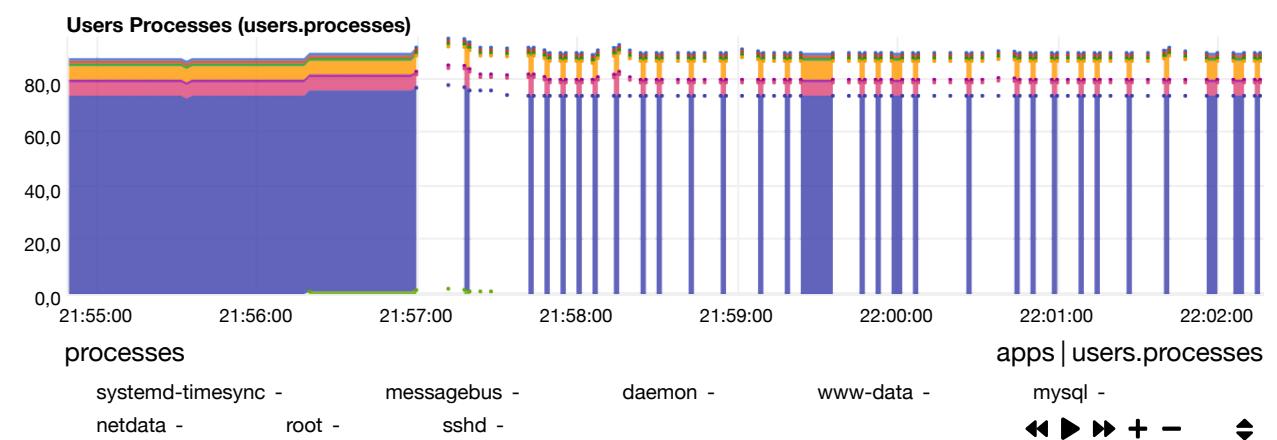
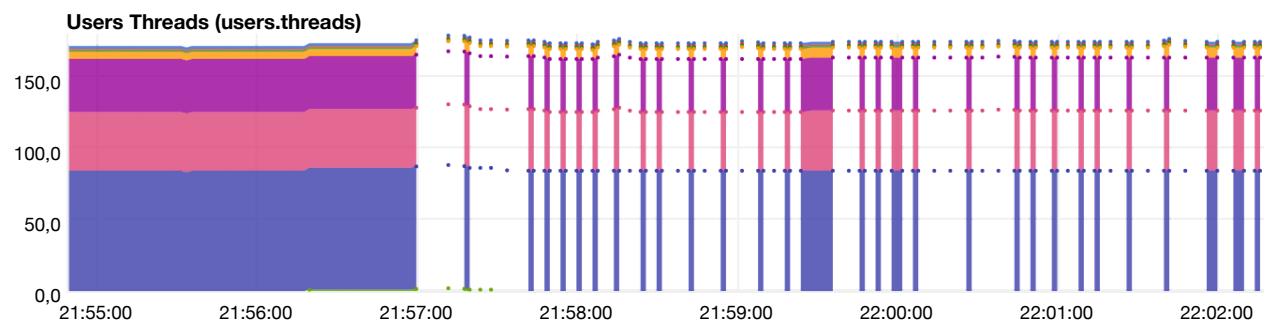
Real memory (RAM) used per user. This does not include shared memory.



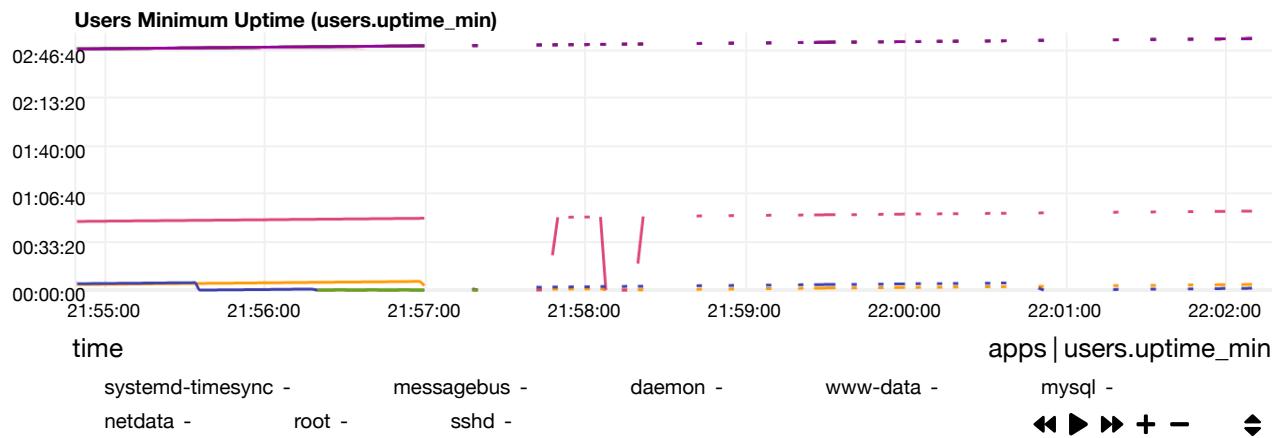
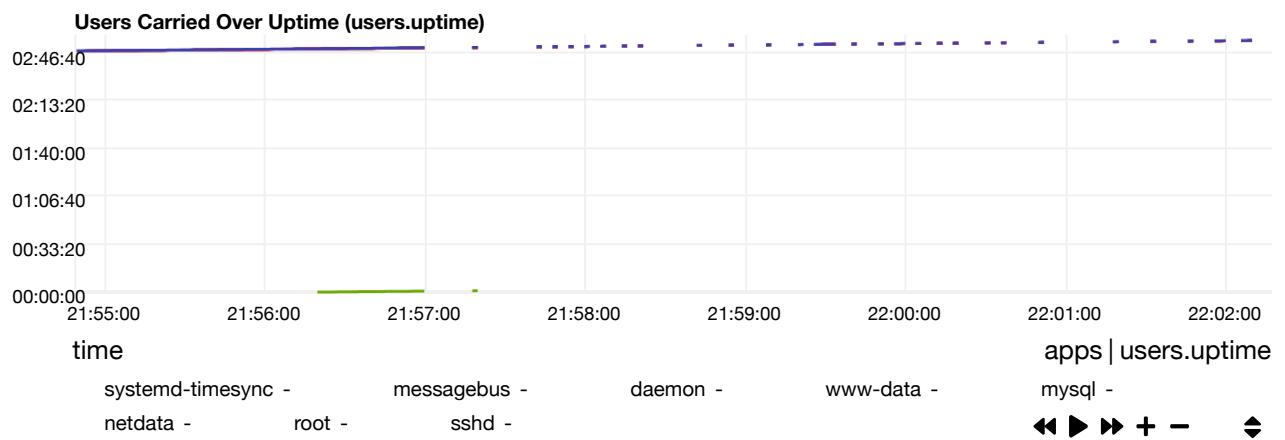
Virtual memory allocated per user. Please check this article (<https://github.com/netdata/netdata/tree/master/daemon#virtual-memory>) for more information.

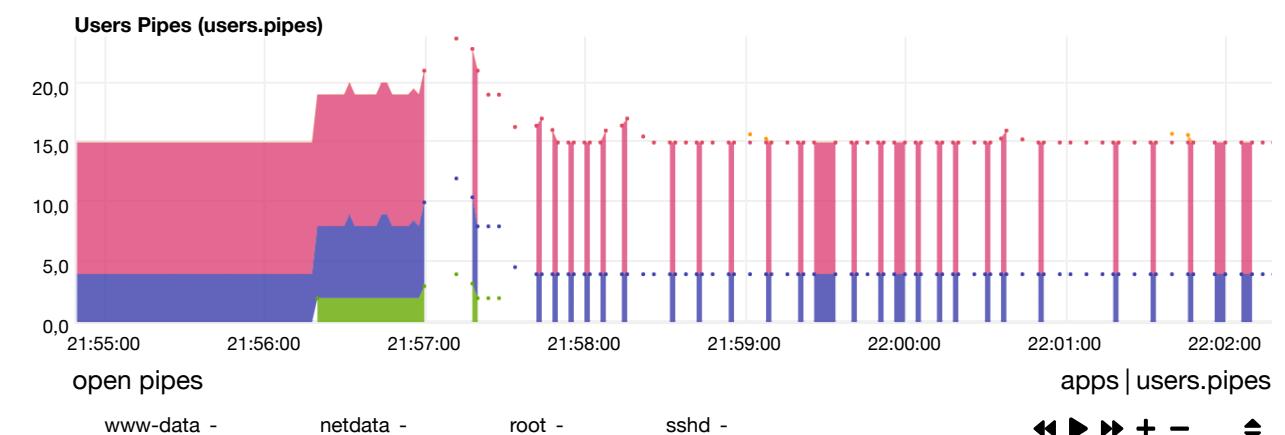
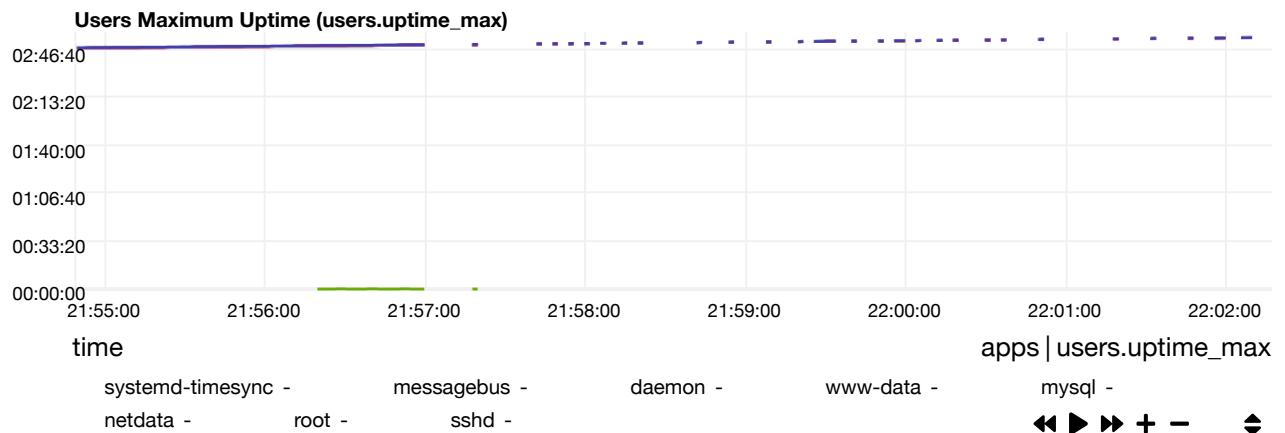
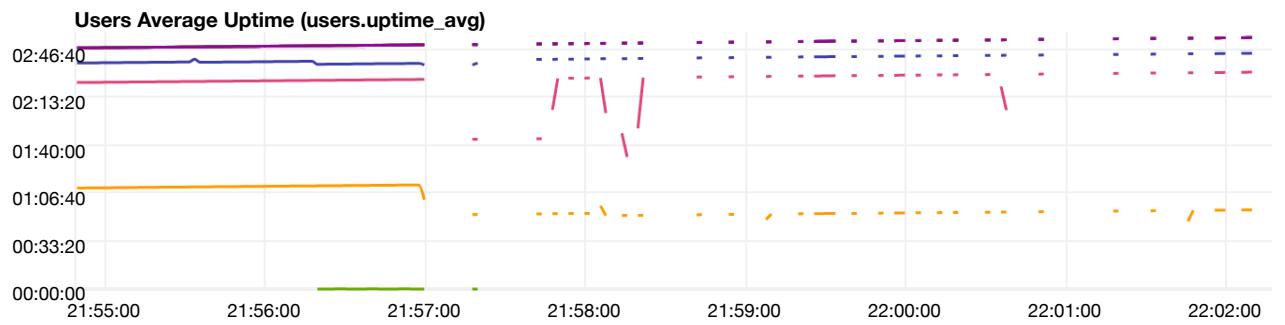


## processes

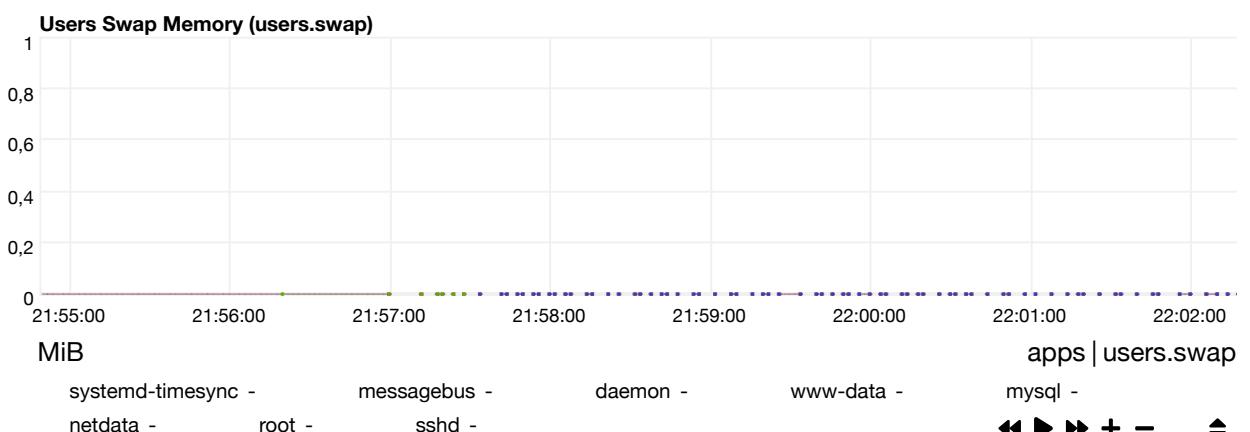


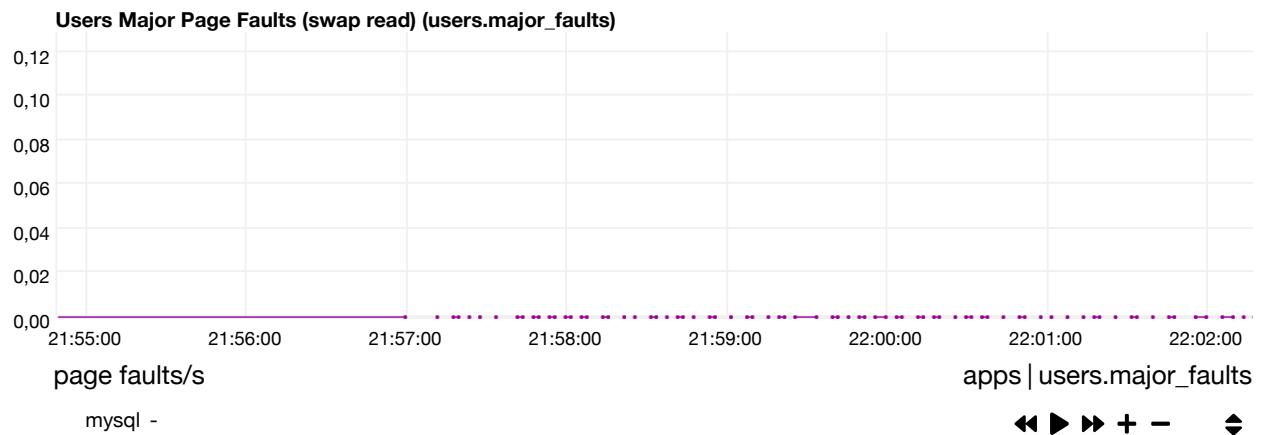
Carried over process group uptime since the Netdata restart. The period of time within which at least one process in the group was running.



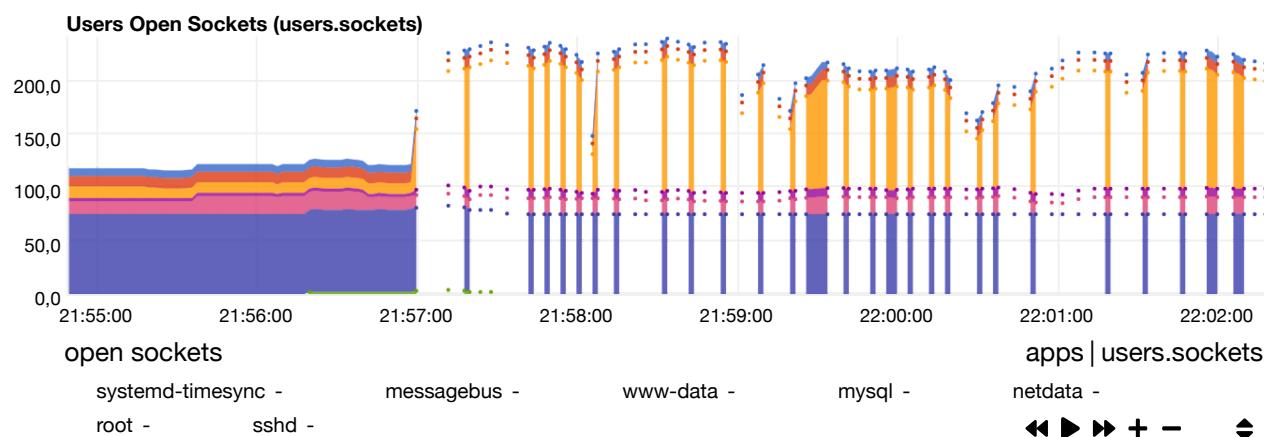


## swap





## net

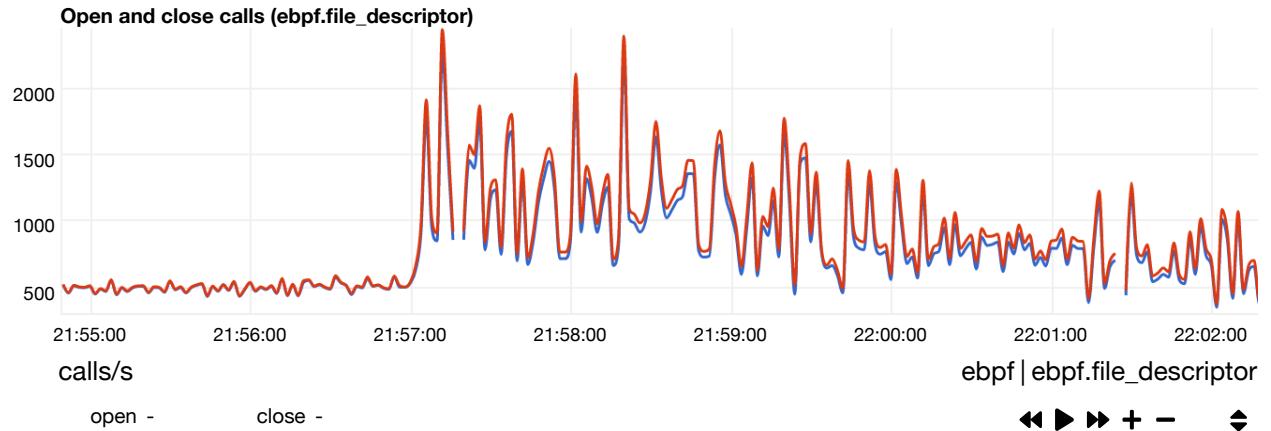


## ❤ eBPF

Monitor system calls, internal functions, bytes read, bytes written and errors using eBPF .

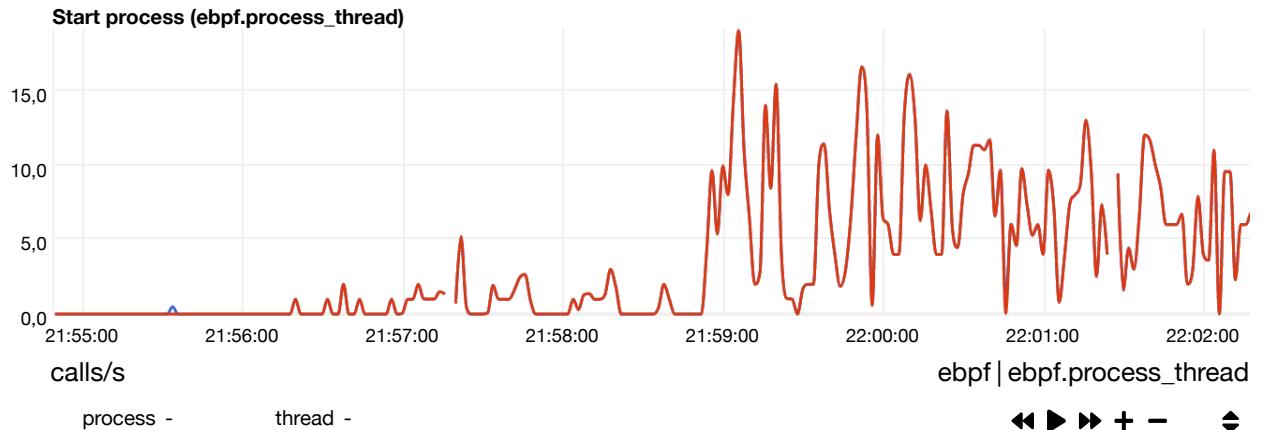
## File

Calls for internal functions on Linux kernel. The open dimension is attached to the kernel internal function `do_sys_open` (For kernels newer than 5.5.19 we add a kprobe to `do_sys_openat2`), which is the common function called from `open(2)` (<https://www.man7.org/linux/man-pages/man2/open.2.html>) and `openat(2)` (<https://www.man7.org/linux/man-pages/man2/openat.2.html>). The close dimension is attached to the function `__close_fd` or `close_fd` according to your kernel version, which is called from system call `close(2)` (<https://www.man7.org/linux/man-pages/man2/close.2.html>).

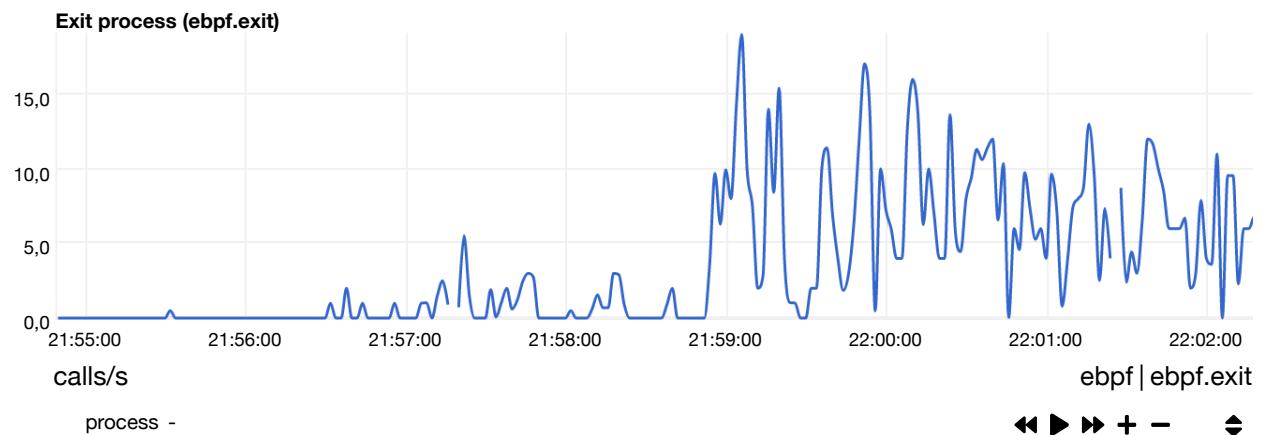


## Process

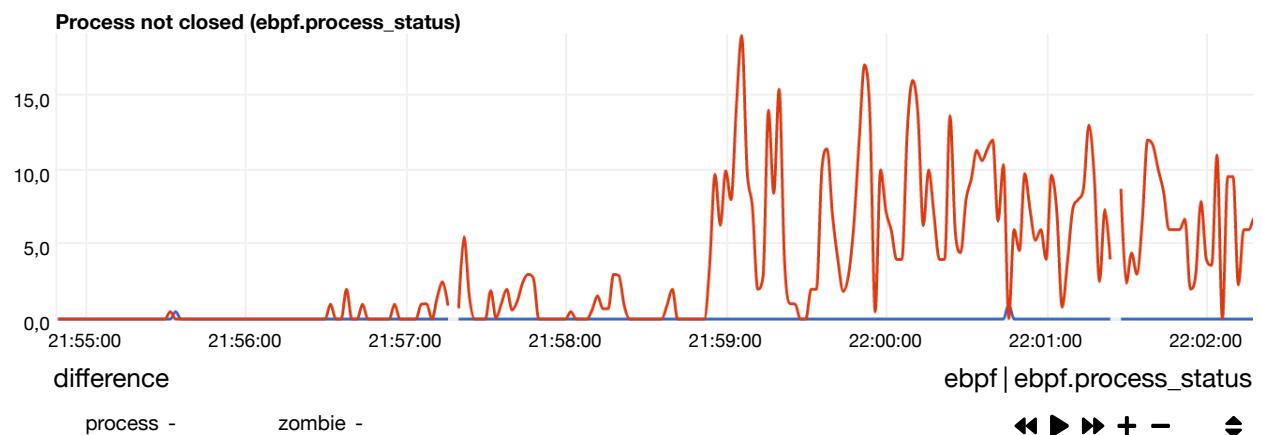
Number of times that either `do_fork` (<https://www.ece.uic.edu/~yshi1/linux/lkse/node4.html#SECTION00421000000000000000>), or `kernel_clone` if you are running kernel newer than 5.9.16, is called to create a new task, which is the common name used to define process and tasks inside the kernel. Netdata identifies the threads by counting the number of calls for `sys_clone` (<https://linux.die.net/man/2/clone>) that has the flag `CLONE_THREAD` set.



Calls for the functions responsible for closing (do\_exit (<https://www.informit.com/articles/article.aspx?p=370047&seqNum=4>)) and releasing (release\_task (<https://www.informit.com/articles/article.aspx?p=370047&seqNum=4>)) tasks.

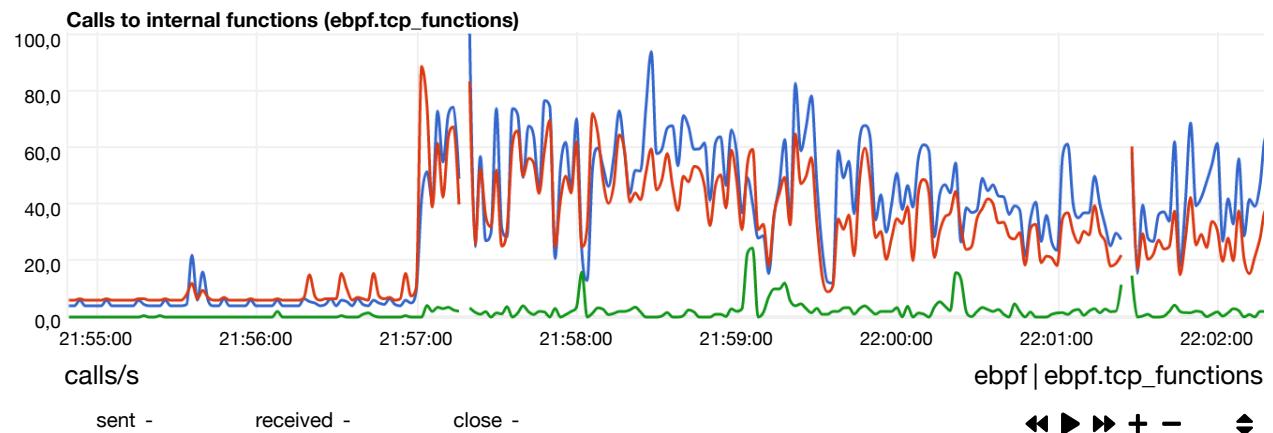


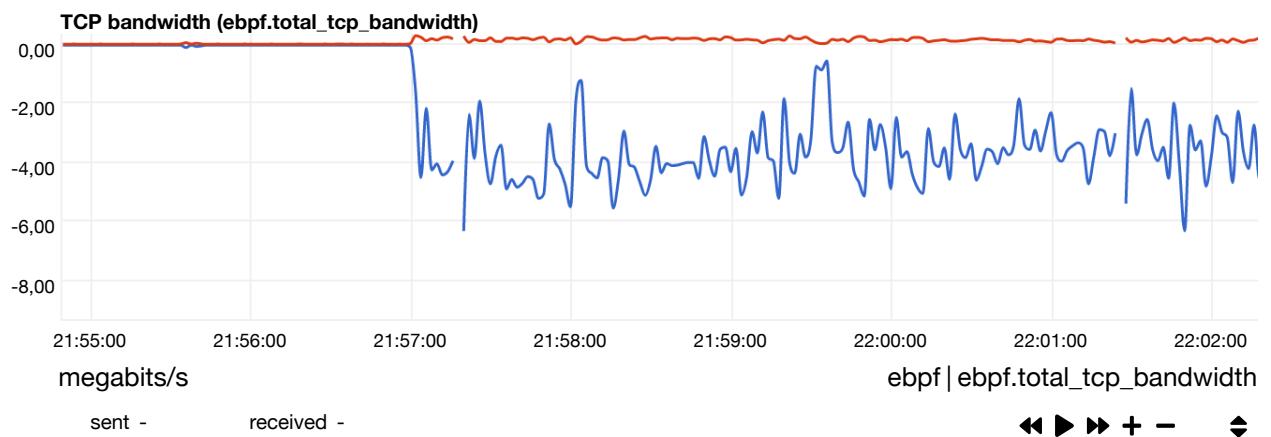
Difference between the number of process created and the number of threads created per period( process dimension), it also shows the number of possible zombie process running on system.



## Socket

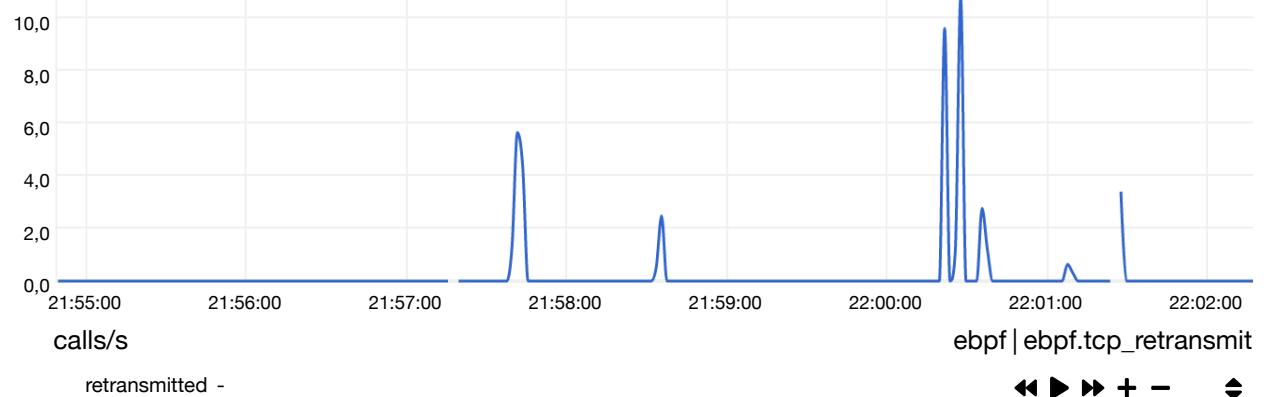
Successful or failed calls to functions `tcp_sendmsg` , `tcp_cleanup_rbuf` and `tcp_close` .





Number of packets retransmitted for function `tcp_retransmit_skb`.

**Packages retransmitted (ebpf.tcp\_retransmit)**

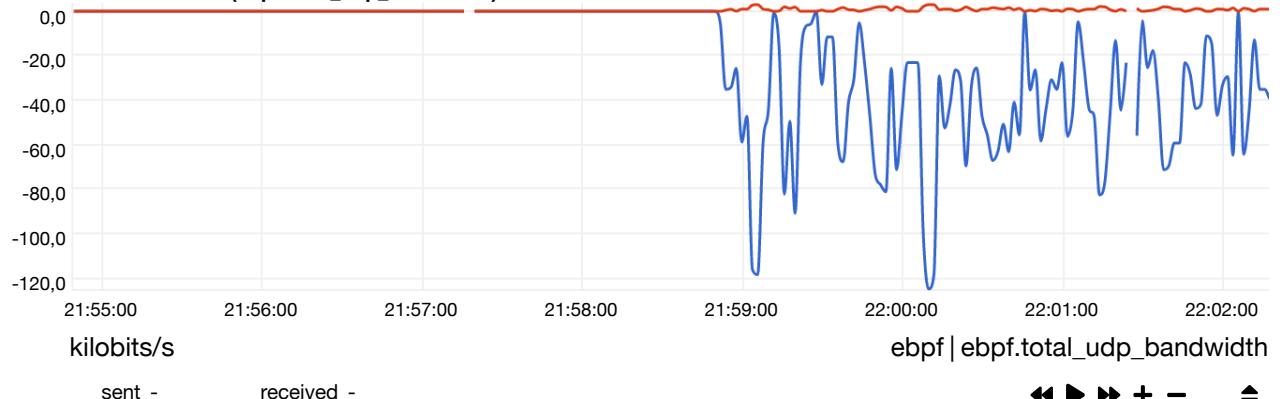


Successful or failed calls to functions `udp_sendmsg` and `udp_recvmsg`.

**UDP calls (ebpf.udp\_functions)**



**UDP bandwidth (ebpf.total\_udp\_bandwidth)**



# MySQL local

Performance metrics for **mysql**, the open-source relational database management system (RDBMS).

## bandwidth

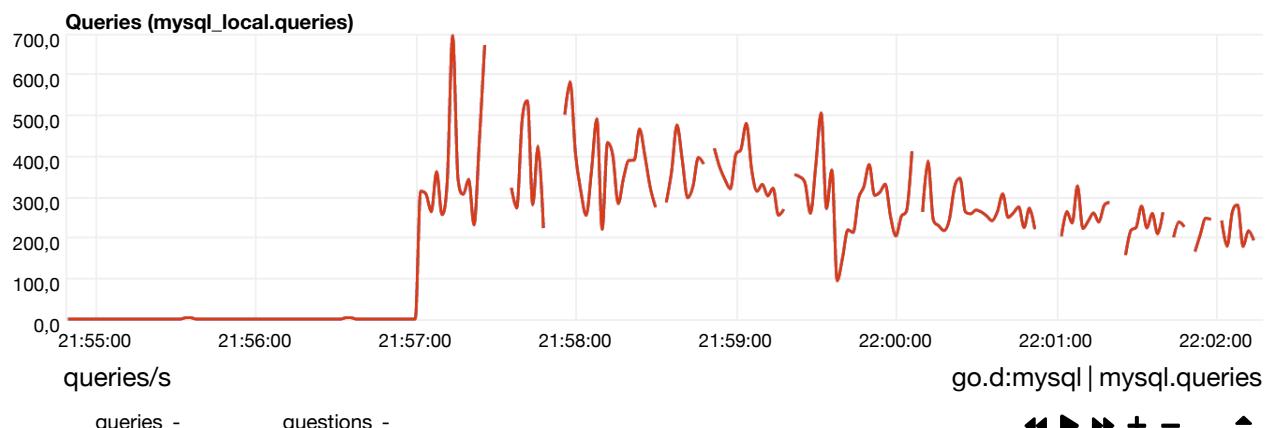
The amount of data sent to mysql clients (**out**) and received from mysql clients (**in**).

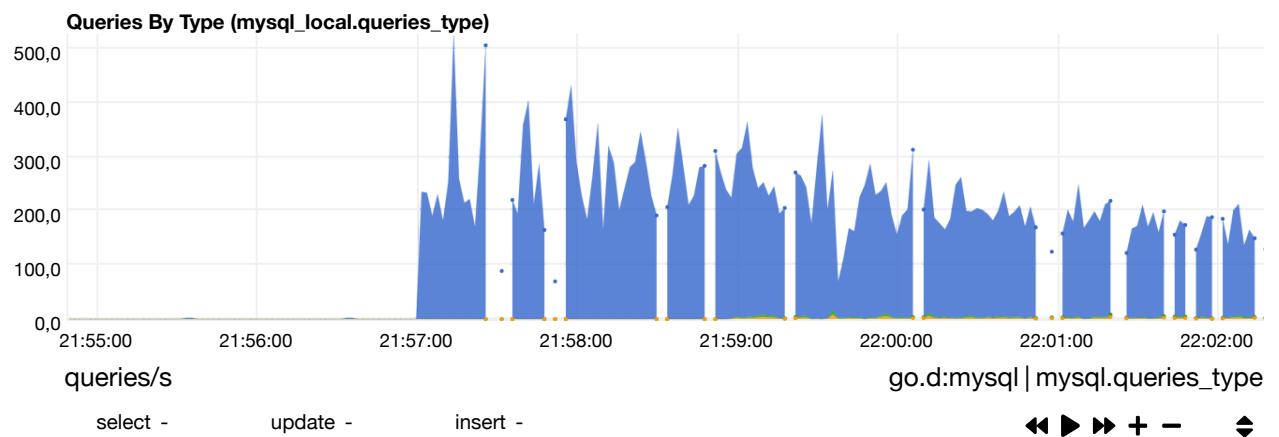


## queries

The number of statements executed by the server.

- **queries** counts the statements executed within stored SQL programs.
- **questions** counts the statements sent to the mysql server by mysql clients.
- **slow queries** counts the number of statements that took more than long\_query\_time ([http://dev.mysql.com/doc/refman/5.7/en/server-system-variables.html#sysvar\\_long\\_query\\_time](http://dev.mysql.com/doc/refman/5.7/en/server-system-variables.html#sysvar_long_query_time)) seconds to be executed. For more information about slow queries check the mysql slow query log (<http://dev.mysql.com/doc/refman/5.7/en/slow-query-log.html>).

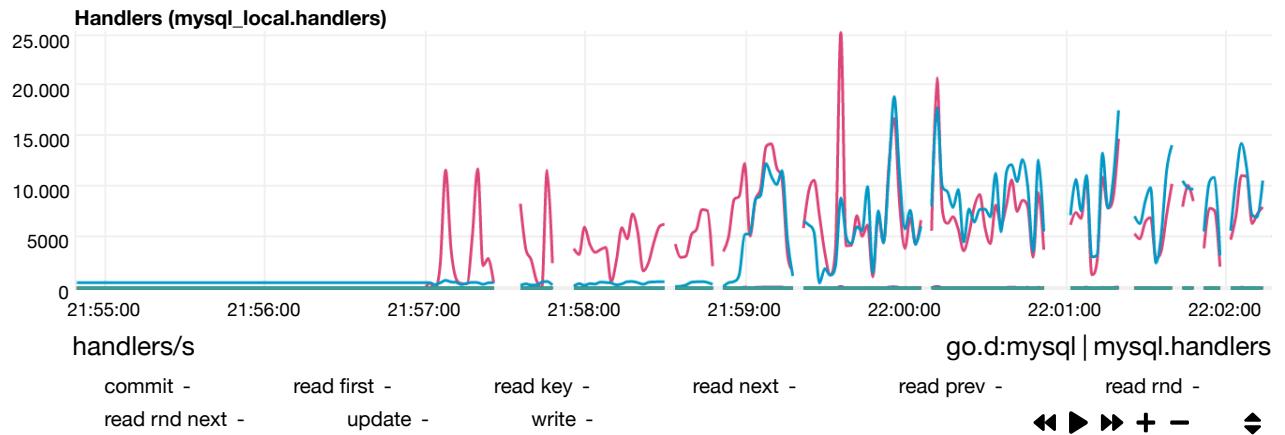




## handlers

Usage of the internal handlers of mysql. This chart provides very good insights of what the mysql server is actually doing. (if the chart is not showing all these dimensions it is because they are zero - set **Which dimensions to show?** to **All** from the dashboard settings, to render even the zero values)

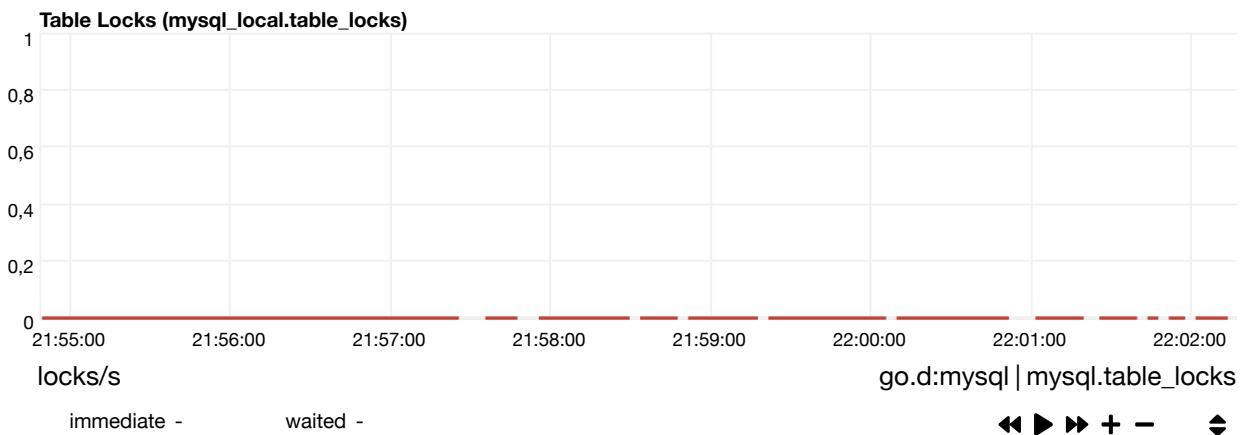
- **commit**, the number of internal COMMIT (<http://dev.mysql.com/doc/refman/5.7/en/commit.html>) statements.
- **delete**, the number of times that rows have been deleted from tables.
- **prepare**, a counter for the prepare phase of two-phase commit operations.
- **read first**, the number of times the first entry in an index was read. A high value suggests that the server is doing a lot of full index scans; e.g. **SELECT col1 FROM foo**, with col1 indexed.
- **read key**, the number of requests to read a row based on a key. If this value is high, it is a good indication that your tables are properly indexed for your queries.
- **read next**, the number of requests to read the next row in key order. This value is incremented if you are querying an index column with a range constraint or if you are doing an index scan.
- **read prev**, the number of requests to read the previous row in key order. This read method is mainly used to optimize **ORDER BY ... DESC**.
- **read rnd**, the number of requests to read a row based on a fixed position. A high value indicates you are doing a lot of queries that require sorting of the result. You probably have a lot of queries that require MySQL to scan entire tables or you have joins that do not use keys properly.
- **read rnd next**, the number of requests to read the next row in the data file. This value is high if you are doing a lot of table scans. Generally this suggests that your tables are not properly indexed or that your queries are not written to take advantage of the indexes you have.
- **rollback**, the number of requests for a storage engine to perform a rollback operation.
- **savepoint**, the number of requests for a storage engine to place a savepoint.
- **savepoint rollback**, the number of requests for a storage engine to roll back to a savepoint.
- **update**, the number of requests to update a row in a table.
- **write**, the number of requests to insert a row in a table.



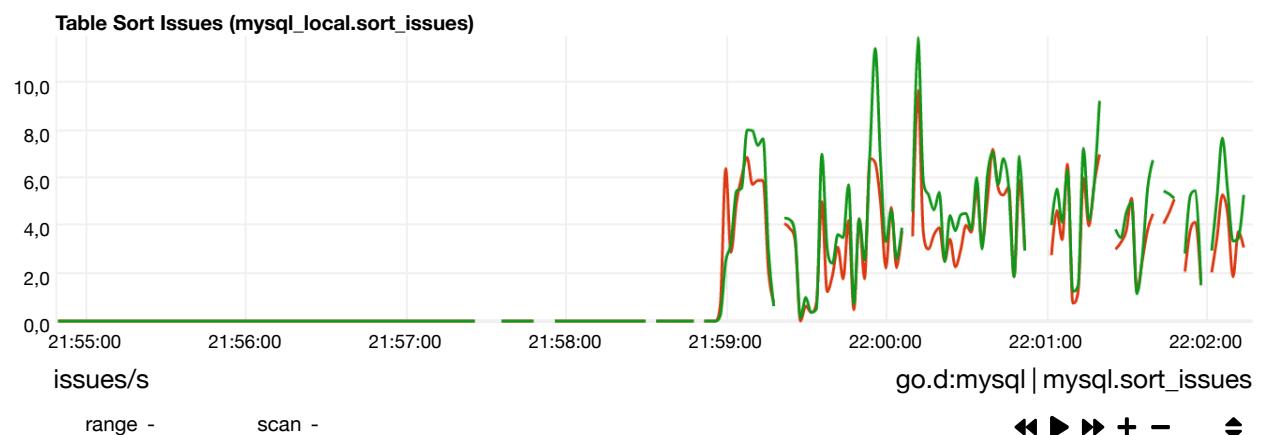
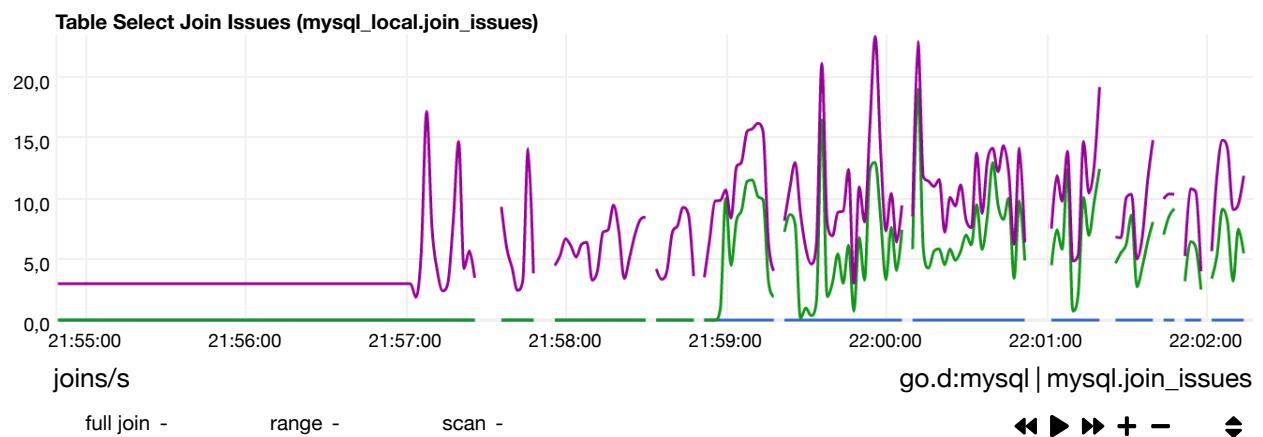
## locks

### MySQL table locks counters:

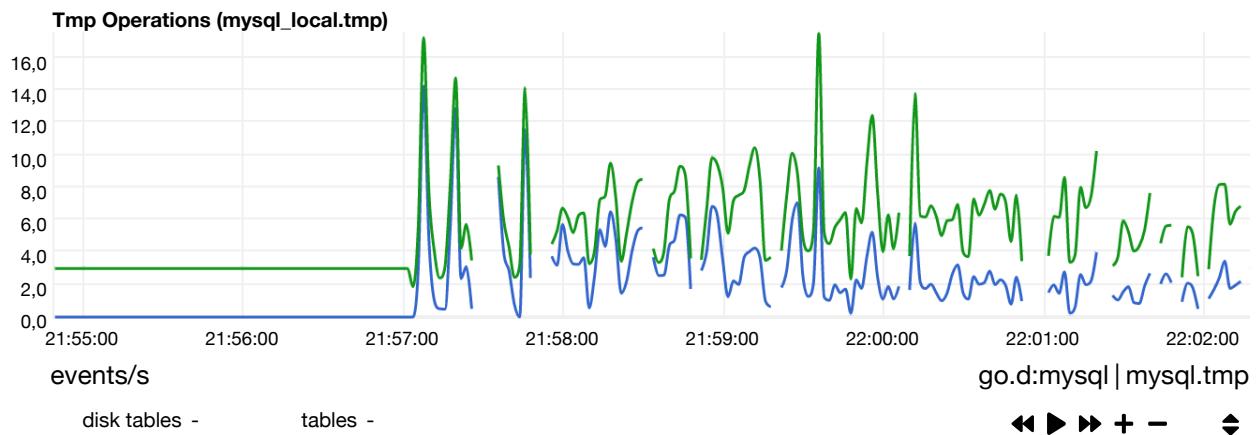
- **immediate**, the number of times that a request for a table lock could be granted immediately.
- **waited**, the number of times that a request for a table lock could not be granted immediately and a wait was needed. If this is high and you have performance problems, you should first optimize your queries, and then either split your table or tables or use replication.



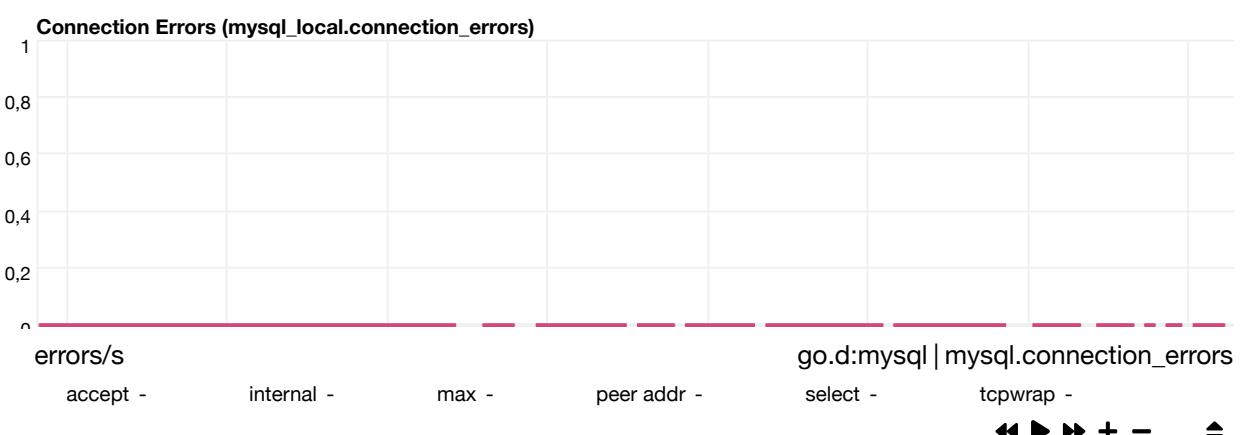
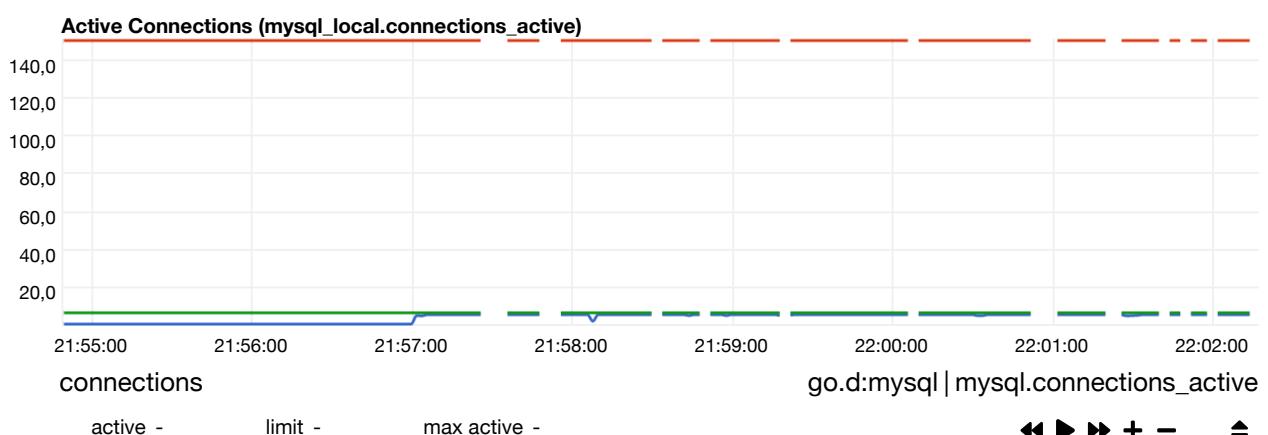
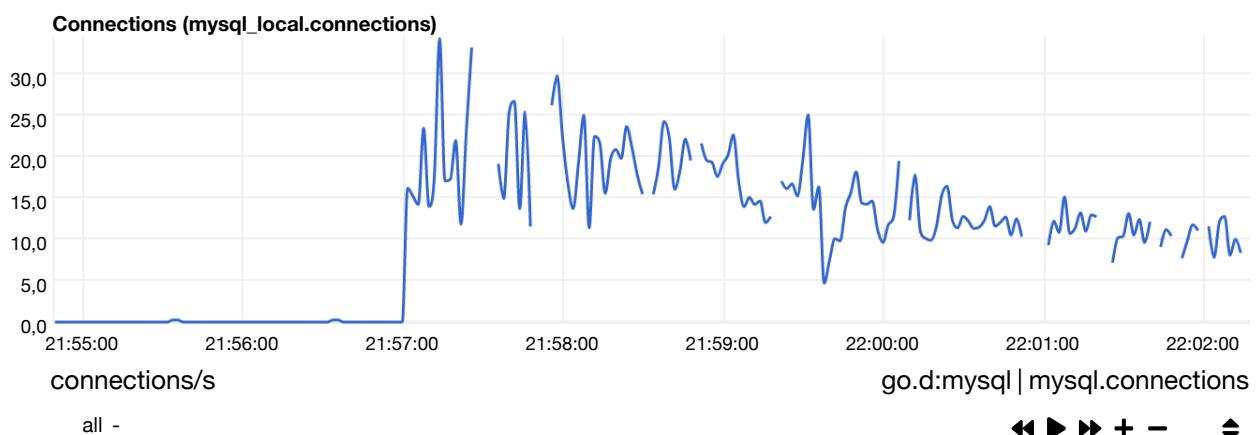
## issues



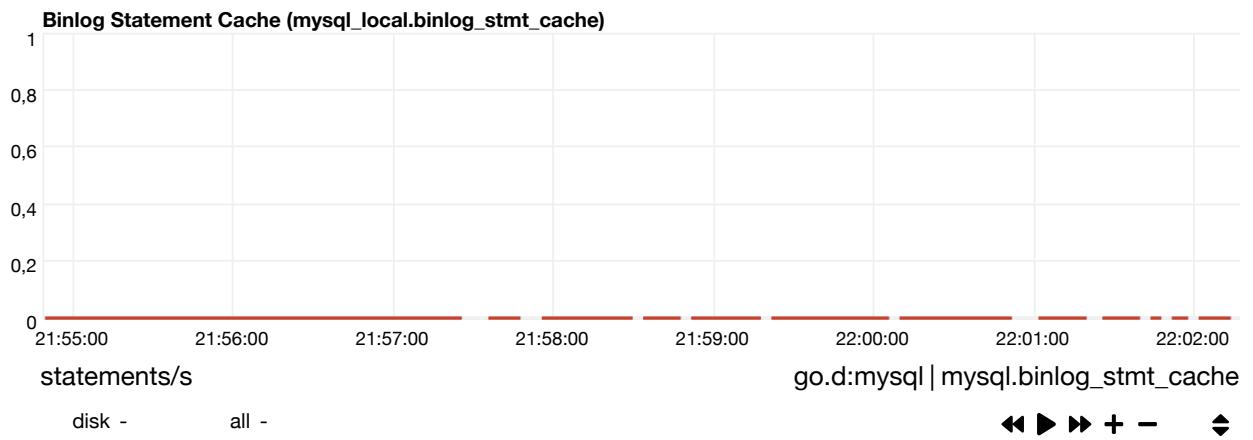
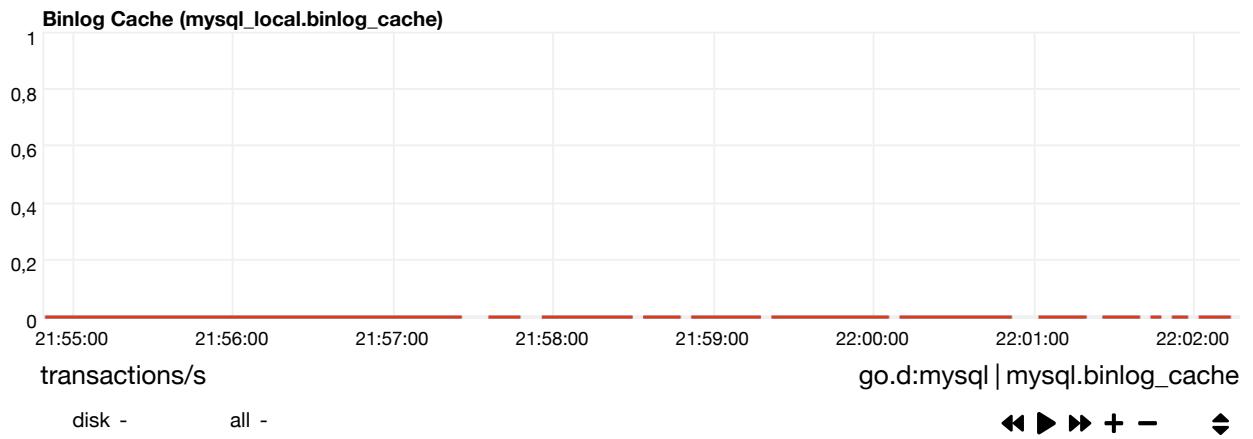
## temporaries



## connections

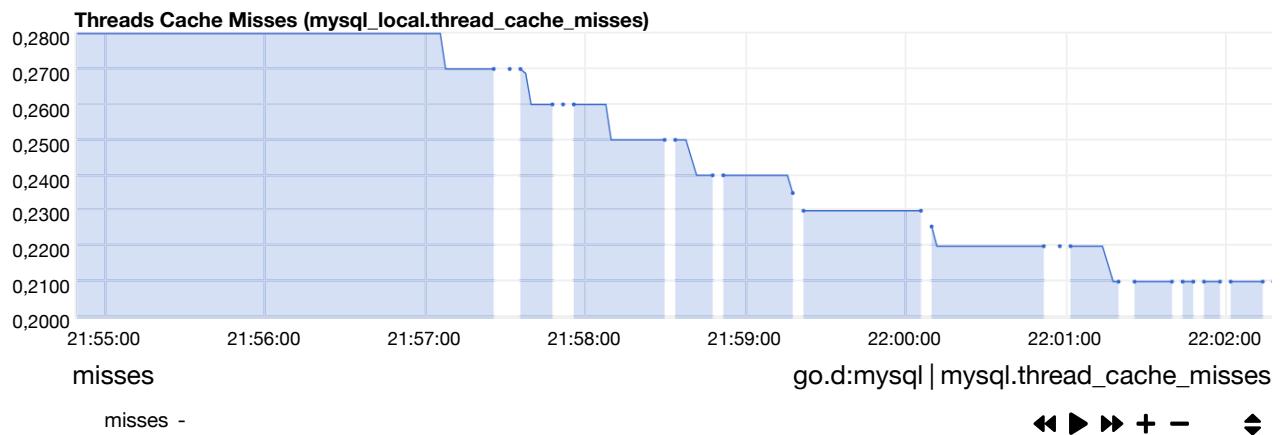
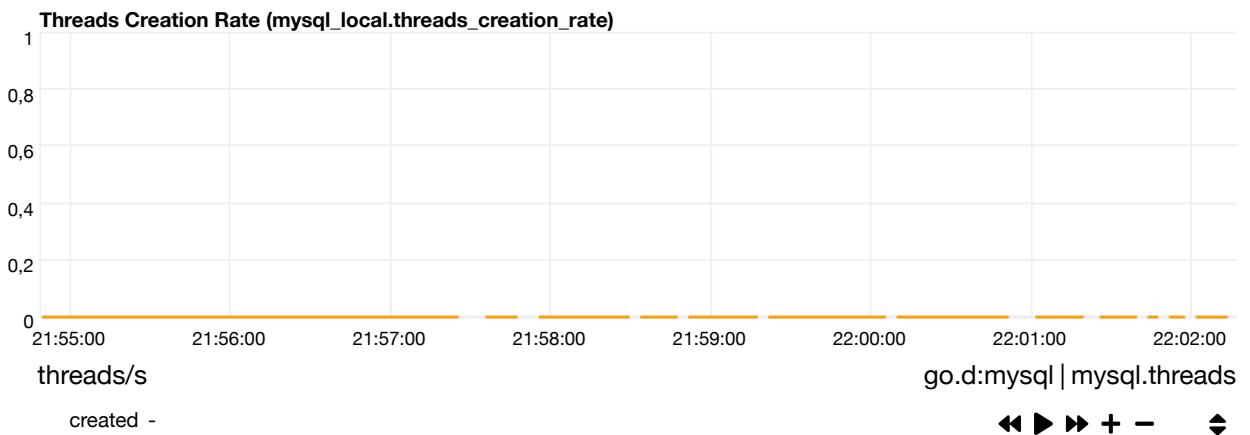


# binlog

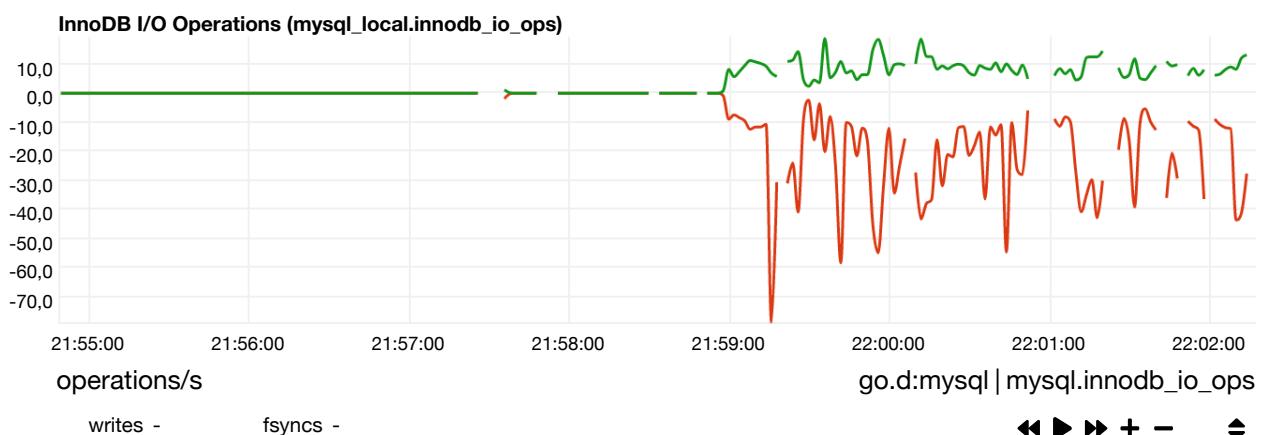
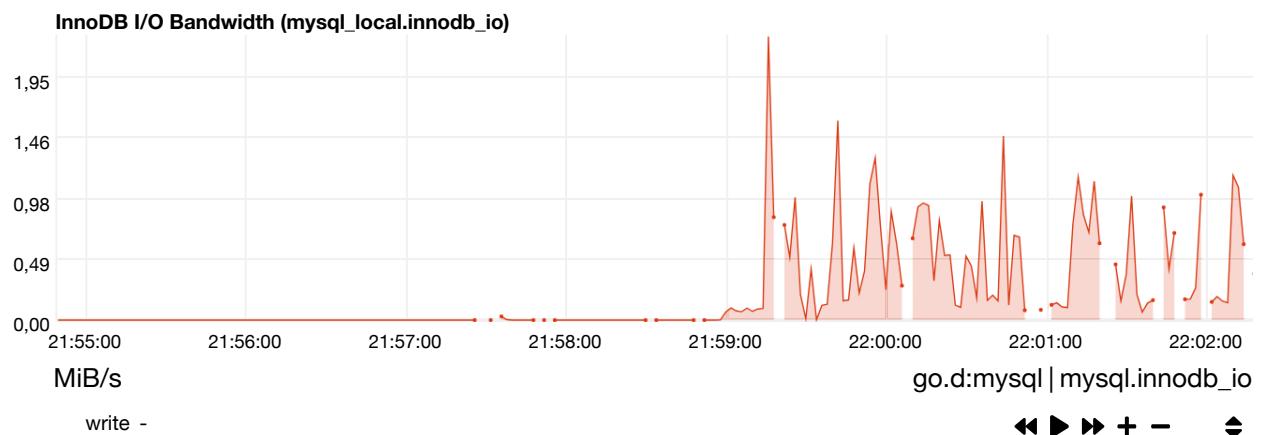


# threads

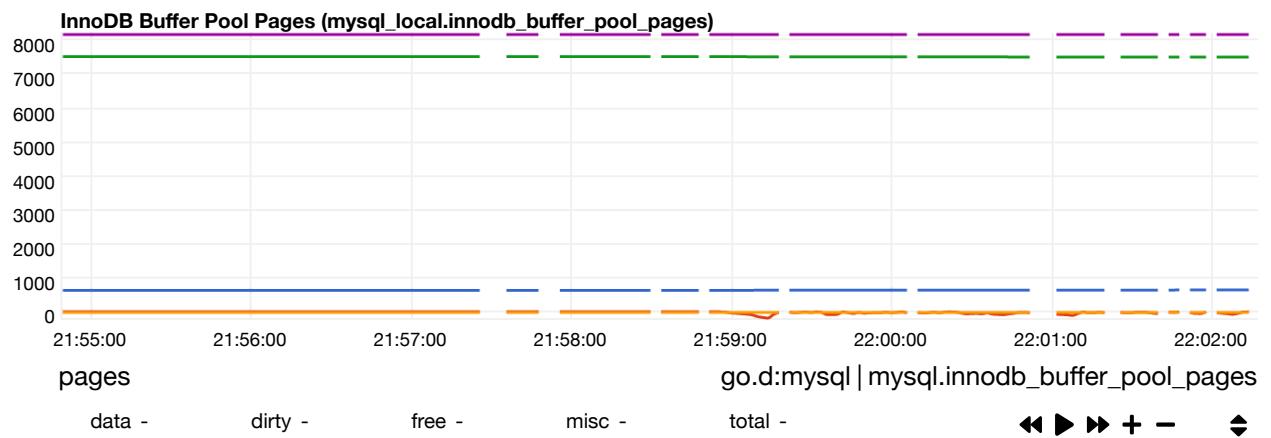
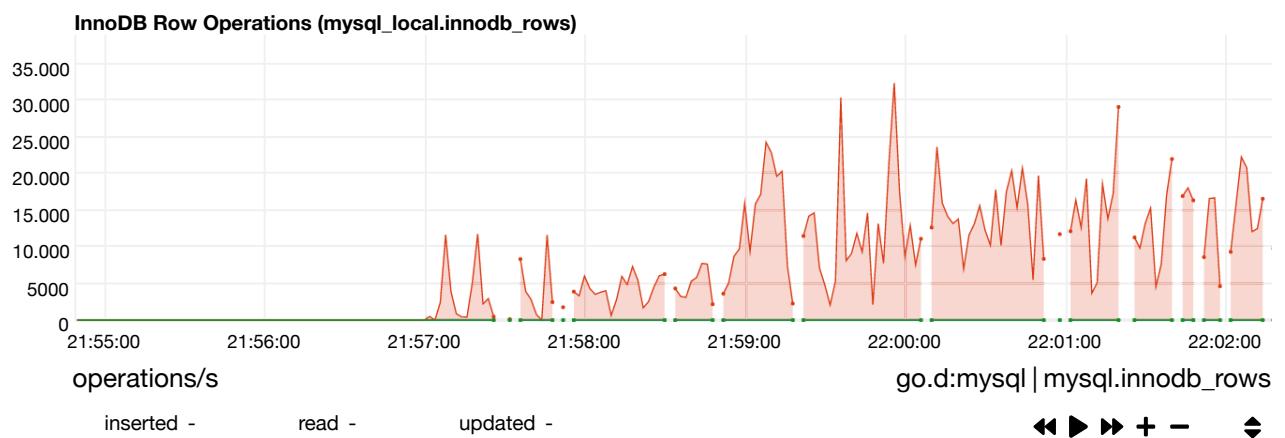
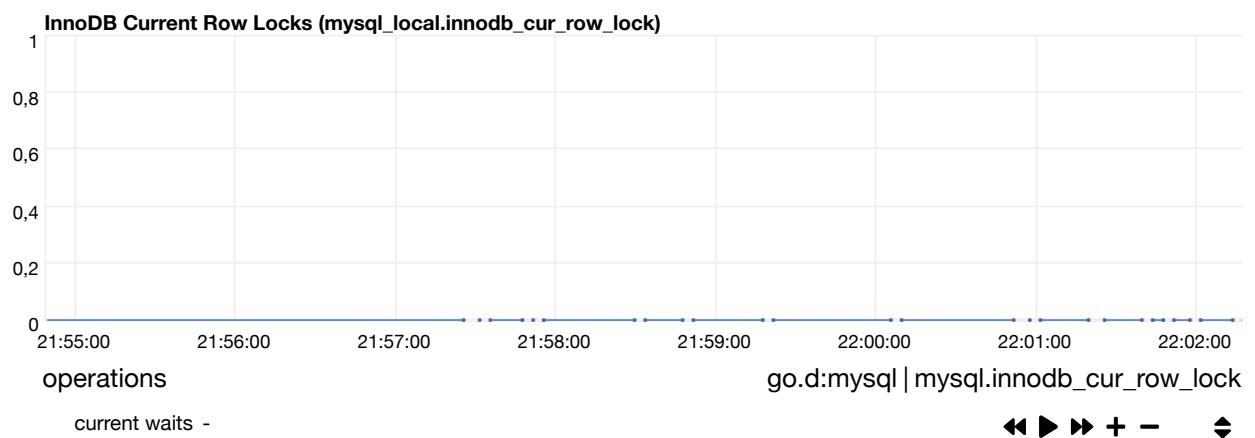
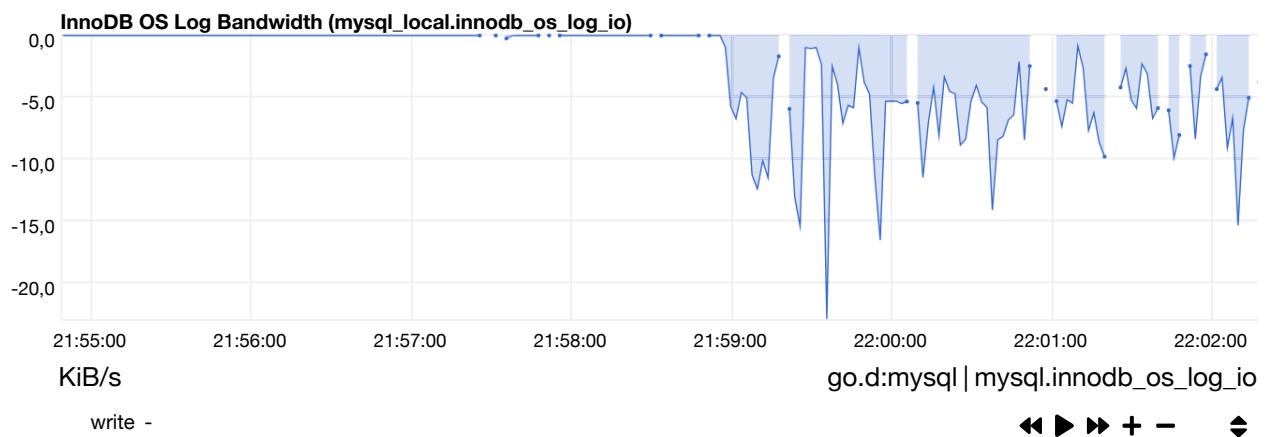




## innodb

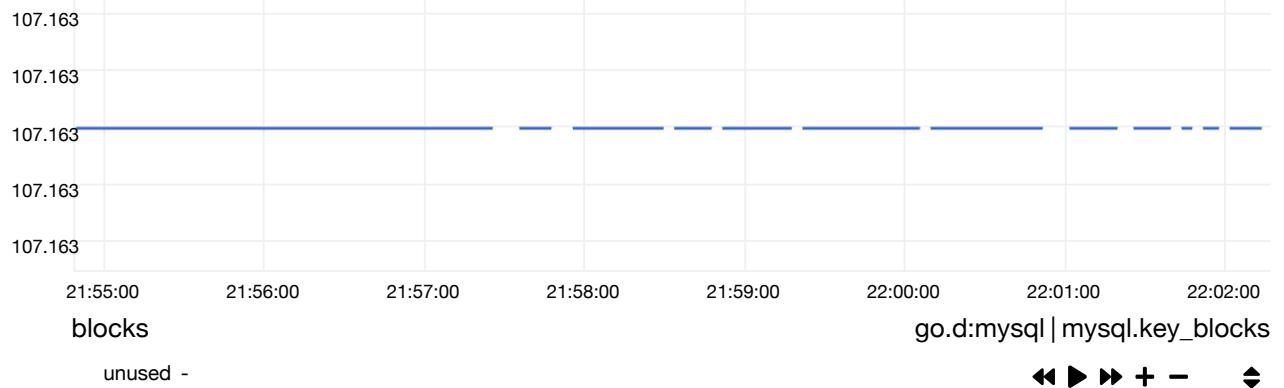




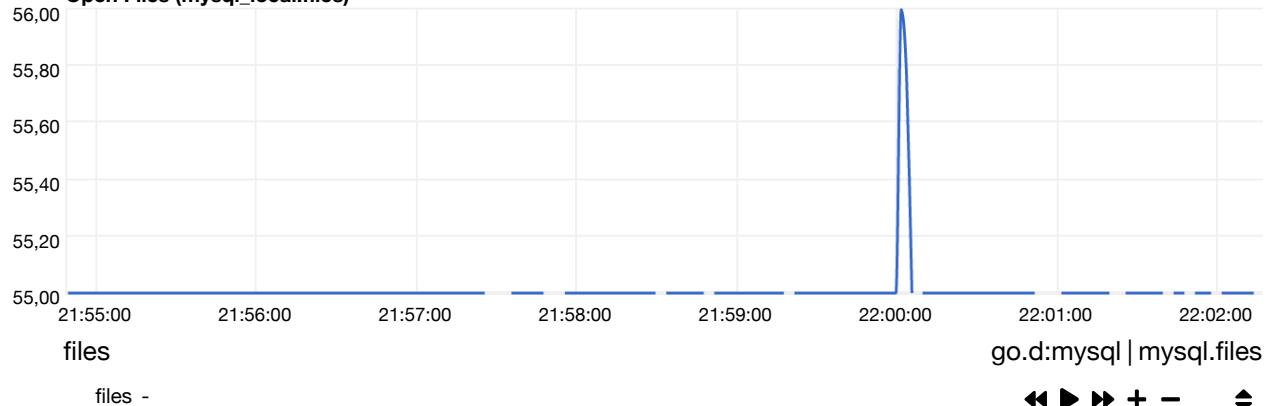


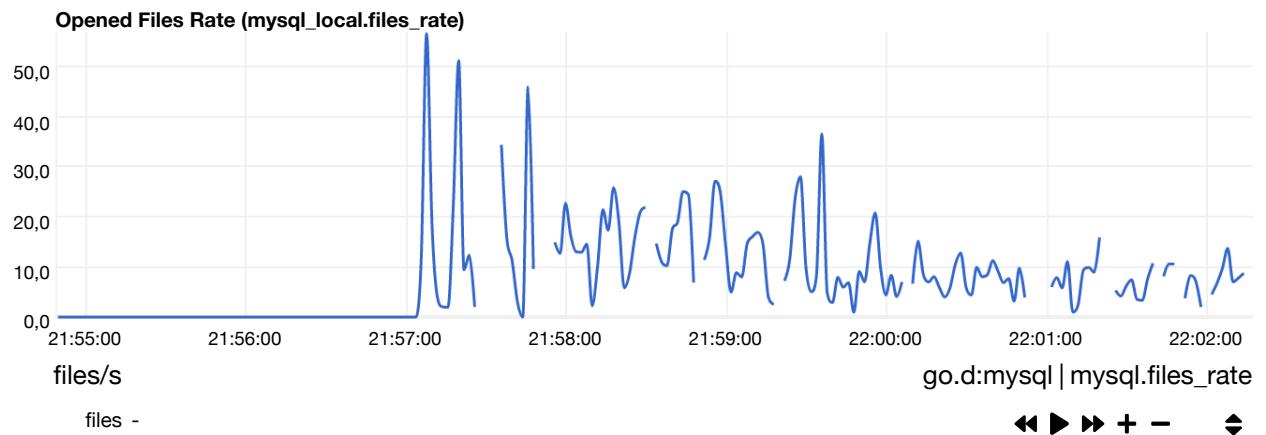


# myisam

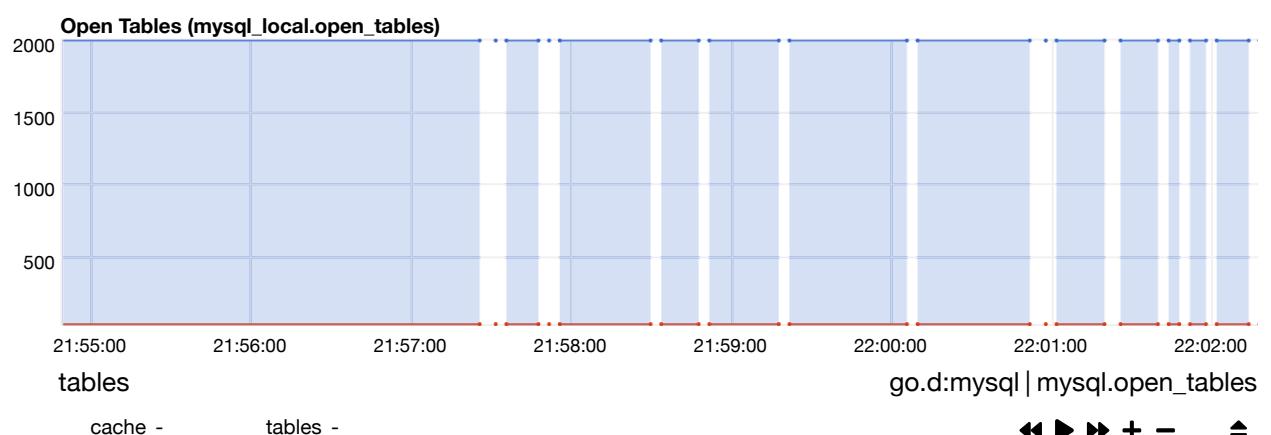
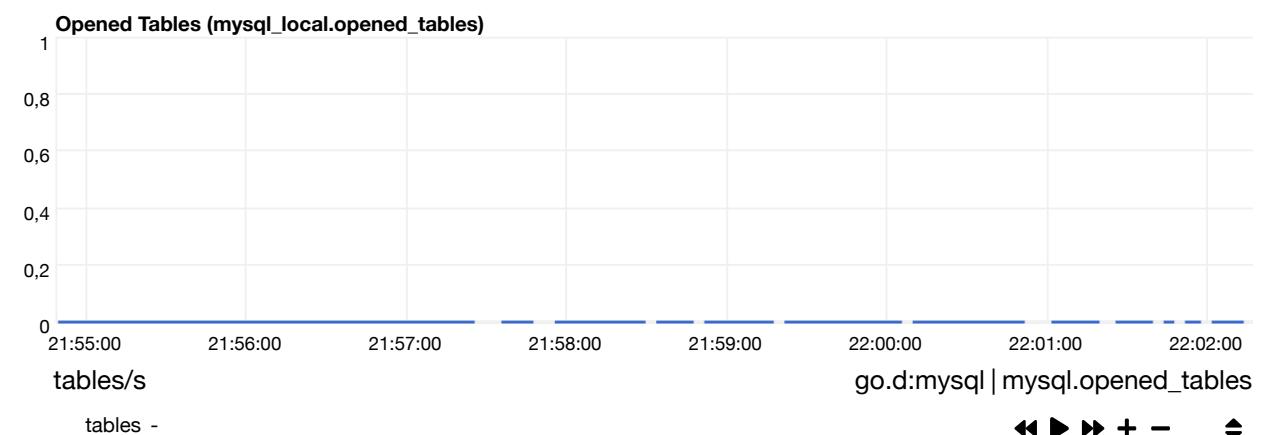
**MyISAM Key Cache Blocks (mysql\_local.key\_blocks)****MyISAM Key Cache Requests (mysql\_local.key\_requests)****MyISAM Key Cache Disk Operations (mysql\_local.key\_disk\_ops)**

# files

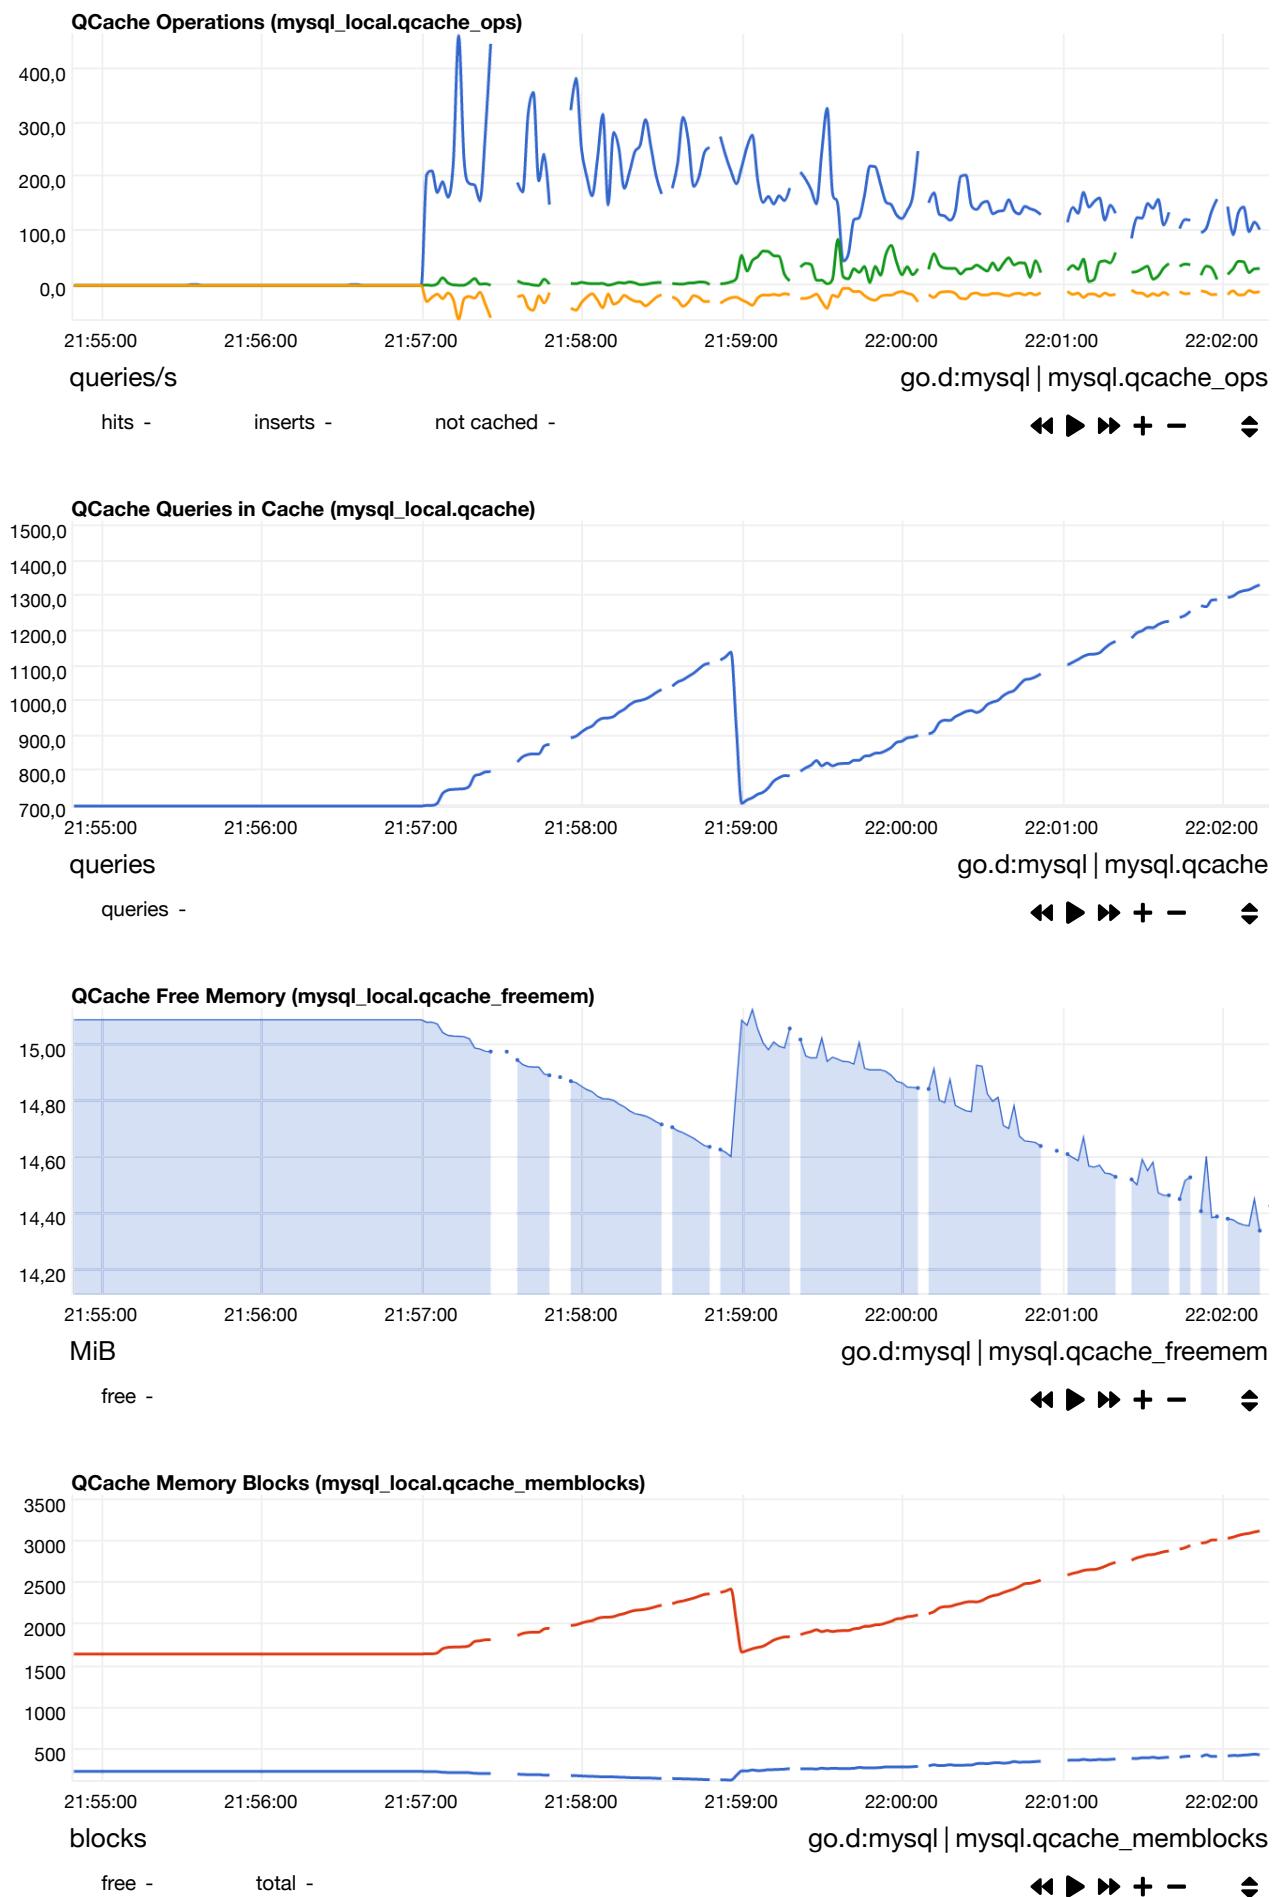
**Open Files (mysql\_local.files)**



## open tables

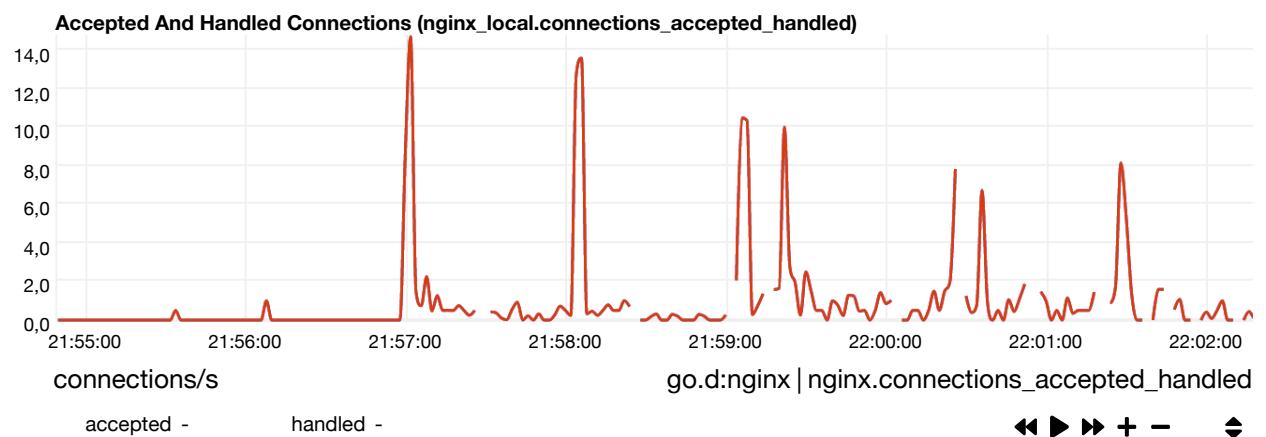
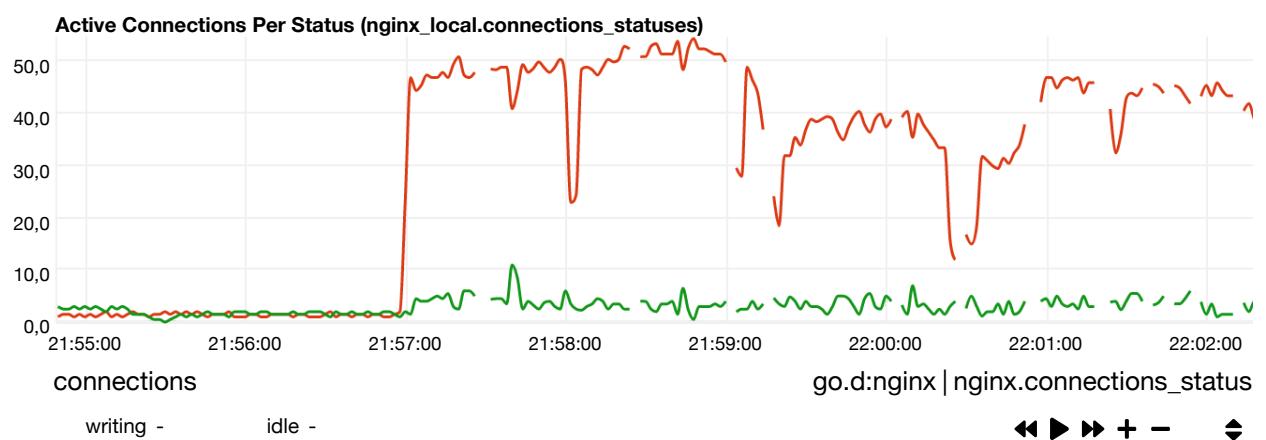
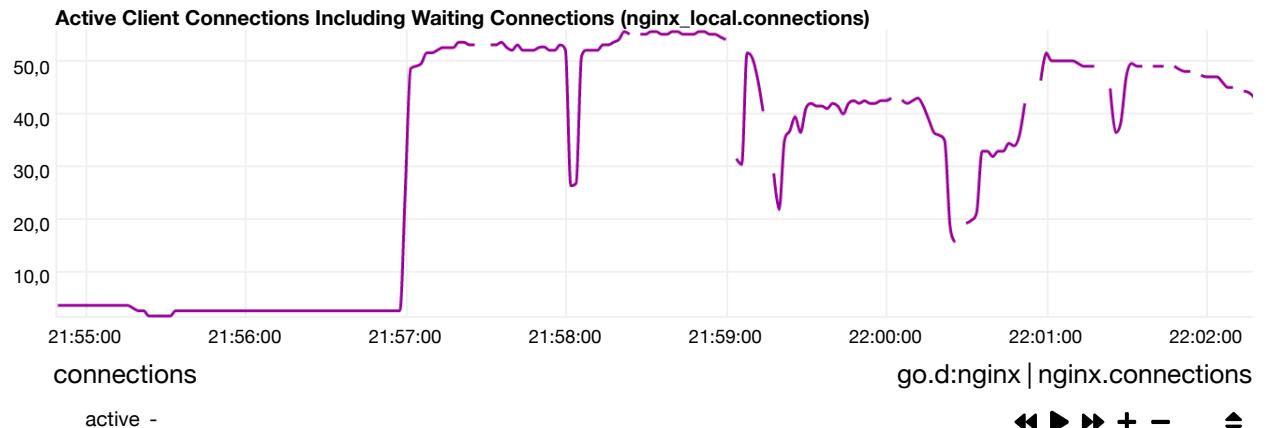


## qcache

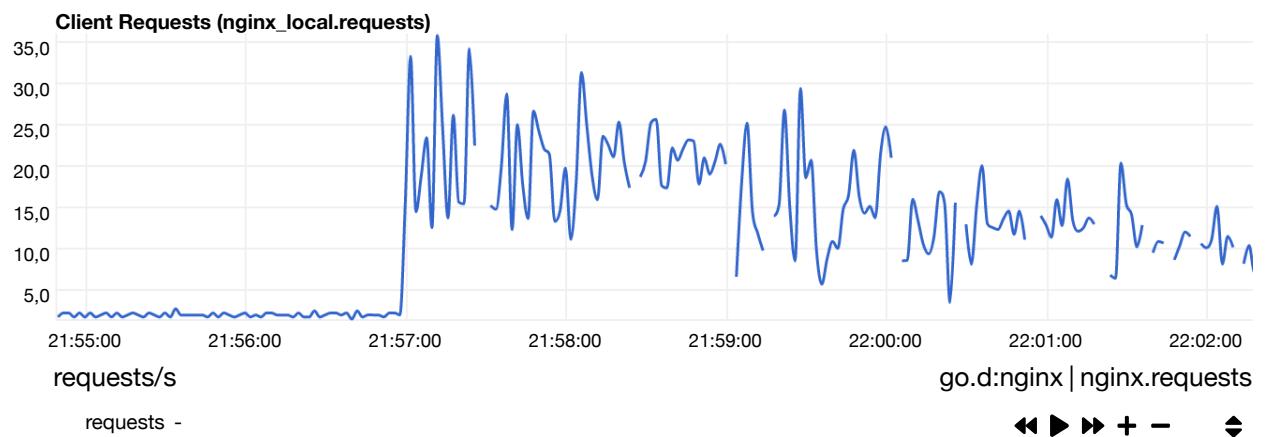


# ⌚ nginx local

## connections



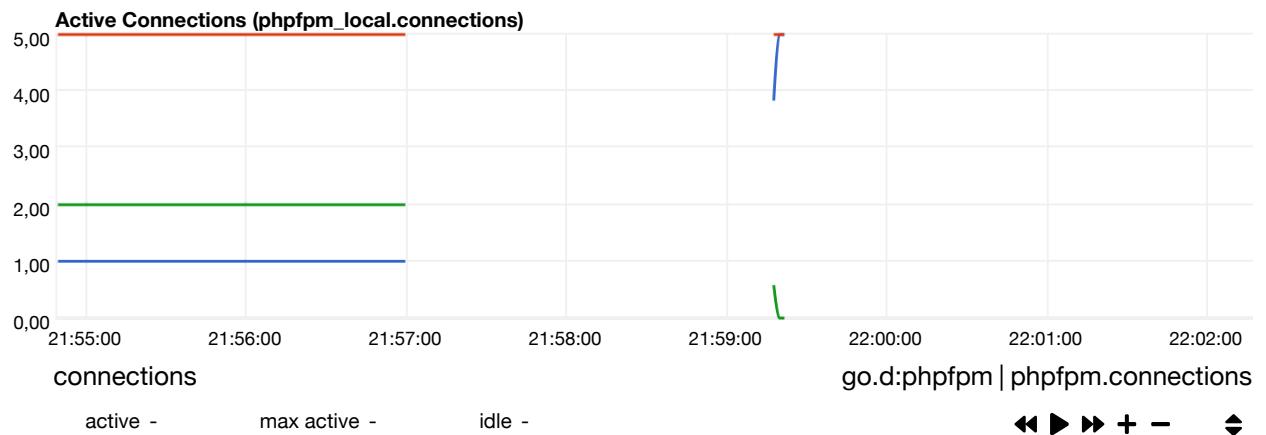
## requests



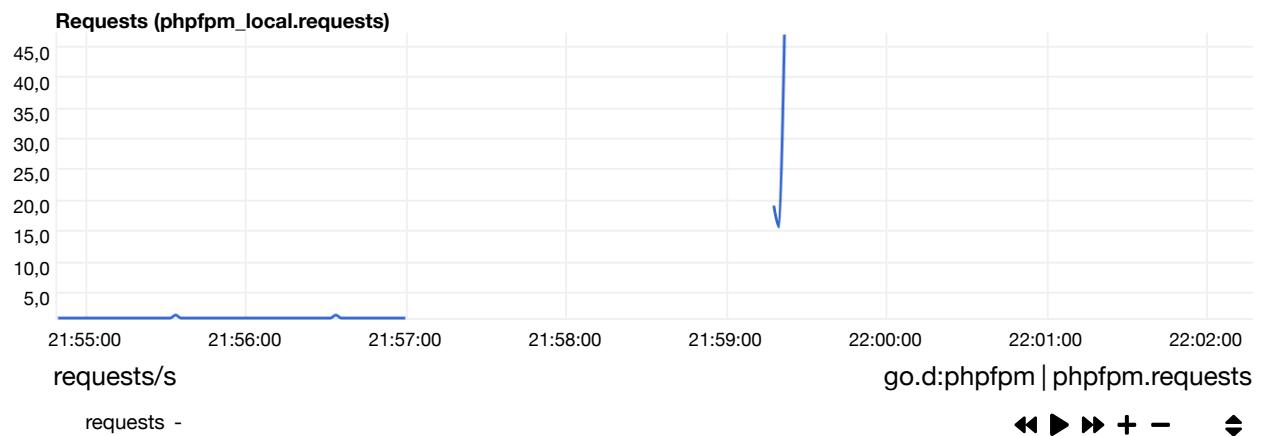
## ⌚ PHP-FPM local

Performance metrics for **PHP-FPM**, an alternative FastCGI implementation for PHP.

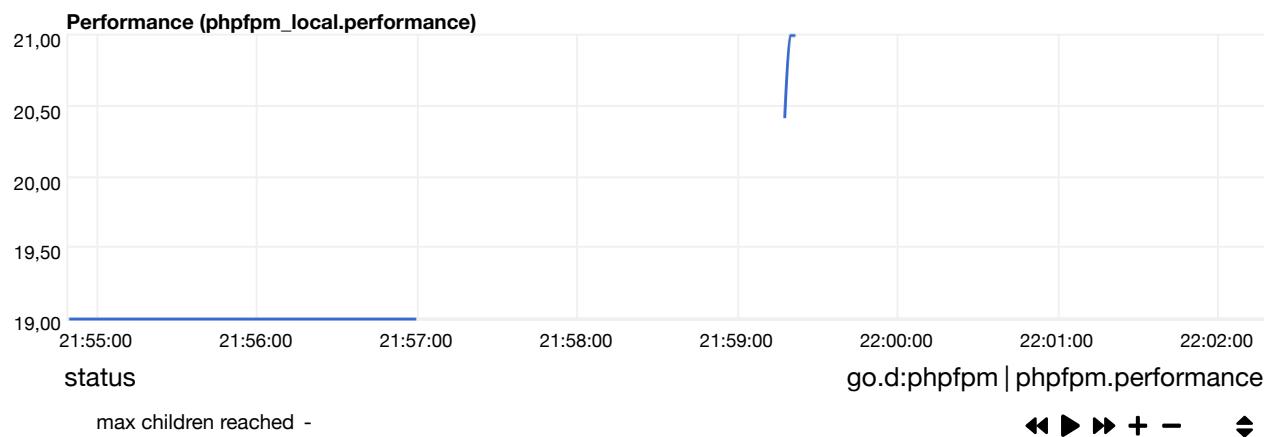
### active connections



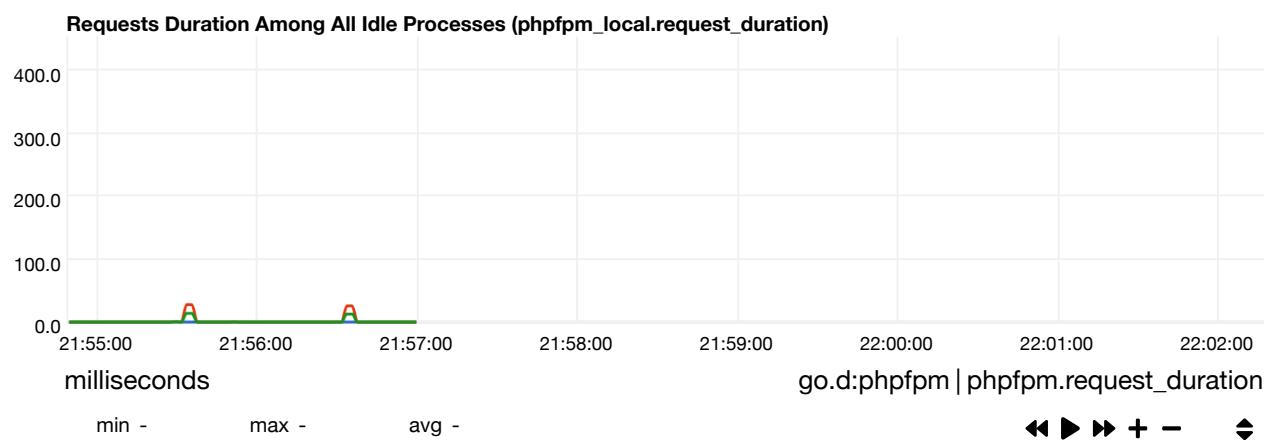
### requests



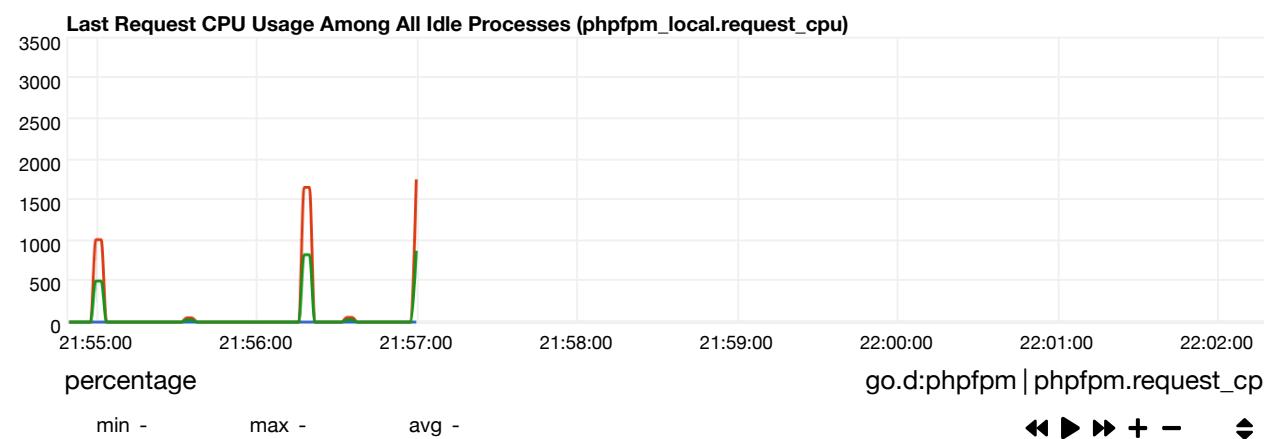
### performance



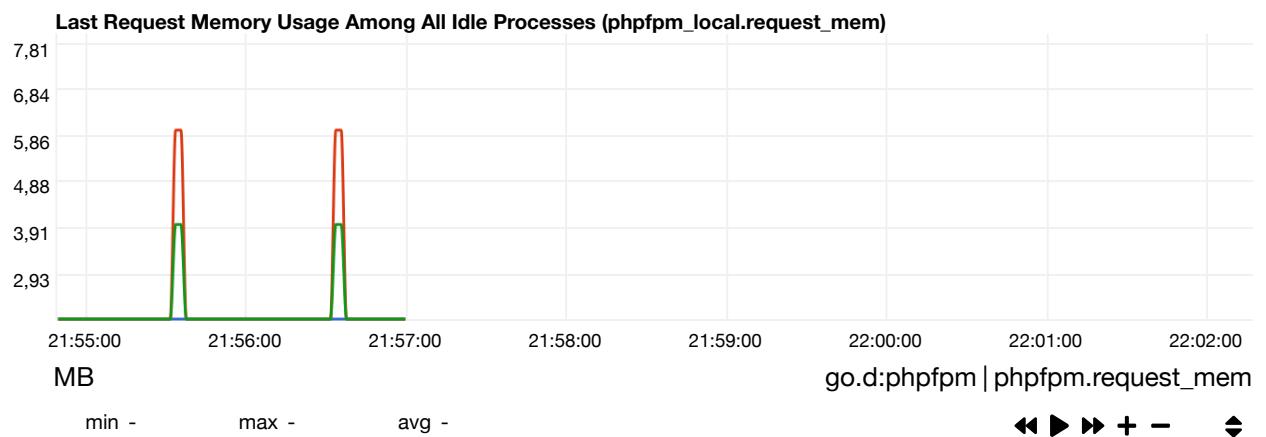
## request duration



## request CPU



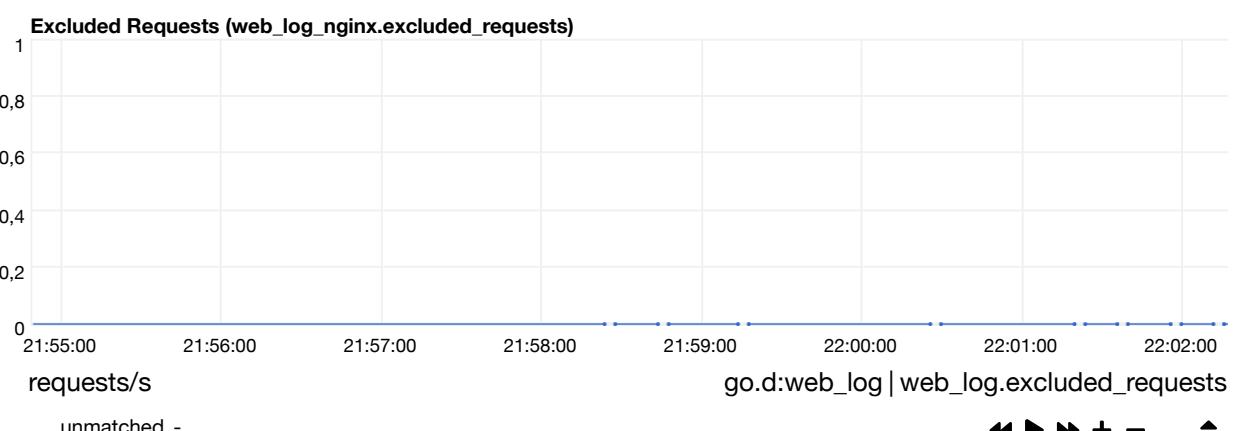
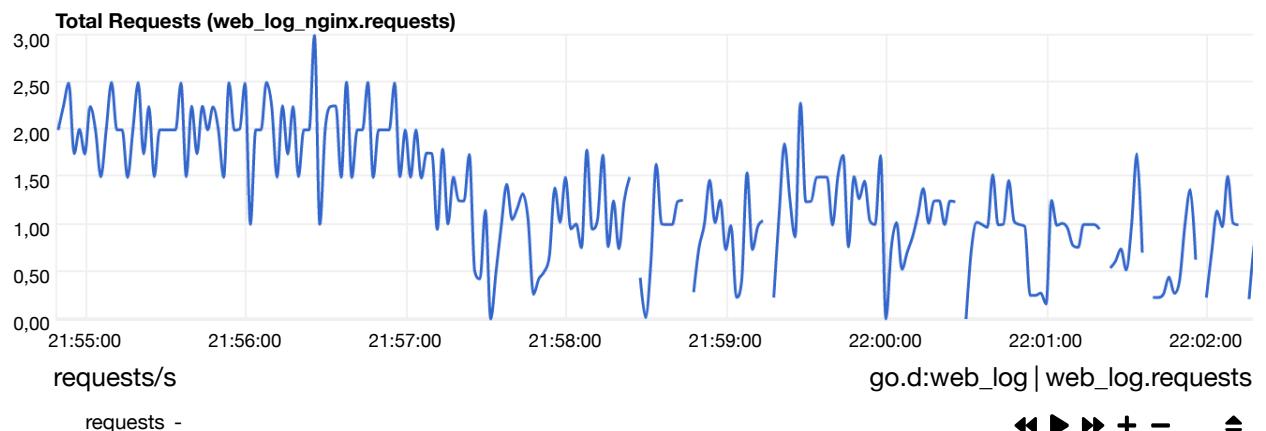
## request memory



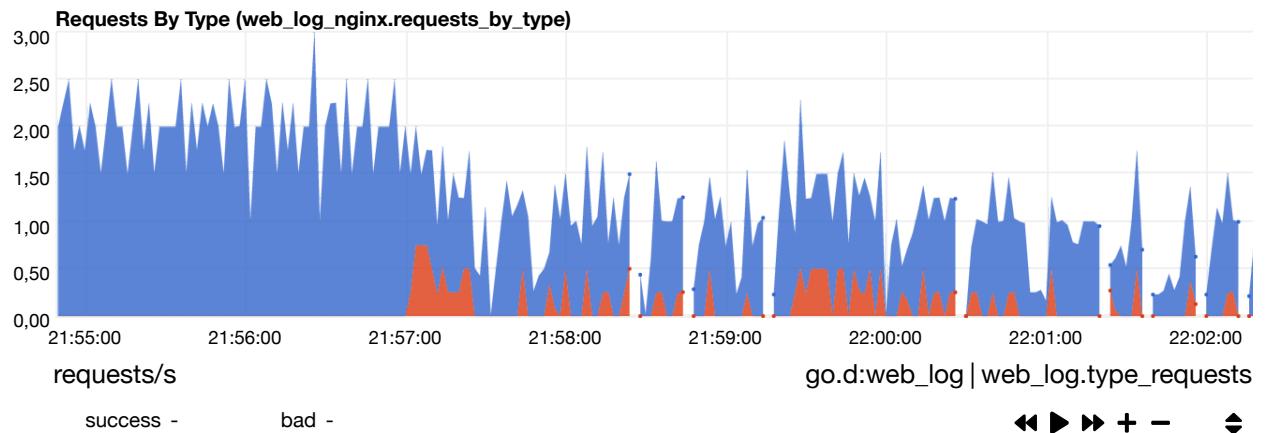
## 📄 web log nginx

Information extracted from a server log file. `web_log` plugin incrementally parses the server log file to provide, in real-time, a break down of key server performance metrics. For web servers, an extended log file format may optionally be used (for `nginx` and `apache`) offering timing information and bandwidth for both requests and responses. `web_log` plugin may also be configured to provide a break down of requests per URL pattern (check `/etc/netdata/python.d/web_log.conf` ([https://github.com/netdata/netdata/blob/master/collectors/python.d.plugin/web\\_log/web\\_log.conf](https://github.com/netdata/netdata/blob/master/collectors/python.d.plugin/web_log/web_log.conf))).

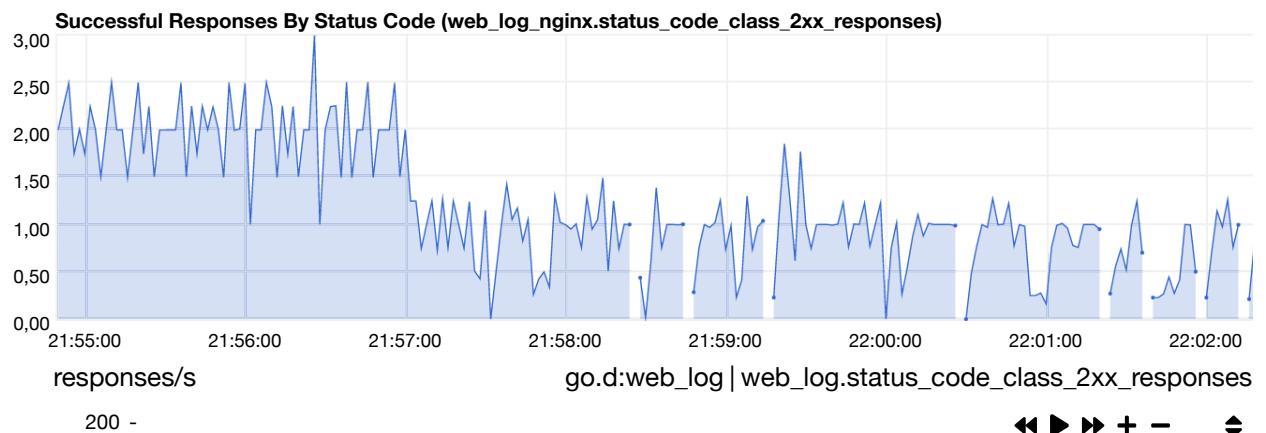
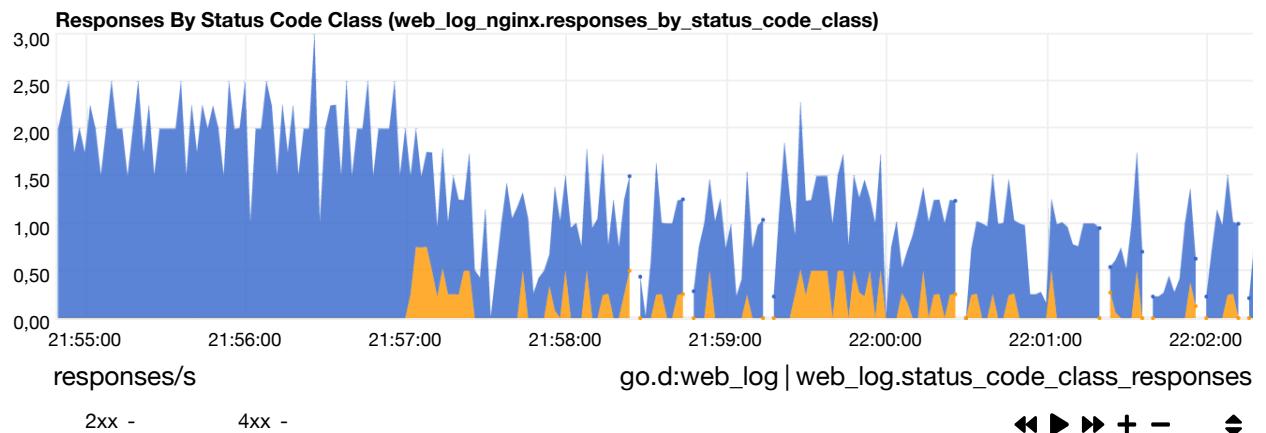
## requests

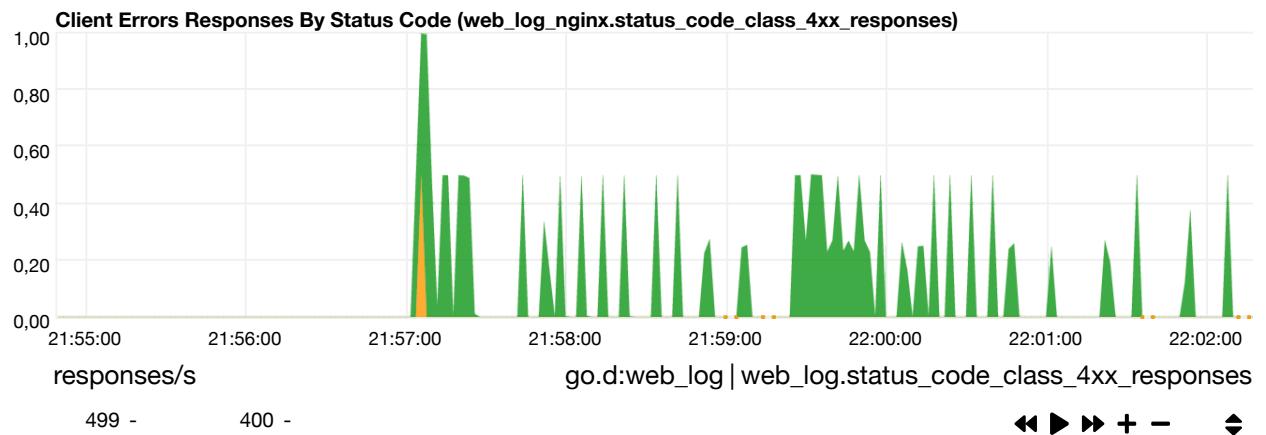


Web server responses by type. success includes **1xx**, **2xx**, **304** and **401**, error includes **5xx**, redirect includes **3xx** except **304**, bad includes **4xx** except **401**, other are all the other responses.



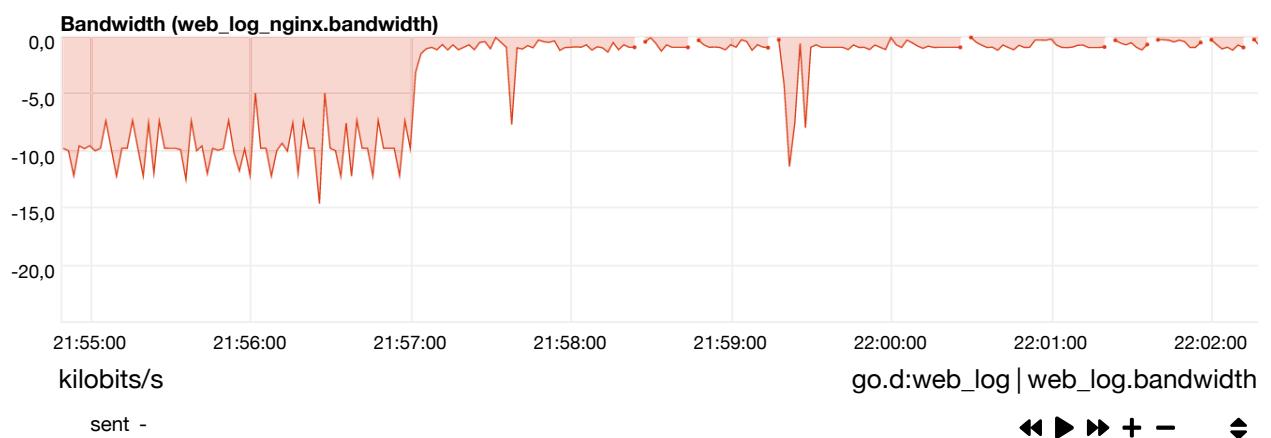
## responses



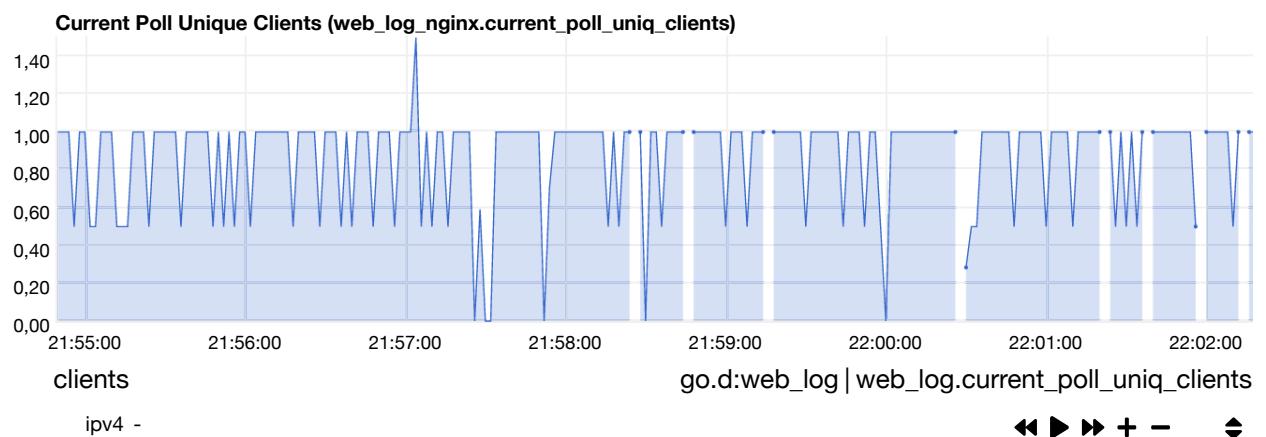


## bandwidth

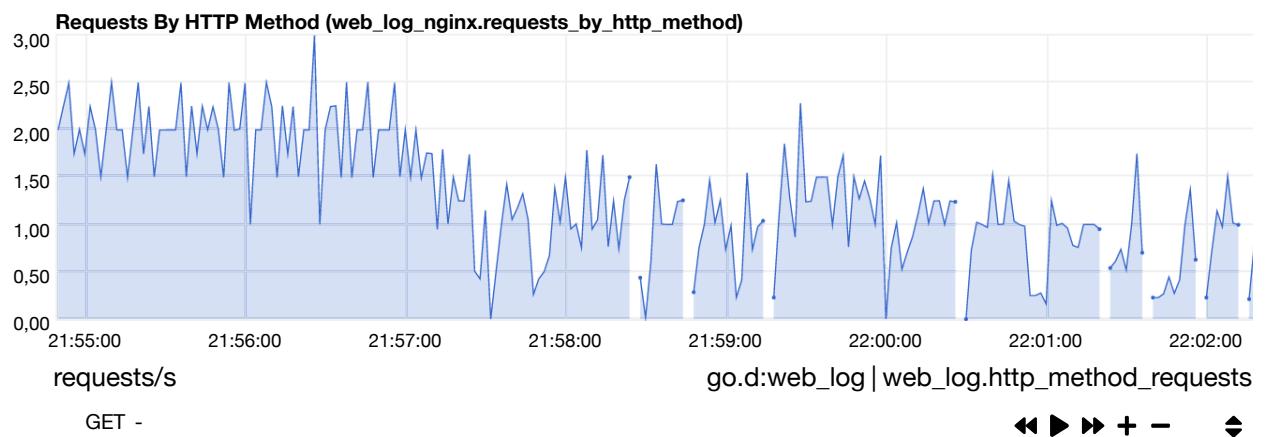
Bandwidth of requests ( received ) and responses ( sent ). received requires an extended log format (without it, the web server log does not have this information). This chart may present unusual spikes, since the bandwidth is accounted at the time the log line is saved by the web server, even if the time needed to serve it spans across a longer duration. We suggest to use QoS (e.g. FireQOS (<http://firehol.org/#fireqos>)) for accurate accounting of the web server bandwidth.



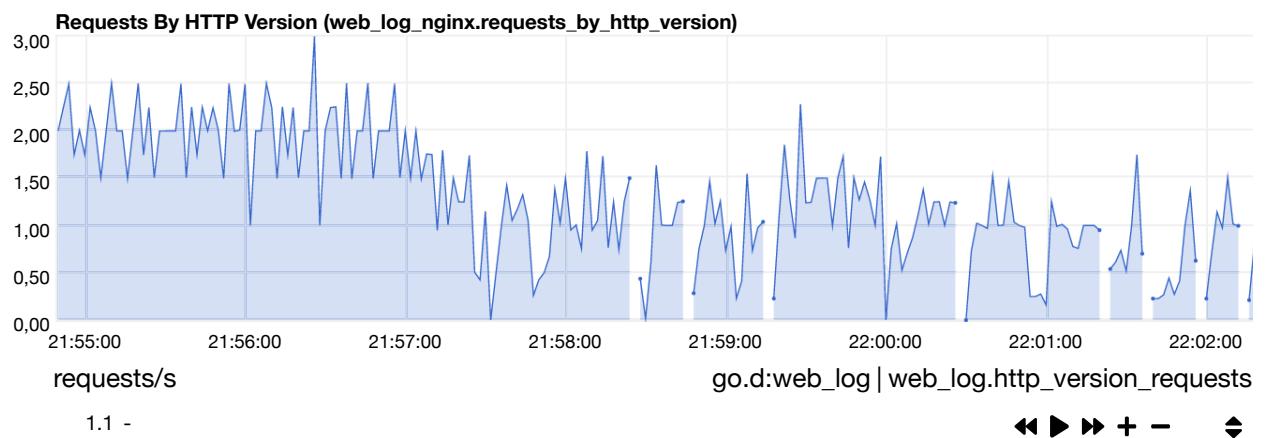
## client



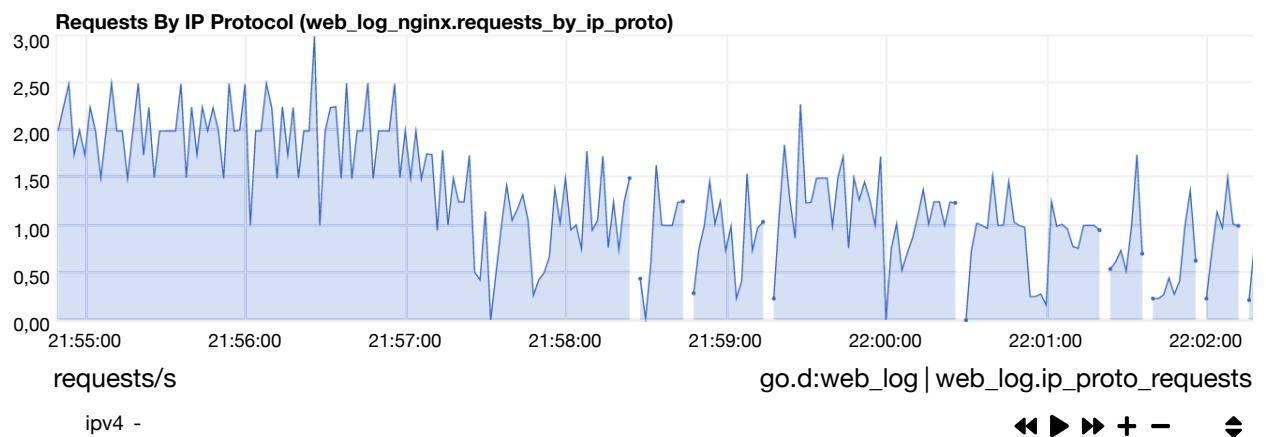
## http method



## http version



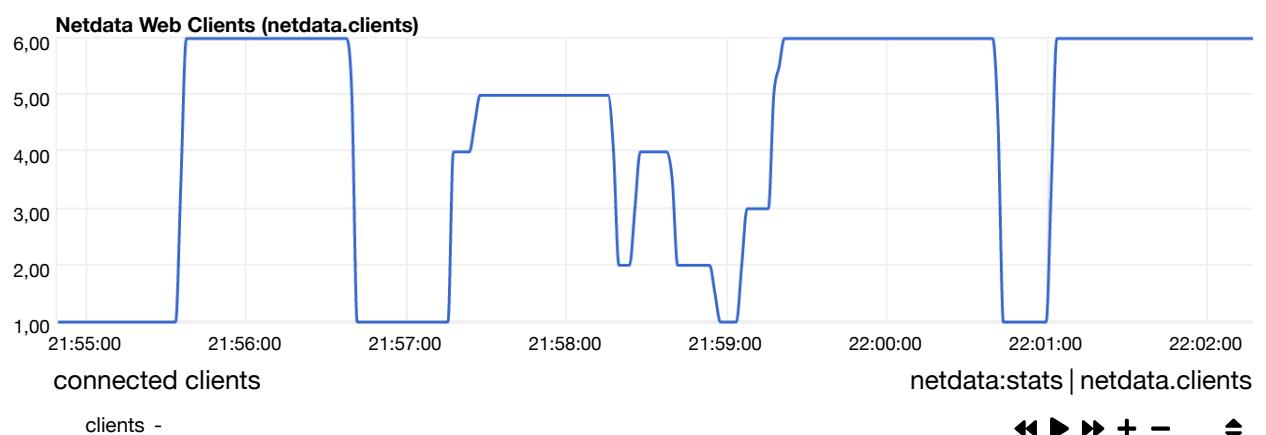
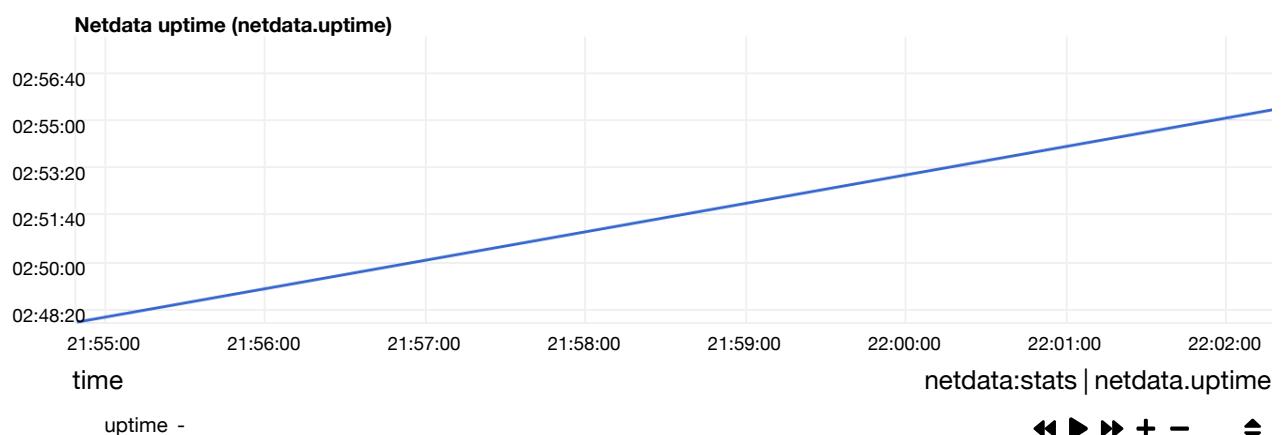
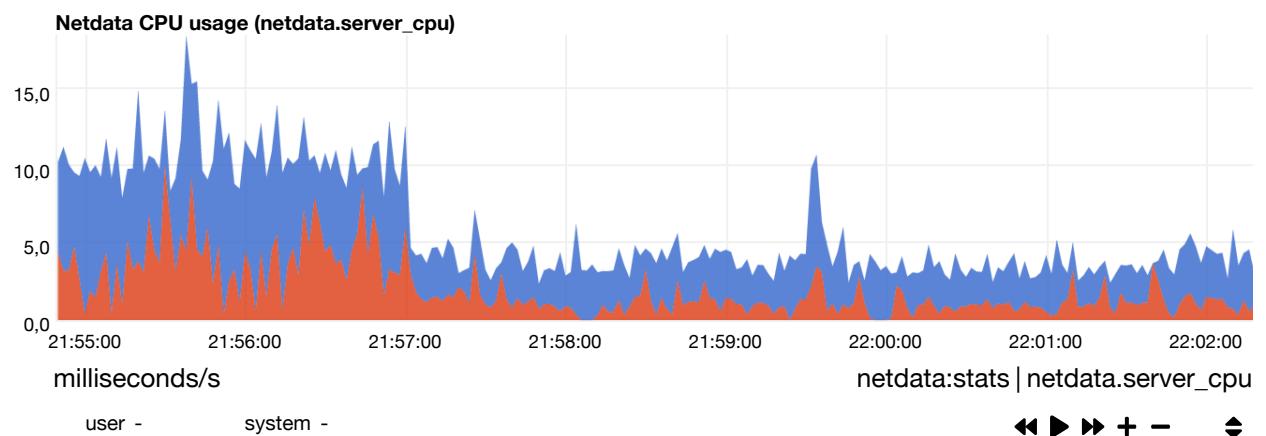
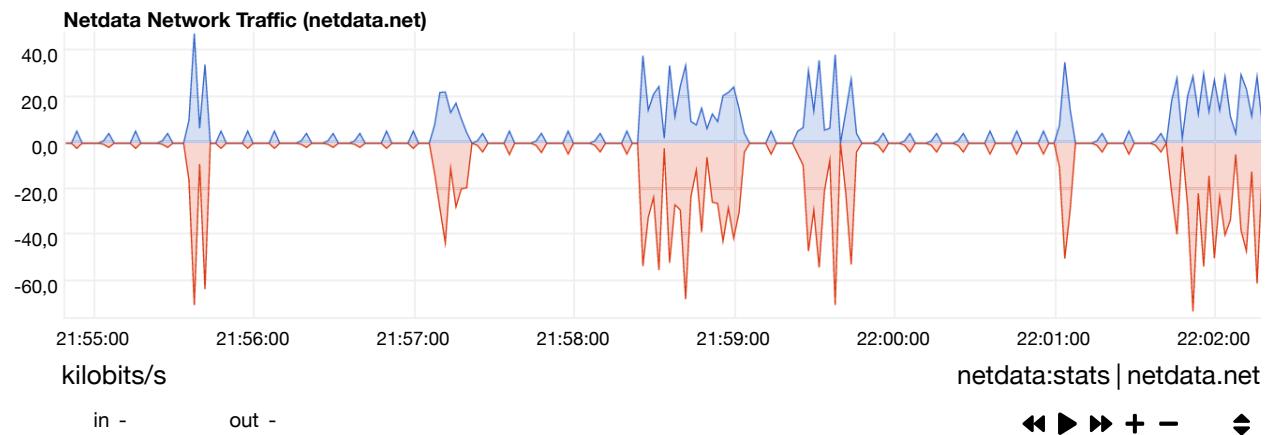
## ip proto

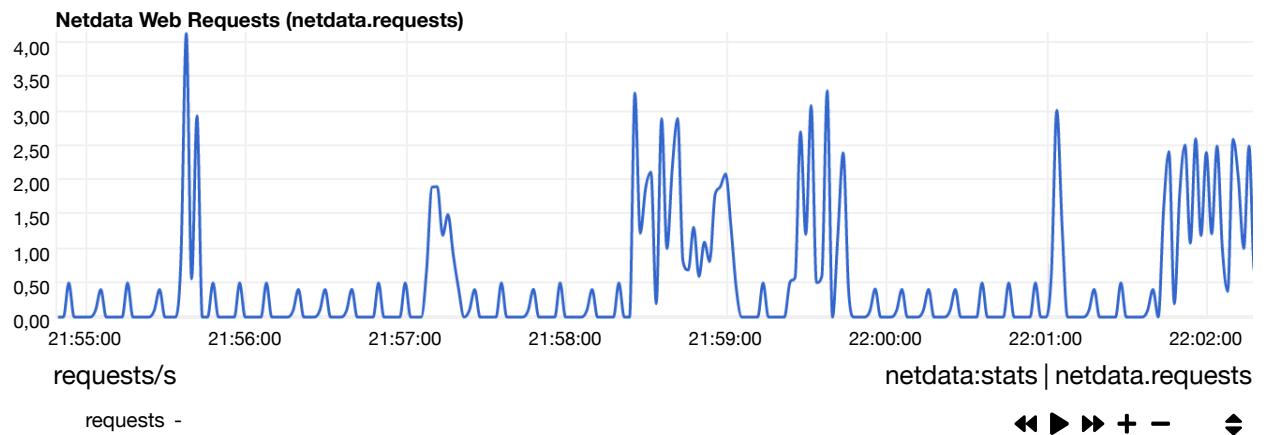


# Netdata Monitoring

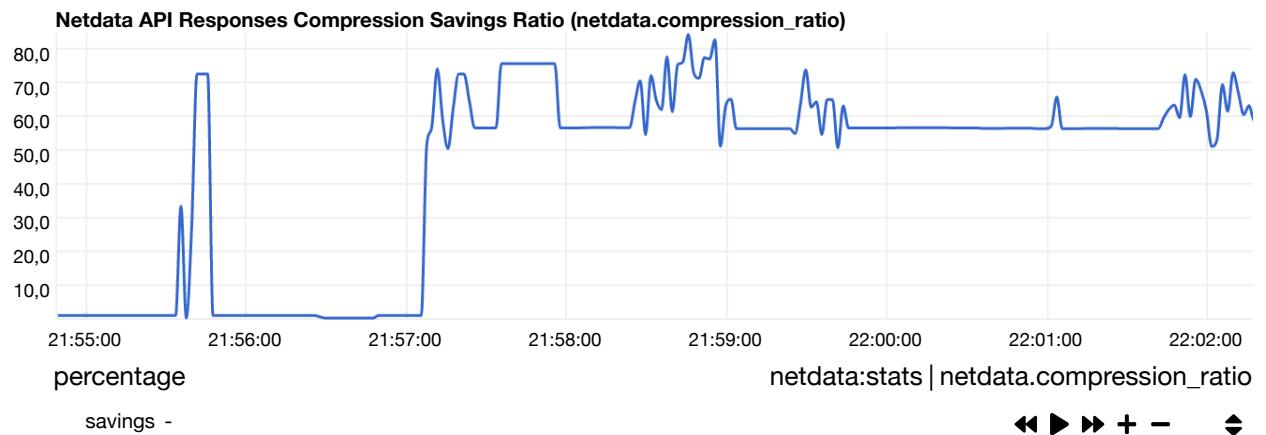
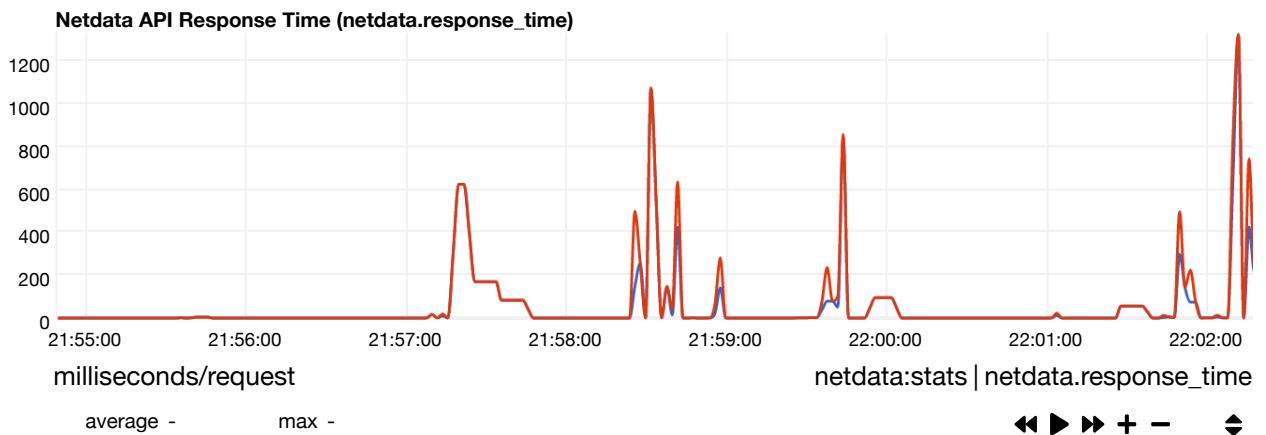
Performance metrics for the operation of netdata itself and its plugins.

## netdata

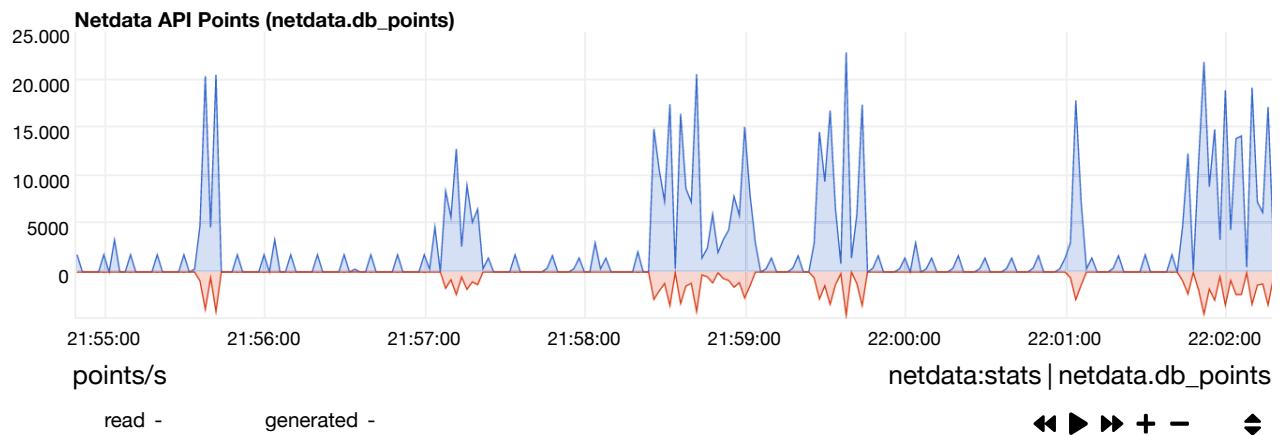
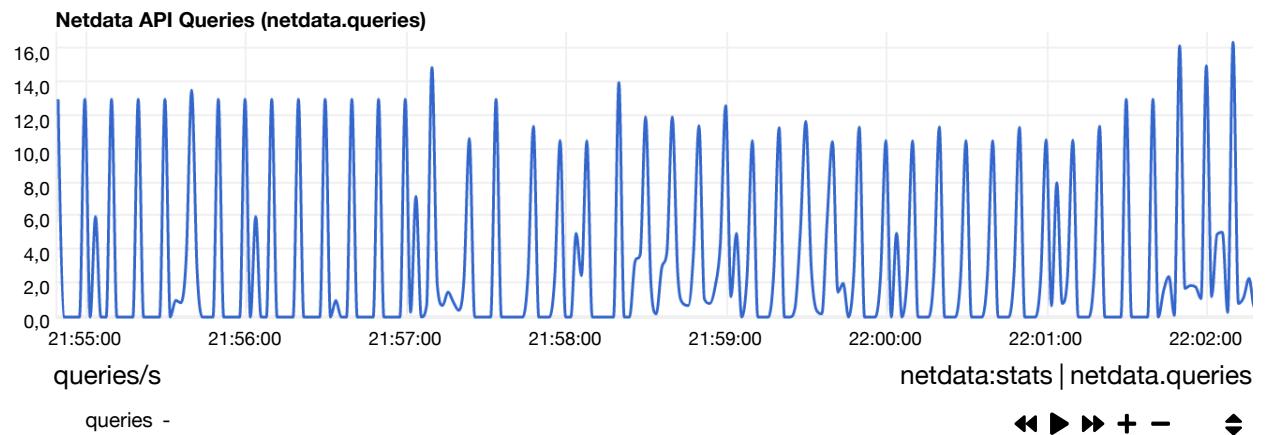




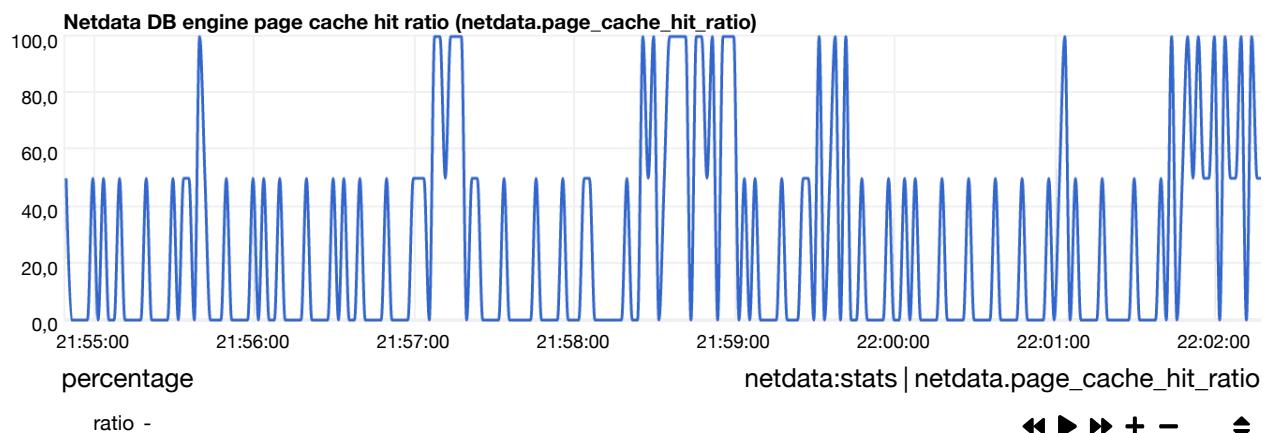
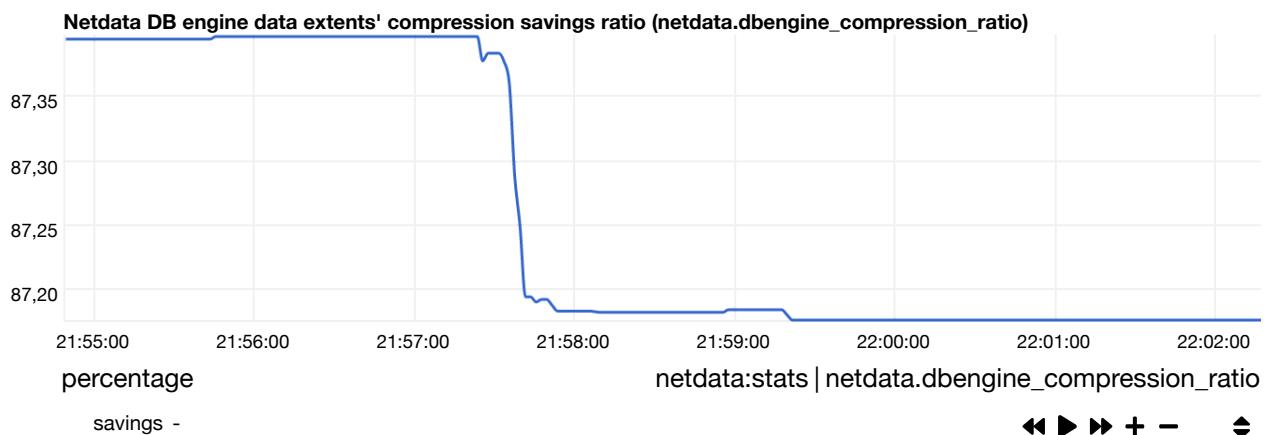
The netdata API response time measures the time netdata needed to serve requests. This time includes everything, from the reception of the first byte of a request, to the dispatch of the last byte of its reply, therefore it includes all network latencies involved (i.e. a client over a slow network will influence these metrics).

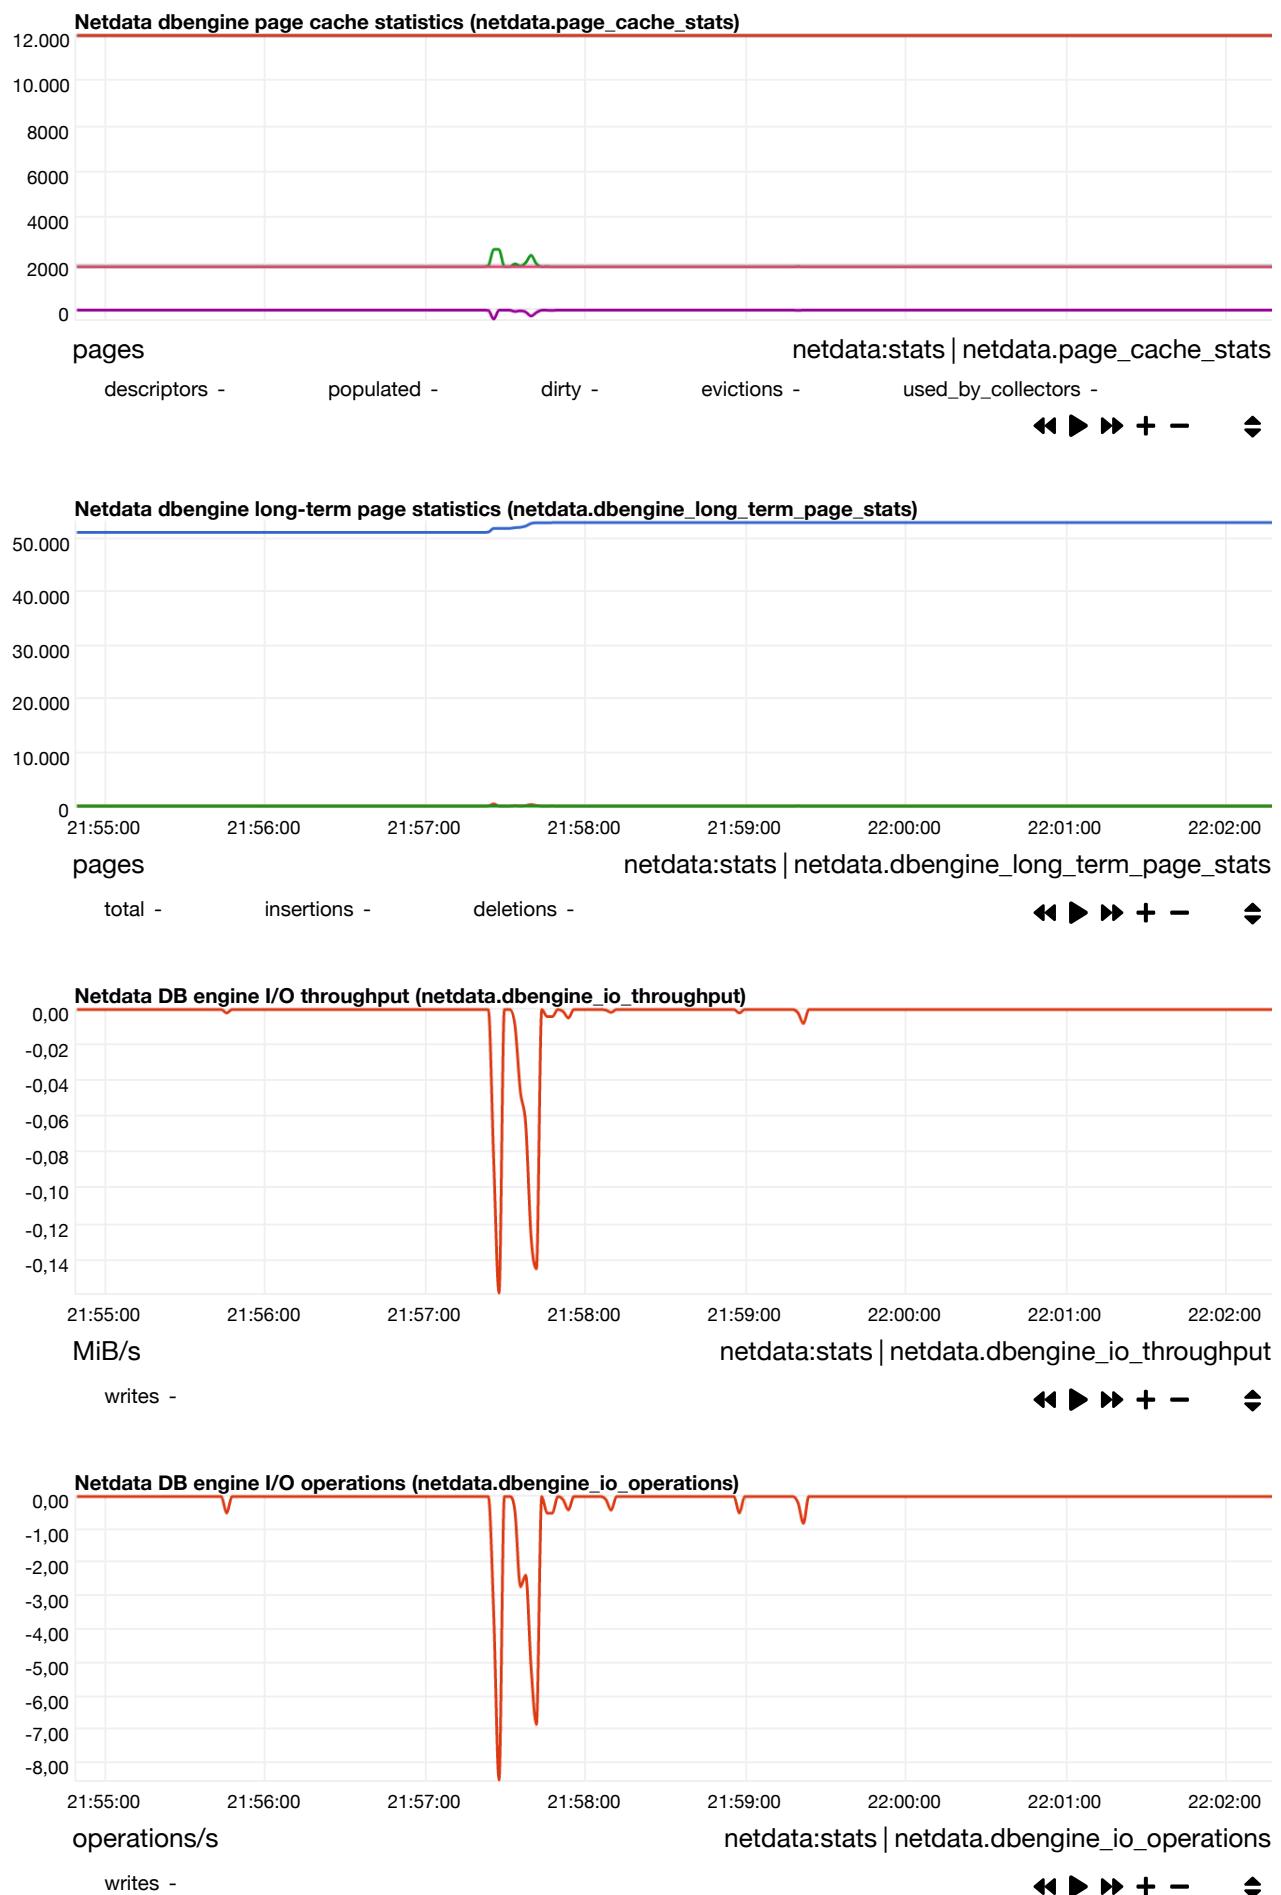


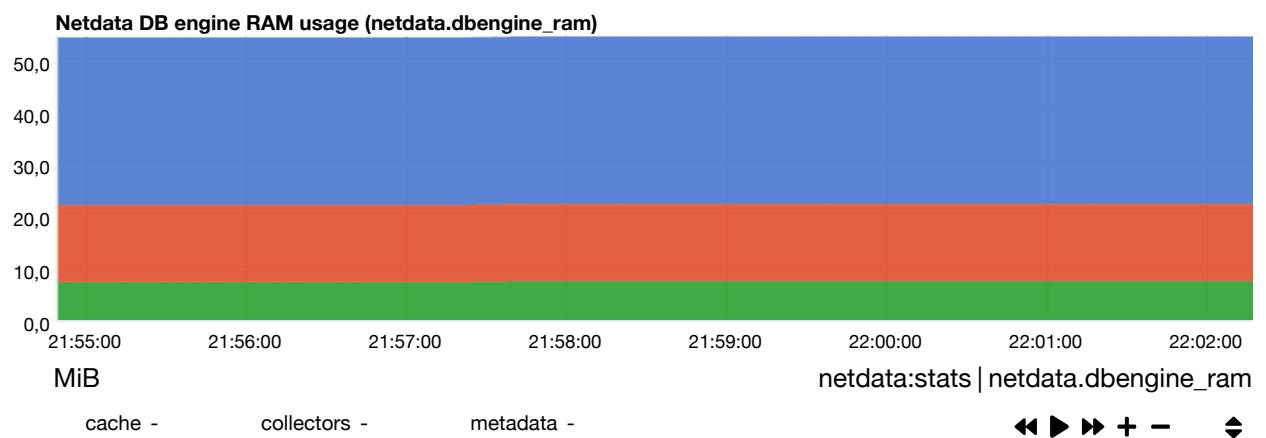
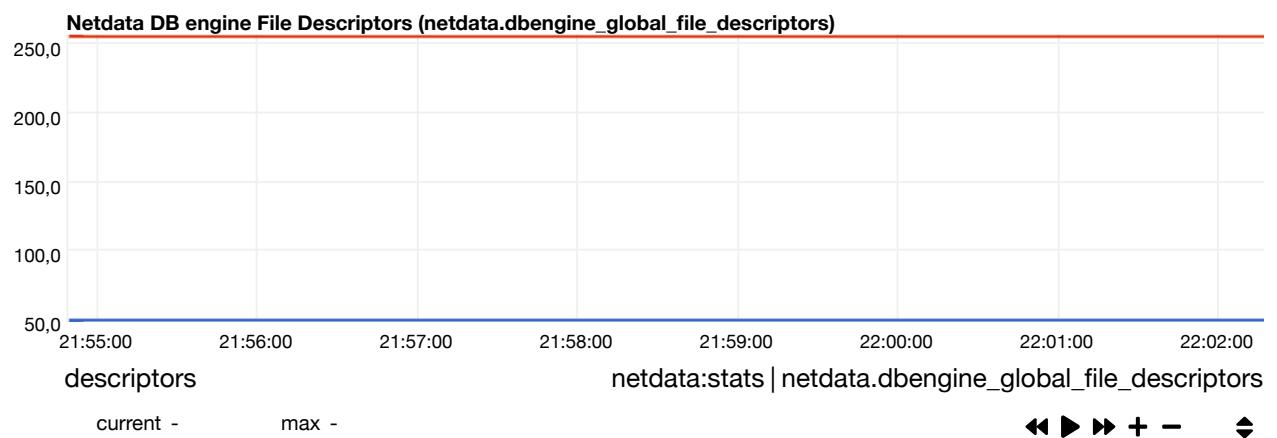
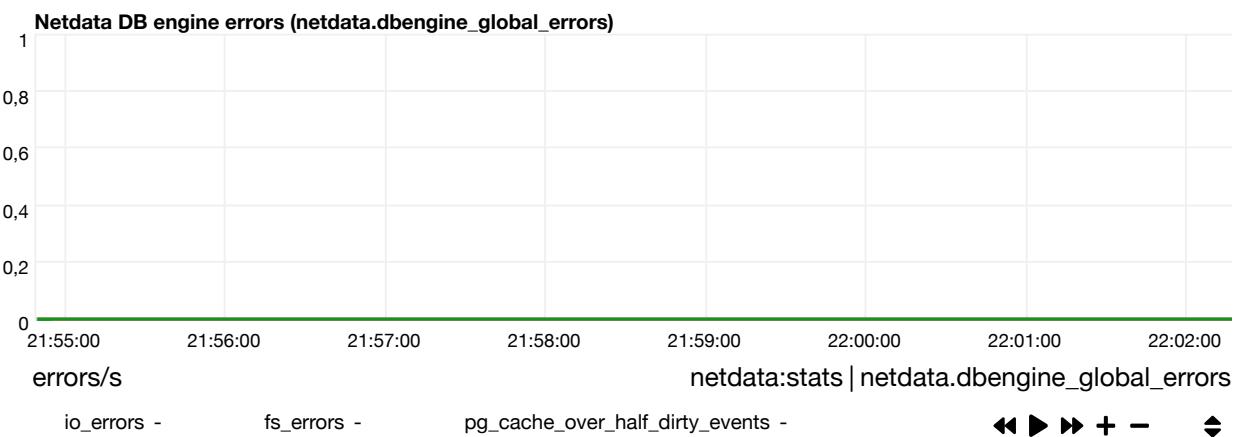
## queries



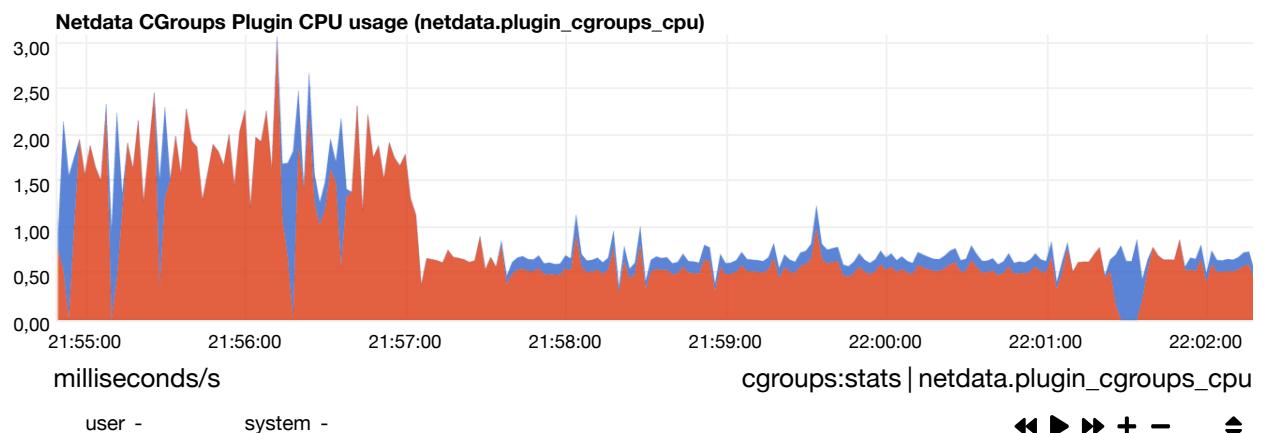
## dbengine



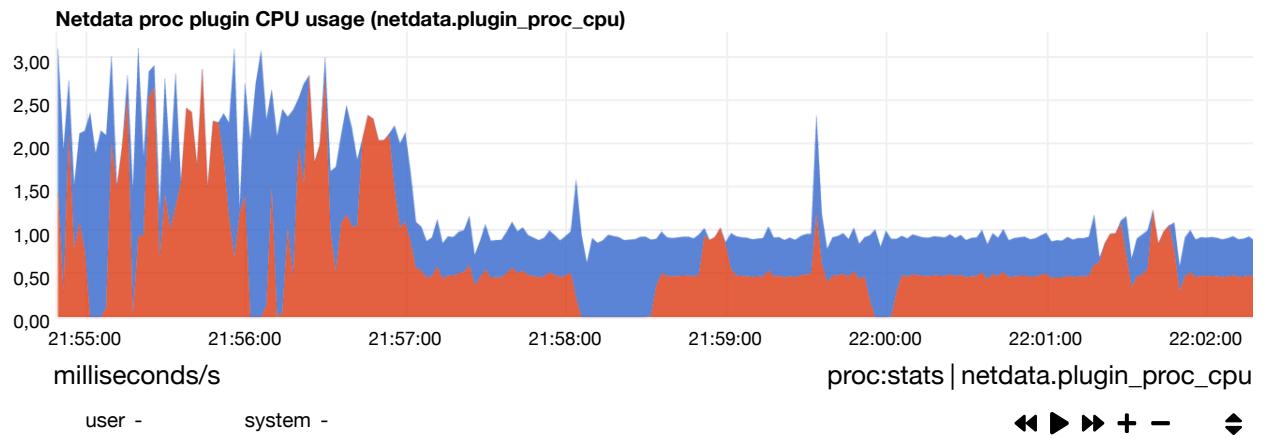




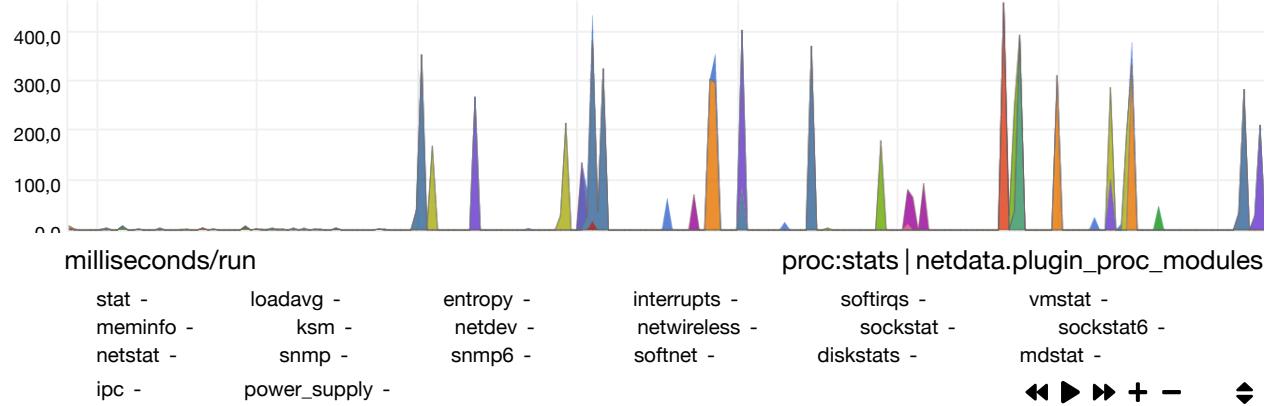
## cgroups



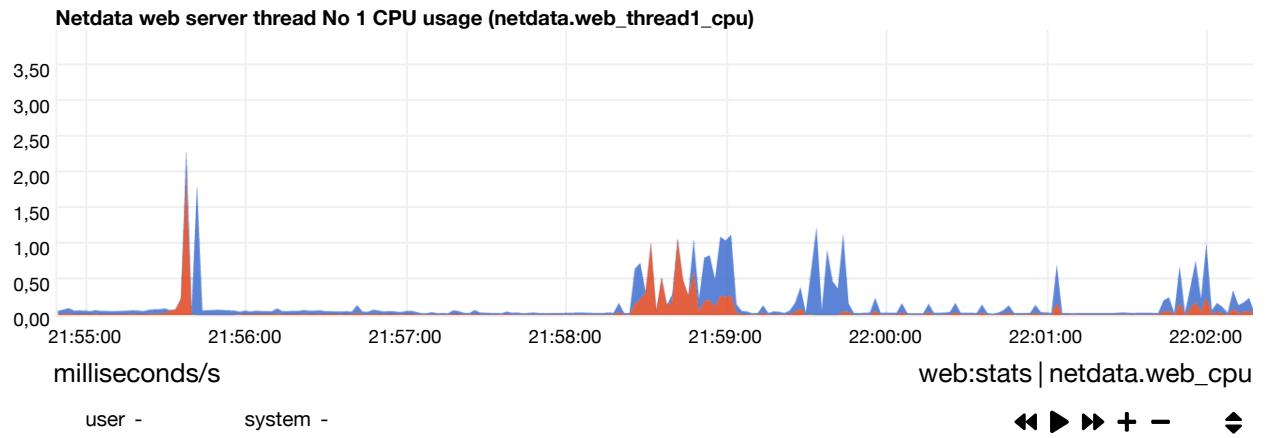
## proc

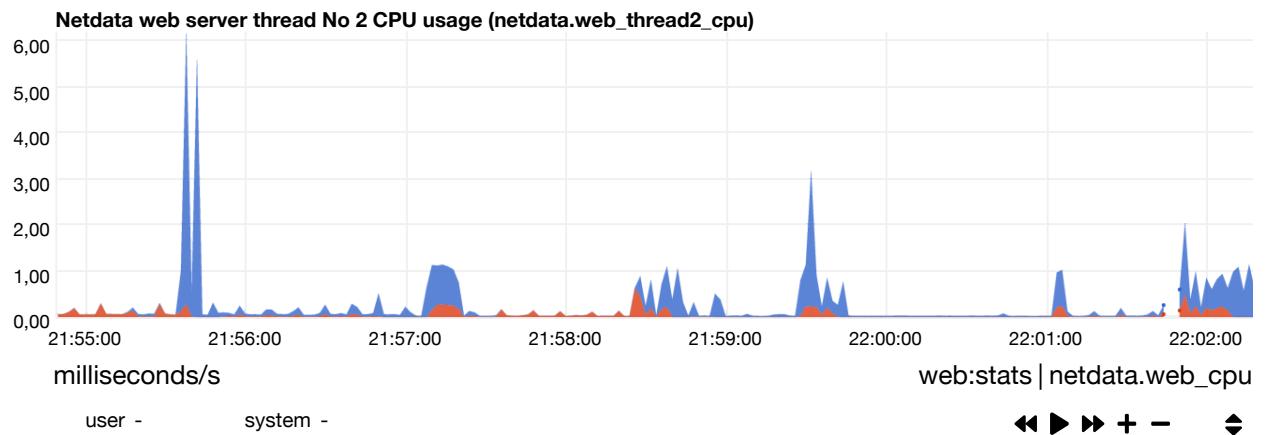


**Netdata proc plugin modules durations (netdata.plugin\_proc\_modules)**

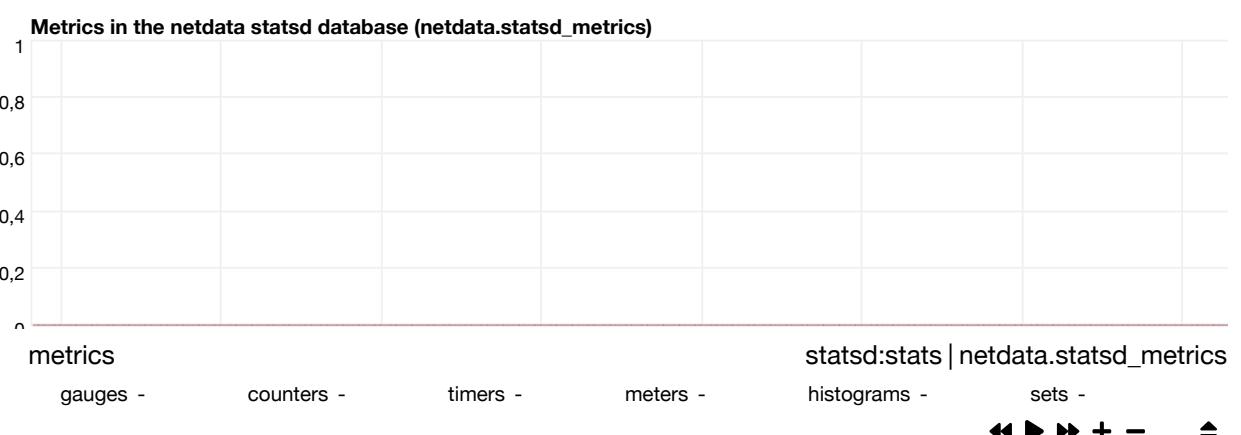
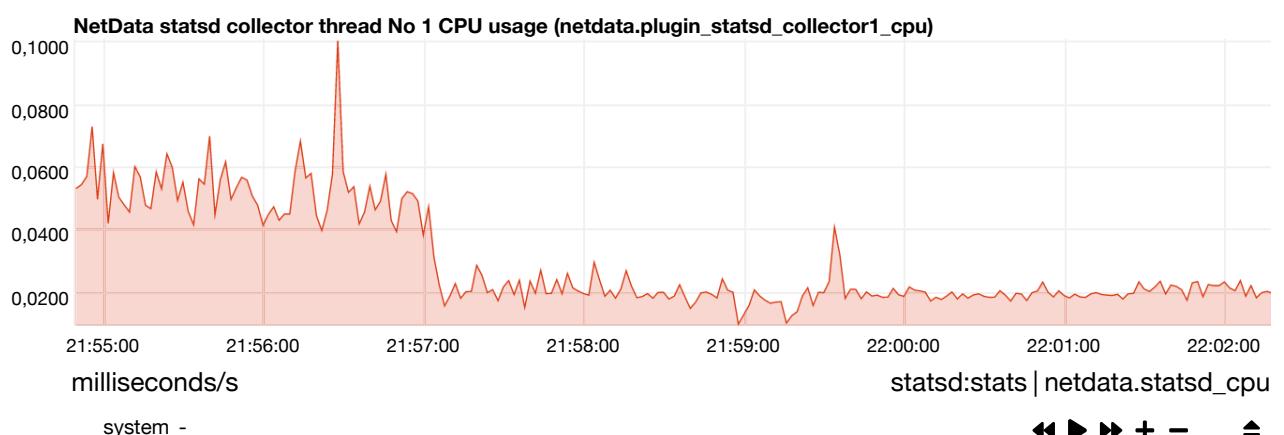
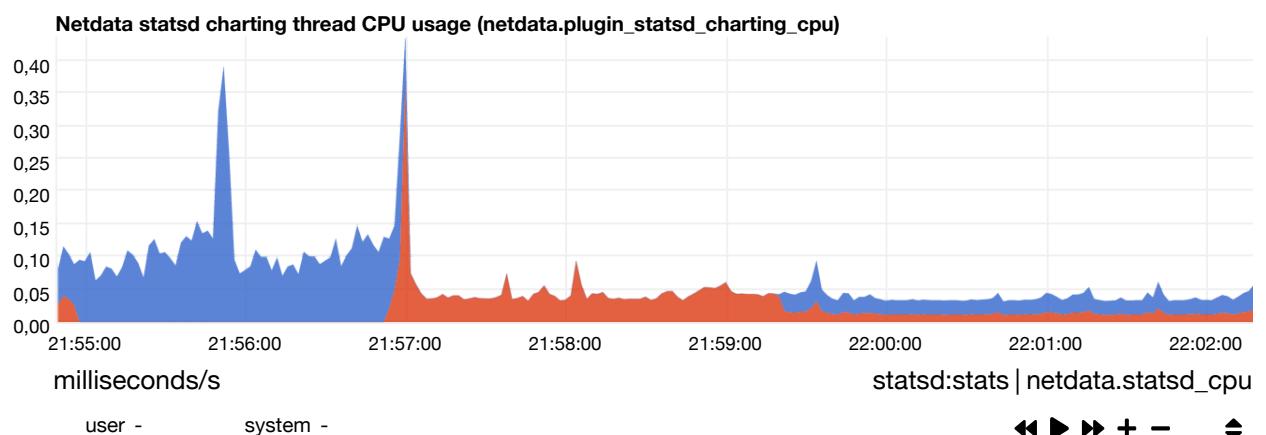


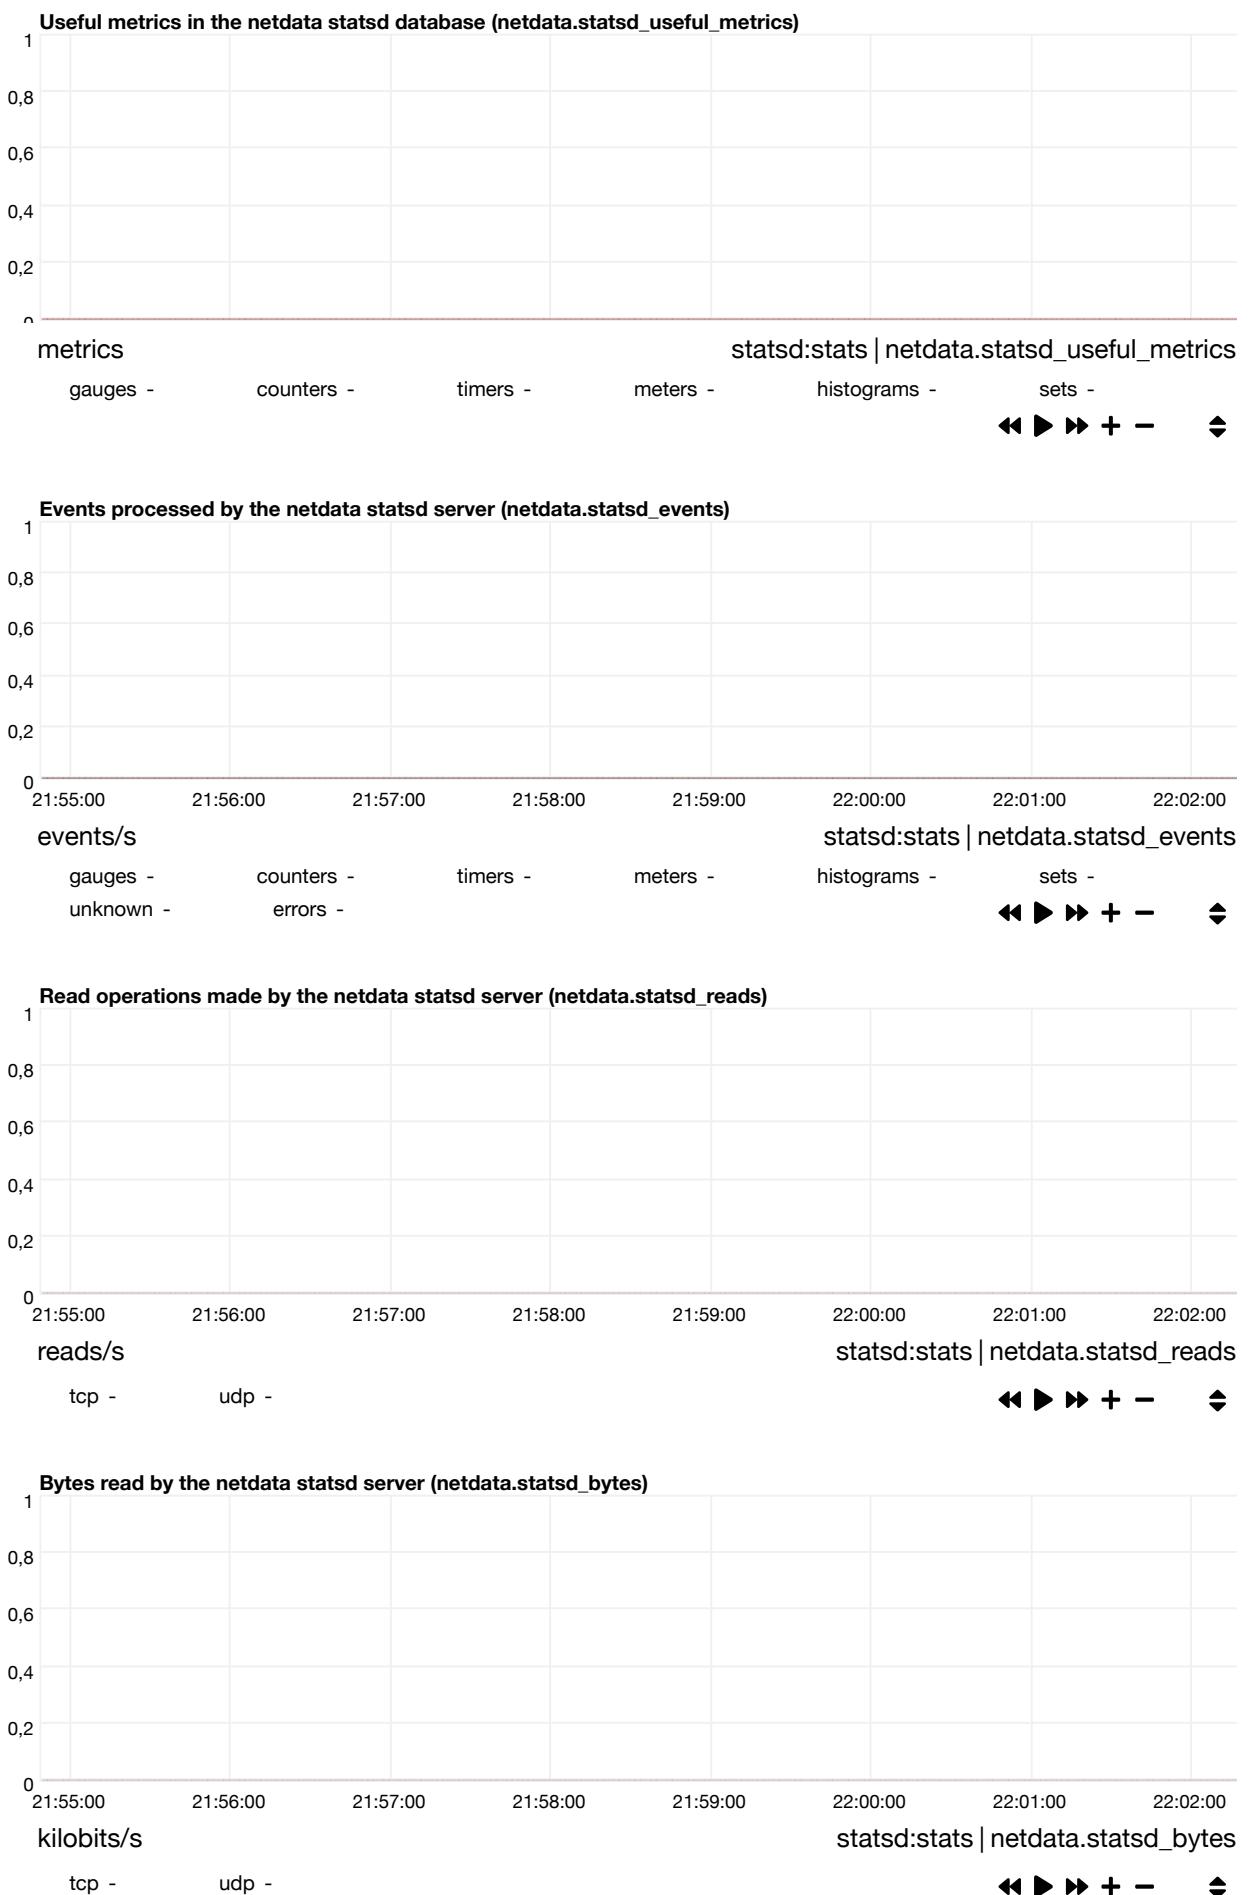
## web

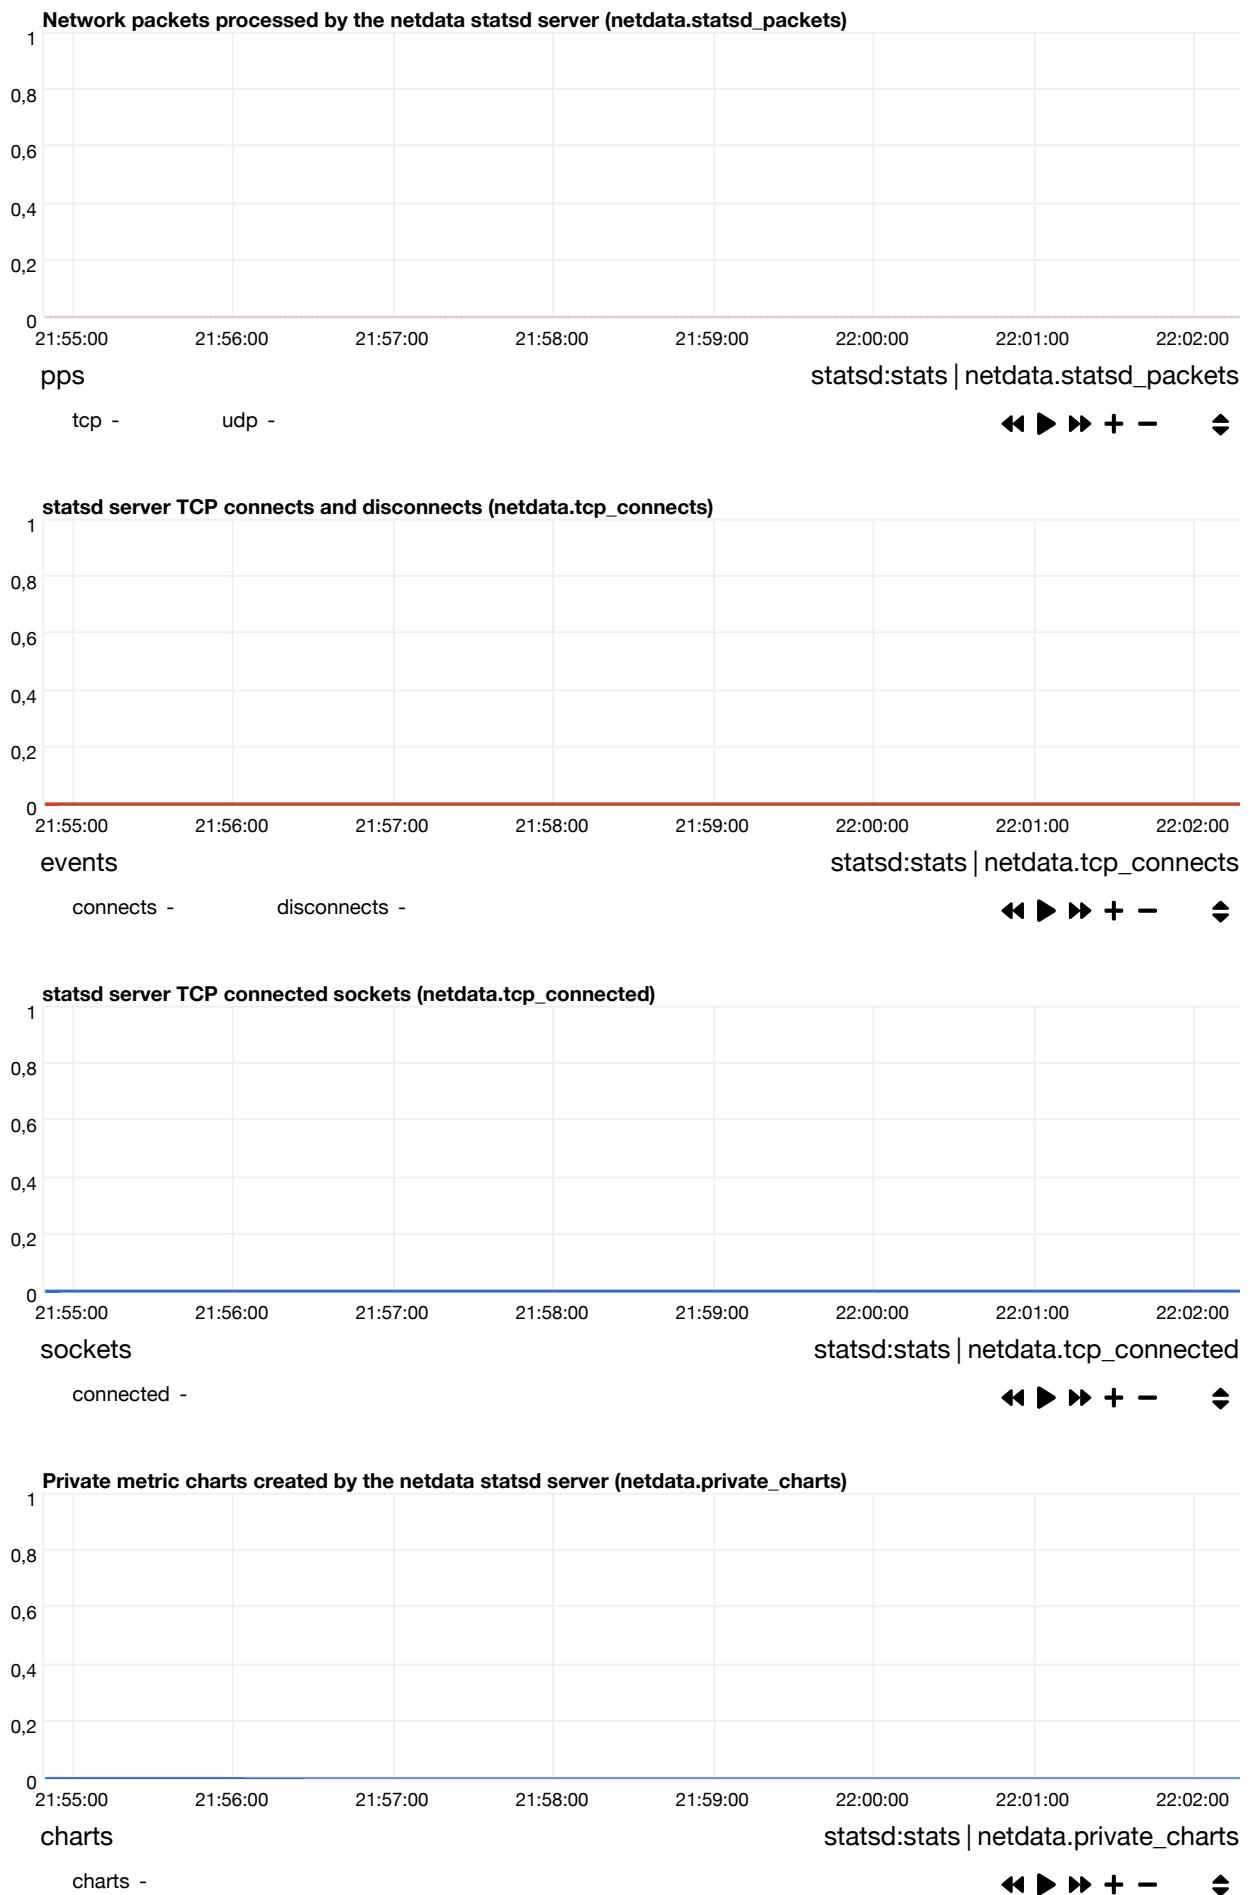




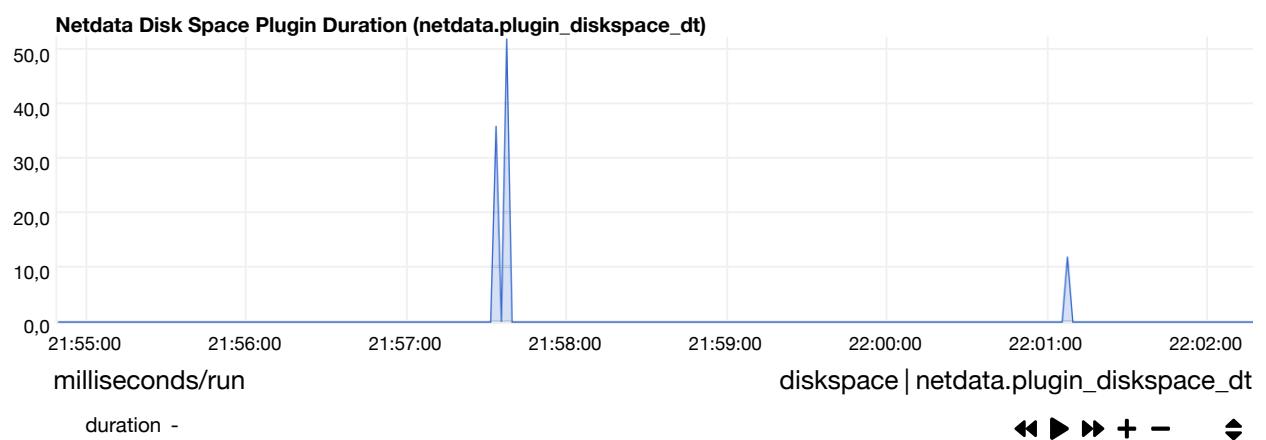
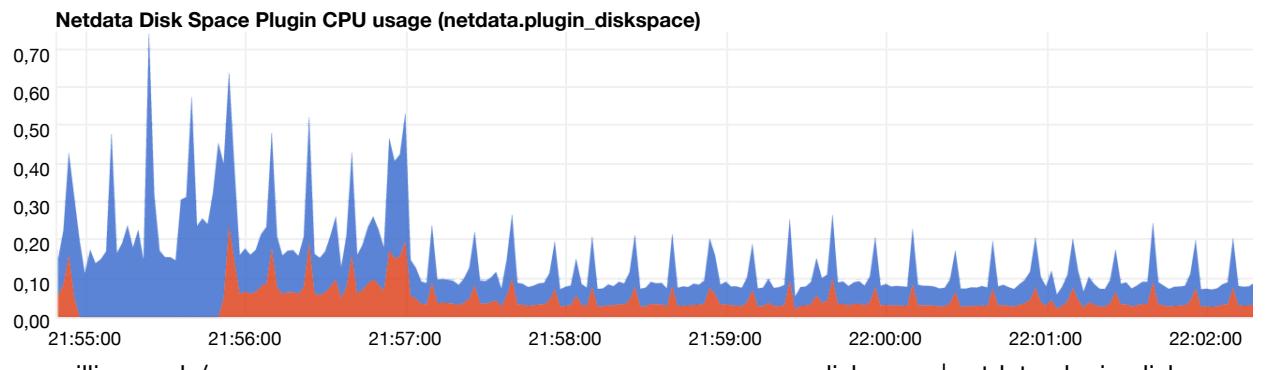
## statsd



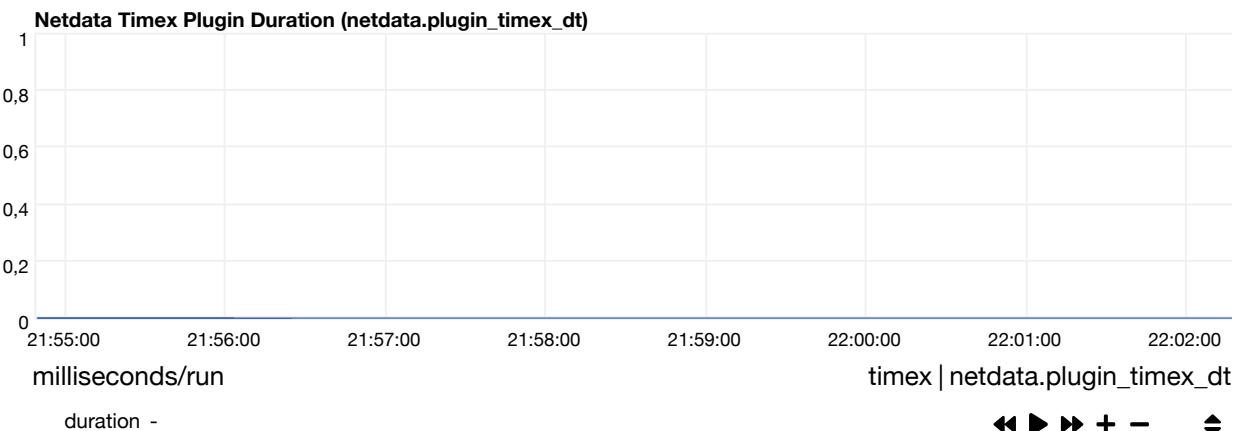
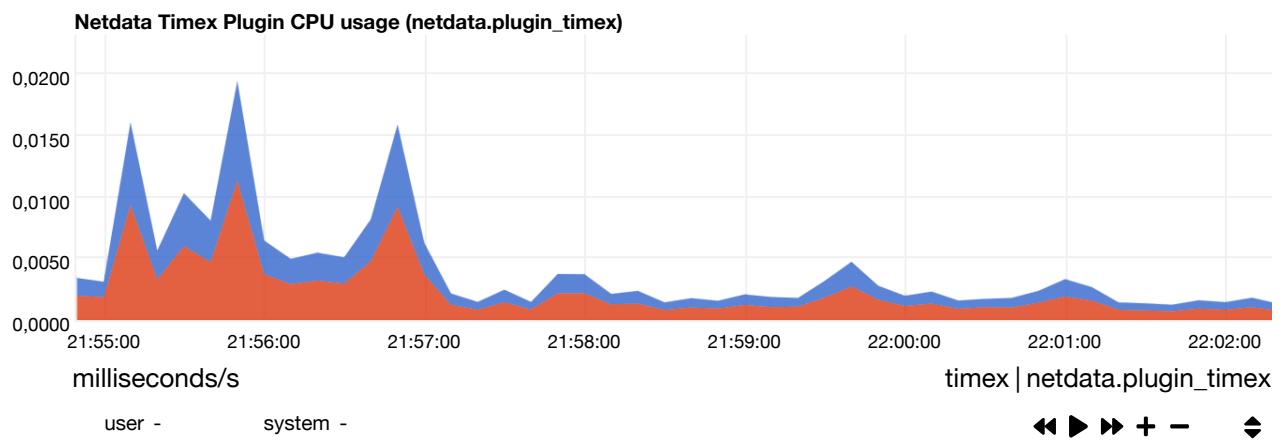




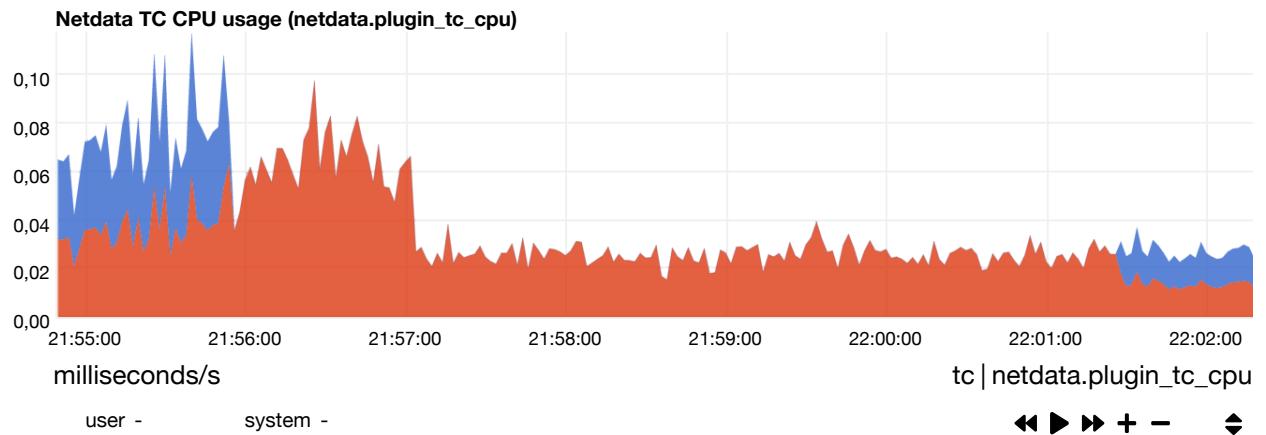
# diskspace



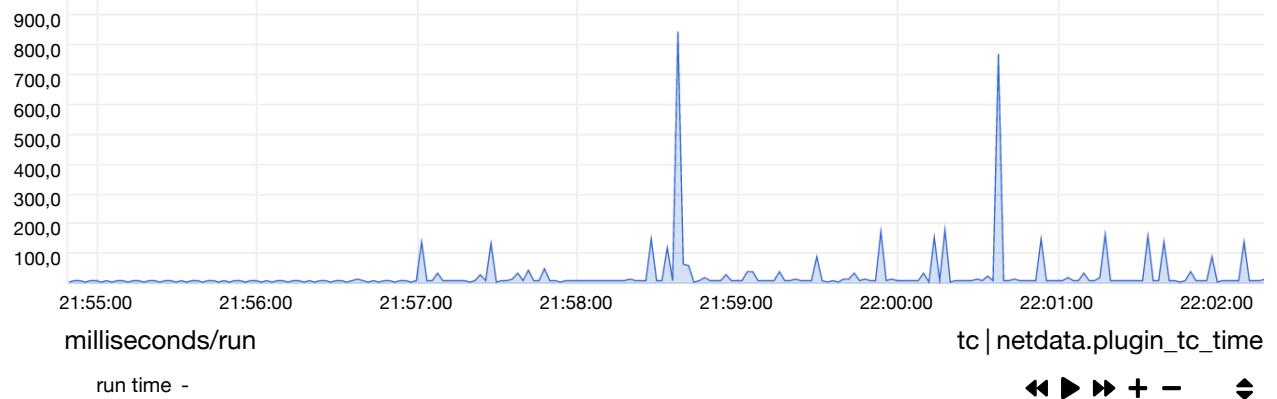
## timex



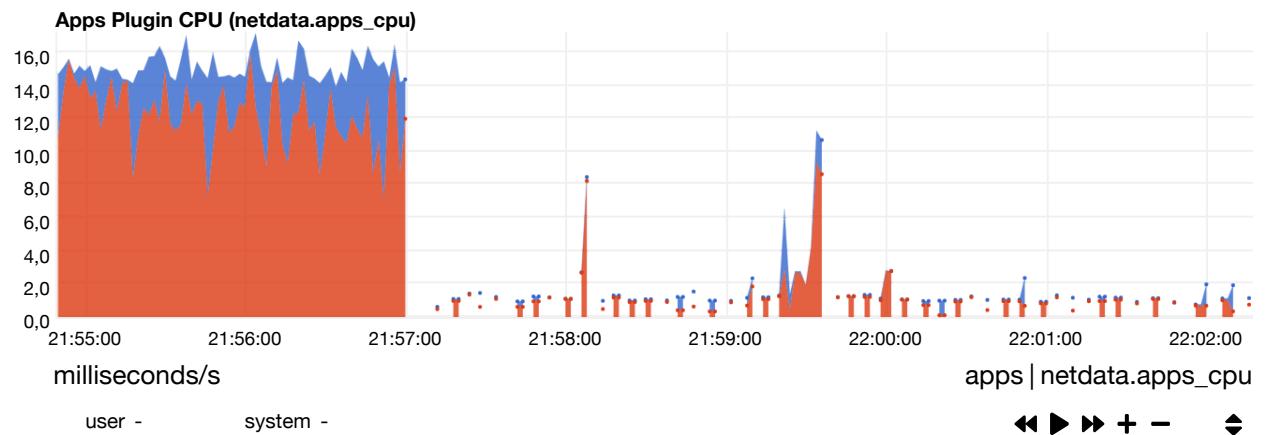
## tc.helper

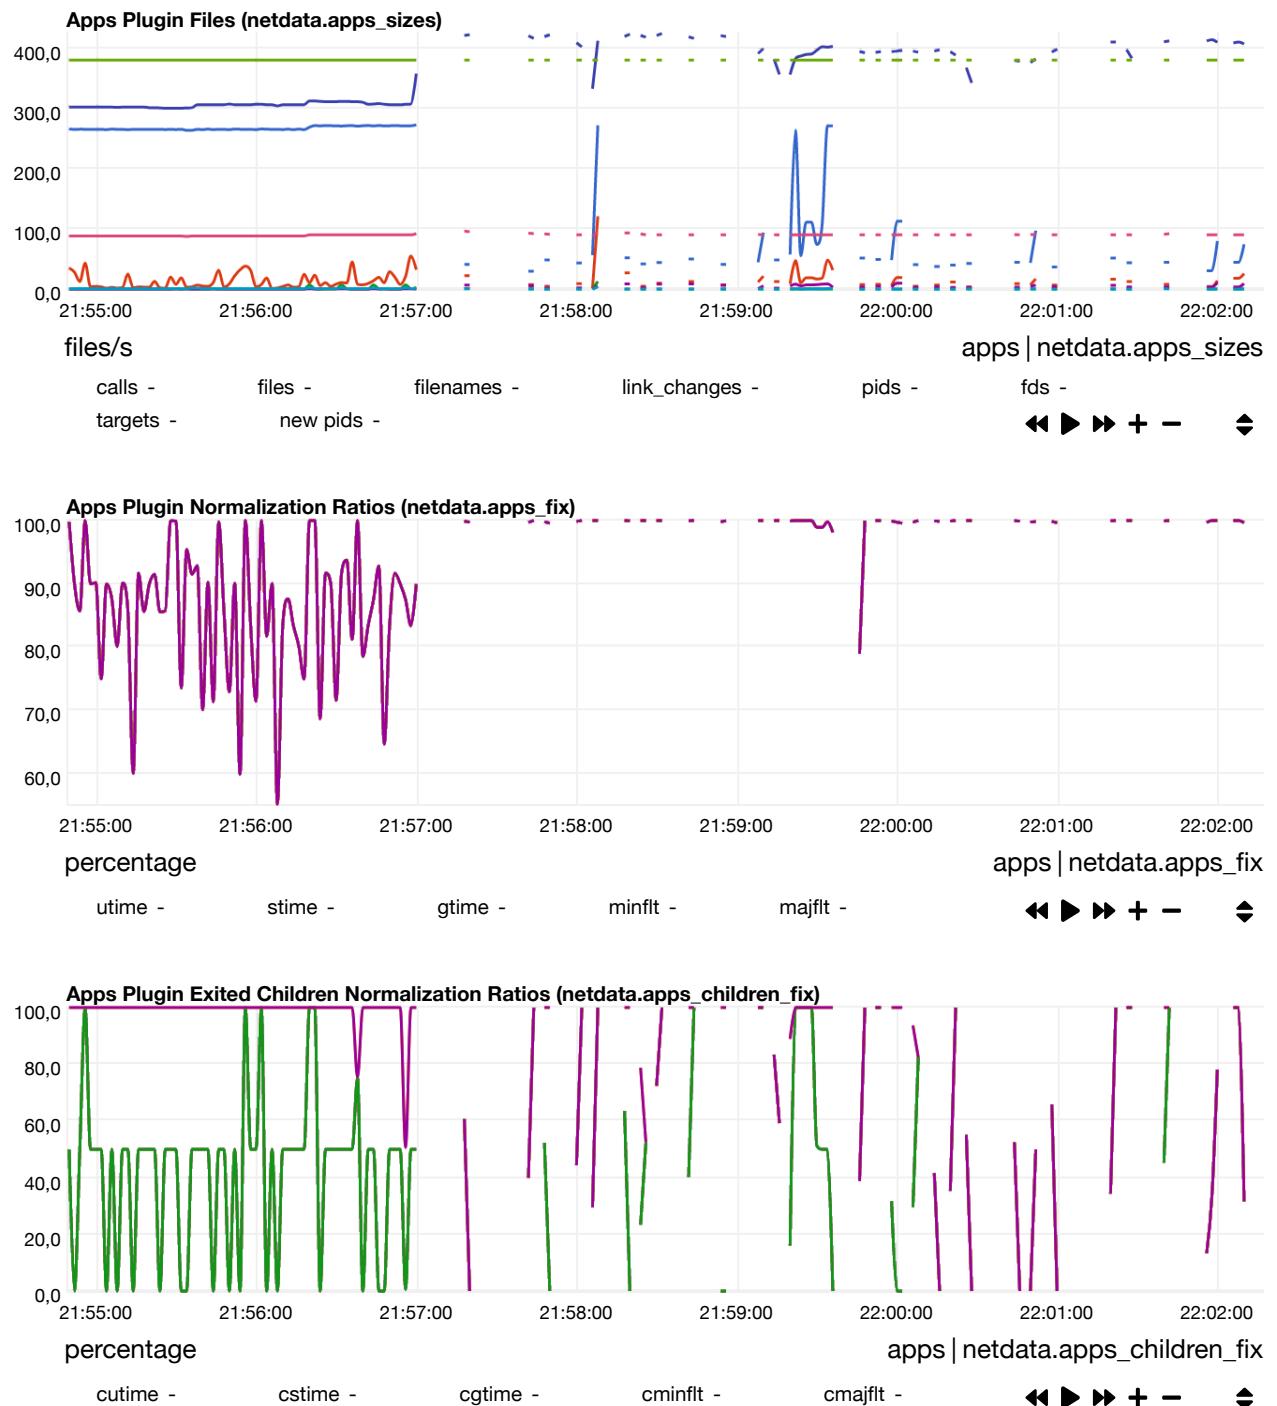


Netdata TC script execution (netdata.plugin\_tc\_time)



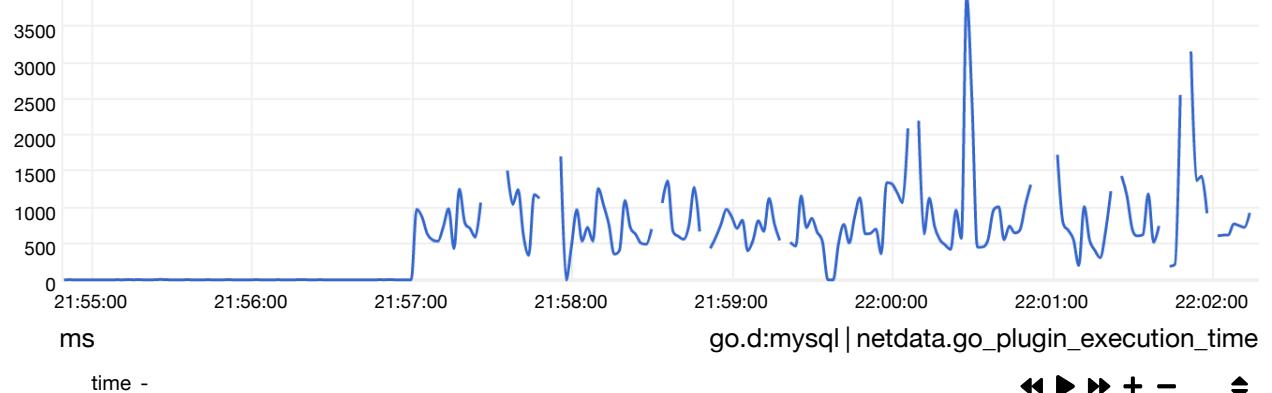
## apps.plugin

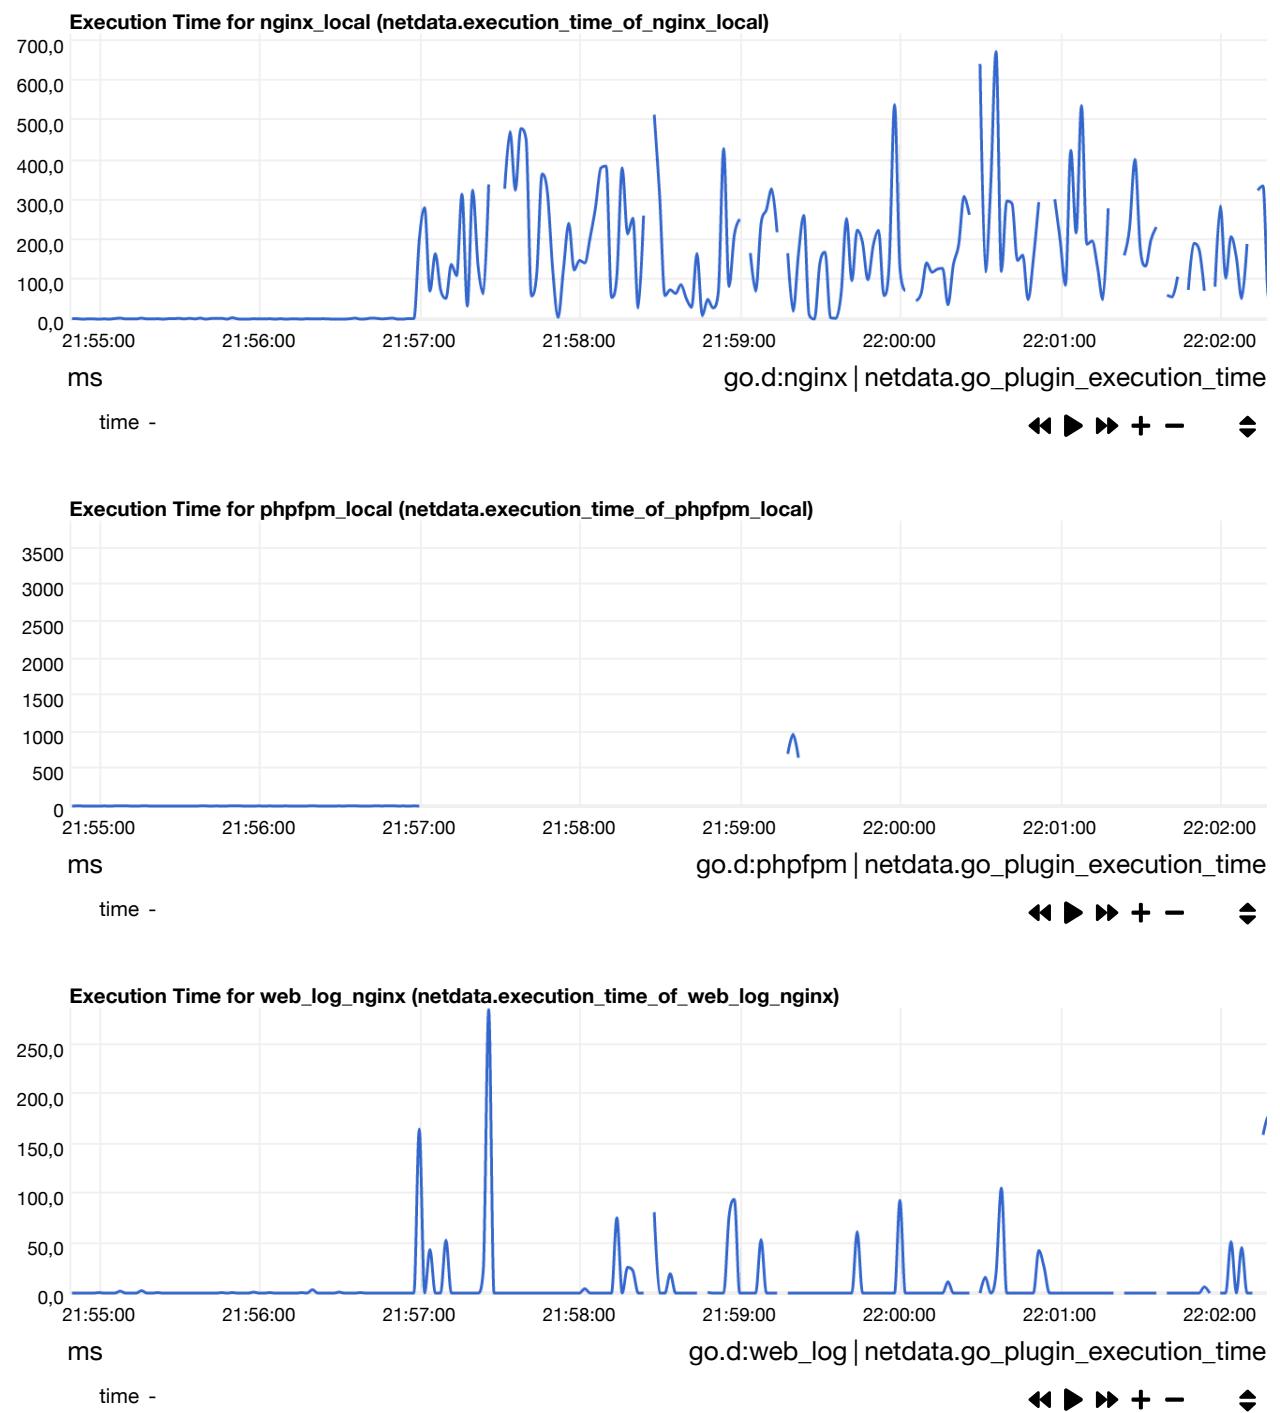




go.d

Execution Time for mysql\_local (netdata.execution\_time\_of\_mysql\_local)





Netdata (<https://github.com/netdata/netdata/wiki>)

© Copyright 2020, Netdata, Inc (<mailto:info@netdata.cloud>).

[Terms and conditions](https://www.netdata.cloud/terms/) (<https://www.netdata.cloud/terms/>).

[Privacy Policy](https://www.netdata.cloud/privacy/) (<https://www.netdata.cloud/privacy/>).

Released under GPL v3 or later (<http://www.gnu.org/licenses/gpl-3.0.en.html>). Netdata uses third party tools (<https://github.com/netdata/netdata/blob/master/REDISTRIBUTED.md>).